

Statistical Practice in Epidemiology with R

SPE practicals with solutions

Krista Fischer¹

Martyn Plummer²

Janne Pitkaniemi³

Bendix Carstensen⁴

Damien Georges⁵

Esa Läärä⁶

1-5 June 2026

¹Estonian Genome Center, University of Tartu, Estonia, Krista.Fischer@ut.ee

²University of Warwick, UK, martyn.plummer@warwick.ac.uk

³Finnish Cancer Registry, Helsinki, Finland, janne.pitkaniemi@cancer.fi

⁴Steno Diabetes Center Copenhagen, Gentofte, Denmark, Dept. of Biostatistics, University of Copenhagen, Denmark, b@bxc.dk, <http://BendixCarstensen.com>

⁵International Agency for Research on Cancer, Lyon, France, georgesd@iarc.who.int

⁶Department of Mathematical Sciences, University of Oulu, Finland, Esa.Laara@oulu.fi, <http://www.oulu.fi/university/researcher/esa-laara>

Contents

About	7
1 Practice with basic R	9
1.1 The working directory	9
1.2 The workspace	9
1.3 Using R as a calculator	10
1.4 Vectors	11
1.5 Sequences	12
1.6 Displaying and changing parts of a vector (indexing)	12
1.7 Lists	14
1.8 Data frames	15
1.9 Working with built-in data frames	17
1.10 Referencing parts of the data frame (indexing)	17
1.11 Summaries	18
1.12 Generating new variables	19
1.13 Turning a variable into a factor	19
1.14 Frequency tables	20
1.15 Grouping the values of a numeric variable	20
2 Reading data into R	23
2.1 Introduction	23
2.2 Data sources	24
2.3 Data in text files	24
2.4 Things that can go wrong	26
2.5 Forgetting the headers	26
2.6 Using the wrong separator	27
2.7 Mis-specifying the representation of missing values	27
2.8 Spreadsheet data	27
2.9 Reading data from the Internet	28
2.10 Reading from the clipboard	29
2.11 Binary data	29
2.12 Summary	31
3 Data manipulation with tidyverse	33
3.1 Introduction	33
3.2 The <code>births</code> data	33
3.3 <code>tibble</code> vs <code>data.frame</code>	33
3.4 Piping functions	35
3.5 <code>mutate</code> columns	35
3.6 <code>select</code> columns, <code>filter</code> and <code>arrange</code> rows	37
3.7 <code>group_by</code> and <code>summarise</code> data	38
3.8 Multiple grouping	40
3.9 Bind and join tables	44
3.10 Data Visualization with <code>ggplot2</code>	46

3.11	pivoting data with <code>tidyr</code>	47
3.12	reading files with <code>readr</code>	48
3.13	String manipulation with <code>stringr</code>	81
3.14	<code>purrr</code> package to apply functions to list	81
3.15	Bonus: Rendering tables	83
4	Tabulation	85
4.1	Introduction	85
4.2	The <code>births</code> data	85
4.3	One-way tables	86
4.4	Improving the Presentation of Tables	88
4.5	Two-way Tables	89
4.6	Printing	90
5	Graphics in R	93
5.1	Simple plot on the screen	93
5.2	Colours	99
5.3	Adding to a plot	101
5.4	Using indexing for plot elements	104
5.5	Generating colours	108
5.6	Saving your graphs for use in other documents	109
5.7	Interacting with a plot	110
6	Analysis of hazard rates, their ratios and differences and binary regression	113
6.1	Hand calculations for a single rate	113
6.2	Poisson model for a single rate with logarithmic link	113
6.3	Poisson model for a single rate with identity link	114
6.4	Poisson model assuming the same rate for several periods	114
6.5	Analysis of rate ratio	115
6.6	Analysis of rate difference	115
6.7	Binary regression	116
6.8	Optional/Homework: Hand calculations and calculations using matrix tools	116
7	Estimation of effects: simple and more complex	119
7.1	Response and explanatory variables	119
7.2	Data set <code>births</code>	120
7.3	Simple estimation with <code>lm()</code> and <code>glm()</code>	120
7.4	Stratified effects, and interaction or effect-measure modification	122
7.5	Controlling or adjusting for the effect of <code>hyp</code> for <code>sex</code>	122
7.6	Numeric exposure, simple linear regression and checking assumptions	123
7.7	Penalized spline model	123
7.8	Analysis of binary outcomes	125
8	Poisson regression & analysis of curved effects	129
8.1	Testis cancer: Data input and housekeeping	129
8.2	Some descriptive analysis	130
8.3	Age and period as categorical factors	131
8.4	Generalized additive model with penalized splines	133
9	Causal inference	137
9.1	Proper adjustment for confounding in regression models	137
9.2	DAG tools in the package <code>dagitty</code>	138
9.3	Identifying the true DAG for the data	139
9.4	Instrumental variables estimation: Mendelian randomization	143
9.5	Why are simulation exercises useful for causal inference?	145
10	Graphics with <code>ggplot2</code>	147
10.1	Data	148

10.2 Building the plot	148
10.3 Self study	157
10.4 Adding the table (advanced topic)	157
10.5 References	159
11 Survival analysis with competing risks: Oral cancer patients	161
11.1 Description of the data	161
11.2 Loading the packages and the data	161
11.3 Total mortality: Kaplan–Meier analyses	162
11.4 Total mortality by stage	165
11.5 Event-specific cumulative mortality curves	168
11.6 Regression modelling of overall mortality.	173
11.7 Modelling event-specific hazards	176
11.8 Lexis object with multi-state set-up	177
12 Time-splitting, time-scales and SMR	179
12.1 Age-specific mortality	181
12.2 Further time scales: period and duration	184
12.3 SMR	188
12.4 SMR modeling	190
13 Nested case-control study and case-cohort study: Risk factors of coronary heart disease	195
13.1 Reading the cohort data, illustrating the study base and risk sets	196
13.2 Nested case-control study	200
13.3 Case-cohort study	203
13.4 Full cohort analysis and comparisons	206
13.5 Further exercises and homework	208
14 Causal inference 2: Model-based estimation of causal estimands	209
14.1 Introduction	209
14.2 Control of confounding	210
14.3 Generation of target population and true models	211
14.4 Factual and counterfactual risks - associational and causal contrasts	211
14.5 Outcome modelling and estimation of causal contrasts by g-formula	212
14.6 Inverse probability weighting (IPW) by propensity scores	214
14.7 Improving IPW estimation and using R package PSweight	215
14.8 Effect of exposure among those exposed	217
14.9 Double robust estimation by augmented IPW	218
14.10 Double robust, targeted maximum likelihood estimation (TMLE)	219
14.11 Double robust estimation with AIPW and tmle packages	219
15 Time-dependent variables and multiple states	221
15.1 The renal failure dataset	221
15.2 Splitting the follow-up time	228

About

This booklet contains a set of exercises to apply concepts learned during the SPE-R course.

Chapter 1

Practice with basic R

The main purpose of this session is to give participants who have not had much (or any) experience with using R a chance to practice the basics and to ask questions. For others, it should serve as a reminder of some of the basic features of the language.

R can be installed on all major platforms (*i.e.* Windows, macOS, Linux). We do not assume in this exercise that you are using any particular platform. Many people like to use the RStudio graphical user interface (GUI), which gives the same look and feel across all platforms.

1.1 The working directory

A key concept in R is the *working directory* (or *folder* in the terminology of Windows). The working directory is where R looks for data files that you may want to read in and where R will write out any files you create. It is a good idea to keep separate working directories for different projects. In this course we recommend that you keep a separate working directory for each exercise.

If you are working on the command line in a terminal, then you can change to the correct working directory and then launch R by typing *R*.

If you are using a GUI then you will typically need to change to the correct working directory after starting R. In RStudio, you can change directory from the *Session* menu. However it is much more useful to create a new *project* to keep your source files and data. When you open a project in the RStudio GUI, your working directory is automatically changed to the directory associated with the project.

You can display the current working directory with the `getwd()` (*get working directory*) function and set it with the `setwd()` (*set working directory*) function. The function `dir()` can be used to see what files you have in the working directory.

1.2 The workspace

You can quit R by typing

```
q()
```

at the R command prompt. You will be asked if you want to save your workspace. We strongly recommend that you answer *no* to this question. If you answer *yes* then R will write a file named `.RData` into the working directory containing all the objects you created during your session. This file will be automatically loaded the next time you start R and this will restore all the objects in your workspace.

It may seem convenient to keep your R objects from one session to another. But this has many disadvantages.

- You may not remember how an object was created. This becomes a problem if you need to redo your analysis after the data has been changed or updated, or if you accidentally delete the object.

- An object might be modified or overwritten. In this case your analysis will give you a different answer but you will not know why. You may not even notice that the answer has changed.
- It becomes impossible to clean your workspace if you cannot remember which objects are required by which analyses. As a result, your workspace will become cluttered with old objects.

We strongly recommend that you follow some basic principles of reproducible research.

- Always start with a clean (empty) workspace.
- Read in the data you need from a pristine source.
- Put your R commands in a script file so that they can be run again in a future session.
- All modifications to the data that you need to make should be done in R using a script and not by editing the data source. This way, if the original data is modified or updated then you can run the script again on the updated data.

1.3 Using R as a calculator

Try using R as an interactive calculator by typing different arithmetic expressions on the command line. Pressing the return key on the command line finishes the expression. R will then evaluate the expression and print out the result.

Note that R allows you to recall previous commands using the vertical arrow key. You can edit a recalled command and then resubmit it by pressing the return key. Keeping that in mind, try the following:

```
12 + 16
(12 + 16) * 5
sqrt((12 + 16) * 5) # square root
round(sqrt((12 + 16) * 5), 2) # round to two decimal places
```

The hash symbol # denotes the start of a *comment*. Anything after the hash is ignored by R.

Round braces are used a lot in R. In the above expressions, they are used in two different ways. Firstly, they can be used to establish the order of operations. In the example

```
(12 + 16) * 5
```

```
## [1] 140
```

they ensure that 12 is added to 16 before the result is multiplied by 5. If you omit the braces then you get a different answer

```
12 + 16 * 5
```

```
## [1] 92
```

because multiplication has higher *precedence* than addition. The second use of round braces is in a function call (e.g. `sqrt`, `round`). To call a function in R, type the name followed by the arguments inside round braces. Some functions take multiple arguments, and in this case they are separated by commas.

You can see that complicated expressions in R can have several levels of nested braces. To keep track of these, it helps to use a syntax-highlighting editor. For example, in RStudio, when you type an opening bracket (, RStudio will automatically add a closing bracket), and when the cursor moves past a closing bracket, RStudio will automatically highlight the corresponding opening bracket. Features like this can make it much easier to write R code free from syntax errors.

Instead of printing the result to the screen, you can store it in an object, say

```
a <- round(sqrt((12 + 16) * 5), 2)
```

In this case R does not print anything to the screen. You can see the results of the calculation, stored in the object `a`, by typing `a` and also use `a` for further calculations, e.g:

```
exp(a)
log(a) # natural logarithm
log10(a) # log to the base 10
```

The left arrow expression `<-`, pronounced *gets*, is called the assignment operator, and is obtained by typing `<` followed by `-` (with no space in between). It is also possible to use the equals sign `=` for assignment.

Note that object names in R are case sensitive. So you can assign different values to objects named `A` and `a`.

1.4 Vectors

All commands in R are *functions* which act on *objects*. One important kind of object is a *vector*, which is an ordered collection of numbers, or character strings (e.g. *Charles Darwin*), or logical values (`TRUE` or `FALSE`). The components of a vector must be of the same type (numeric, character, or logical). The combine function `c()`, together with the assignment operator, is used to create vectors. Thus

```
v <- c(4, 6, 1, 2.2)
```

creates a vector `v` with components 4, 6, 1, 2.2 and assigns the result to the vector `v`.

A key feature of the R language is that many operations are *vectorized*, meaning that you can carry out the same operation on each element of a vector in a single operation. Try

```
v
```

```
## [1] 4.0 6.0 1.0 2.2
```

```
3 + v
```

```
## [1] 7.0 9.0 4.0 5.2
```

```
3 * v
```

```
## [1] 12.0 18.0 3.0 6.6
```

and you will see that R understands what to do in each case.

R extends ordinary arithmetic with the concept of a *missing value* represented by the symbol `NA` (*Not Available*). Any operation on a missing value creates another missing value. You can see this by repeating the same operations on a vector containing a missing value:

```
v <- c(4, 6, NA)
```

```
3 + v
```

```
## [1] 7 9 NA
```

```
3 * v
```

```
## [1] 12 18 NA
```

The fact that every operation on a missing value produces a missing value can be a nuisance when you want to create a summary statistic for a vector:

```
mean(v)
```

```
## [1] NA
```

While it is true that the mean of `v` is unknown because the value of the third element is missing, we normally want the mean of the non-missing elements. Fortunately the `mean` function has an optional argument called `na.rm` which can be used for this.

```
mean(v, na.rm = TRUE)
```

```
## [1] 5
```

Many functions in R have optional arguments that can be omitted, in which case they take their default value (For example, the `mean` function has default `na.rm=FALSE`). You can explicitly values to optional arguments in the function call to override the default behaviour.

You can get a description of the structure of any object using the function `str()`. For example, `str(v)` shows that `v` is numeric with 4 components. If you just want to know the length of a vector then it is much easier to

use the `length` function.

```
length(v)
```

```
## [1] 3
```

1.5 Sequences

There are short-cut functions for creating vectors with a regular structure. For example, if you want a vector containing the sequence of integers from 1 to 10, you can use

```
1:10
```

```
## [1] 1 2 3 4 5 6 7 8 9 10
```

The `seq()` function allows the creation of more general sequences. For example, the vector (15, 20, 25, ... ,85) can be created with

```
seq(from = 15, to = 85, by = 5)
```

```
## [1] 15 20 25 30 35 40 45 50 55 60 65 70 75 80 85
```

The objects created by the `:` operator and the `seq()` function are ordinary vectors, and can be combined with other vectors using the combine function:

```
c(5, seq(from = 20, to = 85, by = 5))
```

```
## [1] 5 20 25 30 35 40 45 50 55 60 65 70 75 80 85
```

You can learn more about functions by typing `?` followed by the function name. For example `?seq` gives information about the syntax and usage of the function `seq()`.

1. Create a vector `w` with components 1, -1, 2, -2
2. Display this vector

```
## [1] 1 -1 2 -2
```

3. Obtain a description of `w` using `str()`

```
## num [1:4] 1 -1 2 -2
```

4. Create the vector `w+1`, and display it.

```
## [1] 2 0 3 -1
```

5. Create the vector `v` with components (5, 10, 15, ... , 75) using `seq()`.

```
## [1] 5 10 15 20 25 30 35 40 45 50 55 60 65 70 75
```

6. Now add the components 0 and 1 to the beginning of `v` using `c()`.

```
## [1] 0 1 5 10 15 20 25 30 35 40 45 50 55 60 65 70 75
```

7. Find the length of this vector.

```
## [1] 17
```

1.6 Displaying and changing parts of a vector (indexing)

Square brackets in R are used to extract parts of vectors. So `x[1]` gives the first element of vector `x`. Since R is vectorized you can also supply a vector of integer index values inside the square brackets. Any expression that creates an integer vector will work.

Try the following commands:

```
x <- c(2, 7, 0, 9, 10, 23, 11, 4, 7, 8, 6, 0)
x[4]
```

```
## [1] 9
```

```
x[3:5]
```

```
## [1] 0 9 10
```

```
x[c(1, 5, 8)]
```

```
## [1] 2 10 4
```

Trying to extract an element that is beyond the end of the vector is, surprisingly, not an error. Instead, this returns a missing value

```
N <- length(x)
```

```
x[N + 1]
```

```
## [1] NA
```

There is a reason for this behaviour, which we will discuss in the recap.

R also allows *logical subscripting*. Try the following

```
x > 10
```

```
## [1] FALSE FALSE FALSE FALSE FALSE TRUE TRUE FALSE FALSE FALSE FALSE
```

```
x[x > 10]
```

```
## [1] 23 11
```

The first expression creates a logical vector of the same length as `x`, where each element has the value `TRUE` or `FALSE` depending on whether or not the corresponding element of `x` is greater than 10. If you supply a logical vector as an index, R selects only those elements for which the conditions is `TRUE`.

You can combine two logical vectors with the operators `&` (logical and) and `|` (logical or). For example, to select elements of `x` that are between 10 and 20 we combine two one-sided logical conditions for $x \geq 10$ and $x \leq 20$:

```
x[x >= 10 & x <= 20]
```

```
## [1] 10 11
```

The remaining elements of `x` that are *either* less than 10 *or* greater than 20 are selected with

```
x[x < 10 | x > 20]
```

```
## [1] 2 7 0 9 23 4 7 8 6 0
```

Indexing can also be used to replace parts of a vector:

```
x[1] <- 1000
```

```
x
```

```
## [1] 1000 7 0 9 10 23 11 4 7 8 6 0
```

This replaces the first element of `x`. Logical subscripting is useful for replacing parts of a vector that satisfy a certain condition. For example to replace all elements that take the value 0 with the value 1:

```
x[x == 0] <- 1
```

```
x
```

```
## [1] 1000 7 1 9 10 23 11 4 7 8 6 1
```

If you want to replace parts of a vector then you need to make sure that the replacement value is either a single value, as in the example above, or a vector equal in length to the number of elements to be replaced. For example, to replace elements 2, 3, and 4 we need to supply a vector of replacement values of length 3.

```
x[2:4] <- c(0, 8, 1)
```

```
x
```

```
## [1] 1000 0 8 1 10 23 11 4 7 8 6 1
```

It is important to remember this when you are using logical subscripting because the number of elements to be replaced is not given explicitly in the R code, and it is easy to get confused about how many values need to be replaced. If we want to add 3 to every element that is less than 3 then we can break the operation down into 3 steps:

```
y <- x[x < 3]
y <- y + 3
x[x < 3] <- y
x
```

```
## [1] 1000 3 8 4 10 23 11 4 7 8 6 4
```

First we extract the values to be modified, then we modify them, then we write back the modified values to the original positions. R experts will normally do this in a single expression.

```
x[x < 3] <- x[x < 3] + 3
```

Remember, if you are confused by a complicated expression you can usually break it down into simpler steps.

If you want to create an entirely new vector based on some logical condition then use the `ifelse()` function. This function takes three arguments: the first is a logical vector; the second is the value taken by elements of the logical vector that are `TRUE`; and the third is the value taken by elements that are `FALSE`.

In this example, we use the remainder operator `%%` to identify elements of `x` that have value 0 when divided by 2 (i.e. the even numbers) and then create a new character vector with the labels *even* and *odd*:

```
x %% 2
```

```
## [1] 0 1 0 0 0 1 1 0 1 0 0 0
```

```
ifelse(x %% 2 == 0, "even", "odd")
```

```
## [1] "even" "odd" "even" "even" "even" "odd" "odd" "even" "odd" "even"
## [11] "even" "even"
```

Now try the following:

1. Display elements that are less than 10, but greater than 4
2. Modify the vector `x`, replacing by 10 all values that are greater than 10
3. Modify the vector `x`, multiplying by 2 all elements that are smaller than 5 (Remember you can do this in steps).

1.7 Lists

Collections of components of different types are called *lists*, and are created with the `list()` function. Thus

```
m <- list(4, TRUE, "name of company")
m
```

```
## [[1]]
## [1] 4
##
## [[2]]
## [1] TRUE
##
## [[3]]
## [1] "name of company"
```

creates a list with 3 components: the first is numeric, the second is logical and the third is character. A list element can be any object, including another list. This flexibility means that functions that need to return a lot of complex information, such as statistical modelling functions, often return a list.

As with vectors, single square brackets are used to take a subset of a list, but the result will always be another list, even if you select only one element

```
m[1:2] # A list containing first two elements of m
```

```
## [[1]]
## [1] 4
##
## [[2]]
## [1] TRUE
```

```
m[3] # A list containing the third element of m
```

```
## [[1]]
## [1] "name of company"
```

If you just want to extract a single element of a list then you must use double square braces:

```
m[[3]] # Extract third element
```

```
## [1] "name of company"
```

Lists are more useful when their elements are named. You can name an element by using the syntax `name=value` in the call to the `list` function:

```
mylist <- list(
  name = c("Joe", "Ann", "Jack", "Tom"),
  age = c(34, 50, 27, 42)
)
mylist
```

```
## $name
## [1] "Joe" "Ann" "Jack" "Tom"
##
## $age
## [1] 34 50 27 42
```

This creates a new list with the elements *name*, a character vector of names, and *age* a numeric vector of ages. The components of the list can be extracted with a dollar sign `$`

```
mylist$name
```

```
## [1] "Joe" "Ann" "Jack" "Tom"
```

```
mylist$age
```

```
## [1] 34 50 27 42
```

1.8 Data frames

Data frames are a special structure used when we want to store several vectors of the same length, and corresponding elements of each vector refer to the same record. For example, here we create a simple data frame containing the names of some individuals along with their age in years, their sex (coded 1 or 2) and their height in cm.

```
mydata <- data.frame(
  name = c("Joe", "Ann", "Jack", "Tom"),
  age = c(34, 50, 27, 42),
  sex = c(1, 2, 1, 1),
  height = c(185, 170, 175, 182)
)
```

The construction of a data frame is just like a named list (except that we use the constructor function `data.frame` instead of `list`). In fact data frames are also lists so, for example, you can extract vectors using the dollar sign:

```
mydata$height
```

```
## [1] 185 170 175 182
```

On the other hand, data frames are also two dimensional objects:

```
mydata
```

```
##   name age sex height
## 1  Joe  34   1    185
## 2  Ann  50   2    170
## 3 Jack  27   1    175
## 4  Tom  42   1    182
```

When you print a data frame, each variable appears in a separate column. You can use square brackets with two comma-separated arguments to take subsets of rows or columns.

```
mydata[1, ]
```

```
##   name age sex height
## 1  Joe  34   1    185
```

```
mydata[, c("age", "height")]
```

```
##   age height
## 1  34    185
## 2  50    170
## 3  27    175
## 4  42    182
```

```
mydata[2, 4]
```

```
## [1] 170
```

We will look into indexing of data frames in more detail below.

Now let's create another data frame with more individuals than the first one:

```
yourdata <- data.frame(
  name = c("Ann", "Peter", "Sue", "Jack", "Tom", "Joe", "Jane"),
  weight = c(67, 81, 56, 90, 72, 79, 69)
)
```

This new data frame contains the weights of the individuals. The two data sets can be joined together with the `merge` function.

```
newdata <- merge(mydata, yourdata)
newdata
```

```
##   name age sex height weight
## 1  Ann  50   2    170     67
## 2 Jack  27   1    175     90
## 3  Joe  34   1    185     79
## 4  Tom  42   1    182     72
```

The `merge` function uses the variables common to both data frames – in this case the variable *name* – to uniquely identify each row. By default, only rows that are in both data frames are preserved, the rest are discarded. In the above example, the records for Peter, Sue, and Jane, which are not in `mydata` are discarded. If you want to keep them, use the optional argument `all=TRUE`.

```
newdata <- merge(mydata, yourdata, all = TRUE)
newdata
```

```
##   name age sex height weight
```

```
## 1  Ann  50  2   170   67
## 2  Jack 27  1   175   90
## 3  Jane NA NA    NA   69
## 4  Joe  34  1   185   79
## 5 Peter NA NA    NA   81
## 6  Sue  NA NA    NA   56
## 7  Tom  42  1   182   72
```

This keeps a row for all individuals but since Peter, Sue and Jane have no recorded age, height, or sex these are missing values.

1.9 Working with built-in data frames

We shall use the `births` data which concern 500 mothers who had singleton births (i.e. no twins) in a large London hospital. The outcome of interest is the birth weight of the baby, also dichotomised as normal or low birth weight. These data are available in the `Epi` package:

```
library(Epi)
data(births)
objects()
```

```
## [1] "births"      "exercise"    "m"           "mydata"
## [5] "mylist"      "N"           "newdata"     "solution"
## [9] "spe_solutions" "v"          "w"           "x"
## [13] "y"           "yourdata"
```

The function `objects()` shows what is in your workspace. To find out a bit more about `births` try

```
help(births)
```

The dataframe `"diet"` in the `Epi` package contains data from a follow-up study with coronary heart disease as the end-point.

1. Load these data with

```
data(diet)
```

and print the contents of the data frame to the screen.. 2. Check that you now have two objects, `births`, and `diet` in your workspace. 3. Get help on the object `diet`. 4. Remove the object `diet` with the command

```
remove(diet)
```

5. Check that the object `diet` is no longer in your workspace.

1.10 Referencing parts of the data frame (indexing)

Typing `births` will list the entire data frame – not usually very helpful. You can use the `head` function to see just the first few rows of a data frame

```
head(births)
```

```
##   id bweight lowbw gestwks preterm matage hyp sex
## 1  1   2974     0   38.52      0     34  0  2
## 2  2   3270     0    NA     NA     30  0  1
## 3  3   2620     0   38.15      0     35  0  2
## 4  4   3751     0   39.80      0     31  0  1
## 5  5   3200     0   38.89      0     33  1  1
## 6  6   3673     0   40.97      0     33  0  2
```

Now try

```
births[1, ]
```

```
##   id bweight lowbw gestwks preterm matage hyp sex
## 1  1   2974     0  38.52      0    34  0  2
```

This will list all the values for the first row. Similarly,

```
births[2, ]
```

```
##   id bweight lowbw gestwks preterm matage hyp sex
## 2  2   3270     0    NA      NA    30  0  1
```

will list the value taken by the second row, and so on. To list the data for the first 10 subjects, try

```
births[1:10, ]
```

```
##   id bweight lowbw gestwks preterm matage hyp sex
## 1  1   2974     0  38.52      0    34  0  2
## 2  2   3270     0    NA      NA    30  0  1
## 3  3   2620     0  38.15      0    35  0  2
## 4  4   3751     0  39.80      0    31  0  1
## 5  5   3200     0  38.89      0    33  1  1
## 6  6   3673     0  40.97      0    33  0  2
## 7  7   3628     0  42.14      0    29  0  2
## 8  8   3773     0  40.21      0    37  0  1
## 9  9   3960     0  42.03      0    36  0  2
## 10 10   3405     0  39.33      0    39  0  1
```

Often we want to extract rows of a data frame based on a condition. To select all subjects with height less than 180 cm from the data frame `mydata` we can use the `subset()` function.

```
subset(mydata, height < 180)
```

```
##   name age sex height
## 2  Ann  50  2   170
## 3  Jack  27  1   175
```

1.11 Summaries

A good way to start an analysis is to ask for a summary of the data by typing

```
summary(births)
```

```
##           id           bweight           lowbw           gestwks           preterm
## Min.      : 1.0   Min.      : 628   Min.      :0.00   Min.      :24.69   Min.      :0.0000
## 1st Qu.:125.8   1st Qu.:2862   1st Qu.:0.00   1st Qu.:37.94   1st Qu.:0.0000
## Median :250.5   Median :3188   Median :0.00   Median :39.12   Median :0.0000
## Mean    :250.5   Mean    :3137   Mean     :0.12   Mean    :38.72   Mean     :0.1286
## 3rd Qu.:375.2   3rd Qu.:3551   3rd Qu.:0.00   3rd Qu.:40.09   3rd Qu.:0.0000
## Max.    :500.0   Max.    :4553   Max.     :1.00   Max.    :43.16   Max.     :1.0000
##
##                NAs      :10           NAs      :10
##           matage           hyp           sex
## Min.      :23.00   Min.      :0.000   Min.      :1.000
## 1st Qu.:31.00   1st Qu.:0.000   1st Qu.:1.000
## Median :34.00   Median :0.000   Median :1.000
## Mean    :34.03   Mean     :0.144   Mean     :1.472
## 3rd Qu.:37.00   3rd Qu.:0.000   3rd Qu.:2.000
## Max.    :43.00   Max.     :1.000   Max.     :2.000
##
```

This prints some summary statistics (minimum, lower quartile, mean, median, upper quartile, maximum). For variables with missing values, the number of NAs is also printed.

To see the names of the variables in the data frame try

```
names(births)
```

```
## [1] "id"      "bweight" "lowbw"   "gestwks" "preterm" "matage"  "hyp"
## [8] "sex"
```

Variables in a data frame can be referred to by name, but to do so it is necessary also to specify the name of the data frame. Thus `births$hyp` refers to the variable `hyp` in the `births` data frame, and typing `births$hyp` will print the data on this variable. To summarize the variable `hyp` try

```
summary(births$hyp)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  0.000   0.000   0.000   0.144   0.000   1.000
```

Alternatively you can use

```
with(births, summary(hyp))
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  0.000   0.000   0.000   0.144   0.000   1.000
```

1.12 Generating new variables

New variables can be produced using assignment together with the usual mathematical operations and functions. For example

```
logbw <- log(births$bweight)
```

produces the variable `logbw` in your workspace, while

```
births$logbw <- log(births$bweight)
```

produces the variable `logbw` in the `births` data frame.

You can also replace existing variables. For example `bweight` measures birth weight in grams. To convert the units to kilograms we replace the original variable with a new one:

```
births$bweight <- births$bweight / 1000
```

1.13 Turning a variable into a factor

In R categorical variables are known as *factors*, and the different categories are called the *levels* of the factor. Variables such as `hyp` and `sex` are originally coded using integer codes, and by default R will interpret these codes as numeric values taken by the variables. Factors will become very important later in the course when we study modelling functions, where factors and numeric variables are treated very differently. For the moment, you can think of factors as *value labels* that are more informative than numeric codes.

For R to recognize that the codes refer to categories it is necessary to convert the variables to be factors, and to label the levels. To convert the variable `hyp` to be a factor, try

```
births$hyp <- factor(births$hyp, labels = c("normal", "hyper"))
```

This takes the original numeric codes (0, 1) and replaces them with informative labels *normal* and *hyper* for normal blood pressure and hypertension, respectively.

1. Convert the variable `sex` into a factor with labels "M" and "F" for values 1 and 2, respectively

1.14 Frequency tables

When starting to look at any new data frame the first step is to check that the values of the variables make sense and correspond to the codes defined in the coding schedule. For categorical variables (factors) this can be done by looking at one-way frequency tables and checking that only the specified codes (levels) occur. The most useful function for making simple frequency tables is `table`. The distribution of the factor `hyp` can be viewed using

```
with(births, table(hyp))
```

```
## hyp
## normal hyper
##    428    72
```

or by specifying the data frame as in

```
table(births$hyp)
```

```
##
## normal hyper
##    428    72
```

For simple expressions the choice is a matter of taste, but `with` is shorter for more complicated expressions.

1. Find the frequency distribution of `sex`.

```
##
## M F
## 264 236

## sex
## M F
## 264 236
```

2. If you give two or more arguments to the `table` function then it produces cross-tabulations. Find the two-way frequency distribution of `sex` and `hyp`.

```
## hyp
## sex normal hyper
## M    221    43
## F    207    29
```

3. Create a logical variable called `early` according to whether `gestwks` is less than 30 or not. Make a frequency table of `early`.

```
## early
## FALSE TRUE
##    485    5
```

1.15 Grouping the values of a numeric variable

For a numeric variable like `matage` it is often useful to group the values and to create a new factor which codes the groups. For example we might cut the values taken by `matage` into the groups 20–29, 30–34, 35–39, 40–44, and then create a factor called `agegrp` with 4 levels corresponding to the four groups. The best way of doing this is with the function `cut`. Try

```
births$agegrp <-
  cut(
    births$matage,
    breaks = c(20, 30, 35, 40, 45),
    right = FALSE
  )
with(births, table(agegrp))
```

```
## agegrp
## [20,30) [30,35) [35,40) [40,45)
##      70      200      194      36
```

By default the factor levels are labelled [20-25), [25-30), etc., where [20-25) refers to the interval which includes the left hand end (20) but not the right hand end (25). This is the reason for `right=FALSE`. When `right=TRUE` (which is the default) the intervals include the right hand end but not the left hand.

Observations which are not inside the range specified by the `breaks` argument result in missing values for the new factor. Hence it is important that the first element in `breaks` is smaller than the smallest value in your data, and the last element is larger than the largest value.

1. Summarize the numeric variable `gestwks`, which records the length of gestation for the baby, and make a note of the range of values.

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.     NAs
## 24.69  37.94   39.12  38.72  40.09   43.16     10
```

2. Create a new factor `gest4` which cuts `gestwks` at 20, 35, 37, 39, and 45 weeks, including the left hand end, but not the right hand. Make a table of the frequencies for the four levels of `gest4`.

Chapter 2

Reading data into R

2.1 Introduction

If you want to have rabbit stew, first catch the rabbit – Old saying, origin unknown

R is a language and environment for data analysis. If you want to do something interesting with it, you need data.

For teaching purposes, data sets are often embedded in R packages. The base R distribution contains a whole package dedicated to data which includes around 100 data sets. This is attached towards the end of the search path, and you can see its contents with

```
objects("package:datasets")
```

```
##      [1] "ability.cov"      "airmiles"         "AirPassengers"
##      [4] "airquality"       "anscombe"         "attenu"
##      [7] "attitude"        "austres"          "beaver1"
##     [10] "beaver2"         "BJsales"          "BJsales.lead"
##     [13] "BOD"             "cars"             "ChickWeight"
##     [16] "chickwts"        "co2"              "CO2"
##     [19] "crimtab"         "discoveries"      "DNase"
##     [22] "esoph"           "euro"             "euro.cross"
##     [25] "eurodist"        "EuStockMarkets"   "faithful"
##     [28] "fdeaths"         "Formaldehyde"     "freeny"
##     [31] "freeny.x"        "freeny.y"         "gait"
##     [34] "HairEyeColor"    "Harman23.cor"     "Harman74.cor"
##     [37] "Indometh"        "infert"           "InsectSprays"
##     [40] "iris"            "iris3"            "islands"
##     [43] "JohnsonJohnson" "LakeHuron"        "ldeaths"
##     [46] "lh"              "LifeCycleSavings" "Loblolly"
##     [49] "longley"         "lynx"             "mdeaths"
##     [52] "morley"          "mtcars"           "nhtemp"
##     [55] "Nile"            "nottem"           "npk"
##     [58] "occupationalStatus" "Orange"          "OrchardSprays"
##     [61] "penguins"        "penguins_raw"     "PlantGrowth"
##     [64] "precip"          "presidents"       "pressure"
##     [67] "Puromycin"       "quakes"           "randu"
##     [70] "rivers"          "rock"             "Seatbelts"
##     [73] "sleep"           "stack.loss"       "stack.x"
##     [76] "stackloss"       "state.abb"        "state.area"
##     [79] "state.center"    "state.division"   "state.name"
##     [82] "state.region"    "state.x77"        "sunspot.m2014"
##     [85] "sunspot.month"   "sunspot.year"     "sunspots"
```

```
## [88] "swiss"           "Theoph"          "Titanic"
## [91] "ToothGrowth"     "treering"        "trees"
## [94] "UCBAdmissions"  "UKDriverDeaths" "UKgas"
## [97] "USAccDeaths"     "USArrests"       "UScitiesD"
## [100] "USJudgeRatings" "USPersonalExpenditure" "uspop"
## [103] "VADeaths"       "volcano"         "warpbreaks"
## [106] "women"          "WorldPhones"     "WWusage"
```

A description of all of these objects is available using the `help()` function. For example

```
help(Titanic)
```

gives an explanation of the `Titanic` data set, along with references giving the source of the data.

The `Epi` package also contains some data sets. These are not available automatically when you load the `Epi` package, but you can make a copy in your workspace using the `data()` function. For example

```
library(Epi)
data(bdendo)
```

will create a data frame called `bdendo` in your workspace containing data from a case-control study of endometrial cancer. Datasets in the `Epi` package also have help pages: type `help(bdendo)` for further information.

To go back to the cooking analogy, these data sets are the equivalent of microwave ready meals, carefully packaged and requiring minimal work by the consumer. Your own data will never be able in this form and you must work harder to read it in to R.

This exercise introduces you to the basics of reading external data into R. It consists of reading the same data from different formats. Although this may appear repetitive, it allows you to see the many options available to you, and should allow you to recognize when things go wrong.

getting the data You will need to download the zip file `data.zip` from the course web site (<https://github.com/SPE-R/SPE/raw/gh-spe-material/SPE-all-material.zip>) and unpack this in your working directory. This will create a sub-directory `data` containing (among other things) the files `fem.dat`, `fem-dot.dat`, `fem.csv`, and `fem.dta` (Reminder: use `setwd()` to set your working directory).

2.2 Data sources

Sources of data can be classified into three groups:

- Data in human readable form, which can be inspected with a text editor.
- Data in binary format, which can only be read by a program that understands that format (SAS, SPSS, Stata, Excel, ...).
- Online data from a database management system (DBMS)

This exercise will deal with the first two forms of data. Epidemiological data sets are rarely large enough to justify being kept in a DBMS. If you want further details on this topic, you can consult the *R Data Import/Export* manual that comes with R.

2.3 Data in text files

Human-readable data files are generally kept in a rectangular format, with individual records in single rows and variables in columns. Such data can be read into a data frame in R.

Before reading in the data, you should inspect the file in a text editor and ask three questions:

- How are columns in the table separated?
- How are missing values represented?
- Are variable names included in the file?

The file `fem.dat` contains data on 118 female psychiatric patients. The data set contains nine variables.

Name	Description
ID	Patient identifier
AGE	Age in years
IQ	Intelligence Quotient (IQ) score
ANXIETY	Anxiety (1=none, 2=mild, 3=moderate,4=severe)
DEPRESS	Depression (1=none, 2=mild, 3=moderate or severe)
SLEEP	Sleeping normally (1=yes, 2=no)
SEX	Lost interest in sex (1=yes, 2=no)
LIFE	Considered suicide (1=yes, 2=no)
WEIGHT	Weight change (kg) in previous 6 months

Inspect the file `fem.dat` with a text editor to answer the questions above.

The most general function for reading in free-format data is `read.table()`. This function reads a text file and returns a data frame. It tries to guess the correct format of each variable in the data frame (integer, double precision, or text).

Read in the table with:

```
fem <- read.table("./data/fem.dat", header = TRUE)
```

Note that you must assign the result of `read.table()` to an object. If this is not done, the data frame will be printed to the screen and then lost.

You can see the names of the variables with

```
names(fem)
```

```
## [1] "ID"      "AGE"      "IQ"      "ANXIETY" "DEPRESS" "SLEEP"   "SEX"
## [8] "LIFE"    "WEIGHT"
```

The structure of the data frame can be seen with

```
str(fem)
```

```
## 'data.frame':   118 obs. of  9 variables:
## $ ID      : int  1 2 3 4 5 6 7 8 9 10 ...
## $ AGE     : int  39 41 42 30 35 44 31 39 35 33 ...
## $ IQ      : int  94 89 83 99 94 90 94 87 -99 92 ...
## $ ANXIETY: int  2 2 3 2 2 NA 2 3 3 2 ...
## $ DEPRESS: int  2 2 3 2 1 1 2 2 2 2 ...
## $ SLEEP   : int  2 2 2 2 1 2 NA 2 2 2 ...
## $ SEX     : int  1 1 1 1 1 2 1 1 1 1 ...
## $ LIFE    : int  1 1 1 1 2 2 1 2 1 1 ...
## $ WEIGHT  : num  2.23 1 1.82 -1.18 -0.14 0.41 -0.68 1.59 -0.55 0.36 ...
```

You can also inspect the top few rows with

```
head(fem)
```

```
##   ID AGE IQ ANXIETY DEPRESS SLEEP SEX LIFE WEIGHT
## 1  1  39 94      2      2      2  1   1   2.23
## 2  2  41 89      2      2      2  1   1   1.00
## 3  3  42 83      3      3      2  1   1   1.82
## 4  4  30 99      2      2      2  1   1  -1.18
## 5  5  35 94      2      1      1  1   2  -0.14
## 6  6  44 90     NA      1      2  2   2   0.41
```

Note that the IQ of subject 9 is -99, which is an illegal value: nobody can have a negative IQ. In fact -99 has been used in this file to represent a missing value. In R the special value `NA` (*Not Available*) is used to represent missing values. All R functions recognize `NA` values and will handle them appropriately, although sometimes the appropriate response is to stop the calculation with an error message.

You can recode the missing values with

```
fem$IQ[fem$IQ == -99] <- NA
```

Of course it is much better to handle special missing value codes when reading in the data. This can be done with the `na.strings` argument of the `read.table()` function. See below.

2.4 Things that can go wrong

Sooner or later when reading data into R, you will make a mistake. The frustrating part of reading data into R is that most mistakes are not fatal: they simply cause the function to return a data frame that is *not what you wanted*. There are three common mistakes, which you should learn to recognize.

2.5 Forgetting the headers

The first row of the file `fem.dat` contains the variable names. The `read.table()` function does not assume this by default so you have to specify this with the argument `header=TRUE`. See what happens when you forget to include this option:

```
fem2 <- read.table("data/fem.dat")
str(fem2)

## 'data.frame':    119 obs. of  9 variables:
## $ V1: chr  "ID" "1" "2" "3" ...
## $ V2: chr  "AGE" "39" "41" "42" ...
## $ V3: chr  "IQ" "94" "89" "83" ...
## $ V4: chr  "ANXIETY" "2" "2" "3" ...
## $ V5: chr  "DEPRESS" "2" "2" "3" ...
## $ V6: chr  "SLEEP" "2" "2" "2" ...
## $ V7: chr  "SEX" "1" "1" "1" ...
## $ V8: chr  "LIFE" "1" "1" "1" ...
## $ V9: chr  "WEIGHT" "2.23" "1" "1.82" ...

head(fem2)

##   V1  V2 V3   V4   V5  V6 V7  V8   V9
## 1 ID AGE IQ ANXIETY DEPRESS SLEEP SEX LIFE WEIGHT
## 2 1  39 94     2     2   2  1   1  2.23
## 3 2  41 89     2     2   2  1   1    1
## 4 3  42 83     3     3   2  1   1  1.82
## 5 4  30 99     2     2   2  1   1 -1.18
## 6 5  35 94     2     1   1  1   2 -0.14
```

and compare the resulting data frame with `fem`. - What are the names of the variables in the data frame? - What is the class of the variables?

Explanation: Remember that `read.table()` tries to guess the mode of the variables in the text file. Without the `header = TRUE` option it reads the first row, containing the variable names, as data, and guesses that all the variables are character, not numeric. By default, all character variables are coerced to factors by `read.table`. The result is a data frame consisting entirely of factors. (You can prevent the conversion of character variables to factors with the argument `as.is = TRUE`).

If the variable names are not specified in the file, then they are given default names `V1`, `V2`, You will soon realise this mistake if you try to access a variable in the data frame by, for example

```
fem2$IQ
```

```
## NULL
```

as the variable will not exist

There is one case where omitting the `header = TRUE` option is harmless (apart from the situation where there is no header line, obviously). When the first row of the file contains **one less** value than subsequent lines, `read.table()` infers that the first row contains the variable names, and the first column of every subsequent row contains its **row name**.

2.6 Using the wrong separator

By default, `read.table` assumes that data values are separated by any amount of white space. Other possibilities can be specified using the `sep` argument. See what happens when you assume the wrong separator, in this case a tab, which is specified using the escape sequence `"\t"`

```
fem3 <- read.table("data/fem.dat", sep = "\t")
str(fem3)
```

```
## 'data.frame':    119 obs. of  1 variable:
## $ V1: chr  "ID AGE IQ ANXIETY DEPRESS SLEEP SEX LIFE WEIGHT" "1 39 94 2 2 2 1 1 2.23" "2 41 89 2 2 2 1
```

- How many variables are there in the data set?

Explanation: If you mis-specify the separator, `read.table()` reads the whole line as a single character variable. Once again, character variables are coerced to factors, so you get a data frame with a single factor variable.

2.7 Mis-specifying the representation of missing values

The file `fem-dot.dat` contains a version of the FEM dataset in which all missing values are represented with a dot. This is a common way of representing missing values, but is not recognized by default by the `read.table()` function, which assumes that missing values are represented by `NA`.

Inspect the file with a text editor, and then see what happens when you read the file in incorrectly:

```
fem4 <- read.table("data/fem-dot.dat", header = TRUE)
str(fem4)
```

```
## 'data.frame':    118 obs. of  9 variables:
## $ ID      : int  1 2 3 4 5 6 7 8 9 10 ...
## $ AGE     : int  39 41 42 30 35 44 31 39 35 33 ...
## $ IQ      : chr  "94" "89" "83" "99" ...
## $ ANXIETY: chr  "2" "2" "3" "2" ...
## $ DEPRESS: chr  "2" "2" "3" "2" ...
## $ SLEEP   : chr  "2" "2" "2" "2" ...
## $ SEX     : chr  "1" "1" "1" "1" ...
## $ LIFE    : chr  "1" "1" "1" "1" ...
## $ WEIGHT  : chr  "2.23" "1" "1.82" "-1.18" ...
```

You should have enough clues by now to work out what went wrong.

You can read the data correctly using the `na.strings` argument

```
fem4 <-
  read.table(
    "data/fem-dot.dat",
    header = TRUE,
    na.strings = "."
  )
```

2.8 Spreadsheet data

Spreadsheets have become a common way of exchanging data. All spreadsheet programs can save a single sheet in *comma-separated variable* (CSV) format, which can then be read into R. There are two functions in R for

reading in CSV data: `read.csv()` and `read.csv2()`.

Both of these are wrappers around the `read.table()` function, *i.e.* the `read.table()` function is still doing the work of reading in the data but the `read.csv()` function provides default argument values for reading in CSV file so all you need to do is specify the file name.

You can see what these default arguments are with the `args()` function.

```
args(read.csv)
```

```
## function (file, header = TRUE, sep = ",", quote = "\"", dec = ".",
##      fill = TRUE, comment.char = "", ...)
## NULL
```

```
args(read.csv2)
```

```
## function (file, header = TRUE, sep = ";", quote = "\"", dec = ",",
##      fill = TRUE, comment.char = "", ...)
## NULL
```

See if you can spot the difference between `read.csv` and `read.csv2`.

Explanation: The CSV format is not a single standard. The file format depends on the *locale* of your computer – the settings that determine how numbers are represented. In some countries, the decimal separator is a point `.` and the variable separator in a CSV file is a comma `,`. In other countries, the decimal separator is a comma `,`, and the variable separator is a semi-colon `;`. This is reflected in the different default values for the arguments `sep` and `dec`. The `read.csv()` function is used for the first format and the `read.csv2()` function is used for the second format.

The file `fem.csv` contains the FEM dataset in CSV format. Inspect the file to work out which format is used, and read it into R.

2.9 Reading data from the Internet

You can also read in data from a remote web site. The `file` argument of `read.table()` does not need to be a local file on your computer; it can be a Uniform Resource Locator (URL), *i.e.* a web address.

A copy of the file `fem.dat` is held at (<https://www.bendixcarstensen.com/SPE/data/fem.dat>). You can read it in with

```
fem6 <-
  read.table(
    "http://www.bendixcarstensen.com/SPE/data/fem.dat",
    header = TRUE
  )
str(fem6)
```

```
## 'data.frame':   118 obs. of  9 variables:
## $ ID      : int  1 2 3 4 5 6 7 8 9 10 ...
## $ AGE     : int  39 41 42 30 35 44 31 39 35 33 ...
## $ IQ      : int  94 89 83 99 94 90 94 87 -99 92 ...
## $ ANXIETY: int  2 2 3 2 2 NA 2 3 3 2 ...
## $ DEPRESS: int  2 2 3 2 1 1 2 2 2 2 ...
## $ SLEEP   : int  2 2 2 2 1 2 NA 2 2 2 ...
## $ SEX     : int  1 1 1 1 1 2 1 1 1 1 ...
## $ LIFE    : int  1 1 1 1 2 2 1 2 1 1 ...
## $ WEIGHT  : num  2.23 1 1.82 -1.18 -0.14 0.41 -0.68 1.59 -0.55 0.36 ...
```

2.10 Reading from the clipboard

On Microsoft Windows, you can copy values directly from an open Excel spreadsheet using the clipboard. Highlight the cells you want to copy in the spread sheet and select copy from the pull-down edit menu. Then type `read.table(file = "clipboard")` to read the data in.

There are two reasons why this is a bad idea

- It is not reproducible. In order to read the data in again you need to complete exactly the same sequence of mouse moves and clicks, and there is no record of what you did before.
- Copying from the clipboard loses precision. If you have a value 1.23456789 in your spreadsheet, but have formatted the cell so it is displayed to two decimal places, then the value read into R will be the truncated value 1.23.

2.11 Binary data

The `foreign` package allows you to read data in binary formats used by other statistical packages. Since R is an open source project, it can only read binary formats that are themselves *open*, in the sense that the standards for reading and writing data are well-documented. For example, there is a function in the `foreign` package for reading SAS XPORT files, a format that has been adopted as a standard by the US Food and Drug Administration (<http://www.sas.com/govedu/fda/faq.html>). However, there is no function in the `foreign` package for reading native SAS binaries (SAS7BDAT files). Other packages are available from CRAN (<http://cran.r-project.org>) that offer the possibility of reading SAS binary files: see the `haven` and `sas7bdat` packages.

The file `fem.dta` contains the FEM dataset in the format used by Stata. Read it into R with

```
library(foreign)
fem5 <- read.dta("data/fem.dta")
head(fem5)
```

```
##   id age iq  anxiety          depress sleep sex life weight
## 1  1  39 94   mild          mild    no yes  yes   2.23
## 2  2  41 89   mild          mild    no yes  yes   1.00
## 3  3  42 83 moderate moderate or severe    no yes  yes   1.82
## 4  4  30 99   mild          mild    no yes  yes  -1.18
## 5  5  35 94   mild          none   yes yes   no  -0.14
## 6  6  44 90   <NA>          none    no  no    no   0.41
```

The Stata data set contains value and variable labels. Stata variables with value labels are automatically converted to factors.

There is no equivalent of variable labels in an R data frame, but the original variable labels are not lost. They are still attached to the data frame as an invisible *attribute*, which you can see with

```
attr(fem5, "var.labels")
```

```
## [1] "Patient ID"
## [2] "Age in years"
## [3] "IQ score"
## [4] "Anxiety"
## [5] "Depression"
## [6] "Sleeping normally"
## [7] "Lost interest in sex"
## [8] "Considered suicide"
## [9] "Weight change (kg) in previous 6 months"
```

A lot of *meta-data* is attached to the data in the form of attributes. You can see the whole list of attributes with

```
attributes(fem5)
```

```
## $datalabel
## [1] "Data on 118 female psychiatric patients"
##
## $time.stamp
## [1] " 1 Jun 2006 12:36"
##
## $names
## [1] "id"      "age"      "iq"      "anxiety" "depress" "sleep"   "sex"
## [8] "life"    "weight"
##
## $formats
## [1] "%8.0g"  "%8.0g"  "%8.0g"  "%8.0g"  "%18.0g" "%8.0g"  "%8.0g"  "%8.0g"
## [9] "%9.0g"
##
## $types
## [1] 252 251 252 251 251 251 251 251 254
##
## $val.labels
## [1] ""      ""      ""      "anxiety" "depress" "yesno"   "yesno"
## [8] "yesno" ""
##
## $var.labels
## [1] "Patient ID"
## [2] "Age in years"
## [3] "IQ score"
## [4] "Anxiety"
## [5] "Depression"
## [6] "Sleeping normally"
## [7] "Lost interest in sex"
## [8] "Considered suicide"
## [9] "Weight change (kg) in previous 6 months"
##
## $row.names
## [1] "1"  "2"  "3"  "4"  "5"  "6"  "7"  "8"  "9"  "10" "11" "12"
## [13] "13" "14" "15" "16" "17" "18" "19" "20" "21" "22" "23" "24"
## [25] "25" "26" "27" "28" "29" "30" "31" "32" "33" "34" "35" "36"
## [37] "37" "38" "39" "40" "41" "42" "43" "44" "45" "46" "47" "48"
## [49] "49" "50" "51" "52" "53" "54" "55" "56" "57" "58" "59" "60"
## [61] "61" "62" "63" "64" "65" "66" "67" "68" "69" "70" "71" "72"
## [73] "73" "74" "75" "76" "77" "78" "79" "80" "81" "82" "83" "84"
## [85] "85" "86" "87" "88" "89" "90" "91" "92" "93" "94" "95" "96"
## [97] "97" "98" "99" "100" "101" "102" "103" "104" "105" "106" "107" "108"
## [109] "109" "110" "111" "112" "113" "114" "115" "116" "117" "118"
##
## $version
## [1] 8
##
## $label.table
## $label.table$anxiety
##      none      mild moderate      severe
##         1         2         3         4
##
## $label.table$depress
##
##              none              mild moderate or severe
```

```
##           1           2           3
##
## $label.table$yesno
## yes  no
##   1   2
##
##
## $class
## [1] "data.frame"

or just the attribute names with
names(attributes(fem5))

## [1] "datalabel"  "time.stamp"  "names"       "formats"     "types"
## [6] "val.labels"  "var.labels"  "row.names"   "version"     "label.table"
## [11] "class"
```

The `read.dta()` function can only read data from Stata versions 5–12. The R Core Team has not been able to keep up with changes in the Stata format. You may wish to try the **haven** package and the **readstata13** package, both available from CRAN.

2.12 Summary

In this exercise we have seen how to create a data frame in R from an external text file. We have also reviewed some common mistakes that result in garbled data.

The capabilities of the **foreign** package for reading binary data have also been demonstrated with a sample Stata data set.

Chapter 3

Data manipulation with tidyverse

3.1 Introduction

In this chapter we will produce *more or less* the same outputs than in chapter 1.2 and 1.4 using **tidyverse** packages framework.

The main objective of this exercise is to get familiar you with some of the main **tidyverse** features.

This is an **optional practical** for participants having already good basic R skills. > All the rest of the course can be done without knowledge of **tidyverse**.

3.2 The births data

We will work with **births** data-set from **Epi** package. First of all, load the **Epi** and **tidyverse** packages. Then load the **births** data-set.

```
library(Epi)
suppressPackageStartupMessages(library(tidyverse))
data(births)
```

You can type `?birth` in the R console to get a description of the **birth** data-set. (Alternatively, you can refer to chapter 1.3.2)

3.3 tibble vs data.frame

Most **dplyr** functions outputs return **tibble** object instead of **data.frame**. Inspect the class and characteristics of the **births** object.

```
class(births)
```

```
## [1] "data.frame"
```

```
head(births)
```

```
##   id bweight lowbw gestwks preterm matage hyp sex
## 1  1   2974     0   38.52      0    34  0  2
## 2  2   3270     0    NA      NA    30  0  1
## 3  3   2620     0   38.15      0    35  0  2
## 4  4   3751     0   39.80      0    31  0  1
## 5  5   3200     0   38.89      0    33  1  1
## 6  6   3673     0   40.97      0    33  0  2
```

Note: As any R object this can be summarized using **str** function.

```
str(births)
```

```
## 'data.frame':    500 obs. of  8 variables:
## $ id      : num  1 2 3 4 5 6 7 8 9 10 ...
## $ bweight: num  2974 3270 2620 3751 3200 ...
## $ lowbw   : num  0 0 0 0 0 0 0 0 0 0 ...
## $ gestwks: num  38.5 NA 38.2 39.8 38.9 ...
## $ preterm: num  0 NA 0 0 0 0 0 0 0 0 ...
## $ matage  : num  34 30 35 31 33 33 29 37 36 39 ...
## $ hyp     : num  0 0 0 0 1 0 0 0 0 0 ...
## $ sex     : num  2 1 2 1 1 2 2 1 2 1 ...
```

births object is a 500 x 8 data.frame. Let's convert births to tibble format with as_tibble function.

```
births_tbl <- as_tibble(births)
```

```
class(births_tbl)
```

```
## [1] "tbl_df"      "tbl"        "data.frame"
births_tbl
```

```
## # A tibble: 500 x 8
##       id bweight lowbw gestwks preterm matage   hyp   sex
##   <dbl> <dbl> <dbl>   <dbl>   <dbl> <dbl> <dbl> <dbl>
## 1     1    2974     0    38.5       0    34     0     2
## 2     2    3270     0     NA       NA    30     0     1
## 3     3    2620     0    38.2       0    35     0     2
## 4     4    3751     0    39.8       0    31     0     1
## 5     5    3200     0    38.9       0    33     1     1
## 6     6    3673     0    41.0       0    33     0     2
## 7     7    3628     0    42.1       0    29     0     2
## 8     8    3773     0    40.2       0    37     0     1
## 9     9    3960     0    42.0       0    36     0     2
## 10    10    3405     0    39.3       0    39     0     1
## # i 490 more rows
```

another way to visualize data set is to use glimpse function

```
glimpse(births_tbl)
```

```
## Rows: 500
## Columns: 8
## $ id      <dbl> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18,~
## $ bweight <dbl> 2974, 3270, 2620, 3751, 3200, 3673, 3628, 3773, 3960, 3405, 40~
## $ lowbw   <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,~
## $ gestwks <dbl> 38.52, NA, 38.15, 39.80, 38.89, 40.97, 42.14, 40.21, 42.03, 39~
## $ preterm <dbl> 0, NA, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,~
## $ matage  <dbl> 34, 30, 35, 31, 33, 33, 29, 37, 36, 39, 37, 35, 38, 34, 28, 38~
## $ hyp     <dbl> 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,~
## $ sex     <dbl> 2, 1, 2, 1, 1, 2, 2, 1, 2, 1, 1, 2, 2, 1, 2, 1, 2, 1, 2, 2,~
```

You can see that tibble objects inherits from data.frame which implies that all functions working with data.frame objects will work with tibble. The opposite is not necessary true. tibble has a couple of extra features compared to classical data.frame. One of them is a slightly more user-friendly console print. The main difference is probably that tibble objects supports grouping/nesting features. Some examples we be done will see latter on.

3.4 Piping functions

This is one of the most popular features of *tidyverse* grammar. It enables function chaining in R. Function output is transparently passed as input to the next function and so on. It can help to make the code more comprehensive and readable. Here is an example of classic vs piped functions.

```
head(births, 4)
```

```
##   id bweight lowbw gestwks preterm matage hyp sex
## 1  1   2974     0   38.52      0    34  0  2
## 2  2   3270     0    NA     NA    30  0  1
## 3  3   2620     0   38.15      0    35  0  2
## 4  4   3751     0   39.80      0    31  0  1
```

```
births |> head(4)
```

```
##   id bweight lowbw gestwks preterm matage hyp sex
## 1  1   2974     0   38.52      0    34  0  2
## 2  2   3270     0    NA     NA    30  0  1
## 3  3   2620     0   38.15      0    35  0  2
## 4  4   3751     0   39.80      0    31  0  1
```

Note: By default the chained object is given as the first argument to the following function. You can use `.` if this is not the case.

Here is a dummy example where we do not give the first argument to `head` function but the second one.

```
4 %>% head(births, .)
```

```
##   id bweight lowbw gestwks preterm matage hyp sex
## 1  1   2974     0   38.52      0    34  0  2
## 2  2   3270     0    NA     NA    30  0  1
## 3  3   2620     0   38.15      0    35  0  2
## 4  4   3751     0   39.80      0    31  0  1
```

This can also be achieved with base pipe `|>` with two differences

- The place holder is an underscore `_` not a dot `.`. This is a deliberate design choice to help people with less acute eyesight
- The argument must be named

```
4 |> head(births, pos=_)
```

```
##   id bweight lowbw gestwks preterm matage hyp sex
## 1  1   2974     0   38.52      0    34  0  2
## 2  2   3270     0    NA     NA    30  0  1
## 3  3   2620     0   38.15      0    35  0  2
## 4  4   3751     0   39.80      0    31  0  1
## 5  5   3200     0   38.89      0    33  1  1
## 6  6   3673     0   40.97      0    33  0  2
```

3.5 mutate columns

`mutate` will allow you to add and or modify columns in a *tibble*. Let's create 2 new variables :

- `agegrp` (5 years mother's age group)
- `gest4` (gestation time split in 4 categories)

And modify 2 others:

- `hyp` (factor version of `hyp`; normal vs hyper)
- `sex` (factor version of `sex`; M vs F)

```

births_tbl <-
  births_tbl |>
  mutate(
    # modify hyp variable (conversion into factor)
    hyp =
      factor(
        hyp,
        levels = c(0, 1),
        labels = c("normal", "hyper")
      ),
    # creating a new variable agegrp
    agegrp =
      cut(
        matage,
        breaks = c(20, 25, 30, 35, 40, 45),
        right = FALSE
      ),
    # modify sex variable (conversion into factor)
    sex =
      factor(
        sex,
        levels = c(1, 2),
        labels = c("M", "F")
      ),
    # creating a new variable gest4 with case_when instead of cut
    gest4 =
      case_when(
        gestwks < 25 ~ "less than 25 weeks",
        gestwks >= 25 & gestwks < 30 ~ "25-30 weeks",
        gestwks >= 30 & gestwks < 35 ~ "30-35 weeks",
        gestwks >= 35 ~ "more than 35 weeks"
      )
  )
births_tbl

```

```

## # A tibble: 500 x 10
##       id bweight lowbw gestwks preterm matage hyp   sex agegrp gest4
##   <dbl> <dbl> <dbl>   <dbl>   <dbl> <dbl> <fct> <fct> <fct> <chr>
## 1     1    2974     0    38.5     0    34 normal F   [30,35) more than 35~
## 2     2    3270     0    NA      NA    30 normal M   [30,35) <NA>
## 3     3    2620     0    38.2     0    35 normal F   [35,40) more than 35~
## 4     4    3751     0    39.8     0    31 normal M   [30,35) more than 35~
## 5     5    3200     0    38.9     0    33 hyper  M   [30,35) more than 35~
## 6     6    3673     0    41.0     0    33 normal F   [30,35) more than 35~
## 7     7    3628     0    42.1     0    29 normal F   [25,30) more than 35~
## 8     8    3773     0    40.2     0    37 normal M   [35,40) more than 35~
## 9     9    3960     0    42.0     0    36 normal F   [35,40) more than 35~
## 10    10    3405     0    39.3     0    39 normal M   [35,40) more than 35~
## # i 490 more rows

```

You can see as header the type of data contained in each column. For instance `<dbl>` stands for double (i.e. numeric value) and `fct` stands for factor. In R `data.frame` (/ `tibble`) data type must be the same within a column (e.g. numeric only) but can be of different type across columns. (note: `matrix` object supports only one type of data)

Note that `case_when` function do not return a `factor` but a `character` variable in this case. You will have to

force the conversion from `character` to `factor` if needed.

3.6 select columns, filter and arrange rows

`select` is used for column sub-setting while `filter` is for row sub-setting. They are equivalent to the `[]` in R base language. Let's display a table where only babies' `id`, `sex`, `bweight` and mothers' `agegrp` are kept for babies with a `bweight` above 4000g.

```
births_tbl |>
  # select only id, women age group, sex
  # and birth weight of the baby
  select(id, agegrp, sex, bweight) |>
  # keep only babies weighing more than 4000g
  filter(bweight > 4000)
```

```
## # A tibble: 30 x 4
##       id agegrp sex   bweight
##   <dbl> <fct> <fct>   <dbl>
## 1    11 [35,40) M      4020
## 2    31 [30,35) F      4022
## 3    57 [25,30) M      4501
## 4    78 [35,40) M      4512
## 5    86 [30,35) M      4141
## 6    87 [30,35) M      4071
## 7    92 [30,35) F      4027
## 8   122 [30,35) M      4092
## 9   139 [35,40) M      4069
## 10  170 [40,45) M      4057
## # i 20 more rows
```

`select` can also be useful to reorder and rename columns. `arrange` is a nice feature to reorder observations according to chosen attributes. Let's rename `agegrp`, `sex` and `bweight` with better looking labels (e.g. Age group, Sex, Birth weight) and reorder the table according to babies' decreasing birth weight.

```
births_tbl |>
  # select only id, women age group, sex
  # and birth weight of the baby
  select(
    id,
    "Age group" = agegrp,
    Sex = sex,
    "Birth weight" = bweight
  ) |>
  # rearrange rows to put the heaviest newborn on top
  arrange(desc(`Birth weight`))
```

```
## # A tibble: 500 x 4
##       id `Age group` Sex   `Birth weight`
##   <dbl> <fct>      <fct>         <dbl>
## 1   319 [30,35)   M      4553
## 2   474 [25,30)   M      4516
## 3    78 [35,40)   M      4512
## 4    57 [25,30)   M      4501
## 5   238 [25,30)   M      4436
## 6   288 [30,35)   M      4423
## 7   206 [40,45)   M      4340
## 8   355 [35,40)   M      4319
## 9   459 [30,35)   M      4304
```

```
## 10 469 [30,35) F 4300
## # i 490 more rows
```

Note: tibble supports blank spaces in the column names which can be handy for final table rendering. When you want to work with columns with blank spaces, do not forget to use the " (back-quote). Try to produce the same table but arranging the rows by decreasing birth weights within each sex.

```
births_tbl |>
  # select only id, women age group, sex
  # and birth weight of the baby
  select(
    id,
    "Age group" = agegrp,
    Sex = sex,
    "Birth weight" = bweight
  ) |>
  # rearrange rows to put the heaviest newborn on top
  arrange(Sex, desc(`Birth weight`))
```

```
## # A tibble: 500 x 4
##       id `Age group` Sex   `Birth weight`
##   <dbl> <fct>    <fct>         <dbl>
## 1  319 [30,35) M       4553
## 2  474 [25,30) M       4516
## 3   78 [35,40) M       4512
## 4   57 [25,30) M       4501
## 5  238 [25,30) M       4436
## 6  288 [30,35) M       4423
## 7  206 [40,45) M       4340
## 8  355 [35,40) M       4319
## 9  459 [30,35) M       4304
## 10 387 [35,40) M       4287
## # i 490 more rows
```

You can arrange the tibble according to more than one column.

3.7 group_by and summarise data

One greatest features of dplyr is the ability to aggregate data sharing a common attribute to process per group operations. Here we want to compute the number of boys and girls in the data-set. The idea here is to split the births table in two groups. One with the boys, the other with the girls and to count the number of rows in each group.

```
births.01 <-
  births_tbl |>
  # group the data according to the sex attribute
  group_by(sex) |>
  # count the number of rows/individuals in each group
  summarise(
    count = n()
  )
births.01
```

```
## # A tibble: 2 x 2
##   sex   count
##   <fct> <int>
## 1 M       264
## 2 F       236
```

Note: `n` function is equivalent to `nrow`

Now we have the number of boys and girls, we can compute the distribution (in percentage) of newborns per sex.

```
births.02 <-
  births.01 |>
  mutate(
    percent = count / sum(count) * 100
  )
```

Trick: most of `dplyr` functions can be combined with a column selection execution statement using `across` function. This can be very handy in some cases. As an example below a code to compute the `sum` of every `birth.02` numerical columns (numerical columns only)

```
births.03 <-
  births_tbl |>
  select(gest4, sex, gestwks, bweight, matage) |>
  group_by(gest4, sex) |>
  summarise(
    across(
      where(is.numeric),
      ~ mean(.x, na.rm = TRUE)
    ),
    .groups = "drop"
  )
births.03
```

```
## # A tibble: 9 x 5
##   gest4      sex gestwks bweight matage
##   <chr>    <fct>   <dbl>   <dbl>   <dbl>
## 1 25-30 weeks M      28.0   1000    30
## 2 25-30 weeks F      27.1    668   29.5
## 3 30-35 weeks M      33.3  2224.   32.8
## 4 30-35 weeks F      32.8  1672.   33
## 5 less than 25 weeks F      24.7   864    40
## 6 more than 35 weeks M      39.1  3300.   34.3
## 7 more than 35 weeks F      39.3  3158.   34.0
## 8 <NA>      M      NaN   3056.   33
## 9 <NA>      F      NaN   3188.   31
```

`across` function supports the `purrr`-style lambda format, e.g. `~ mean(.x, na.rm = TRUE)` where `.x` refers to the values from the data set to be passed to the function. This is a common notation you will find across several `tidyverse` functions.

Some other functions ending by `_with` can be used conditionally within `dplyr`. As an example we can rename only columns which are not numeric at once (here we want to code all column names using upper characters) using the combination of `rename_with` and `where`.

```
births.03 |>
  rename_with(toupper, where(~ !is.numeric(.x)))
```

```
## # A tibble: 9 x 5
##   GEST4      SEX gestwks bweight matage
##   <chr>    <fct>   <dbl>   <dbl>   <dbl>
## 1 25-30 weeks M      28.0   1000    30
## 2 25-30 weeks F      27.1    668   29.5
## 3 30-35 weeks M      33.3  2224.   32.8
## 4 30-35 weeks F      32.8  1672.   33
## 5 less than 25 weeks F      24.7   864    40
## 6 more than 35 weeks M      39.1  3300.   34.3
```

```
## 7 more than 35 weeks F      39.3   3158.   34.0
## 8 <NA>                  M      NaN    3056.   33
## 9 <NA>                  F      NaN    3188.   31
```

Let's now compute the number of births and the mean birth weight according to newborn gender.

```
births.05 <-
  births_tbl |>
  group_by(sex) |>
  summarise(
    count = n(),
    bweight.mean = mean(bweight)
  )
births.05
```

```
## # A tibble: 2 x 3
##   sex    count bweight.mean
##   <fct> <int>      <dbl>
## 1 M       264      3230.
## 2 F       236      3033.
```

With `births.05` table, compute the global mean birth weight.

Note: with such a table the mean baby's birth weight have to be weighted by number of boys and girls (see. `?weighted.mean`).

```
births.05 |>
  summarise(
    count.tot = sum(count),
    bweight.mean.tot = weighted.mean(bweight.mean, count)
  )
```

```
## # A tibble: 1 x 2
##   count.tot bweight.mean.tot
##   <int>      <dbl>
## 1      500      3137.
```

this is equivalent to

```
births_tbl |>
  summarise(
    count.tot = n(),
    bweight.mean.tot = mean(bweight)
  )
```

```
## # A tibble: 1 x 2
##   count.tot bweight.mean.tot
##   <int>      <dbl>
## 1      500      3137.
```

3.8 Multiple grouping

In some cases, we can be interested in looking at more than a single strata. This can be achieved using multiple grouping. Let's count the number of births per gender and birth weight class (low vs not low)

```
births.06 <-
  births_tbl |>
  group_by(sex, lowbw) |>
  summarise(
    count = n()
  )
```

```
## `summarise()` has regrouped the output.
## i Summaries were computed grouped by sex and lowbw.
## i Output is grouped by sex.
## i Use `summarise(.groups = "drop_last")` to silence this message.
## i Use `summarise(.by = c(sex, lowbw))` for per-operation grouping
## (`?dplyr::dplyr_by`) instead.
```

```
births.06
```

```
## # A tibble: 4 x 3
## # Groups:   sex [2]
##   sex    lowbw count
##   <fct> <dbl> <int>
## 1 M         0   237
## 2 M         1    27
## 3 F         0   203
## 4 F         1    33
```

Try then to compute the percentage of babies in each group. Look at the difference between the 2 following command lines:

```
births.06 |>
  mutate(
    percent = count / sum(count) * 100
  )
```

```
## # A tibble: 4 x 4
## # Groups:   sex [2]
##   sex    lowbw count percent
##   <fct> <dbl> <int>   <dbl>
## 1 M         0   237    89.8
## 2 M         1    27    10.2
## 3 F         0   203    86.0
## 4 F         1    33    14.0
```

```
births.06 |>
  ungroup() |>
  mutate(
    percent = count / sum(count) * 100
  )
```

```
## # A tibble: 4 x 4
##   sex    lowbw count percent
##   <fct> <dbl> <int>   <dbl>
## 1 M         0   237    47.4
## 2 M         1    27     5.4
## 3 F         0   203   40.6
## 4 F         1    33     6.6
```

Are the results the same?

Note: summarizing a data-set will remove the last level of grouping but not the other ones if multiple grouping has been performed. In some cases you might have to explicitly ungroup your `data.frame` before doing further calculations. In the previous examples, if you do not ungroup the data-set, percentages are computed per gender. Ungrouping will let you compute the overall percentages.

Trick: a good practice is to always ungroup the summarized dataset in order to prevent form confusion. You can do it using the `.group = 'drop'` option in `summarize()`.

```
# this tibble will still be grouped by sex
births_tbl |>
  group_by(sex, lowbw) |>
```

```

summarise(
  count = n()
)

## # A tibble: 4 x 3
## # Groups:   sex [2]
##   sex    lowbw count
##   <fct> <dbl> <int>
## 1 M         0   237
## 2 M         1    27
## 3 F         0   203
## 4 F         1    33
# this tibble will be group free
births_tbl |>
  group_by(sex, lowbw) |>
  summarise(
    count = n(),
    .groups = "drop"
  )

## # A tibble: 4 x 3
##   sex    lowbw count
##   <fct> <dbl> <int>
## 1 M         0   237
## 2 M         1    27
## 3 F         0   203
## 4 F         1    33

```

The same exercise can be done using gestation time group (`gest4`) as stratifying variable. Lets compute number of births and mean birth weights according to gestation time category.

```

births_tbl |>
  group_by(gest4) |>
  summarise(
    count = n(),
    bweight.mean = mean(bweight)
  )

## # A tibble: 5 x 3
##   gest4          count bweight.mean
##   <chr>         <int>         <dbl>
## 1 25-30 weeks         4          834
## 2 30-35 weeks        26         1906.
## 3 less than 25 weeks  1          864
## 4 more than 35 weeks 459         3233.
## 5 <NA>              10         3082.

```

Any trend? It seems that birth weight increases with gestation time. We can also spot that in our data-set the gestation time is missing for 10 newborns. We will do not consider this observation for the rest of the exercise. Lets cross-tabulate the birth weight category and the gestation time groups.

```

births_tbl |>
  # keep only the newborn with defined gestational time category
  filter(
    !is.na(gest4)
  ) |>
  group_by(lowbw, gest4) |>
  # compute the number of babies in each cross category
  summarise(

```

```

  count = n()
) |>
# compute the percentage of babies in each gestational
# time category per birth weight category
mutate(
  percent = count / sum(count, na.rm = TRUE)
)

```

```

## `summarise()` has regrouped the output.
## i Summaries were computed grouped by lowbw and gest4.
## i Output is grouped by lowbw.
## i Use `summarise(.groups = "drop_last")` to silence this message.
## i Use `summarise(.by = c(lowbw, gest4))` for per-operation grouping
##   (`?dplyr::dplyr_by`) instead.

```

```

## # A tibble: 6 x 4
## # Groups:   lowbw [2]
##   lowbw gest4          count percent
##   <dbl> <chr>          <int>   <dbl>
## 1     0 30-35 weeks           6  0.0139
## 2     0 more than 35 weeks    425  0.986
## 3     1 25-30 weeks           4  0.0678
## 4     1 30-35 weeks          20  0.339
## 5     1 less than 25 weeks     1  0.0169
## 6     1 more than 35 weeks    34  0.576

```

Similarly we can be interested in the birth weight distribution per gestational time.

```

births_tbl |>
  filter(
    !is.na(gest4)
  ) |>
  group_by(gest4, lowbw) |>
  summarise(
    count = n()
  ) |>
  # compute the percentage of babies in each birth weight category
  # per gestational time category
  mutate(
    percent = count / sum(count, na.rm = TRUE)
  )

```

```

## `summarise()` has regrouped the output.
## i Summaries were computed grouped by gest4 and lowbw.
## i Output is grouped by gest4.
## i Use `summarise(.groups = "drop_last")` to silence this message.
## i Use `summarise(.by = c(gest4, lowbw))` for per-operation grouping
##   (`?dplyr::dplyr_by`) instead.

```

```

## # A tibble: 6 x 4
## # Groups:   gest4 [4]
##   gest4          lowbw count percent
##   <chr>          <dbl> <int>   <dbl>
## 1 25-30 weeks           1     4     1
## 2 30-35 weeks           0     6  0.231
## 3 30-35 weeks           1    20  0.769
## 4 less than 25 weeks     1     1     1
## 5 more than 35 weeks     0   425  0.926
## 6 more than 35 weeks     1    34  0.0741

```

Note: grouping order matters! and can be confusing so think about ungrouping intermediate tables.

3.9 Bind and join tables

Another nice feature of `dplyr` is tables binding and joining. To practice we will create two `tibbles`:

- `age` an individual database which contains `pid` (unique individuals id) and their `age` in year
- `center` an study center database which contains `pid` (unique individuals id) and `center` (the center where an individual is registered coded as a letter)

```
age <-
  tibble(
    pid = 1:6,
    age = sample(15:25, size = 6, replace = TRUE)
  )

center <-
  tibble(
    pid = c(1, 2, 3, 4, 10),
    center = c("A", "B", "A", "B", "C")
  )

age
```

```
## # A tibble: 6 x 2
##   pid  age
##   <int> <int>
## 1     1    20
## 2     2    23
## 3     3    19
## 4     4    23
## 5     5    24
## 6     6    15
```

```
center
```

```
## # A tibble: 5 x 2
##   pid center
##   <dbl> <chr>
## 1     1 A
## 2     2 B
## 3     3 A
## 4     4 B
## 5    10 C
```

Now the tables are define we will try to make the linkage between individuals ages and the center they belong to. First of all let's have a look to `bind_rows` function.

```
bind_rows(age, center)
```

```
## # A tibble: 11 x 3
##   pid  age center
##   <dbl> <int> <chr>
## 1     1    20 <NA>
## 2     2    23 <NA>
## 3     3    19 <NA>
## 4     4    23 <NA>
## 5     5    24 <NA>
## 6     6    15 <NA>
## 7     1     NA A
```

```
## 8      2      NA B
## 9      3      NA A
## 10     4      NA B
## 11    10     NA C
```

Is it useful? Here not really because we do not want to *bind* the data-set (but *join* them instead) but that can be in other situations (e.g. several individuals data base to merge..).

Note: in `bind_rows`, if columns names do not match, they are fill with NA.

Here we want to join the 2 `tibble` according to their common attribute `pid`. Depending on the context you can be interested in joining tables differently. Have a look at the differences between `left_join`, `full_join` and `inner_join`.

```
# all individuals from ages are kept
left_join(age, center, by = c("pid"))
```

```
## # A tibble: 6 x 3
##   pid  age center
##   <dbl> <int> <chr>
## 1     1    20 A
## 2     2    23 B
## 3     3    19 A
## 4     4    23 B
## 5     5    24 <NA>
## 6     6    15 <NA>
```

```
# everithing is kept
full_join(age, center, by = c("pid"))
```

```
## # A tibble: 7 x 3
##   pid  age center
##   <dbl> <int> <chr>
## 1     1    20 A
## 2     2    23 B
## 3     3    19 A
## 4     4    23 B
## 5     5    24 <NA>
## 6     6    15 <NA>
## 7    10     NA C
```

```
# only the individuals present in both dataset are kept
inner_join(age, center, by = c("pid"))
```

```
## # A tibble: 4 x 3
##   pid  age center
##   <dbl> <int> <chr>
## 1     1    20 A
## 2     2    23 B
## 3     3    19 A
## 4     4    23 B
```

Can you spot the differences between the commands above? As an exercise, you can try to compute the individuals' mean age per center.

```
inner_join(age, center, by = c("pid")) |>
  group_by(center) |>
  summarise(
    mean_age = mean(age)
  )
```

```
## # A tibble: 2 x 2
```

```
##   center mean_age
##   <chr>      <dbl>
## 1 A          19.5
## 2 B          23
```

Note: the `by` argument indicates which column should be used to make the *join*. In some cases, you might have to use several columns to match (e.g. per sex and age group), this can be easily done specifying a vector of column names.


From now on, we will consider other packages than `dplyr` from the `tidyverse` suite.

3.10 Data Visualization with ggplot2

One of the packages that have contributed to `tidyverse` success is for sure `ggplot2`. We will go more into the details on how to produce advanced graphs with `ggplot2` in another practical. Let's just have a quick example of graphic creation using `ggplot2`. Let's draw a bar plot to visualize the number of births by women age group. First you have to create a table with the number of birth per age group.

```
birth_per_ageg <- births_tbl |>
  group_by(agegrp) |>
  summarise(total_births = n())
```


```
(gg.01 <-
  ggplot(birth_per_ageg, aes(x = agegrp, y = total_births)) +
  geom_bar(stat = "identity"))
```



This graph can be customized by adding labels and a title to the plot:

```
(gg.02 <-
  gg.01 +
  xlab("Women Age Group") +
```

```
ylab("Total Births") +
ggtitle("Number of Births per Women Age Group"))
```



As

you can see, plots from `ggplot` family are built incrementally using the `+` operator for each additional element.

3.11 pivoting data with `tidyr`

`dplyr` often comes with its good friend `tidyr` when we are performing data manipulation. `tidyr` main features is to reshape tables from long to wide format and vis-versa. Let's have an example. Let's transform in wide format the previously created `birth_per_ageg` table. We want to have a table with one column per age group containing the `total_births` numbers.

```
birth_per_ageg
```

```
## # A tibble: 5 x 2
##   agegrp total_births
##   <fct>      <int>
## 1 [20,25)         2
## 2 [25,30)        68
## 3 [30,35)       200
## 4 [35,40)       194
## 5 [40,45)        36

birth_per_ageg_wide <-
  birth_per_ageg |>
  pivot_wider(
    names_from = "agegrp",
    values_from = "total_births"
  )
```

```
birth_per_ageg_wide
```

```
## # A tibble: 1 x 5
##   `[20,25)` `[25,30)` `[30,35)` `[35,40)` `[40,45)`
##   <int>      <int>      <int>      <int>      <int>
## 1         2        68        200        194         36
```

This table can easily be formatted back in long format using `pivot_longer` function:

```
birth_per_ageg_long <-
  birth_per_ageg_wide |>
  pivot_longer(
    cols = 1:5,
    names_to = "agegrp",
    values_to = "total_births"
  )
```

```
birth_per_ageg_long
```

```
## # A tibble: 5 x 2
##   agegrp total_births
##   <chr>      <int>
## 1 [20,25)         2
## 2 [25,30)        68
## 3 [30,35)       200
## 4 [35,40)       194
## 5 [40,45)        36
```

Are the tables `birth_per_ageg` and `birth_per_ageg_long` identical?

```
identical(birth_per_ageg, birth_per_ageg_long)
```

```
## [1] FALSE
```

Not really because the factor type of `agegrp` column has been lost during the transformation. Let's convert `agegrp` column into a factor. Is the new table identical to `birth_per_ageg`?

```
birth_per_ageg_long_02 <-
  birth_per_ageg_long |>
  mutate(agegrp = as.factor(agegrp))
```

```
identical(birth_per_ageg, birth_per_ageg_long_02)
```

```
## [1] TRUE
```

Here we have seen the simplest example you can have of table reshaping with `tidyr`. If you are interested check the dedicated vignette (`vignette("pivot")`) to learn how to perform more advanced tables reshaping.

3.12 reading files with readr

Another package from `tidyverse` that can be introduced here is `readr` that contains a set of functions equivalent to the core R data.frame reading functions (e.g. `read.table()`, `read.csv()`, `read.delim()`, ...). The main change is that data are loaded in R as `tibble` instead of `data.frame`, type of variables (columns) are *guessed* if possible, and some extra data checking tests are performed.

Let's explore this differences with `fem` dataset available in `data` directory.

```
# read a csv using core R
fem.csv.core <- read.csv("data/fem.csv")
# read a csv using tidyverse
fem.csv.tidy <- read_csv("data/fem.csv")
```

```
## Rows: 118 Columns: 9
## -- Column specification -----
## Delimiter: ","
## dbl (9): ID, AGE, IQ, ANXIETY, DEPRESS, SLEEP, SEX, LIFE, WEIGHT
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
# compare
fem.csv.core
```

	ID	AGE	IQ	ANXIETY	DEPRESS	SLEEP	SEX	LIFE	WEIGHT
## 1	1	39	94	2	2	2	1	1	2.23
## 2	2	41	89	2	2	2	1	1	1.00
## 3	3	42	83	3	3	2	1	1	1.82
## 4	4	30	99	2	2	2	1	1	-1.18
## 5	5	35	94	2	1	1	1	2	-0.14
## 6	6	44	90	NA	1	2	2	2	0.41
## 7	7	31	94	2	2	NA	1	1	-0.68
## 8	8	39	87	3	2	2	1	2	1.59
## 9	9	35	NA	3	2	2	1	1	-0.55
## 10	10	33	92	2	2	2	1	1	0.36
## 11	11	38	92	2	1	1	2	2	-0.86
## 12	12	31	94	2	2	2	NA	2	2.50
## 13	13	40	91	3	2	2	1	2	1.23
## 14	14	44	86	2	2	2	1	1	2.00
## 15	15	43	90	3	2	2	1	1	1.45
## 16	16	32	NA	1	1	1	1	2	-0.68
## 17	17	32	91	1	2	2	NA	2	-0.86
## 18	18	43	82	4	3	2	1	1	3.77
## 19	19	46	86	3	2	2	1	1	1.64
## 20	20	30	88	2	2	2	1	2	0.64
## 21	21	34	97	3	3	NA	1	1	NA
## 22	22	37	96	3	2	2	1	2	NA
## 23	23	35	95	2	1	2	1	2	-0.45
## 24	24	45	87	2	2	2	1	1	2.95
## 25	25	35	103	2	2	2	1	2	-0.95
## 26	26	31	NA	2	2	2	1	2	-0.18
## 27	27	32	91	2	2	2	1	2	-0.86
## 28	28	44	87	2	2	2	1	1	1.68
## 29	29	40	91	3	3	2	1	1	2.05
## 30	30	42	89	3	3	2	1	1	1.91
## 31	31	36	92	3	NA	2	1	1	NA
## 32	32	42	84	3	3	2	1	1	0.77
## 33	33	46	94	2	NA	2	1	1	2.18
## 34	34	41	92	2	1	2	1	2	0.77
## 35	35	30	96	NA	2	2	1	1	-1.36
## 36	36	39	96	2	2	2	2	2	0.36
## 37	37	40	86	2	3	2	1	1	0.68
## 38	38	42	92	3	2	2	1	2	0.59
## 39	39	35	102	2	2	2	1	1	1.36
## 40	40	31	82	2	2	2	1	2	0.45
## 41	41	33	92	3	3	2	1	1	0.68
## 42	42	43	90	NA	NA	2	1	1	1.55
## 43	43	37	92	2	1	1	2	2	NA
## 44	44	32	88	4	2	2	1	2	NA
## 45	45	34	98	2	2	2	1	NA	0.27
## 46	46	34	93	3	2	2	1	1	0.27

## 47	47	42	90	2	1	1	1	2	1.50
## 48	48	41	91	2	1	1	2	2	2.18
## 49	49	31	NA	3	1	2	1	2	-1.00
## 50	50	32	92	3	2	2	1	1	0.45
## 51	51	29	92	2	2	2	2	1	-0.55
## 52	52	41	91	2	2	2	1	1	1.82
## 53	53	39	91	2	2	2	1	1	2.68
## 54	54	41	86	2	1	1	1	2	0.09
## 55	55	34	95	2	1	1	1	2	1.59
## 56	56	39	91	1	1	2	2	2	1.32
## 57	57	35	96	3	2	2	2	2	-0.27
## 58	58	31	100	2	2	2	1	1	-0.27
## 59	59	32	99	4	3	2	1	1	-1.14
## 60	60	41	89	2	1	2	2	2	1.45
## 61	61	41	89	3	2	2	1	1	0.95
## 62	62	44	98	3	2	2	1	1	1.73
## 63	63	35	98	2	2	2	1	2	-1.09
## 64	64	41	103	2	2	2	1	1	-0.36
## 65	65	41	91	3	1	2	1	2	2.64
## 66	66	42	91	4	3	NA	NA	1	1.14
## 67	67	33	94	2	2	2	1	2	-0.82
## 68	68	41	91	2	1	2	1	2	1.95
## 69	69	43	85	2	2	2	2	2	NA
## 70	70	37	92	1	1	2	1	2	0.45
## 71	71	36	96	3	3	2	1	1	1.59
## 72	72	44	90	2	NA	2	1	1	1.50
## 73	73	42	87	2	2	2	2	1	-0.32
## 74	74	31	95	2	3	2	1	1	-0.73
## 75	75	29	95	3	3	2	1	1	-0.09
## 76	76	32	87	1	1	2	1	2	-1.68
## 77	77	35	95	2	2	2	1	1	1.73
## 78	78	42	88	1	1	1	1	2	-0.45
## 79	79	32	94	2	2	2	1	2	2.14
## 80	80	39	NA	3	2	2	1	1	-2.23
## 81	81	34	NA	3	NA	2	1	2	NA
## 82	82	34	87	3	3	2	1	2	1.00
## 83	83	42	92	1	1	2	2	2	2.27
## 84	84	43	86	2	3	2	1	1	0.18
## 85	85	31	93	NA	2	2	1	1	-1.91
## 86	86	31	92	2	2	2	1	2	-0.50
## 87	87	36	106	2	2	2	2	1	-0.45
## 88	88	37	93	2	2	2	1	1	1.91
## 89	89	43	95	2	2	2	1	2	1.09
## 90	90	32	95	3	2	2	1	1	2.23
## 91	91	32	92	NA	NA	NA	1	1	1.36
## 92	92	32	98	2	2	2	1	1	-0.14
## 93	93	43	92	2	2	2	1	1	0.55
## 94	94	41	88	2	2	2	1	2	1.18
## 95	95	43	85	1	1	2	1	2	0.86
## 96	96	39	92	2	2	2	1	2	1.59
## 97	97	41	84	2	2	2	1	1	-0.27
## 98	98	41	92	2	1	2	1	2	0.64
## 99	99	32	91	2	2	2	1	1	2.59
## 100	100	44	86	3	2	2	1	1	2.09
## 101	101	42	92	3	2	2	1	2	NA
## 102	102	39	89	2	2	2	1	2	0.91
## 103	103	45	NA	2	2	2	1	1	0.27

```
## 104 104 39 96      3      NA      2 1 1      NA
## 105 105 31 97      2      NA     NA NA 1 1.27
## 106 106 34 92      3      2      2 1 1 -0.95
## 107 107 41 92      2      2      2 1 1 -1.14
## 108 108 33 98      3      2      2 1 1 1.14
## 109 109 34 91      2      1      1 1 2 2.59
## 110 100 42 91      3      3      2 1 1 1.09
## 111 111 40 89      3      1      1 2 2 0.68
## 112 112 35 94      3      3      2 1 1 0.77
## 113 113 41 90      3      2      2 1 1 1.14
## 114 114 32 96      2      1      1 1 2      NA
## 115 115 39 87      2      2      2 2 1      NA
## 116 116 41 86      3      2      1 2 1 -0.45
## 117 117 33 89      1      1      1 2 2 2.95
## 118 118 42 NA      3      2      2 1 1 2.23
```

```
fem.csv.tidy
```

```
## # A tibble: 118 x 9
##       ID   AGE   IQ ANXIETY DEPRESS SLEEP  SEX  LIFE WEIGHT
##   <dbl> <dbl> <dbl>   <dbl>   <dbl> <dbl> <dbl> <dbl> <dbl>
## 1     1     39    94     2       2     2     1     1  2.23
## 2     2     41    89     2       2     2     1     1     1
## 3     3     42    83     3       3     2     1     1  1.82
## 4     4     30    99     2       2     2     1     1 -1.18
## 5     5     35    94     2       1     1     1     2 -0.14
## 6     6     44    90    NA       1     2     2     2  0.41
## 7     7     31    94     2       2    NA     1     1 -0.68
## 8     8     39    87     3       2     2     1     2  1.59
## 9     9     35    NA     3       2     2     1     1 -0.55
## 10    10     33    92     2       2     2     1     1  0.36
## # i 108 more rows
```

```
# table dimensions
```

```
dim(fem.csv.core)
```

```
## [1] 118 9
```

```
dim(fem.csv.tidy)
```

```
## [1] 118 9
```

```
# compare column types
```

```
map(fem.csv.core, class)
```

```
## $ID
```

```
## [1] "integer"
```

```
##
```

```
## $AGE
```

```
## [1] "integer"
```

```
##
```

```
## $IQ
```

```
## [1] "integer"
```

```
##
```

```
## $ANXIETY
```

```
## [1] "integer"
```

```
##
```

```
## $DEPRESS
```

```
## [1] "integer"
```

```
##
```

```
## $SLEEP
## [1] "integer"
##
## $SEX
## [1] "integer"
##
## $LIFE
## [1] "integer"
##
## $WEIGHT
## [1] "numeric"
map(fem.csv.tidy, class)
```

```
## $ID
## [1] "numeric"
##
## $AGE
## [1] "numeric"
##
## $IQ
## [1] "numeric"
##
## $ANXIETY
## [1] "numeric"
##
## $DEPRESS
## [1] "numeric"
##
## $SLEEP
## [1] "numeric"
##
## $SEX
## [1] "numeric"
##
## $LIFE
## [1] "numeric"
##
## $WEIGHT
## [1] "numeric"
```

note: in case you do not fully get the last lines and the `map()` call, it will be explained in the next section on `purrr` package.

Here we see that the only difference is the type of object loaded `data.frame` vs `tibble` and the default type chosen to cast numeric values (`integer` vs `numeric`).

What about loading `occoh.txt` you will be using in some other practical in the coming days.

```
# read a csv using core R
occoh.txt.core <- read.table("data/occoh.txt")
# read a csv using tidyverse
occoh.txt.tidy <- read_table("data/occoh.txt")
```

```
##
## -- Column specification -----
## cols(
##   id = col_double(),
##   birth = col_double(),
##   entry = col_date(format = ""),
```

```
## exit = col_date(format = ""),
## death = col_date(format = ""),
## chdeath = col_double()
## )
```

```
## Warning: 1501 parsing failures.
```

```
## row col expected actual file
## 1 -- 6 columns 7 columns 'data/occoh.txt'
## 2 -- 6 columns 7 columns 'data/occoh.txt'
## 3 -- 6 columns 7 columns 'data/occoh.txt'
## 4 -- 6 columns 7 columns 'data/occoh.txt'
## 5 -- 6 columns 7 columns 'data/occoh.txt'
## ... ..
## See problems(...) for more details.
```

```
occoh.txt.tidy <- read_table("data/occoh.txt")
```

```
##
## -- Column specification -----
## cols(
##   id = col_double(),
##   birth = col_double(),
##   entry = col_date(format = ""),
##   exit = col_date(format = ""),
##   death = col_date(format = ""),
##   chdeath = col_double()
## )
```

```
## Warning: 1501 parsing failures.
```

```
## row col expected actual file
## 1 -- 6 columns 7 columns 'data/occoh.txt'
## 2 -- 6 columns 7 columns 'data/occoh.txt'
## 3 -- 6 columns 7 columns 'data/occoh.txt'
## 4 -- 6 columns 7 columns 'data/occoh.txt'
## 5 -- 6 columns 7 columns 'data/occoh.txt'
## ... ..
## See problems(...) for more details.
```

```
# compare
```

```
occoh.txt.core
```

```
##      id      birth      entry      exit death chdeath
## 1      1 1943-02-19 1990-08-14 2009-12-31      0      0
## 2      2 1934-07-06 1990-08-14 2009-12-31      0      0
## 3      3 1939-03-05 1990-08-14 2009-12-31      0      0
## 4      4 1939-07-03 1990-08-14 2009-12-31      0      0
## 5      5 1935-02-18 1990-08-14 2006-03-13      1      0
## 6      6 1936-03-07 1990-08-14 2007-06-10      1      0
## 7      7 1944-03-30 1990-08-15 2007-04-14      1      0
## 8      8 1942-11-24 1990-08-15 2006-10-30      1      1
## 9      9 1942-09-11 1990-08-15 2009-12-31      0      0
## 10     10 1931-03-01 1990-08-15 2009-12-31      0      0
## 11     11 1943-02-20 1990-08-15 2009-12-31      0      0
## 12     12 1934-07-26 1990-08-15 2009-12-31      0      0
## 13     13 1935-10-04 1990-08-18 2009-12-31      0      0
## 14     14 1931-12-09 1990-08-18 2009-12-31      0      0
## 15     15 1931-08-24 1990-08-18 2009-12-31      0      0
## 17     16 1947-03-29 1990-08-18 2009-12-31      0      0
## 18     17 1939-03-19 1990-08-18 2009-12-31      0      0
```

## 19	18	1940-07-19	1990-08-18	2009-12-31	0	0
## 20	19	1945-08-10	1990-08-18	2005-07-30	1	0
## 21	20	1945-01-13	1990-08-18	2009-12-31	0	0
## 22	21	1942-02-08	1990-08-18	2009-12-31	0	0
## 23	22	1941-04-11	1990-08-18	2009-12-31	0	0
## 24	23	1948-07-07	1990-08-18	2009-12-31	0	0
## 25	24	1944-02-12	1990-08-18	2009-12-31	0	0
## 26	25	1947-11-19	1990-08-19	2009-12-31	0	0
## 27	26	1945-11-24	1990-08-19	2009-12-31	0	0
## 28	27	1941-05-01	1990-08-19	2009-12-31	0	0
## 29	28	1932-03-19	1990-08-19	2009-12-31	0	0
## 30	29	1944-02-20	1990-08-19	2009-12-31	0	0
## 31	30	1943-09-26	1990-08-19	2009-12-31	0	0
## 32	31	1948-08-29	1990-08-20	2009-12-31	0	0
## 33	32	1949-06-18	1990-08-20	2009-12-31	0	0
## 34	33	1937-08-15	1990-08-20	2009-12-31	0	0
## 35	34	1938-09-14	1990-08-20	2009-12-31	0	0
## 36	35	1943-11-13	1990-08-20	2009-12-31	0	0
## 37	36	1949-01-22	1990-08-20	2009-12-31	0	0
## 38	37	1950-04-30	1990-08-21	2009-12-31	0	0
## 39	38	1943-09-09	1990-08-21	2009-12-31	0	0
## 40	39	1943-11-25	1990-08-21	2009-12-31	0	0
## 41	40	1945-06-22	1990-08-21	1995-03-26	1	0
## 42	41	1950-06-11	1990-08-21	2009-12-14	1	1
## 43	42	1933-10-06	1990-08-21	2004-01-14	1	1
## 44	43	1936-04-28	1990-08-21	2003-06-13	1	1
## 45	44	1943-06-29	1990-08-22	2009-12-31	0	0
## 46	45	1945-06-18	1990-08-22	2009-12-31	0	0
## 47	46	1942-01-05	1990-08-22	2009-12-31	0	0
## 48	47	1930-03-01	1990-08-22	2009-12-31	0	0
## 49	48	1932-01-10	1990-08-22	2005-12-08	1	0
## 50	49	1939-01-31	1990-08-22	2009-12-31	0	0
## 51	50	1947-03-10	1990-08-25	2009-12-31	0	0
## 52	51	1935-04-08	1990-08-26	2006-04-26	1	1
## 54	52	1950-05-10	1990-08-27	2009-12-31	0	0
## 55	53	1939-02-14	1990-08-27	2009-12-31	0	0
## 56	54	1943-10-20	1990-08-27	2009-12-31	0	0
## 57	55	1935-08-21	1990-08-27	2009-12-31	0	0
## 59	56	1935-07-21	1990-08-29	2009-12-31	0	0
## 60	57	1942-01-01	1990-08-29	2009-12-31	0	0
## 61	58	1944-05-29	1990-09-05	2009-12-31	0	0
## 62	59	1936-08-17	1990-09-05	2009-12-31	0	0
## 63	60	1945-01-20	1990-09-12	2009-12-31	0	0
## 64	61	1944-01-14	1990-09-12	2009-12-31	0	0
## 65	62	1932-09-19	1990-09-29	1999-02-01	1	1
## 66	63	1941-02-28	1990-09-29	2004-06-06	1	1
## 67	64	1944-09-21	1990-09-29	2009-12-31	0	0
## 68	65	1940-07-04	1990-09-29	2009-12-31	0	0
## 69	66	1945-04-23	1990-09-29	2006-09-15	1	0
## 70	67	1944-08-21	1990-09-29	2009-12-31	0	0
## 71	68	1937-02-06	1990-09-29	2009-12-19	1	0
## 72	69	1948-04-28	1990-09-29	2009-12-31	0	0
## 73	70	1937-06-11	1990-09-29	2009-12-31	0	0
## 74	71	1948-06-28	1990-09-30	2009-12-31	0	0
## 75	72	1945-07-22	1990-09-30	2009-12-31	0	0
## 76	73	1939-08-14	1990-09-30	2009-12-31	0	0
## 77	74	1946-05-02	1990-09-30	2009-12-31	0	0

## 78	75	1931-09-05	1990-09-30	2000-01-27	1	0
## 79	76	1943-07-03	1990-09-30	2009-12-31	0	0
## 80	77	1938-06-29	1990-10-01	2009-12-31	0	0
## 81	78	1936-06-27	1990-10-01	2009-12-31	0	0
## 82	79	1944-11-11	1990-10-01	2009-12-31	0	0
## 83	80	1944-05-22	1990-10-01	2009-12-31	0	0
## 84	81	1947-06-07	1990-10-01	2009-12-31	0	0
## 85	82	1943-09-12	1990-10-01	2009-12-31	0	0
## 86	83	1933-10-14	1990-10-01	2009-12-31	0	0
## 87	84	1943-01-23	1990-10-01	2009-12-31	0	0
## 88	85	1945-07-12	1990-10-01	2008-02-09	1	1
## 89	86	1947-10-15	1990-10-02	2009-12-31	0	0
## 90	87	1935-12-23	1990-10-02	2009-12-31	0	0
## 91	88	1931-09-09	1990-10-02	2009-12-31	0	0
## 92	89	1938-11-18	1990-10-03	2009-12-31	0	0
## 93	90	1943-04-10	1990-10-03	2009-12-31	0	0
## 94	91	1930-06-26	1990-10-03	2009-12-31	0	0
## 95	92	1941-12-15	1990-10-03	2009-12-31	0	0
## 96	93	1947-07-29	1990-10-03	2009-12-31	0	0
## 97	94	1948-10-30	1990-10-03	2009-12-31	0	0
## 98	95	1943-03-19	1990-10-03	2009-11-21	1	1
## 99	96	1945-09-13	1990-10-06	2009-12-31	0	0
## 100	97	1933-01-29	1990-10-06	1993-11-28	1	1
## 101	98	1935-05-07	1990-10-06	2004-12-11	1	0
## 102	99	1940-07-18	1990-10-06	1997-08-19	1	0
## 104	100	1946-01-16	1990-10-06	2009-12-31	0	0
## 105	101	1940-02-12	1990-10-06	2004-08-02	1	0
## 106	102	1950-05-13	1990-10-07	2009-12-31	0	0
## 107	103	1937-04-11	1990-10-07	2000-01-15	1	1
## 108	104	1936-08-28	1990-10-07	2001-04-12	1	1
## 109	105	1935-12-04	1990-10-07	2009-12-31	0	0
## 110	106	1943-03-06	1990-10-07	2009-12-31	0	0
## 111	107	1950-08-30	1990-10-07	2009-12-31	0	0
## 112	108	1942-04-16	1990-10-08	2009-12-31	0	0
## 113	109	1941-04-16	1990-10-08	2009-12-31	0	0
## 114	110	1946-02-18	1990-10-08	2009-12-31	0	0
## 115	111	1946-01-28	1990-10-08	2009-12-31	0	0
## 116	112	1938-08-28	1990-10-08	2009-12-31	0	0
## 117	113	1934-01-08	1990-10-09	2009-12-31	0	0
## 118	114	1935-02-04	1990-10-09	2009-12-31	0	0
## 119	115	1943-01-06	1990-10-09	2007-12-19	1	1
## 120	116	1943-07-27	1990-10-09	2009-12-31	0	0
## 121	117	1933-02-07	1990-10-09	2009-12-31	0	0
## 122	118	1934-04-28	1990-10-09	2009-06-29	1	0
## 123	119	1938-12-11	1990-10-09	2009-12-31	0	0
## 124	120	1931-06-21	1990-10-09	1993-05-12	1	1
## 125	121	1949-09-12	1990-10-10	2009-12-31	0	0
## 126	122	1949-01-27	1990-10-10	2009-12-31	0	0
## 127	123	1947-01-22	1990-10-10	2009-12-31	0	0
## 128	124	1929-12-23	1990-10-10	1993-02-03	1	0
## 129	125	1931-02-19	1990-10-15	2006-06-08	1	0
## 130	126	1938-02-24	1990-10-15	2009-12-31	0	0
## 131	127	1942-10-26	1990-10-15	2009-12-31	0	0
## 132	128	1940-10-14	1990-10-15	2009-12-31	0	0
## 133	129	1934-12-28	1990-10-15	2009-12-31	0	0
## 134	130	1945-02-23	1990-10-27	2009-12-31	0	0
## 135	131	1932-05-12	1990-10-27	2009-12-31	0	0

## 136	132	1942-07-25	1990-10-27	2009-12-31	0	0
## 137	133	1933-06-25	1990-10-28	1996-12-15	1	0
## 138	134	1937-12-05	1990-10-28	2009-12-31	0	0
## 139	135	1947-06-23	1990-10-28	2009-12-31	0	0
## 140	136	1945-09-04	1990-10-29	1999-05-23	1	1
## 142	137	1931-08-29	1990-10-29	2002-06-01	1	0
## 143	138	1930-02-11	1990-10-29	2009-12-31	0	0
## 144	139	1944-05-12	1990-10-29	2009-12-31	0	0
## 146	140	1936-07-01	1990-10-29	2009-12-31	0	0
## 147	141	1937-10-10	1990-10-30	2009-12-31	0	0
## 148	142	1945-01-25	1990-10-30	2009-12-31	0	0
## 149	143	1939-06-11	1990-10-30	2009-12-31	0	0
## 150	144	1941-06-15	1990-10-30	2009-12-31	0	0
## 151	145	1941-08-01	1990-10-30	2009-12-31	0	0
## 152	146	1938-09-15	1990-10-30	2009-12-31	0	0
## 153	147	1939-11-30	1990-10-30	2009-12-31	0	0
## 154	148	1936-11-27	1990-10-30	1992-07-05	1	1
## 155	149	1935-04-05	1990-10-30	1992-04-06	1	1
## 156	150	1943-03-08	1990-10-31	2009-12-31	0	0
## 160	151	1942-03-03	1990-10-31	2009-12-31	0	0
## 161	152	1937-03-12	1990-11-03	2009-12-31	0	0
## 162	153	1941-07-27	1990-11-03	2009-12-31	0	0
## 163	154	1940-01-23	1990-11-03	2006-06-01	1	1
## 164	155	1936-02-24	1990-11-03	2000-01-23	1	0
## 165	156	1949-08-08	1990-11-03	2009-12-31	0	0
## 167	157	1936-01-13	1990-11-04	2009-12-31	0	0
## 168	158	1949-05-16	1990-11-04	2009-12-31	0	0
## 169	159	1939-08-06	1990-11-04	2009-12-31	0	0
## 170	160	1948-06-04	1990-11-04	2009-12-31	0	0
## 173	161	1948-12-14	1990-11-04	2009-12-31	0	0
## 174	162	1946-05-02	1990-11-04	2009-12-31	0	0
## 175	163	1931-05-16	1990-11-05	2009-12-31	0	0
## 176	164	1942-09-16	1990-11-05	2009-12-31	0	0
## 177	165	1937-09-16	1990-11-05	2009-12-31	0	0
## 178	166	1950-06-14	1990-11-05	2009-12-31	0	0
## 180	167	1950-10-17	1990-11-05	2009-12-31	0	0
## 181	168	1936-11-17	1990-11-05	2000-06-19	1	1
## 182	169	1942-05-15	1990-11-05	2009-12-31	0	0
## 183	170	1943-01-18	1990-11-06	2002-06-20	1	0
## 184	171	1938-09-10	1990-11-06	2009-12-31	0	0
## 185	172	1948-05-20	1990-11-06	2009-12-31	0	0
## 186	173	1944-07-30	1990-11-06	2009-12-31	0	0
## 187	174	1938-10-06	1990-11-06	2009-12-31	0	0
## 188	175	1944-07-11	1990-11-06	2009-12-31	0	0
## 189	176	1944-03-31	1990-11-06	2009-12-31	0	0
## 190	177	1944-08-21	1990-11-07	2009-12-31	0	0
## 191	178	1940-11-12	1990-11-07	2009-12-31	0	0
## 192	179	1939-01-24	1990-11-07	2009-10-05	1	0
## 193	180	1932-01-04	1990-11-07	2009-12-31	0	0
## 194	181	1935-01-12	1990-11-07	2009-12-31	0	0
## 195	182	1949-05-10	1990-11-07	2009-12-31	0	0
## 196	183	1946-06-17	1990-11-07	2009-12-31	0	0
## 197	184	1943-12-27	1990-11-07	1997-12-14	1	1
## 198	185	1949-09-05	1990-11-07	2009-12-31	0	0
## 199	186	1943-02-21	1990-11-07	2009-12-31	0	0
## 200	187	1939-02-20	1990-11-10	2009-12-31	0	0
## 201	188	1946-12-26	1990-11-10	2009-12-31	0	0

##	203	189	1939-01-03	1990-11-10	2009-12-31	0	0
##	204	190	1946-12-31	1990-11-10	2007-01-17	1	0
##	205	191	1942-05-09	1990-11-10	2000-07-01	1	1
##	206	192	1942-10-14	1990-11-11	2009-12-31	0	0
##	207	193	1942-05-24	1990-11-11	2009-12-31	0	0
##	208	194	1945-03-13	1990-11-11	2009-12-31	0	0
##	209	195	1929-08-01	1990-11-11	2009-12-31	0	0
##	210	196	1949-11-25	1990-11-11	2009-12-31	0	0
##	211	197	1938-04-25	1990-11-12	2009-12-31	0	0
##	212	198	1937-05-10	1990-11-12	2009-12-31	0	0
##	213	199	1936-04-29	1990-11-12	2006-11-26	1	1
##	214	200	1931-05-21	1990-11-12	1999-02-02	1	0
##	215	201	1932-02-14	1990-11-13	1996-07-24	1	1
##	216	202	1947-06-21	1990-11-13	2009-05-13	1	0
##	217	203	1942-10-31	1990-11-13	2009-12-31	0	0
##	218	204	1943-04-22	1990-11-13	2009-12-31	0	0
##	219	205	1945-01-25	1990-11-14	2009-12-31	0	0
##	220	206	1945-05-28	1990-11-14	2009-12-31	0	0
##	221	207	1931-05-13	1990-11-14	2009-12-31	0	0
##	222	208	1940-08-17	1990-11-14	2000-11-05	1	0
##	223	209	1948-03-17	1990-11-14	2002-05-27	1	0
##	225	210	1930-09-17	1990-11-17	1998-12-06	1	0
##	226	211	1949-10-03	1990-11-17	2009-12-31	0	0
##	227	212	1948-01-16	1990-11-17	2009-12-31	0	0
##	228	213	1948-04-20	1990-11-17	2009-12-31	0	0
##	229	214	1944-03-09	1990-11-17	2009-12-31	0	0
##	230	215	1933-03-21	1990-11-17	2009-12-31	0	0
##	231	216	1946-10-08	1990-11-17	2009-12-31	0	0
##	232	217	1936-04-19	1990-11-18	2009-12-31	0	0
##	233	218	1949-12-24	1990-11-18	2006-03-28	1	0
##	234	219	1934-09-24	1990-11-18	2009-12-31	0	0
##	235	220	1944-12-13	1990-11-18	2009-12-31	0	0
##	236	221	1944-02-28	1990-11-18	2009-12-31	0	0
##	237	222	1948-07-20	1990-11-18	2009-12-31	0	0
##	238	223	1941-04-28	1990-11-18	2009-12-31	0	0
##	239	224	1946-07-04	1990-11-19	2009-12-31	0	0
##	240	225	1939-04-25	1990-11-19	2009-12-31	0	0
##	241	226	1934-02-11	1990-11-19	2009-12-31	0	0
##	242	227	1943-11-21	1990-12-08	2009-12-31	0	0
##	243	228	1935-04-25	1990-12-08	2009-12-31	0	0
##	244	229	1938-01-03	1990-12-08	2009-12-31	0	0
##	245	230	1947-02-23	1990-12-08	2009-12-31	0	0
##	246	231	1943-04-06	1990-12-08	2009-12-31	0	0
##	247	232	1944-10-03	1990-12-08	2009-12-31	0	0
##	248	233	1948-06-14	1990-12-09	2009-12-31	0	0
##	249	234	1936-04-23	1990-12-09	2009-12-31	0	0
##	250	235	1933-03-31	1990-12-09	2009-12-31	0	0
##	251	236	1934-09-30	1990-12-09	2008-10-14	1	0
##	252	237	1933-03-06	1990-12-09	2005-02-18	1	0
##	253	238	1942-05-10	1990-12-09	2009-12-31	0	0
##	255	239	1933-09-08	1990-12-10	2009-12-31	0	0
##	256	240	1945-11-17	1990-12-10	2009-12-31	0	0
##	257	241	1943-09-15	1990-12-10	2009-12-31	0	0
##	258	242	1933-07-27	1990-12-10	2009-12-31	0	0
##	259	243	1942-10-05	1990-12-10	2009-12-31	0	0
##	260	244	1950-02-06	1990-12-10	2009-12-31	0	0
##	261	245	1937-03-18	1990-12-10	2009-12-31	0	0

## 262	246	1930-10-03	1990-12-29	2003-03-07	1	0
## 263	247	1941-07-07	1990-12-29	2009-12-31	0	0
## 264	248	1935-11-16	1990-12-30	2000-04-13	1	0
## 266	249	1933-03-27	1990-12-30	2009-12-31	0	0
## 267	250	1933-07-05	1990-12-30	1995-05-19	1	0
## 268	251	1945-08-18	1990-12-31	2009-12-31	0	0
## 269	252	1939-04-04	1990-12-31	2005-11-08	1	1
## 270	253	1945-01-21	1991-01-01	2009-12-31	0	0
## 271	254	1944-04-28	1991-01-01	2009-12-31	0	0
## 272	255	1934-08-30	1991-01-01	1992-02-25	1	1
## 273	256	1933-12-29	1991-01-01	2009-12-31	0	0
## 274	257	1930-09-08	1991-01-01	2006-12-15	1	1
## 275	258	1930-12-31	1991-01-02	1995-07-16	1	0
## 276	259	1941-05-18	1991-01-02	2009-12-31	0	0
## 277	260	1944-12-02	1991-01-02	2004-05-28	1	0
## 278	261	1943-06-14	1991-01-02	2009-12-31	0	0
## 279	262	1940-04-27	1991-01-02	2009-12-31	0	0
## 281	263	1938-01-10	1991-01-05	2009-12-31	0	0
## 282	264	1931-01-28	1991-01-05	2002-06-04	1	0
## 283	265	1945-10-18	1991-01-05	2009-12-31	0	0
## 285	266	1948-03-09	1991-01-05	2009-12-31	0	0
## 286	267	1944-10-18	1991-01-06	2009-12-31	0	0
## 288	268	1949-11-02	1991-01-06	2009-12-31	0	0
## 289	269	1935-03-31	1991-01-06	2009-12-31	0	0
## 290	270	1943-04-01	1991-01-06	2009-12-31	0	0
## 291	271	1930-10-24	1991-01-06	2009-12-31	0	0
## 292	272	1944-04-16	1991-01-06	2009-12-31	0	0
## 293	273	1941-06-13	1991-01-07	2009-12-31	0	0
## 294	274	1946-05-09	1991-01-07	2009-12-31	0	0
## 295	275	1933-02-28	1991-01-07	2009-12-31	0	0
## 296	276	1935-07-05	1991-01-08	2009-12-31	0	0
## 297	277	1949-12-24	1991-01-08	2009-12-31	0	0
## 298	278	1943-03-01	1991-01-08	2009-12-31	0	0
## 299	279	1930-02-24	1991-01-08	2009-12-31	0	0
## 300	280	1948-09-02	1991-01-09	2009-12-31	0	0
## 301	281	1940-04-04	1991-01-09	2009-12-31	0	0
## 302	282	1950-06-23	1991-01-09	2003-02-07	1	0
## 303	283	1942-10-16	1991-01-09	2009-12-31	0	0
## 304	284	1949-08-27	1991-01-09	2009-12-31	0	0
## 305	285	1930-07-19	1991-01-12	2003-03-06	1	0
## 306	286	1944-08-01	1991-01-12	2009-12-31	0	0
## 307	287	1942-04-19	1991-01-12	2009-12-31	0	0
## 308	288	1944-03-28	1991-01-12	2007-09-12	1	1
## 309	289	1945-01-09	1991-01-12	2009-12-31	0	0
## 310	290	1934-05-25	1991-01-12	2009-12-31	0	0
## 311	291	1939-04-19	1991-01-13	1997-05-27	1	0
## 312	292	1943-02-13	1991-01-13	2009-12-31	0	0
## 313	293	1943-01-02	1991-01-13	2009-12-31	0	0
## 314	294	1937-01-01	1991-01-13	2008-01-30	1	0
## 316	295	1935-02-18	1991-01-13	1995-04-12	1	0
## 317	296	1930-01-25	1991-01-13	1995-12-03	1	1
## 319	297	1937-12-17	1991-01-14	2009-12-31	0	0
## 320	298	1935-01-02	1991-01-14	2008-07-17	1	0
## 321	299	1946-05-19	1991-01-14	2009-12-31	0	0
## 322	300	1948-05-03	1991-01-14	2009-12-31	0	0
## 323	301	1942-09-01	1991-01-14	2009-12-31	0	0
## 324	302	1931-04-06	1991-01-14	2009-12-31	0	0

## 325	303	1945-02-20	1991-01-15	2009-12-31	0	0
## 327	304	1942-04-08	1991-01-15	2001-10-13	1	0
## 328	305	1943-03-22	1991-01-15	2009-12-31	0	0
## 329	306	1939-05-28	1991-01-15	2009-12-31	0	0
## 330	307	1934-06-20	1991-01-30	2009-12-31	0	0
## 331	308	1942-06-24	1991-01-30	2008-03-23	1	1
## 333	309	1942-02-14	1991-01-30	2009-12-31	0	0
## 334	310	1940-01-07	1991-03-30	2009-12-31	0	0
## 335	311	1945-11-18	1991-03-30	2009-12-31	0	0
## 337	312	1942-09-30	1991-03-30	1996-08-23	1	0
## 339	313	1940-02-21	1991-03-30	2009-12-31	0	0
## 340	314	1943-09-22	1991-03-30	2009-12-31	0	0
## 341	315	1945-10-05	1991-03-30	2009-12-31	0	0
## 342	316	1945-09-04	1991-03-30	2009-12-31	0	0
## 343	317	1947-04-18	1991-03-30	2009-12-31	0	0
## 344	318	1949-07-29	1991-03-30	2007-12-06	1	1
## 345	319	1941-06-16	1991-03-30	1995-12-17	1	1
## 346	320	1941-04-23	1991-03-31	2009-12-31	0	0
## 347	321	1931-03-31	1991-03-31	2009-12-31	0	0
## 348	322	1932-07-27	1991-03-31	1996-05-23	1	1
## 349	323	1935-02-05	1991-03-31	2009-12-31	0	0
## 350	324	1947-06-27	1991-03-31	2009-12-31	0	0
## 351	325	1937-12-18	1991-03-31	2009-12-31	0	0
## 352	326	1948-11-10	1991-03-31	2009-12-31	0	0
## 353	327	1944-06-11	1991-04-01	2009-12-31	0	0
## 354	328	1943-04-28	1991-04-01	2009-12-31	0	0
## 355	329	1938-03-09	1991-04-01	2009-12-31	0	0
## 356	330	1945-12-04	1991-04-01	2009-12-31	0	0
## 357	331	1942-10-21	1991-04-01	2009-12-31	0	0
## 358	332	1944-07-31	1991-04-01	2009-12-31	0	0
## 359	333	1940-03-11	1991-04-01	2009-12-31	0	0
## 360	334	1932-08-28	1991-04-01	2009-12-31	0	0
## 361	335	1933-02-16	1991-04-02	1996-05-13	1	1
## 362	336	1930-07-05	1991-04-02	1998-08-28	1	1
## 363	337	1937-01-18	1991-04-02	2009-12-31	0	0
## 364	338	1947-08-29	1991-04-03	2001-04-13	1	0
## 365	339	1938-12-17	1991-04-03	2000-03-04	1	1
## 366	340	1945-09-26	1991-04-03	2009-12-31	0	0
## 369	341	1943-05-15	1991-04-03	2009-12-31	0	0
## 370	342	1937-07-09	1991-04-06	2009-12-31	0	0
## 371	343	1939-04-16	1991-04-06	2009-12-31	0	0
## 372	344	1946-03-10	1991-04-06	2009-12-31	0	0
## 373	345	1935-07-01	1991-04-06	2009-12-31	0	0
## 374	346	1946-02-12	1991-04-07	2009-12-31	0	0
## 375	347	1950-06-09	1991-04-07	2009-12-31	0	0
## 376	348	1940-08-13	1991-04-07	2009-12-31	0	0
## 377	349	1939-11-28	1991-04-07	2009-12-31	0	0
## 378	350	1936-08-21	1991-04-07	2009-12-31	0	0
## 379	351	1943-04-16	1991-04-08	2009-12-31	0	0
## 380	352	1933-10-25	1991-04-08	2007-01-24	1	1
## 381	353	1934-06-13	1991-04-08	2009-12-31	0	0
## 382	354	1933-06-20	1991-04-08	2009-12-31	0	0
## 384	355	1943-04-21	1991-04-09	2009-12-31	0	0
## 385	356	1936-09-14	1991-04-09	2009-12-31	0	0
## 386	357	1933-02-28	1991-04-09	2009-12-31	0	0
## 387	358	1946-02-15	1991-04-09	2009-04-04	1	0
## 388	359	1950-11-22	1991-04-10	2002-06-12	1	1

## 389	360	1947-10-07	1991-04-10	2009-12-31	0	0
## 390	361	1947-10-18	1991-04-10	2009-12-31	0	0
## 391	362	1945-07-17	1991-04-10	2009-12-31	0	0
## 392	363	1942-07-28	1991-04-10	2009-12-31	0	0
## 393	364	1938-11-15	1991-04-10	2009-12-31	0	0
## 394	365	1944-03-13	1991-04-10	2009-12-31	0	0
## 395	366	1946-11-22	1991-04-10	2004-06-22	1	1
## 396	367	1944-04-05	1991-04-10	1993-08-06	1	0
## 397	368	1941-04-02	1991-04-10	2009-12-31	0	0
## 399	369	1945-08-08	1991-04-10	2009-12-31	0	0
## 400	370	1944-10-07	1991-04-10	2009-12-31	0	0
## 401	371	1932-04-09	1991-04-13	2005-04-05	1	0
## 402	372	1947-07-05	1991-04-13	2009-12-31	0	0
## 403	373	1944-03-15	1991-04-13	2009-12-31	0	0
## 404	374	1937-01-18	1991-04-13	2005-03-10	1	0
## 405	375	1939-04-25	1991-04-13	2009-12-31	0	0
## 406	376	1943-01-02	1991-04-14	2009-12-31	0	0
## 407	377	1945-12-05	1991-04-14	2009-12-31	0	0
## 408	378	1949-02-23	1991-04-14	2009-12-31	0	0
## 409	379	1939-03-28	1991-04-14	2009-12-31	0	0
## 410	380	1942-04-27	1991-04-14	2009-12-31	0	0
## 411	381	1943-04-16	1991-04-14	2009-12-31	0	0
## 413	382	1947-09-09	1991-04-14	2009-12-31	0	0
## 414	383	1943-05-20	1991-04-15	2009-12-31	0	0
## 415	384	1943-04-04	1991-04-15	2009-12-31	0	0
## 417	385	1944-10-25	1991-04-15	2009-12-31	0	0
## 418	386	1931-07-09	1991-04-15	2009-12-31	0	0
## 419	387	1949-09-25	1991-04-16	2009-12-31	0	0
## 420	388	1939-05-20	1991-04-16	2009-12-31	0	0
## 421	389	1935-12-02	1991-04-16	2009-12-31	0	0
## 422	390	1940-05-23	1991-04-16	2007-08-11	1	0
## 423	391	1938-01-18	1991-04-16	2009-12-31	0	0
## 424	392	1935-02-28	1991-04-17	2009-12-31	0	0
## 425	393	1945-06-03	1991-04-17	2009-12-31	0	0
## 426	394	1942-03-28	1991-04-17	2009-12-31	0	0
## 427	395	1943-05-31	1991-04-17	2002-08-04	1	0
## 429	396	1930-08-18	1991-04-17	2009-12-31	0	0
## 431	397	1936-12-08	1991-04-20	2005-01-01	1	0
## 432	398	1949-11-30	1991-04-20	2009-12-31	0	0
## 433	399	1932-03-03	1991-04-20	1996-02-14	1	0
## 434	400	1943-02-06	1991-04-21	2009-12-31	0	0
## 437	401	1934-08-19	1991-04-21	1995-12-21	1	1
## 439	402	1942-12-31	1991-04-22	2009-12-31	0	0
## 440	403	1945-09-02	1991-04-22	2009-12-31	0	0
## 442	404	1934-11-16	1991-04-22	2009-12-31	0	0
## 443	405	1944-09-20	1991-04-23	2009-12-31	0	0
## 444	406	1931-12-03	1991-04-23	2009-12-31	0	0
## 445	407	1938-07-16	1991-04-23	2009-12-31	0	0
## 446	408	1941-08-29	1991-04-23	2009-12-31	0	0
## 447	409	1930-03-04	1991-04-23	2009-12-31	0	0
## 448	410	1937-10-15	1991-04-23	2009-12-31	0	0
## 449	411	1933-01-14	1991-04-24	1999-11-24	1	0
## 450	412	1932-02-08	1991-04-24	2009-12-31	0	0
## 451	413	1936-01-31	1991-04-24	2006-07-09	1	0
## 452	414	1951-03-02	1991-04-24	2009-12-31	0	0
## 453	415	1936-12-23	1991-04-24	2009-12-31	0	0
## 454	416	1935-11-17	1991-04-24	2009-12-31	0	0

## 455	417	1930-07-14	1991-04-24	2009-12-31	0	0
## 456	418	1932-08-21	1991-04-24	2004-09-08	1	1
## 457	419	1933-12-01	1991-04-24	1997-08-21	1	1
## 458	420	1940-05-14	1991-04-24	2009-12-31	0	0
## 459	421	1945-01-01	1991-04-24	2009-12-31	0	0
## 460	422	1930-03-17	1991-04-30	2009-12-31	0	0
## 461	423	1946-02-04	1991-04-30	2009-12-31	0	0
## 462	424	1940-10-22	1991-04-30	2009-12-31	0	0
## 463	425	1932-10-10	1991-04-30	2007-12-11	1	1
## 464	426	1936-04-23	1991-05-11	1999-12-22	1	1
## 465	427	1947-10-28	1991-05-11	2009-12-31	0	0
## 466	428	1938-07-07	1991-05-11	2008-03-23	1	0
## 468	429	1947-12-31	1991-05-11	2009-12-31	0	0
## 469	430	1931-05-10	1991-05-12	1997-09-03	1	0
## 470	431	1943-04-29	1991-05-12	2009-12-31	0	0
## 471	432	1931-02-28	1991-05-12	2009-12-31	0	0
## 473	433	1940-01-28	1991-05-12	2009-12-31	0	0
## 474	434	1940-10-14	1991-05-12	2009-12-31	0	0
## 475	435	1939-12-30	1991-05-12	2009-12-31	0	0
## 476	436	1938-06-23	1991-05-13	2009-12-31	0	0
## 477	437	1932-03-08	1991-05-13	2009-12-31	0	0
## 478	438	1949-10-29	1991-05-13	2009-12-31	0	0
## 479	439	1947-11-07	1991-05-13	2009-12-31	0	0
## 480	440	1948-12-02	1991-05-13	2009-12-31	0	0
## 481	441	1946-04-11	1991-05-14	2009-12-31	0	0
## 482	442	1949-08-10	1991-05-14	2009-12-31	0	0
## 483	443	1938-04-14	1991-05-15	2009-12-31	0	0
## 484	444	1942-11-05	1991-05-15	2009-12-31	0	0
## 485	445	1940-07-29	1991-05-15	2009-12-31	0	0
## 486	446	1937-12-10	1991-05-15	2009-12-31	0	0
## 487	447	1932-11-19	1991-05-15	1993-11-22	1	1
## 488	448	1941-10-15	1991-05-25	2001-09-30	1	1
## 489	449	1941-04-20	1991-05-25	2009-12-31	0	0
## 490	450	1949-03-23	1991-05-25	2009-12-31	0	0
## 491	451	1935-12-13	1991-05-25	2009-12-31	0	0
## 492	452	1944-03-26	1991-05-25	2009-12-31	0	0
## 493	453	1949-05-13	1991-05-25	2008-07-29	1	0
## 494	454	1935-04-18	1991-05-25	2009-12-31	0	0
## 495	455	1941-01-02	1991-05-26	2009-12-31	0	0
## 497	456	1947-10-16	1991-05-26	2008-05-26	1	0
## 498	457	1935-02-16	1991-05-26	2009-12-31	0	0
## 499	458	1940-04-19	1991-05-26	2009-12-31	0	0
## 500	459	1942-04-16	1991-05-27	2009-12-31	0	0
## 501	460	1935-03-23	1991-05-27	2002-10-09	1	0
## 502	461	1938-05-23	1991-05-27	2009-12-31	0	0
## 503	462	1943-06-13	1991-05-27	2009-12-31	0	0
## 505	463	1946-09-10	1991-05-27	1992-11-13	1	1
## 506	464	1932-02-10	1991-05-27	2008-03-03	1	1
## 507	465	1946-02-25	1991-05-28	2009-12-31	0	0
## 508	466	1937-12-21	1991-05-28	2009-12-31	0	0
## 509	467	1931-03-20	1991-05-28	2009-12-31	0	0
## 510	468	1938-04-13	1991-05-29	2008-03-12	1	1
## 511	469	1936-11-28	1991-05-29	1999-08-05	1	1
## 513	470	1947-12-12	1991-06-01	2009-12-31	0	0
## 514	471	1945-03-29	1991-06-01	2009-12-31	0	0
## 515	472	1937-10-21	1991-06-02	2009-12-31	0	0
## 517	473	1948-07-28	1991-06-02	2009-12-31	0	0

## 518	474	1947-04-07	1991-06-02	2009-12-31	0	0
## 519	475	1934-07-14	1991-06-02	2009-12-31	0	0
## 520	476	1941-02-04	1991-06-02	2009-12-31	0	0
## 521	477	1944-05-20	1991-06-03	2009-12-31	0	0
## 522	478	1936-07-21	1991-06-03	2009-12-31	0	0
## 523	479	1940-06-28	1991-06-03	2009-12-31	0	0
## 524	480	1934-12-30	1991-06-03	2009-12-31	0	0
## 525	481	1948-02-25	1991-06-03	2009-12-31	0	0
## 526	482	1931-09-19	1991-06-04	2009-12-31	0	0
## 527	483	1944-06-17	1991-06-04	2009-12-31	0	0
## 528	484	1948-02-27	1991-06-04	2009-12-31	0	0
## 529	485	1944-09-06	1991-06-05	2009-12-31	0	0
## 530	486	1935-04-16	1991-06-05	2009-12-31	0	0
## 531	487	1943-03-24	1991-06-05	2009-12-31	0	0
## 533	488	1950-04-09	1991-06-05	2009-12-31	0	0
## 534	489	1951-04-27	1991-06-05	2009-12-31	0	0
## 536	490	1948-05-18	1991-06-05	2009-12-31	0	0
## 537	491	1940-06-01	1991-06-05	2009-07-10	1	1
## 538	492	1943-05-17	1991-06-05	2009-12-31	0	0
## 539	493	1944-02-19	1991-06-08	2009-12-31	0	0
## 540	494	1945-10-15	1991-06-08	2009-12-31	0	0
## 541	495	1937-03-03	1991-06-08	2009-12-31	0	0
## 542	496	1932-09-04	1991-06-08	2009-12-31	0	0
## 543	497	1946-11-26	1991-06-08	2009-12-31	0	0
## 544	498	1947-12-03	1991-06-08	2009-12-31	0	0
## 545	499	1940-11-16	1991-06-08	2009-12-31	0	0
## 546	500	1938-11-13	1991-06-08	2009-12-31	0	0
## 547	501	1935-08-02	1991-06-09	2004-04-04	1	1
## 548	502	1942-12-02	1991-06-09	2009-12-31	0	0
## 549	503	1936-11-18	1991-06-09	2009-12-31	0	0
## 550	504	1944-06-01	1991-06-09	2007-07-24	1	1
## 551	505	1947-11-29	1991-06-09	2009-12-31	0	0
## 552	506	1947-08-29	1991-06-09	2009-12-31	0	0
## 553	507	1937-05-26	1991-06-09	2009-12-31	0	0
## 555	508	1945-08-17	1991-06-10	2009-12-31	0	0
## 556	509	1943-12-07	1991-06-11	2009-12-31	0	0
## 557	510	1945-09-18	1991-06-11	2009-12-31	0	0
## 558	511	1946-06-23	1991-06-11	2009-12-31	0	0
## 559	512	1949-07-18	1991-06-11	2009-12-31	0	0
## 560	513	1950-04-10	1991-06-11	2009-12-31	0	0
## 561	514	1933-05-24	1991-06-11	2009-12-31	0	0
## 562	515	1944-08-28	1991-06-11	2009-12-31	0	0
## 563	516	1948-12-06	1991-06-11	2009-12-31	0	0
## 564	517	1937-05-28	1991-06-11	2008-03-01	1	0
## 565	518	1942-08-27	1991-06-12	2009-12-31	0	0
## 566	519	1934-02-11	1991-06-12	2009-12-31	0	0
## 567	520	1941-07-01	1991-06-12	2009-12-31	0	0
## 568	521	1934-06-12	1991-06-12	2003-02-25	1	1
## 569	522	1938-08-25	1991-06-12	2009-12-31	0	0
## 570	523	1947-08-24	1991-06-12	2009-12-31	0	0
## 571	524	1941-02-17	1991-06-12	2009-12-31	0	0
## 572	525	1933-01-23	1991-06-12	2009-12-31	0	0
## 573	526	1944-03-23	1991-06-12	2009-12-31	0	0
## 574	527	1940-01-23	1991-06-15	2009-12-31	0	0
## 575	528	1934-06-04	1991-06-15	2009-12-31	0	0
## 576	529	1932-01-04	1991-06-15	2004-12-30	1	1
## 577	530	1943-01-28	1991-06-15	2009-12-31	0	0

## 578	531	1942-08-21	1991-06-15	2009-12-31	0	0
## 579	532	1944-09-04	1991-06-15	2009-12-31	0	0
## 580	533	1932-06-15	1991-06-15	2009-12-31	0	0
## 581	534	1931-02-19	1991-06-16	2007-12-10	1	1
## 582	535	1932-02-19	1991-06-16	2009-12-31	0	0
## 583	536	1945-09-30	1991-06-16	2009-12-31	0	0
## 584	537	1940-01-13	1991-06-16	2000-04-06	1	0
## 585	538	1944-12-23	1991-06-16	2009-12-31	0	0
## 586	539	1942-10-01	1991-06-16	2009-12-31	0	0
## 587	540	1942-08-06	1991-06-17	2009-12-31	0	0
## 588	541	1931-04-09	1991-06-17	2009-12-31	0	0
## 589	542	1939-11-03	1991-06-17	2009-12-31	0	0
## 590	543	1948-01-24	1991-06-17	2009-12-31	0	0
## 591	544	1936-05-06	1991-06-18	2009-12-31	0	0
## 592	545	1943-08-28	1991-06-18	1993-02-08	1	0
## 593	546	1940-09-22	1991-06-18	2009-12-31	0	0
## 594	547	1934-01-15	1991-06-18	2009-12-31	0	0
## 595	548	1941-12-06	1991-06-18	2009-12-31	0	0
## 596	549	1947-04-28	1991-06-18	2009-12-31	0	0
## 597	550	1943-12-02	1991-06-18	2009-12-31	0	0
## 599	551	1943-04-16	1991-06-19	2009-12-31	0	0
## 600	552	1931-09-15	1991-06-19	2009-12-31	0	0
## 601	553	1943-05-26	1991-06-19	2009-12-31	0	0
## 602	554	1932-10-22	1991-06-19	2009-12-31	0	0
## 604	555	1931-11-20	1991-06-22	2001-11-27	1	1
## 605	556	1935-03-29	1991-06-22	2009-12-31	0	0
## 606	557	1932-09-06	1991-06-22	2009-12-31	0	0
## 607	558	1946-02-04	1991-06-22	2009-12-31	0	0
## 608	559	1940-03-29	1991-06-23	2009-12-31	0	0
## 609	560	1944-07-10	1991-06-24	2009-12-31	0	0
## 610	561	1946-02-16	1991-06-24	2009-12-31	0	0
## 611	562	1947-04-10	1991-06-24	2009-12-31	0	0
## 612	563	1933-05-28	1991-06-24	2009-12-31	0	0
## 613	564	1931-11-28	1991-06-25	1994-04-07	1	1
## 614	565	1944-02-15	1991-06-25	2009-12-31	0	0
## 615	566	1939-06-27	1991-06-25	2009-12-31	0	0
## 616	567	1949-09-09	1991-06-25	2009-12-31	0	0
## 617	568	1942-03-28	1991-06-26	2009-12-31	0	0
## 618	569	1947-04-28	1991-06-26	2009-12-31	0	0
## 619	570	1936-03-08	1991-06-26	1996-09-06	1	1
## 620	571	1937-07-13	1991-06-26	2009-12-31	0	0
## 621	572	1933-06-24	1991-06-26	1999-02-24	1	1
## 622	573	1933-11-04	1991-06-29	2008-05-08	1	0
## 623	574	1939-06-14	1991-06-29	2009-12-31	0	0
## 624	575	1934-08-18	1991-06-29	2009-12-31	0	0
## 625	576	1935-04-03	1991-06-29	2009-12-31	0	0
## 626	577	1939-02-27	1991-06-30	1997-01-04	1	1
## 627	578	1932-06-19	1991-06-30	2009-12-31	0	0
## 628	579	1944-03-13	1991-06-30	2009-12-31	0	0
## 629	580	1940-09-22	1991-06-30	2009-12-31	0	0
## 630	581	1949-06-27	1991-06-30	2009-12-31	0	0
## 631	582	1940-12-05	1991-07-01	2009-12-31	0	0
## 632	583	1931-10-08	1991-07-01	2009-12-31	0	0
## 633	584	1934-09-20	1991-07-01	2001-10-07	1	0
## 634	585	1946-08-13	1991-07-01	2009-12-31	0	0
## 636	586	1936-05-21	1991-07-01	2009-12-31	0	0
## 637	587	1937-03-23	1991-07-01	2009-12-31	0	0

## 638	588	1937-10-09	1991-07-02	2009-12-31	0	0
## 639	589	1934-01-20	1991-07-02	2009-12-31	0	0
## 642	590	1932-05-25	1991-07-03	2009-12-31	0	0
## 645	591	1939-09-17	1991-07-03	2009-12-31	0	0
## 646	592	1932-11-25	1991-07-03	2009-12-31	0	0
## 647	593	1948-03-21	1991-07-03	2009-12-31	0	0
## 648	594	1941-12-20	1991-08-05	2009-12-31	0	0
## 649	595	1937-10-18	1991-08-05	2009-12-31	0	0
## 650	596	1942-05-31	1991-08-05	2009-12-31	0	0
## 651	597	1933-04-07	1991-08-05	2005-05-22	1	0
## 652	598	1942-07-13	1991-08-05	2009-12-31	0	0
## 654	599	1932-03-14	1991-08-06	2003-07-17	1	1
## 655	600	1933-06-13	1991-08-06	2009-12-31	0	0
## 656	601	1948-04-08	1991-08-06	2009-12-31	0	0
## 657	602	1931-11-26	1991-08-06	1999-10-19	1	0
## 658	603	1949-04-15	1991-08-06	2009-12-31	0	0
## 659	604	1936-01-04	1991-08-06	2009-12-31	0	0
## 660	605	1946-12-26	1991-08-06	2009-12-31	0	0
## 661	606	1946-10-09	1991-08-06	2009-12-31	0	0
## 662	607	1935-05-19	1991-08-07	2009-12-31	0	0
## 664	608	1942-07-24	1991-08-07	2009-12-31	0	0
## 665	609	1937-10-29	1991-08-07	2009-12-31	0	0
## 666	610	1936-10-04	1991-08-07	2000-09-21	1	1
## 667	611	1933-06-08	1991-08-07	2009-12-31	0	0
## 668	612	1947-07-07	1991-08-08	2000-06-18	1	0
## 670	613	1945-08-27	1991-08-08	2009-12-31	0	0
## 671	614	1944-04-03	1991-08-11	2009-12-31	0	0
## 672	615	1935-08-25	1991-08-11	2007-02-08	1	0
## 673	616	1942-11-18	1991-08-12	2000-10-14	1	0
## 674	617	1944-07-28	1991-08-12	2009-12-31	0	0
## 675	618	1939-02-02	1991-08-12	2009-12-31	0	0
## 676	619	1936-05-06	1991-08-12	2003-04-29	1	1
## 677	620	1932-01-11	1991-08-13	2009-12-31	0	0
## 678	621	1930-07-11	1991-08-13	1995-10-25	1	0
## 679	622	1945-07-27	1991-08-13	2009-12-31	0	0
## 680	623	1941-04-15	1991-08-14	2009-12-31	0	0
## 681	624	1947-01-02	1991-08-14	2009-12-31	0	0
## 682	625	1946-10-09	1991-08-14	2009-12-31	0	0
## 683	626	1948-06-04	1991-08-14	2009-12-31	0	0
## 684	627	1951-06-03	1991-08-14	2009-12-31	0	0
## 686	628	1940-03-27	1991-08-14	2009-12-31	0	0
## 687	629	1940-03-19	1991-08-17	2003-12-06	1	1
## 688	630	1935-12-26	1991-08-17	2009-12-31	0	0
## 689	631	1945-03-17	1991-08-17	2009-12-31	0	0
## 690	632	1936-11-12	1991-08-17	2009-12-31	0	0
## 692	633	1945-11-28	1991-08-17	2009-12-31	0	0
## 693	634	1940-11-10	1991-08-18	2009-12-31	0	0
## 694	635	1934-07-05	1991-08-18	2007-02-28	1	1
## 695	636	1939-08-15	1991-08-18	2005-10-18	1	0
## 696	637	1937-01-26	1991-08-19	2009-12-31	0	0
## 697	638	1938-07-17	1991-08-19	2009-12-31	0	0
## 698	639	1948-03-11	1991-08-19	2009-12-31	0	0
## 699	640	1934-09-14	1991-08-19	2002-08-31	1	0
## 700	641	1934-10-30	1991-08-19	2009-12-31	0	0
## 701	642	1931-06-02	1991-08-19	1996-03-08	1	0
## 702	643	1938-01-05	1991-08-20	2009-12-31	0	0
## 703	644	1942-03-31	1991-08-20	2009-12-31	0	0

##	704	645	1936-10-19	1991-08-20	2009-12-31	0	0
##	706	646	1934-10-12	1991-08-20	2009-03-31	1	0
##	707	647	1942-12-07	1991-08-20	2009-12-31	0	0
##	708	648	1938-02-03	1991-08-20	2009-02-16	1	1
##	709	649	1948-10-31	1991-08-21	2009-12-31	0	0
##	711	650	1935-10-24	1991-08-21	2009-12-31	0	0
##	712	651	1947-08-18	1991-08-21	2009-12-31	0	0
##	713	652	1938-04-01	1991-08-24	2009-12-30	1	0
##	714	653	1939-10-26	1991-08-24	2009-12-31	0	0
##	715	654	1933-05-13	1991-08-24	1998-10-20	1	1
##	717	655	1936-10-21	1991-08-24	2007-09-13	1	0
##	719	656	1933-09-21	1991-08-25	1994-12-27	1	0
##	721	657	1933-05-21	1991-08-25	2008-12-08	1	1
##	722	658	1941-03-17	1991-08-25	2009-12-31	0	0
##	723	659	1931-03-03	1991-08-26	2001-12-22	1	0
##	724	660	1944-02-14	1991-08-26	1997-03-03	1	0
##	725	661	1940-01-31	1991-08-26	2009-12-31	0	0
##	726	662	1931-02-15	1991-08-26	1998-10-22	1	0
##	727	663	1947-08-31	1991-08-27	2009-12-31	0	0
##	728	664	1936-07-05	1991-08-27	1996-09-07	1	0
##	729	665	1951-06-15	1991-08-27	2009-12-31	0	0
##	731	666	1947-04-19	1991-08-27	2009-12-31	0	0
##	732	667	1931-09-08	1991-08-27	2009-12-31	0	0
##	733	668	1933-08-24	1991-08-27	2009-12-31	0	0
##	734	669	1932-06-16	1991-08-27	2009-12-16	1	0
##	735	670	1947-07-10	1991-08-28	2009-12-31	0	0
##	736	671	1940-10-11	1991-08-28	2009-12-31	0	0
##	737	672	1938-01-26	1991-08-28	2009-12-31	0	0
##	738	673	1944-04-07	1991-08-31	1995-04-08	1	0
##	739	674	1937-10-10	1991-08-31	2009-12-31	0	0
##	741	675	1939-03-10	1991-09-01	2009-12-31	0	0
##	742	676	1950-01-07	1991-09-02	2009-12-31	0	0
##	743	677	1939-02-15	1991-09-02	2009-12-31	0	0
##	744	678	1935-09-02	1991-09-03	2005-06-12	1	1
##	745	679	1942-07-05	1991-09-03	2009-12-31	0	0
##	747	680	1938-10-31	1991-09-03	2009-12-31	0	0
##	748	681	1937-09-05	1991-09-03	2009-12-31	0	0
##	749	682	1937-08-21	1991-09-03	2009-12-31	0	0
##	750	683	1935-07-21	1991-09-03	2001-05-14	1	0
##	751	684	1935-06-24	1991-09-03	2009-12-31	0	0
##	752	685	1937-02-05	1991-09-03	2009-12-31	0	0
##	753	686	1949-08-21	1991-09-04	2009-12-31	0	0
##	754	687	1942-06-27	1991-09-04	2009-12-31	0	0
##	755	688	1941-01-03	1991-09-04	2009-12-31	0	0
##	756	689	1939-04-20	1991-09-04	2009-12-31	0	0
##	757	690	1949-03-15	1991-09-04	2009-12-31	0	0
##	758	691	1930-08-23	1991-09-04	2009-12-31	0	0
##	759	692	1948-11-02	1991-09-04	2009-12-31	0	0
##	760	693	1932-07-14	1991-09-04	2009-12-31	0	0
##	761	694	1951-07-26	1991-09-04	2009-12-31	0	0
##	764	695	1931-12-19	1991-09-07	2009-12-31	0	0
##	765	696	1935-05-04	1991-09-07	2009-12-31	0	0
##	766	697	1938-01-19	1991-09-07	2009-12-31	0	0
##	767	698	1936-03-13	1991-09-07	2009-12-31	0	0
##	768	699	1931-12-01	1991-09-07	2007-07-01	1	1
##	769	700	1939-02-26	1991-09-07	2009-12-31	0	0
##	770	701	1939-01-30	1991-09-08	2009-12-31	0	0

## 771	702	1945-08-05	1991-09-08	2009-12-31	0	0
## 772	703	1948-02-15	1991-09-08	2009-12-31	0	0
## 773	704	1944-09-29	1991-09-08	2009-12-31	0	0
## 774	705	1948-11-16	1991-09-08	2009-12-31	0	0
## 776	706	1936-09-11	1991-09-09	2008-01-16	1	1
## 777	707	1937-04-19	1991-09-09	2009-12-31	0	0
## 778	708	1931-02-02	1991-09-09	2009-05-31	1	0
## 779	709	1940-12-28	1991-09-09	2009-12-31	0	0
## 780	710	1936-04-16	1991-09-09	2009-12-31	0	0
## 781	711	1941-07-21	1991-09-09	2009-12-31	0	0
## 783	712	1934-10-02	1991-09-09	2009-12-31	0	0
## 787	713	1941-09-14	1991-09-10	2009-12-31	0	0
## 788	714	1944-02-19	1991-09-10	2009-12-31	0	0
## 789	715	1937-06-07	1991-09-10	2009-12-31	0	0
## 790	716	1931-09-02	1991-09-10	2002-02-22	1	1
## 791	717	1946-06-03	1991-09-10	2009-12-31	0	0
## 793	718	1946-03-31	1991-09-11	2009-12-31	0	0
## 794	719	1936-06-15	1991-09-11	2008-01-11	1	0
## 795	720	1933-08-14	1991-09-11	2009-12-31	0	0
## 796	721	1937-11-22	1991-09-11	2009-12-31	0	0
## 797	722	1944-01-30	1991-09-11	2009-12-31	0	0
## 798	723	1931-10-01	1991-09-11	2004-02-27	1	0
## 799	724	1934-11-08	1991-09-11	2002-03-01	1	1
## 800	725	1946-01-05	1991-09-11	2009-12-31	0	0
## 801	726	1941-06-16	1991-09-11	2009-09-11	1	0
## 802	727	1944-09-04	1991-09-11	2009-12-31	0	0
## 803	728	1937-05-25	1991-09-21	2009-12-31	0	0
## 804	729	1944-11-05	1991-09-21	2009-12-31	0	0
## 805	730	1937-03-24	1991-09-21	2009-12-31	0	0
## 806	731	1944-07-16	1991-09-21	2007-04-21	1	0
## 808	732	1949-09-08	1991-09-21	2009-12-31	0	0
## 809	733	1939-08-14	1991-09-22	2007-01-21	1	1
## 810	734	1943-05-10	1991-09-22	2009-12-31	0	0
## 811	735	1949-07-09	1991-09-22	2009-12-31	0	0
## 812	736	1944-06-03	1991-09-22	2006-05-27	1	0
## 813	737	1944-08-21	1991-09-23	2009-12-31	0	0
## 814	738	1947-05-30	1991-09-23	2009-12-31	0	0
## 815	739	1951-04-03	1991-09-23	2009-12-31	0	0
## 816	740	1931-01-29	1991-09-24	2009-12-31	0	0
## 817	741	1940-09-11	1991-09-24	2009-12-31	0	0
## 818	742	1943-08-15	1991-09-24	2009-12-31	0	0
## 819	743	1942-03-03	1991-09-24	2009-12-31	0	0
## 820	744	1935-10-13	1991-09-24	2009-12-31	0	0
## 821	745	1939-04-13	1991-09-25	2009-12-31	0	0
## 822	746	1943-04-12	1991-09-25	2008-12-25	1	0
## 823	747	1933-05-27	1991-09-25	2009-12-31	0	0
## 824	748	1939-06-30	1991-09-25	2009-12-31	0	0
## 826	749	1936-06-24	1991-09-25	2009-12-31	0	0
## 827	750	1948-12-02	1991-09-25	2009-12-31	0	0
## 828	751	1946-05-28	1991-09-25	2009-12-31	0	0
## 829	752	1937-06-19	1991-09-28	1999-12-07	1	0
## 831	753	1948-04-13	1991-09-28	2009-12-31	0	0
## 832	754	1950-03-24	1991-09-28	2009-12-31	0	0
## 833	755	1933-12-29	1991-09-28	2009-12-31	0	0
## 834	756	1948-05-15	1991-09-28	2009-12-31	0	0
## 835	757	1941-07-01	1991-09-28	2009-12-31	0	0
## 836	758	1931-09-08	1991-09-29	2009-12-31	0	0

## 837	759	1937-06-29	1991-09-29	2009-12-31	0	0
## 838	760	1942-11-12	1991-09-29	2009-12-31	0	0
## 839	761	1939-05-27	1991-09-29	2009-12-31	0	0
## 840	762	1933-06-07	1991-09-29	2006-08-28	1	0
## 841	763	1939-01-21	1991-09-29	2007-11-13	1	1
## 842	764	1941-01-21	1991-09-29	2009-12-31	0	0
## 843	765	1937-09-01	1991-09-30	2009-12-31	0	0
## 844	766	1949-07-05	1991-09-30	2009-12-31	0	0
## 845	767	1939-01-27	1991-09-30	2009-12-31	0	0
## 846	768	1940-10-20	1991-09-30	2009-12-31	0	0
## 848	769	1933-11-18	1991-09-30	2009-12-31	0	0
## 849	770	1940-06-02	1991-09-30	2009-12-31	0	0
## 850	771	1938-06-22	1991-09-30	2009-12-31	0	0
## 851	772	1948-07-18	1991-10-01	2009-12-31	0	0
## 852	773	1938-08-17	1991-10-01	2009-12-31	0	0
## 853	774	1940-01-20	1991-10-01	2009-12-31	0	0
## 854	775	1935-12-03	1991-10-01	2009-12-31	0	0
## 855	776	1940-08-02	1991-10-01	2009-12-31	0	0
## 856	777	1942-08-10	1991-10-01	2009-12-31	0	0
## 857	778	1942-07-29	1991-10-01	2009-12-31	0	0
## 858	779	1945-04-21	1991-10-02	2009-12-31	0	0
## 859	780	1942-11-28	1991-10-02	2009-12-31	0	0
## 860	781	1950-08-08	1991-10-05	2009-12-31	0	0
## 861	782	1938-11-28	1991-10-05	1999-05-31	1	0
## 862	783	1944-02-05	1991-10-05	2009-12-31	0	0
## 863	784	1943-09-06	1991-10-05	2009-12-31	0	0
## 865	785	1943-09-20	1991-10-05	2009-12-31	0	0
## 866	786	1950-07-14	1991-10-05	2009-12-31	0	0
## 867	787	1937-12-20	1991-10-06	2009-12-31	0	0
## 868	788	1935-09-10	1991-10-07	2009-12-31	0	0
## 869	789	1943-02-27	1991-10-07	2009-12-31	0	0
## 870	790	1930-12-23	1991-10-07	2009-12-31	0	0
## 871	791	1935-08-21	1991-10-07	2009-12-31	0	0
## 872	792	1943-06-06	1991-10-07	2009-12-31	0	0
## 873	793	1938-11-22	1991-10-07	2009-12-31	0	0
## 874	794	1938-08-27	1991-10-08	2009-12-31	0	0
## 875	795	1949-01-19	1991-10-08	2009-12-31	0	0
## 876	796	1939-04-23	1991-10-08	2009-12-31	0	0
## 877	797	1934-01-07	1991-10-09	2009-11-06	1	0
## 878	798	1933-12-23	1991-10-09	2009-12-31	0	0
## 879	799	1933-05-07	1991-10-09	2009-12-31	0	0
## 880	800	1932-06-03	1991-10-09	2009-12-31	0	0
## 881	801	1933-07-18	1991-10-09	2009-12-31	0	0
## 882	802	1936-04-30	1991-10-09	2003-11-10	1	0
## 883	803	1932-07-22	1991-10-09	2003-12-14	1	1
## 884	804	1950-02-17	1991-10-12	2009-12-31	0	0
## 885	805	1940-07-11	1991-10-12	2002-01-11	1	0
## 886	806	1942-12-28	1991-10-12	2008-01-26	1	0
## 887	807	1934-07-03	1991-10-12	2009-12-31	0	0
## 888	808	1934-01-23	1991-10-15	1994-01-10	1	0
## 889	809	1931-07-11	1991-10-15	2009-12-31	0	0
## 890	810	1937-01-19	1991-10-15	2008-02-09	1	0
## 891	811	1932-08-31	1991-10-15	2002-03-05	1	1
## 892	812	1946-12-20	1991-10-16	2009-12-31	0	0
## 893	813	1950-11-29	1991-10-16	2009-12-31	0	0
## 895	814	1936-12-30	1991-10-16	2009-12-31	0	0
## 896	815	1947-01-11	1991-10-16	2009-12-31	0	0

## 897	816	1943-05-14	1991-10-16	2009-12-31	0	0
## 900	817	1946-09-12	1991-10-19	2009-12-31	0	0
## 901	818	1936-01-21	1991-10-19	1999-01-22	1	0
## 902	819	1940-05-27	1991-10-19	2009-12-31	0	0
## 903	820	1942-10-13	1991-10-19	2009-12-31	0	0
## 904	821	1941-04-30	1991-10-19	2009-12-31	0	0
## 905	822	1944-11-23	1991-10-19	2009-12-31	0	0
## 906	823	1935-07-08	1991-10-20	2009-12-31	0	0
## 907	824	1936-09-22	1991-10-20	2009-12-31	0	0
## 908	825	1943-01-22	1991-10-20	2009-12-31	0	0
## 909	826	1932-10-02	1991-10-20	2009-12-31	0	0
## 910	827	1945-02-09	1991-10-21	2009-12-31	0	0
## 912	828	1939-04-04	1991-10-21	2001-08-19	1	1
## 913	829	1950-02-04	1991-10-21	2009-12-31	0	0
## 914	830	1948-01-27	1991-10-21	1993-01-06	1	0
## 915	831	1940-07-23	1991-10-21	2002-10-21	1	0
## 916	832	1938-02-25	1991-10-21	2006-06-10	1	1
## 917	833	1938-11-29	1991-10-21	2009-12-31	0	0
## 918	834	1949-04-08	1991-10-22	2009-12-31	0	0
## 919	835	1940-08-16	1991-10-22	1995-08-31	1	1
## 920	836	1932-09-10	1991-10-22	2009-12-31	0	0
## 921	837	1932-08-28	1991-10-22	2009-12-31	0	0
## 923	838	1934-06-07	1991-10-22	2009-12-31	0	0
## 924	839	1943-01-12	1991-10-22	2009-12-31	0	0
## 926	840	1945-01-06	1991-10-23	2009-12-31	0	0
## 927	841	1933-05-08	1991-10-23	2003-05-10	1	1
## 928	842	1937-04-02	1991-10-23	2009-12-31	0	0
## 929	843	1934-01-04	1991-10-23	2009-12-31	0	0
## 930	844	1930-08-11	1991-10-23	2009-12-31	0	0
## 932	845	1950-11-19	1991-10-26	2003-05-28	1	1
## 933	846	1935-06-08	1991-10-26	2009-12-31	0	0
## 934	847	1944-07-03	1991-10-26	2009-12-31	0	0
## 935	848	1941-10-13	1991-10-26	2009-12-31	0	0
## 936	849	1949-02-15	1991-10-26	2009-12-31	0	0
## 938	850	1942-09-14	1991-10-27	2009-12-31	0	0
## 939	851	1937-09-29	1991-10-27	1997-01-03	1	1
## 941	852	1939-06-15	1991-10-27	2009-12-31	0	0
## 942	853	1942-05-11	1991-10-27	2009-12-31	0	0
## 944	854	1946-03-18	1991-10-27	2009-12-31	0	0
## 945	855	1933-02-28	1991-10-28	2009-12-31	0	0
## 946	856	1937-08-11	1991-10-28	2009-12-31	0	0
## 947	857	1949-04-16	1991-10-28	2009-12-31	0	0
## 948	858	1936-10-06	1991-10-28	2009-12-31	0	0
## 949	859	1942-05-07	1991-10-28	2009-12-31	0	0
## 950	860	1950-12-04	1991-10-28	2009-12-31	0	0
## 951	861	1946-06-20	1991-10-29	2009-12-31	0	0
## 952	862	1933-12-18	1991-10-29	2009-12-31	0	0
## 953	863	1931-10-25	1991-10-29	2005-01-30	1	0
## 955	864	1944-07-04	1991-10-29	2009-12-31	0	0
## 958	865	1931-02-11	1991-10-29	2009-12-31	0	0
## 959	866	1940-02-24	1991-10-30	2009-12-31	0	0
## 960	867	1932-06-17	1991-10-30	2009-12-31	0	0
## 961	868	1938-09-15	1991-11-02	2009-12-31	0	0
## 962	869	1940-06-13	1991-11-02	2009-12-31	0	0
## 964	870	1950-11-20	1991-11-03	2003-05-23	1	1
## 965	871	1936-03-26	1991-11-03	2000-08-26	1	1
## 966	872	1940-07-27	1991-11-03	1998-03-10	1	0

## 967	873	1947-06-17	1991-11-03	2009-12-31	0	0
## 968	874	1948-08-31	1991-11-03	2006-04-02	1	1
## 969	875	1937-07-15	1991-11-03	2009-12-31	0	0
## 971	876	1942-07-30	1991-11-04	2009-12-31	0	0
## 972	877	1932-02-05	1991-11-04	2009-12-31	0	0
## 973	878	1937-09-18	1991-11-04	2009-12-31	0	0
## 974	879	1936-11-05	1991-11-04	2001-05-15	1	0
## 975	880	1949-09-13	1991-11-04	2009-12-31	0	0
## 976	881	1935-10-06	1991-11-04	2009-12-31	0	0
## 977	882	1938-09-25	1991-11-05	2009-12-31	0	0
## 978	883	1947-07-07	1991-11-05	2009-12-31	0	0
## 979	884	1933-12-11	1991-11-05	2009-12-31	0	0
## 980	885	1938-02-18	1991-11-05	2009-12-31	0	0
## 981	886	1947-01-26	1991-11-09	2009-12-31	0	0
## 983	887	1937-11-22	1991-11-09	2009-12-31	0	0
## 984	888	1933-04-08	1991-11-09	2007-09-16	1	0
## 985	889	1950-11-24	1991-11-09	2009-12-31	0	0
## 986	890	1931-10-29	1991-11-09	2009-12-31	0	0
## 987	891	1932-12-17	1991-11-09	2008-07-10	1	0
## 988	892	1934-08-30	1991-11-09	2009-12-31	0	0
## 989	893	1948-04-11	1991-11-09	2009-12-31	0	0
## 990	894	1932-12-15	1991-11-10	2009-12-31	0	0
## 991	895	1934-06-15	1991-11-10	2009-12-31	0	0
## 993	896	1945-03-17	1991-11-11	2009-12-31	0	0
## 994	897	1941-03-13	1991-11-11	2004-10-23	1	0
## 995	898	1949-03-21	1991-11-11	2009-12-31	0	0
## 996	899	1933-03-02	1991-11-11	1996-12-10	1	0
## 997	900	1932-04-01	1991-11-11	2009-12-31	0	0
## 998	901	1939-12-30	1991-11-11	2009-12-31	0	0
## 1000	902	1936-01-24	1991-11-12	2002-02-14	1	1
## 1001	903	1938-10-11	1991-12-07	2009-12-31	0	0
## 1002	904	1932-02-02	1991-12-07	2006-04-25	1	1
## 1003	905	1939-12-01	1991-12-07	2009-12-31	0	0
## 1004	906	1950-11-22	1991-12-07	2009-12-31	0	0
## 1005	907	1935-03-13	1991-12-07	2009-12-31	0	0
## 1006	908	1930-12-09	1991-12-08	2009-12-31	0	0
## 1007	909	1933-07-25	1991-12-08	2009-12-31	0	0
## 1008	910	1945-02-17	1991-12-08	2009-12-31	0	0
## 1009	911	1940-01-04	1991-12-08	2007-09-09	1	0
## 1010	912	1943-03-13	1991-12-09	2009-12-31	0	0
## 1011	913	1937-03-30	1991-12-09	2009-08-22	1	0
## 1012	914	1935-04-09	1991-12-09	2009-12-31	0	0
## 1013	915	1941-06-18	1991-12-09	2009-05-11	1	0
## 1014	916	1942-02-28	1991-12-10	2009-12-31	0	0
## 1015	917	1942-06-23	1991-12-10	2009-12-31	0	0
## 1016	918	1938-10-26	1991-12-10	2000-10-04	1	0
## 1017	919	1943-03-14	1991-12-11	2009-12-31	0	0
## 1018	920	1941-06-26	1991-12-11	2009-12-31	0	0
## 1019	921	1945-12-18	1991-12-11	2009-12-31	0	0
## 1021	922	1936-04-09	1991-12-14	2009-12-31	0	0
## 1022	923	1938-12-02	1991-12-14	2008-05-08	1	0
## 1024	924	1936-09-10	1991-12-15	2009-12-31	0	0
## 1025	925	1944-12-21	1991-12-15	2009-12-31	0	0
## 1026	926	1938-07-16	1991-12-15	2009-12-31	0	0
## 1027	927	1942-01-16	1991-12-15	2009-12-31	0	0
## 1028	928	1947-08-06	1991-12-15	2009-12-31	0	0
## 1030	929	1932-11-07	1991-12-16	2000-01-23	1	0

##	1032	930	1942-10-19	1991-12-16	2002-11-16	1	0
##	1033	931	1948-08-07	1991-12-16	2009-12-31	0	0
##	1034	932	1945-08-16	1991-12-17	2009-12-31	0	0
##	1035	933	1941-04-17	1991-12-17	2009-12-31	0	0
##	1036	934	1941-07-18	1991-12-17	2009-12-31	0	0
##	1037	935	1951-11-16	1991-12-17	2009-12-31	0	0
##	1039	936	1934-02-05	1991-12-17	2009-12-31	0	0
##	1040	937	1939-12-25	1991-12-17	2009-12-31	0	0
##	1041	938	1946-11-30	1991-12-17	2009-12-31	0	0
##	1043	939	1944-08-14	1991-12-18	2009-12-31	0	0
##	1044	940	1942-01-27	1991-12-18	2009-12-31	0	0
##	1045	941	1934-06-26	1991-12-18	2003-04-30	1	1
##	1046	942	1943-03-12	1991-12-18	2009-12-31	0	0
##	1047	943	1947-03-29	1991-12-18	2009-12-31	0	0
##	1048	944	1947-10-10	1991-12-18	2009-12-31	0	0
##	1049	945	1937-05-15	1991-12-18	2009-12-31	0	0
##	1050	946	1940-02-17	1991-12-18	1998-02-14	1	0
##	1051	947	1941-12-21	1991-12-18	2009-12-31	0	0
##	1052	948	1936-04-02	1991-12-18	2009-12-31	0	0
##	1053	949	1943-01-29	1991-12-18	1995-05-24	1	0
##	1054	950	1941-03-23	1991-12-21	2009-12-31	0	0
##	1057	951	1942-04-01	1991-12-21	2009-12-31	0	0
##	1058	952	1942-02-22	1991-12-21	2009-12-31	0	0
##	1059	953	1945-11-08	1991-12-21	2009-12-31	0	0
##	1060	954	1941-05-05	1991-12-21	2009-12-31	0	0
##	1061	955	1943-03-12	1991-12-22	2009-12-31	0	0
##	1062	956	1947-11-08	1991-12-22	2009-12-31	0	0
##	1063	957	1931-05-13	1991-12-22	2005-02-13	1	0
##	1064	958	1931-12-07	1991-12-22	2009-12-31	0	0
##	1065	959	1933-09-06	1991-12-22	2001-08-24	1	0
##	1066	960	1946-12-29	1991-12-22	2009-12-31	0	0
##	1067	961	1935-09-30	1991-12-22	2000-04-19	1	1
##	1068	962	1944-08-03	1991-12-22	2009-12-31	0	0
##	1069	963	1936-03-17	1991-12-23	2009-12-31	0	0
##	1070	964	1946-01-23	1991-12-23	2009-12-31	0	0
##	1071	965	1933-04-04	1991-12-23	2009-12-31	0	0
##	1072	966	1931-04-02	1991-12-23	2009-12-31	0	0
##	1073	967	1947-06-03	1991-12-23	2009-12-31	0	0
##	1074	968	1941-05-07	1991-12-23	2009-12-31	0	0
##	1075	969	1944-08-14	1991-12-24	2009-12-31	0	0
##	1076	970	1932-01-26	1991-12-24	2009-12-31	0	0
##	1077	971	1948-02-13	1991-12-24	2009-12-31	0	0
##	1079	972	1938-06-30	1991-12-31	2009-12-31	0	0
##	1080	973	1939-01-07	1991-12-31	2009-12-31	0	0
##	1081	974	1943-12-05	1991-12-31	2009-12-31	0	0
##	1082	975	1947-01-29	1991-12-31	2009-12-31	0	0
##	1083	976	1934-10-12	1992-01-11	2009-12-31	0	0
##	1086	977	1945-08-13	1992-01-11	2008-05-01	1	0
##	1088	978	1935-02-03	1992-01-12	1995-05-13	1	0
##	1089	979	1937-03-13	1992-01-12	2009-12-31	0	0
##	1090	980	1941-08-29	1992-01-12	2009-12-31	0	0
##	1091	981	1940-03-15	1992-01-12	2009-12-31	0	0
##	1092	982	1944-12-13	1992-01-12	2007-12-17	1	0
##	1094	983	1944-01-05	1992-01-13	2009-12-31	0	0
##	1095	984	1943-07-03	1992-01-13	2009-12-31	0	0
##	1096	985	1945-11-24	1992-01-13	2009-12-31	0	0
##	1097	986	1944-01-25	1992-01-13	2004-07-23	1	0

##	1098	987	1935-05-17	1992-01-13	2009-12-31	0	0
##	1099	988	1931-08-24	1992-01-13	2003-09-14	1	0
##	1100	989	1950-01-19	1992-01-14	2009-12-31	0	0
##	1101	990	1940-03-06	1992-01-14	2006-09-16	1	0
##	1102	991	1946-11-22	1992-01-14	2009-12-31	0	0
##	1103	992	1951-08-24	1992-01-14	2009-12-31	0	0
##	1104	993	1934-12-24	1992-01-14	2009-12-31	0	0
##	1105	994	1946-05-21	1992-01-14	2009-12-31	0	0
##	1106	995	1949-12-07	1992-01-14	1994-04-29	1	0
##	1107	996	1934-09-06	1992-01-15	2009-12-31	0	0
##	1109	997	1937-01-25	1992-01-15	2009-12-31	0	0
##	1110	998	1944-10-12	1992-01-15	2009-12-31	0	0
##	1111	999	1935-07-19	1992-01-15	2009-12-15	1	0
##	1112	1000	1946-04-30	1992-01-15	2009-12-31	0	0
##	1113	1001	1950-08-20	1992-01-15	2009-12-31	0	0
##	1114	1002	1946-07-26	1992-01-15	2009-12-31	0	0
##	1115	1003	1944-09-12	1992-01-18	2009-12-31	0	0
##	1116	1004	1943-08-04	1992-01-18	2009-12-31	0	0
##	1117	1005	1943-01-07	1992-01-18	2009-12-31	0	0
##	1118	1006	1950-08-30	1992-01-19	2009-12-31	0	0
##	1119	1007	1950-09-14	1992-01-19	2009-12-31	0	0
##	1120	1008	1942-02-12	1992-01-19	2009-12-31	0	0
##	1121	1009	1942-03-22	1992-01-19	2009-12-31	0	0
##	1122	1010	1943-08-14	1992-01-19	2009-12-31	0	0
##	1123	1011	1948-10-05	1992-01-19	2009-12-31	0	0
##	1124	1012	1939-11-14	1992-01-19	2009-12-31	0	0
##	1125	1013	1946-05-29	1992-01-19	2009-12-31	0	0
##	1126	1014	1944-01-20	1992-01-19	2009-12-31	0	0
##	1129	1015	1951-01-19	1992-01-20	2009-12-31	0	0
##	1130	1016	1946-09-15	1992-01-20	2009-12-31	0	0
##	1131	1017	1951-04-26	1992-01-20	2009-12-31	0	0
##	1132	1018	1942-03-07	1992-01-20	2009-12-31	0	0
##	1133	1019	1949-01-03	1992-01-20	2009-12-31	0	0
##	1134	1020	1948-08-07	1992-01-21	2009-12-31	0	0
##	1135	1021	1943-03-14	1992-01-21	2009-12-31	0	0
##	1136	1022	1933-12-23	1992-01-21	2009-12-31	0	0
##	1137	1023	1937-05-27	1992-01-21	1999-03-16	1	0
##	1138	1024	1947-07-14	1992-01-21	2009-12-31	0	0
##	1139	1025	1942-10-30	1992-01-21	2009-12-31	0	0
##	1140	1026	1941-10-18	1992-01-21	2009-12-31	0	0
##	1141	1027	1948-02-25	1992-01-21	2009-12-31	0	0
##	1142	1028	1935-03-07	1992-01-21	2009-12-31	0	0
##	1143	1029	1939-05-22	1992-01-22	2009-12-31	0	0
##	1144	1030	1944-11-20	1992-01-22	2009-12-31	0	0
##	1145	1031	1948-03-10	1992-01-22	2003-07-31	1	0
##	1146	1032	1942-05-01	1992-01-22	2009-12-31	0	0
##	1147	1033	1948-08-18	1992-01-22	2009-12-31	0	0
##	1148	1034	1940-01-15	1992-01-22	2001-03-05	1	0
##	1149	1035	1936-03-22	1992-01-22	2009-12-31	0	0
##	1150	1036	1930-10-24	1992-01-25	2009-12-31	0	0
##	1151	1037	1950-07-16	1992-01-25	2009-12-31	0	0
##	1152	1038	1938-05-19	1992-01-25	2009-12-31	0	0
##	1153	1039	1938-07-08	1992-01-25	2009-12-31	0	0
##	1154	1040	1936-09-01	1992-01-26	2007-11-05	1	1
##	1155	1041	1943-10-14	1992-01-26	2009-12-31	0	0
##	1157	1042	1951-10-30	1992-01-26	2009-12-31	0	0
##	1158	1043	1939-09-04	1992-01-26	2009-12-31	0	0

##	1159	1044	1944-02-11	1992-01-26	2009-12-31	0	0
##	1160	1045	1941-09-21	1992-01-26	2009-12-31	0	0
##	1161	1046	1942-12-02	1992-01-26	2009-12-31	0	0
##	1162	1047	1945-08-22	1992-01-27	2009-12-31	0	0
##	1163	1048	1932-09-17	1992-01-27	2009-12-31	0	0
##	1164	1049	1943-09-29	1992-01-27	2009-12-31	0	0
##	1165	1050	1946-12-03	1992-01-27	2009-12-31	0	0
##	1166	1051	1936-02-23	1992-01-27	2009-12-31	0	0
##	1167	1052	1945-11-08	1992-01-27	2009-12-31	0	0
##	1168	1053	1933-10-07	1992-01-27	2009-07-26	1	0
##	1169	1054	1938-07-08	1992-01-27	1997-06-04	1	1
##	1170	1055	1931-04-26	1992-01-28	2009-12-31	0	0
##	1171	1056	1941-08-28	1992-01-28	2009-12-31	0	0
##	1172	1057	1937-01-19	1992-01-29	2009-12-31	0	0
##	1173	1058	1946-03-06	1992-01-29	2009-12-31	0	0
##	1174	1059	1931-10-24	1992-01-29	2009-12-31	0	0
##	1175	1060	1944-07-22	1992-01-29	2009-12-31	0	0
##	1176	1061	1947-11-11	1992-01-29	2009-12-31	0	0
##	1177	1062	1949-09-30	1992-01-29	2009-12-31	0	0
##	1178	1063	1945-07-10	1992-01-29	2009-12-31	0	0
##	1179	1064	1947-09-12	1992-01-29	2009-12-31	0	0
##	1180	1065	1947-03-28	1992-01-29	2009-12-31	0	0
##	1181	1066	1938-05-26	1992-02-01	2009-12-31	0	0
##	1182	1067	1935-02-24	1992-02-02	1994-02-25	1	1
##	1183	1068	1931-04-28	1992-02-02	2009-12-31	0	0
##	1184	1069	1937-06-20	1992-02-02	2009-12-31	0	0
##	1185	1070	1935-01-24	1992-02-08	2009-12-31	0	0
##	1186	1071	1938-08-22	1992-02-08	2009-12-31	0	0
##	1188	1072	1942-12-23	1992-02-08	2009-12-31	0	0
##	1189	1073	1931-05-24	1992-02-08	2005-08-28	1	0
##	1190	1074	1936-04-25	1992-02-08	2009-12-31	0	0
##	1191	1075	1943-01-14	1992-02-08	2009-12-31	0	0
##	1193	1076	1931-08-05	1992-02-09	2009-12-31	0	0
##	1194	1077	1951-09-10	1992-02-09	2009-12-31	0	0
##	1195	1078	1932-01-26	1992-02-09	2009-12-31	0	0
##	1196	1079	1949-06-24	1992-02-09	2009-12-31	0	0
##	1197	1080	1948-07-15	1992-02-10	2009-12-31	0	0
##	1198	1081	1939-08-13	1992-02-10	2009-12-31	0	0
##	1200	1082	1933-07-16	1992-02-10	2009-12-31	0	0
##	1201	1083	1934-11-19	1992-02-10	1992-07-24	1	0
##	1202	1084	1939-01-10	1992-02-10	2009-12-31	0	0
##	1203	1085	1945-09-05	1992-02-10	2009-12-31	0	0
##	1204	1086	1937-08-30	1992-02-10	2009-12-31	0	0
##	1205	1087	1933-01-27	1992-02-11	2009-12-31	0	0
##	1206	1088	1935-11-25	1992-02-11	2009-12-31	0	0
##	1208	1089	1939-08-15	1992-02-11	1994-09-27	1	0
##	1209	1090	1940-11-09	1992-02-12	2009-12-31	0	0
##	1210	1091	1931-09-10	1992-02-12	2009-12-31	0	0
##	1211	1092	1931-11-27	1992-02-12	2009-12-31	0	0
##	1212	1093	1945-10-30	1992-02-15	2009-12-31	0	0
##	1213	1094	1945-03-21	1992-02-15	2009-12-31	0	0
##	1214	1095	1945-08-04	1992-02-15	2009-12-31	0	0
##	1216	1096	1946-03-22	1992-02-16	2009-12-31	0	0
##	1218	1097	1939-07-24	1992-02-16	2009-12-31	0	0
##	1219	1098	1937-09-12	1992-02-16	2009-12-31	0	0
##	1220	1099	1951-04-08	1992-02-16	2009-12-31	0	0
##	1221	1100	1940-01-02	1992-02-16	2009-12-31	0	0

##	1222	1101	1936-05-12	1992-02-16	2009-12-31	0	0
##	1223	1102	1942-07-18	1992-02-16	2007-08-24	1	0
##	1225	1103	1934-11-29	1992-02-17	2009-12-31	0	0
##	1226	1104	1931-09-21	1992-02-17	1994-11-13	1	1
##	1227	1105	1942-07-16	1992-02-17	2009-12-31	0	0
##	1228	1106	1945-02-03	1992-02-17	2009-12-31	0	0
##	1230	1107	1936-07-21	1992-02-17	2009-12-31	0	0
##	1231	1108	1939-05-22	1992-02-17	2006-11-18	1	1
##	1233	1109	1951-07-04	1992-02-17	2009-12-31	0	0
##	1234	1110	1932-02-10	1992-02-18	2009-12-31	0	0
##	1235	1111	1951-11-14	1992-02-18	2009-12-31	0	0
##	1236	1112	1944-08-29	1992-02-19	2009-12-31	0	0
##	1237	1113	1942-09-04	1992-02-19	2009-12-31	0	0
##	1238	1114	1937-10-30	1992-02-19	2009-12-31	0	0
##	1240	1115	1937-07-13	1992-02-19	2009-12-31	0	0
##	1241	1116	1939-06-16	1992-02-19	2009-12-31	0	0
##	1242	1117	1940-01-18	1992-02-19	2009-12-31	0	0
##	1243	1118	1942-06-01	1992-02-19	2009-12-31	0	0
##	1244	1119	1933-10-15	1992-02-22	2009-12-31	0	0
##	1245	1120	1933-06-06	1992-02-22	1999-07-24	1	0
##	1246	1121	1940-03-15	1992-02-22	2009-12-31	0	0
##	1247	1122	1937-06-20	1992-02-22	2009-12-31	0	0
##	1248	1123	1934-10-21	1992-02-23	1993-07-06	1	1
##	1249	1124	1938-03-13	1992-02-23	1998-07-23	1	1
##	1250	1125	1947-10-03	1992-02-23	2009-12-31	0	0
##	1251	1126	1939-09-05	1992-02-23	2009-12-31	0	0
##	1252	1127	1941-07-09	1992-02-23	2009-12-31	0	0
##	1253	1128	1932-12-29	1992-02-24	2009-12-31	0	0
##	1254	1129	1934-04-18	1992-02-24	2009-12-31	0	0
##	1255	1130	1932-10-05	1992-02-24	1998-06-09	1	0
##	1256	1131	1941-10-23	1992-02-24	2009-12-31	0	0
##	1257	1132	1939-10-19	1992-02-24	1998-08-04	1	0
##	1258	1133	1943-05-27	1992-02-24	2009-12-31	0	0
##	1259	1134	1935-12-27	1992-02-25	2009-12-31	0	0
##	1261	1135	1937-01-27	1992-02-25	2009-12-31	0	0
##	1262	1136	1930-12-30	1992-02-25	1992-11-28	1	0
##	1263	1137	1936-05-11	1992-02-25	2009-12-31	0	0
##	1264	1138	1936-09-10	1992-02-25	2009-12-31	0	0
##	1265	1139	1939-02-25	1992-02-25	2009-12-31	0	0
##	1266	1140	1948-07-19	1992-02-25	2009-12-31	0	0
##	1267	1141	1932-08-24	1992-02-26	2009-12-31	0	0
##	1268	1142	1933-01-06	1992-02-26	2009-12-31	0	0
##	1269	1143	1944-08-29	1992-02-26	2009-12-31	0	0
##	1270	1144	1943-04-26	1992-02-26	2009-12-31	0	0
##	1271	1145	1942-11-20	1992-02-26	2009-12-31	0	0
##	1272	1146	1943-09-12	1992-02-26	2009-12-31	0	0
##	1274	1147	1943-02-12	1992-02-29	2009-12-31	0	0
##	1275	1148	1935-02-17	1992-02-29	2009-12-31	0	0
##	1276	1149	1949-07-25	1992-02-29	2009-12-31	0	0
##	1277	1150	1936-09-21	1992-02-29	2009-12-31	0	0
##	1278	1151	1935-09-01	1992-02-29	2006-06-30	1	0
##	1279	1152	1933-09-21	1992-02-29	2001-01-24	1	0
##	1280	1153	1939-01-28	1992-02-29	2009-12-31	0	0
##	1281	1154	1944-05-13	1992-03-01	2009-12-31	0	0
##	1282	1155	1945-02-11	1992-03-01	2009-12-31	0	0
##	1283	1156	1936-04-02	1992-03-01	2009-12-31	0	0
##	1285	1157	1940-10-22	1992-03-01	2009-12-31	0	0

##	1286	1158	1933-05-13	1992-03-01	1994-02-13	1	1
##	1287	1159	1935-11-06	1992-03-02	2009-12-31	0	0
##	1288	1160	1936-02-27	1992-03-02	2009-12-31	0	0
##	1289	1161	1941-05-27	1992-03-02	2009-12-31	0	0
##	1290	1162	1940-05-22	1992-03-02	2009-12-31	0	0
##	1291	1163	1948-02-11	1992-03-02	2009-12-31	0	0
##	1293	1164	1945-12-18	1992-03-10	2009-12-31	0	0
##	1294	1165	1939-09-09	1992-03-10	2000-11-01	1	0
##	1295	1166	1939-12-07	1992-03-10	2009-12-31	0	0
##	1296	1167	1946-07-15	1992-03-10	2009-12-31	0	0
##	1297	1168	1947-12-12	1992-03-11	2009-12-31	0	0
##	1298	1169	1944-06-28	1992-03-11	2009-12-31	0	0
##	1299	1170	1937-11-26	1992-03-15	2008-02-24	1	1
##	1301	1171	1934-12-28	1992-03-16	2009-12-31	0	0
##	1303	1172	1945-08-02	1992-03-16	2003-12-31	1	0
##	1304	1173	1932-11-07	1992-03-16	2000-03-09	1	0
##	1306	1174	1931-02-19	1992-03-17	2001-11-04	1	0
##	1307	1175	1937-05-22	1992-03-17	1999-06-17	1	1
##	1308	1176	1931-10-01	1992-03-17	2009-12-31	0	0
##	1309	1177	1934-01-20	1992-03-17	2008-11-16	1	1
##	1310	1178	1938-01-21	1992-03-17	2009-12-31	0	0
##	1311	1179	1931-08-29	1992-03-17	2009-12-31	0	0
##	1312	1180	1945-08-05	1992-03-17	2009-12-31	0	0
##	1313	1181	1936-10-10	1992-03-17	2001-10-13	1	0
##	1314	1182	1933-09-11	1992-03-18	2009-12-31	0	0
##	1315	1183	1931-10-05	1992-03-18	2009-12-31	0	0
##	1316	1184	1951-10-01	1992-03-18	2009-12-31	0	0
##	1317	1185	1936-10-15	1992-03-18	2009-12-31	0	0
##	1318	1186	1944-10-22	1992-03-18	2009-12-31	0	0
##	1320	1187	1934-01-30	1992-03-21	2004-12-22	1	0
##	1321	1188	1944-12-16	1992-03-21	2009-12-31	0	0
##	1322	1189	1935-05-10	1992-03-21	2009-12-31	0	0
##	1326	1190	1939-06-29	1992-03-21	2009-12-31	0	0
##	1327	1191	1931-02-27	1992-03-22	2005-12-12	1	1
##	1330	1192	1937-03-02	1992-03-22	2009-12-31	0	0
##	1331	1193	1931-07-14	1992-03-22	2009-12-31	0	0
##	1332	1194	1945-01-31	1992-03-22	2009-12-31	0	0
##	1333	1195	1947-11-09	1992-03-22	2009-12-31	0	0
##	1334	1196	1944-01-10	1992-03-23	2009-12-31	0	0
##	1337	1197	1931-03-27	1992-03-23	2005-04-14	1	1
##	1339	1198	1934-07-03	1992-03-24	1998-09-24	1	0
##	1340	1199	1951-06-09	1992-03-24	2009-12-31	0	0
##	1341	1200	1945-01-16	1992-03-24	2009-12-31	0	0
##	1342	1201	1945-04-16	1992-03-24	2009-12-31	0	0
##	1343	1202	1933-04-30	1992-03-24	2008-05-24	1	0
##	1344	1203	1934-04-09	1992-03-25	2009-12-31	0	0
##	1345	1204	1932-09-15	1992-03-25	2002-10-04	1	0
##	1346	1205	1936-11-23	1992-03-25	2009-12-31	0	0
##	1348	1206	1940-03-27	1992-03-25	2009-12-31	0	0
##	1349	1207	1935-08-11	1992-03-25	2009-12-31	0	0
##	1350	1208	1932-03-15	1992-03-25	2009-12-31	0	0
##	1351	1209	1944-01-03	1992-03-25	2009-12-31	0	0
##	1352	1210	1944-03-01	1992-03-25	2009-12-31	0	0
##	1353	1211	1941-09-03	1992-03-25	2009-12-31	0	0
##	1354	1212	1948-03-29	1992-03-28	2009-12-31	0	0
##	1355	1213	1944-01-24	1992-03-28	2009-12-31	0	0
##	1356	1214	1942-07-30	1992-03-28	2009-12-31	0	0

##	1357	1215	1932-11-14	1992-03-28	2009-12-31	0	0
##	1358	1216	1946-11-19	1992-03-28	2009-12-31	0	0
##	1359	1217	1948-06-02	1992-03-28	2009-12-31	0	0
##	1360	1218	1943-10-22	1992-03-28	2009-12-31	0	0
##	1361	1219	1949-08-15	1992-03-29	2009-12-31	0	0
##	1362	1220	1946-07-23	1992-03-29	2009-12-31	0	0
##	1363	1221	1945-05-22	1992-03-29	2009-12-31	0	0
##	1364	1222	1947-03-21	1992-03-29	2009-12-31	0	0
##	1365	1223	1951-08-29	1992-03-30	2009-12-31	0	0
##	1366	1224	1947-01-27	1992-03-30	2009-12-31	0	0
##	1367	1225	1934-08-18	1992-03-30	2009-12-31	0	0
##	1368	1226	1944-09-20	1992-03-30	2009-12-31	0	0
##	1369	1227	1944-01-01	1992-03-30	2009-12-31	0	0
##	1370	1228	1944-11-27	1992-03-30	2009-12-31	0	0
##	1371	1229	1948-02-07	1992-03-30	2009-12-31	0	0
##	1372	1230	1942-08-12	1992-03-30	1998-11-23	1	1
##	1373	1231	1951-02-23	1992-03-30	2009-12-31	0	0
##	1374	1232	1944-09-23	1992-03-30	2009-12-31	0	0
##	1375	1233	1935-07-03	1992-03-31	2009-12-31	0	0
##	1376	1234	1941-07-18	1992-03-31	2003-01-13	1	0
##	1377	1235	1938-07-28	1992-03-31	1999-05-08	1	0
##	1378	1236	1936-07-16	1992-03-31	2009-12-31	0	0
##	1379	1237	1936-01-05	1992-03-31	2009-12-31	0	0
##	1380	1238	1941-07-01	1992-03-31	2009-12-31	0	0
##	1381	1239	1939-12-10	1992-03-31	2009-12-31	0	0
##	1382	1240	1932-12-11	1992-04-01	2009-12-31	0	0
##	1383	1241	1944-02-16	1992-04-01	2009-12-31	0	0
##	1384	1242	1937-02-20	1992-04-01	2009-12-31	0	0
##	1385	1243	1941-09-24	1992-04-04	2009-12-31	0	0
##	1386	1244	1933-06-21	1992-04-04	1999-01-19	1	0
##	1387	1245	1942-10-07	1992-04-04	2008-11-02	1	1
##	1388	1246	1937-07-22	1992-04-04	2005-02-09	1	1
##	1390	1247	1948-12-05	1992-04-05	2009-12-31	0	0
##	1391	1248	1949-10-16	1992-04-05	2000-09-01	1	1
##	1392	1249	1933-12-19	1992-04-05	2009-12-31	0	0
##	1393	1250	1951-10-25	1992-04-05	2009-12-31	0	0
##	1395	1251	1934-08-13	1992-04-08	2009-12-31	0	0
##	1396	1252	1948-03-15	1992-04-08	2009-12-31	0	0
##	1397	1253	1944-04-17	1992-04-08	2009-12-31	0	0
##	1398	1254	1948-05-04	1992-04-08	2009-12-31	0	0
##	1400	1255	1931-11-15	1992-04-08	2009-12-31	0	0
##	1401	1256	1943-04-12	1992-04-26	2009-12-31	0	0
##	1403	1257	1931-10-12	1992-04-26	2009-12-31	0	0
##	1404	1258	1935-11-14	1992-04-26	2009-12-31	0	0
##	1405	1259	1934-03-29	1992-04-26	2009-03-16	1	0
##	1407	1260	1947-10-19	1992-04-27	2001-10-05	1	1
##	1410	1261	1939-07-18	1992-04-28	1992-10-28	1	1
##	1411	1262	1950-10-17	1992-04-28	2009-12-31	0	0
##	1412	1263	1941-06-01	1992-04-28	2009-12-31	0	0
##	1413	1264	1951-08-29	1992-04-28	2009-12-31	0	0
##	1414	1265	1947-05-02	1992-04-28	2009-12-31	0	0
##	1415	1266	1935-02-20	1992-04-29	2009-12-31	0	0
##	1416	1267	1934-01-07	1992-04-29	1999-04-25	1	0
##	1417	1268	1932-03-18	1992-04-29	2009-12-31	0	0
##	1418	1269	1946-04-22	1992-04-29	2009-12-31	0	0
##	1419	1270	1951-11-14	1992-04-29	2009-12-31	0	0
##	1421	1271	1935-05-16	1992-04-29	2009-12-31	0	0

##	1422	1272	1948-02-03	1992-04-29	2009-12-31	0	0
##	1425	1273	1937-08-17	1992-05-02	1999-06-11	1	1
##	1426	1274	1931-07-24	1992-05-02	2009-12-31	0	0
##	1427	1275	1951-06-27	1992-05-02	2009-12-31	0	0
##	1429	1276	1938-06-17	1992-05-03	2009-12-31	0	0
##	1431	1277	1933-06-18	1992-05-05	2009-12-31	0	0
##	1432	1278	1945-12-19	1992-05-16	2009-12-31	0	0
##	1433	1279	1939-08-11	1992-05-16	2009-12-31	0	0
##	1434	1280	1935-10-15	1992-05-16	2009-12-31	0	0
##	1435	1281	1939-04-30	1992-05-16	2009-12-31	0	0
##	1436	1282	1936-01-12	1992-05-16	2009-12-31	0	0
##	1437	1283	1947-02-07	1992-05-16	2009-12-31	0	0
##	1438	1284	1936-12-02	1992-05-17	2009-12-31	0	0
##	1439	1285	1944-02-17	1992-05-17	2009-12-31	0	0
##	1440	1286	1934-04-29	1992-05-17	2008-09-17	1	0
##	1441	1287	1941-09-25	1992-05-17	2009-12-31	0	0
##	1442	1288	1932-03-29	1992-05-17	2009-12-31	0	0
##	1443	1289	1936-12-03	1992-05-18	2009-12-31	0	0
##	1444	1290	1937-10-04	1992-05-18	2009-12-31	0	0
##	1445	1291	1939-04-08	1992-05-18	2009-12-31	0	0
##	1446	1292	1944-01-20	1992-05-18	2009-12-31	0	0
##	1447	1293	1942-01-19	1992-05-18	2009-12-31	0	0
##	1448	1294	1948-08-26	1992-05-19	2009-12-31	0	0
##	1449	1295	1944-10-27	1992-05-19	2009-12-31	0	0
##	1450	1296	1939-11-17	1992-05-19	2009-12-31	0	0
##	1451	1297	1939-06-02	1992-05-19	2009-12-31	0	0
##	1452	1298	1943-05-05	1992-05-19	2009-12-31	0	0
##	1453	1299	1935-05-25	1992-05-20	2009-12-31	0	0
##	1454	1300	1934-08-29	1992-05-20	2009-12-31	0	0
##	1455	1301	1940-01-11	1992-05-20	2009-12-31	0	0
##	1456	1302	1947-05-22	1992-05-20	2009-12-31	0	0
##	1457	1303	1942-02-19	1992-05-20	2009-12-31	0	0
##	1458	1304	1948-09-03	1992-05-20	2009-12-31	0	0
##	1459	1305	1944-09-03	1992-05-20	2009-12-31	0	0
##	1461	1306	1943-10-20	1992-05-23	2009-12-31	0	0
##	1462	1307	1943-05-03	1992-05-23	2009-12-31	0	0
##	1464	1308	1934-02-02	1992-05-23	2009-12-31	0	0
##	1465	1309	1939-03-05	1992-05-23	2009-12-31	0	0
##	1466	1310	1951-08-05	1992-05-23	2009-12-31	0	0
##	1467	1311	1940-03-17	1992-05-23	2009-12-31	0	0
##	1468	1312	1932-08-01	1992-05-24	1996-06-18	1	0
##	1469	1313	1945-12-17	1992-05-24	2009-12-31	0	0
##	1470	1314	1939-05-15	1992-05-24	2009-12-31	0	0
##	1471	1315	1947-02-12	1992-05-24	2002-11-27	1	1
##	1472	1316	1950-12-19	1992-05-25	2009-12-31	0	0
##	1473	1317	1941-12-16	1992-05-25	2009-12-31	0	0
##	1474	1318	1942-03-08	1992-05-25	2009-12-31	0	0
##	1475	1319	1937-12-16	1992-05-26	2009-12-31	0	0
##	1476	1320	1937-08-25	1992-05-26	2009-12-31	0	0
##	1477	1321	1946-09-03	1992-05-26	2009-12-31	0	0
##	1479	1322	1941-03-05	1992-05-27	2005-02-01	1	1
##	1480	1323	1948-02-06	1992-05-27	2009-12-31	0	0
##	1481	1324	1946-06-23	1992-05-27	2009-12-31	0	0
##	1482	1325	1948-09-14	1992-05-27	2009-12-31	0	0
##	1483	1326	1935-08-15	1992-05-27	2009-12-31	0	0
##	1484	1327	1941-03-05	1992-05-30	2009-12-31	0	0
##	1485	1328	1940-08-31	1992-05-30	2009-12-31	0	0

##	1486	1329	1942-11-05	1992-05-30	2009-12-31	0	0
##	1487	1330	1944-07-20	1992-05-30	2009-12-31	0	0
##	1488	1331	1944-12-07	1992-05-30	2009-12-31	0	0
##	1489	1332	1945-08-03	1992-05-30	2009-12-31	0	0
##	1490	1333	1934-10-05	1992-05-30	2009-12-31	0	0
##	1491	1334	1948-07-10	1992-05-31	2009-12-31	0	0
##	1492	1335	1945-07-09	1992-05-31	2009-12-31	0	0
##	1493	1336	1937-01-02	1992-05-31	2005-05-04	1	1
##	1494	1337	1943-04-02	1992-05-31	2009-12-31	0	0
##	1495	1338	1946-05-11	1992-05-31	2009-12-31	0	0
##	1496	1339	1945-05-05	1992-05-31	2009-12-31	0	0
##	1497	1340	1949-03-27	1992-06-01	2009-12-31	0	0
##	1498	1341	1941-11-16	1992-06-01	2009-12-31	0	0
##	1499	1342	1946-01-26	1992-06-01	2009-12-31	0	0
##	1500	1343	1944-11-26	1992-06-01	2009-12-31	0	0
##	1501	1344	1943-04-14	1992-06-02	2009-12-31	0	0
##	1502	1345	1939-02-12	1992-06-02	2009-12-31	0	0
##	1504	1346	1945-06-03	1992-06-02	2009-12-31	0	0
##	1505	1347	1935-10-07	1992-06-02	2005-03-25	1	1
##	1506	1348	1945-08-27	1992-06-02	2009-12-31	0	0
##	1507	1349	1942-03-13	1992-06-02	2009-12-31	0	0
##	1508	1350	1936-06-10	1992-06-03	2009-12-31	0	0
##	1509	1351	1935-03-03	1992-06-03	2009-12-31	0	0
##	1510	1352	1943-09-26	1992-06-03	2009-12-31	0	0
##	1511	1353	1940-03-16	1992-06-03	2009-12-31	0	0
##	1512	1354	1947-03-16	1992-06-03	2009-12-31	0	0
##	1513	1355	1950-12-11	1992-06-06	2009-12-31	0	0
##	1514	1356	1938-11-15	1992-06-06	1994-01-25	1	0
##	1515	1357	1936-06-09	1992-06-13	2009-12-31	0	0
##	1516	1358	1942-09-18	1992-06-13	2009-12-31	0	0
##	1518	1359	1936-12-03	1992-06-13	2009-12-31	0	0
##	1519	1360	1951-02-22	1992-06-14	2009-12-31	0	0
##	1520	1361	1936-07-10	1992-06-14	2009-12-31	0	0
##	1521	1362	1931-03-15	1992-06-14	2009-12-31	0	0
##	1522	1363	1931-09-11	1992-06-14	2009-12-31	0	0
##	1524	1364	1938-04-30	1992-06-14	2005-11-28	1	0
##	1525	1365	1951-02-26	1992-06-15	2009-12-31	0	0
##	1526	1366	1945-09-27	1992-06-15	2009-12-31	0	0
##	1527	1367	1939-07-06	1992-06-15	2009-12-31	0	0
##	1528	1368	1951-02-10	1992-06-15	2009-12-31	0	0
##	1529	1369	1944-03-12	1992-06-15	2009-12-31	0	0
##	1530	1370	1947-12-06	1992-06-15	2009-12-31	0	0
##	1531	1371	1949-09-09	1992-06-15	2009-12-31	0	0
##	1532	1372	1943-06-02	1992-06-15	2009-12-31	0	0
##	1533	1373	1945-09-24	1992-06-15	2009-12-31	0	0
##	1534	1374	1945-07-13	1992-06-16	2009-12-31	0	0
##	1535	1375	1943-03-03	1992-06-16	2009-12-31	0	0
##	1536	1376	1943-05-27	1992-06-16	2009-12-31	0	0
##	1537	1377	1946-07-11	1992-06-17	2009-12-31	0	0
##	1538	1378	1942-01-19	1992-06-17	2009-12-31	0	0
##	1539	1379	1950-08-12	1992-06-17	2009-12-31	0	0
##	1540	1380	1942-11-02	1992-06-17	2009-12-31	0	0
##	1542	1381	1934-12-04	1992-06-17	2009-12-31	0	0
##	1543	1382	1942-08-31	1992-06-17	2009-12-31	0	0
##	1544	1383	1945-02-25	1992-06-17	2009-12-31	0	0
##	1545	1384	1937-05-11	1992-06-20	2009-12-31	0	0
##	1546	1385	1941-05-10	1992-06-20	2009-12-31	0	0

##	1547	1386	1943-12-20	1992-06-20	2009-12-31	0	0
##	1548	1387	1944-11-06	1992-06-20	2009-12-31	0	0
##	1550	1388	1932-10-02	1992-06-21	2009-12-31	0	0
##	1551	1389	1943-04-07	1992-06-21	2009-12-31	0	0
##	1552	1390	1932-09-19	1992-06-21	2009-12-31	0	0
##	1553	1391	1951-05-30	1992-06-21	2009-12-31	0	0
##	1554	1392	1945-03-25	1992-06-22	2009-12-31	0	0
##	1555	1393	1939-02-13	1992-06-23	2009-12-31	0	0
##	1556	1394	1935-03-20	1992-06-23	2009-12-31	0	0
##	1557	1395	1940-04-21	1992-06-23	2002-05-28	1	1
##	1558	1396	1945-05-28	1992-06-23	2009-12-31	0	0
##	1559	1397	1943-03-19	1992-06-23	2009-12-31	0	0
##	1560	1398	1945-04-27	1992-08-08	2009-12-31	0	0
##	1561	1399	1952-03-05	1992-08-08	2009-12-31	0	0
##	1563	1400	1951-11-29	1992-08-08	2009-12-31	0	0
##	1564	1401	1949-02-06	1992-08-08	2009-12-31	0	0
##	1565	1402	1931-04-30	1992-08-08	2009-12-31	0	0
##	1566	1403	1932-07-22	1992-08-08	2009-12-31	0	0
##	1567	1404	1935-01-30	1992-08-08	2009-12-31	0	0
##	1569	1405	1936-08-13	1992-08-09	2009-12-31	0	0
##	1570	1406	1941-03-25	1992-08-09	2009-12-31	0	0
##	1571	1407	1934-10-17	1992-08-09	2000-10-04	1	0
##	1572	1408	1947-12-27	1992-08-09	2009-02-01	1	1
##	1573	1409	1934-06-23	1992-08-09	2002-09-16	1	0
##	1574	1410	1939-03-31	1992-08-09	2002-04-03	1	0
##	1576	1411	1946-09-01	1992-08-09	2009-12-31	0	0
##	1577	1412	1946-03-23	1992-08-09	2009-12-31	0	0
##	1578	1413	1939-12-21	1992-08-10	2009-12-31	0	0
##	1579	1414	1945-10-30	1992-08-10	2009-12-31	0	0
##	1581	1415	1940-02-21	1992-08-10	2002-06-02	1	1
##	1582	1416	1943-03-07	1992-08-11	2009-12-31	0	0
##	1583	1417	1933-02-05	1992-08-11	2009-12-31	0	0
##	1584	1418	1937-08-25	1992-08-11	2009-12-31	0	0
##	1585	1419	1939-07-04	1992-08-11	2009-12-31	0	0
##	1586	1420	1948-02-27	1992-08-11	2009-12-31	0	0
##	1587	1421	1943-12-04	1992-08-11	2006-02-06	1	0
##	1591	1422	1938-08-07	1992-08-12	2006-06-12	1	0
##	1592	1423	1944-09-15	1992-08-12	2009-12-31	0	0
##	1593	1424	1938-01-17	1992-08-12	2009-12-31	0	0
##	1594	1425	1949-03-10	1992-08-12	2009-12-31	0	0
##	1595	1426	1939-08-26	1992-08-12	2009-12-31	0	0
##	1597	1427	1949-04-29	1992-08-16	2009-12-31	0	0
##	1598	1428	1942-09-01	1992-08-16	2009-12-31	0	0
##	1599	1429	1950-07-07	1992-08-16	2009-12-31	0	0
##	1600	1430	1952-02-05	1992-08-16	2009-05-11	1	0
##	1601	1431	1940-07-29	1992-08-16	2009-12-31	0	0
##	1602	1432	1940-07-18	1992-08-16	2005-08-12	1	0
##	1603	1433	1943-12-02	1992-10-31	2009-12-31	0	0
##	1604	1434	1944-04-09	1992-10-31	2009-12-31	0	0
##	1606	1435	1932-01-16	1992-10-31	2009-12-31	0	0
##	1607	1436	1937-10-09	1992-10-31	2009-12-31	0	0
##	1608	1437	1952-01-09	1992-11-01	2009-12-31	0	0
##	1610	1438	1939-11-12	1992-11-01	2009-12-31	0	0
##	1611	1439	1944-05-12	1992-11-01	2002-05-14	1	0
##	1612	1440	1947-02-10	1992-11-01	2009-12-31	0	0
##	1613	1441	1949-03-30	1992-11-01	2009-12-31	0	0
##	1614	1442	1941-11-19	1992-11-01	2009-12-31	0	0

##	1615	1443	1947-08-08	1992-11-02	2009-12-31	0	0
##	1616	1444	1936-03-10	1992-11-02	2009-12-31	0	0
##	1617	1445	1944-01-12	1992-11-02	2009-12-31	0	0
##	1618	1446	1942-11-30	1992-11-02	2009-12-31	0	0
##	1620	1447	1934-03-27	1992-11-02	2009-12-31	0	0
##	1621	1448	1952-07-23	1992-11-04	2009-12-31	0	0
##	1622	1449	1933-07-07	1992-11-04	1997-04-02	1	0
##	1623	1450	1937-09-19	1992-11-04	2009-12-31	0	0
##	1624	1451	1946-05-07	1992-11-04	2009-12-31	0	0
##	1625	1452	1932-09-14	1992-11-04	2009-12-31	0	0
##	1626	1453	1946-07-09	1992-11-04	2009-12-31	0	0
##	1627	1454	1932-06-28	1992-11-08	2007-09-28	1	0
##	1629	1455	1943-10-25	1992-11-08	2009-12-31	0	0
##	1630	1456	1935-10-24	1992-11-08	2002-04-21	1	1
##	1631	1457	1951-03-27	1992-11-08	2009-12-31	0	0
##	1633	1458	1939-09-27	1992-11-09	2009-12-31	0	0
##	1634	1459	1931-08-24	1992-11-09	2000-08-15	1	1
##	1635	1460	1933-11-13	1992-11-09	2001-10-27	1	0
##	1636	1461	1939-03-27	1992-11-09	1999-11-29	1	0
##	1637	1462	1935-05-07	1992-11-09	2005-04-11	1	0
##	1638	1463	1950-08-30	1992-11-09	2009-12-31	0	0
##	1639	1464	1940-04-24	1992-11-10	2009-12-31	0	0
##	1640	1465	1950-05-07	1992-11-10	2009-12-31	0	0
##	1641	1466	1944-10-01	1992-11-10	2009-12-31	0	0
##	1642	1467	1943-07-31	1992-11-10	2009-12-31	0	0
##	1643	1468	1937-08-28	1992-11-10	2009-12-31	0	0
##	1644	1469	1951-10-30	1992-11-10	2009-12-31	0	0
##	1645	1470	1944-06-25	1992-11-10	2009-12-31	0	0
##	1646	1471	1948-06-05	1992-11-10	2009-12-31	0	0
##	1647	1472	1933-12-08	1992-11-10	2009-12-31	0	0
##	1648	1473	1952-08-17	1992-11-11	2009-12-31	0	0
##	1649	1474	1938-09-08	1992-11-11	2009-12-31	0	0
##	1650	1475	1940-03-31	1992-11-15	2009-12-31	0	0
##	1651	1476	1951-06-20	1992-11-15	2009-12-31	0	0
##	1652	1477	1952-06-20	1992-11-16	2009-12-31	0	0
##	1653	1478	1940-04-06	1992-11-16	1997-08-01	1	0
##	1654	1479	1952-05-12	1992-11-17	2009-12-31	0	0
##	1655	1480	1952-03-08	1992-11-17	2009-12-31	0	0
##	1656	1481	1944-09-29	1992-11-17	2009-12-31	0	0
##	1657	1482	1951-12-18	1992-11-17	2009-12-31	0	0
##	1658	1483	1943-11-14	1992-11-17	2009-12-31	0	0
##	1659	1484	1951-04-08	1992-11-17	2009-12-31	0	0
##	1660	1485	1944-12-09	1992-11-17	1996-10-14	1	0
##	1661	1486	1952-03-22	1992-11-17	2009-12-31	0	0
##	1662	1487	1952-02-05	1992-12-05	2009-12-31	0	0
##	1663	1488	1947-06-23	1992-12-05	2009-12-31	0	0
##	1664	1489	1939-12-13	1992-12-05	2009-12-31	0	0
##	1665	1490	1951-10-15	1992-12-05	2009-12-31	0	0
##	1666	1491	1934-05-19	1992-12-05	2009-12-31	0	0
##	1667	1492	1933-11-02	1992-12-07	2009-12-31	0	0
##	1668	1493	1939-03-14	1992-12-07	2009-12-31	0	0
##	1669	1494	1943-05-18	1992-12-08	2009-12-31	0	0
##	1670	1495	1933-06-06	1992-12-08	2009-12-31	0	0
##	1671	1496	1946-05-24	1992-12-08	2009-12-31	0	0
##	1672	1497	1946-09-13	1992-12-08	2009-12-31	0	0
##	1673	1498	1939-09-02	1992-12-09	2009-12-31	0	0
##	1675	1499	1933-10-25	1992-12-09	2003-01-31	1	1

```
## 1676 1500 1939-04-25 1992-12-09 2009-12-31    0    0
## 1677 1501 1933-05-07 1992-12-09 2009-12-31    0    0
```

```
occoh.txt.tidy
```

```
## # A tibble: 1,501 x 6
##       id birth entry      exit      death      chdeath
##   <dbl> <dbl> <date>    <date>    <date>    <dbl>
## 1     1     1     1 1943-02-19 1990-08-14 2009-12-31      0
## 2     2     2     2 1934-07-06 1990-08-14 2009-12-31      0
## 3     3     3     3 1939-03-05 1990-08-14 2009-12-31      0
## 4     4     4     4 1939-07-03 1990-08-14 2009-12-31      0
## 5     5     5     5 1935-02-18 1990-08-14 2006-03-13      1
## 6     6     6     6 1936-03-07 1990-08-14 2007-06-10      1
## 7     7     7     7 1944-03-30 1990-08-15 2007-04-14      1
## 8     8     8     8 1942-11-24 1990-08-15 2006-10-30      1
## 9     9     9     9 1942-09-11 1990-08-15 2009-12-31      0
## 10    10    10    10 1931-03-01 1990-08-15 2009-12-31      0
## # i 1,491 more rows
```

```
# table dimensions
```

```
dim(occoh.txt.core)
```

```
## [1] 1501     6
```

```
dim(occoh.txt.tidy)
```

```
## [1] 1501     6
```

```
# compare column types
```

```
map(occoh.txt.core, class)
```

```
## $id
## [1] "integer"
##
## $birth
## [1] "character"
##
## $entry
## [1] "character"
##
## $exit
## [1] "character"
##
## $death
## [1] "integer"
##
## $chdeath
## [1] "integer"
```

```
map(occoh.txt.tidy, class)
```

```
## $id
## [1] "numeric"
##
## $birth
## [1] "numeric"
##
## $entry
## [1] "Date"
##
```

```
## $exit
## [1] "Date"
##
## $death
## [1] "Date"
##
## $chdeath
## [1] "numeric"
```

As you can see, in addition to inferring the type of columns in the input data (here some dates), using `readr` to load your data-set can help you to detect inconsistencies in input data formatting (there are no true problem here).

If you are interested, you can explore the other functions of `readr` and see how you can tune it.

3.13 String manipulation with `stringr`

Another popular `tidyverse` popular package is `stringr` package. This package is specialized in the string manipulation. Here are couple of examples.

Let's create a character vector with the following elements representing country names: "Estonia", "Finland", "Denmark", "United Kingdom", "France".

```
countries <-
  c("Estonia", "Finland", "Denmark", "United Kingdom", "France")
```

With `stringr` functions perform the following actions.

Extract the first three characters from each country name:

```
country_initials <- str_sub(countries, start = 1, end = 3)
```

Convert all country names to uppercase:

```
countries_upper <- str_to_upper(countries)
```

Replace "United" with "Utd" in each country name:

```
countries_modified <- str_replace(countries, "United", "Utd")
```

Find the positions of the letter "n" in each country name:

```
a_positions <- str_locate_all(countries, "n")
```

As you can see, the output of `str_locate_all` is a list (one element per character string) containing a 2 column table with one line for each match. The first column (start) being the position of the beginning of the match and the second one (end) being the end of the match. In our case, since we are searching for a single character match, this 2 indexes are always the same.

Count the number of characters in each country name:

```
character_counts <- str_length(countries)
```

These examples demonstrate various string manipulation operations using the `stringr` package. You can modify the exercises, combine several operations or explore other string manipulation functions provided by `stringr` to further practice and enhance your skills in manipulating and analyzing text data.

3.14 `purrr` package to apply functions to list

Among my favorite `tidyverse` packages, you will find `purrr`. This package contains several functions that are very similar to `lapply` function.

Apply a function to each element of the vector using `map()`. Here producing the mean of some grades per class:

```
# define the grade dataset
grades <-
  list(
    c1 = c(80, 85, 90),
    c2 = c(75, 70, 85, 88),
    c3 = c(90, 85, 95)
  )
# compute grades
mean_grades <- map(grades, mean)
```

By default `map()` return a list. One of the nice feature of **purrr** functions is to be able to specify the type of output you want (e.g. `_dbl` for numeric, `_chr` for characters, ...). Check and try to explain the differences between the following command lines:

```
map(grades, mean)
```

```
## $c1
## [1] 85
##
## $c2
## [1] 79.5
##
## $c3
## [1] 90
```

```
map_dbl(grades, mean)
```

```
##   c1   c2   c3
## 85.0 79.5 90.0
```

`map_chr()` requires the function to return a character. Since `mean()` returns a number, the call below errors with a clear type-mismatch message (in older versions of **purrr** the result was silently coerced to a string).

```
map_chr(grades, mean)
```

```
## Error in `map_chr()`:
## i In index: 1.
## i With name: c1.
## Caused by error:
## ! Can't coerce from a number to a string.
```

```
map_df(grades, mean)
```

```
## # A tibble: 1 x 3
##       c1     c2     c3
##   <dbl> <dbl> <dbl>
## 1    85   79.5    90
```

Other nice features of `map` like functions is the availability to support more than one argument. `map2()` for 2 arguments and `pmap()` for more than 2. This can be very handy in some conditions. If you are interested you can have a look to this function help file and play with the examples.

purrr package has also a set of functions that can be used to apply iteratively a function using **reduce** and/or **accumulate**. The 2 functions behave the same way, it takes the 2 first element of a list, apply a function taking 2 arguments. The results is combined with the third element of the list and given as input to the same function and so on.. The only difference is that **accumulate** return intermediate results while **reduce** return only the final results.

Here an example of the cumulative product of the 10 first numbers.

```
1:10 |> purrr::reduce(`*`)
```

```
## [1] 3628800
1:10 |> purrr::accumulate(`*`)

## [1]      1      2      6     24    120    720   5040  40320 362880
## [10] 3628800
```

`purrr` have many of others useful features. Please check the dedicated documentation if you want to go further with this package.

3.15 Bonus: Rendering tables

Once you have produced a nice data-set we can be interested in rendering it in a nice format that can meet presentation/publication expectations. The `kableExtra` table can be useful to achieve this goal.

```
# if(!require(kableExtra)) install.packages('kableExtra')
library(kableExtra)

births.08 <-
  births_tbl |>
  filter(
    !is.na(gest4)
  ) |>
  group_by(gest4) |>
  summarise(
    N = n()
  ) |>
  mutate(
    `(`)` = (N / sum(N)) |> scales::percent()
  )

# default
births.08

# create an html version of the table and save it on the hard drive
births.08 |>
  kable() |>
  kable_styling(
    bootstrap_options =
      c("striped", "hover", "condensed", "responsive"),
    full_width = FALSE
  ) |>
  save_kable(file = "births.08.html", self_contained = TRUE)
```

note: One other very cool package to produce advance formatted Excel spreadsheet I am using more and more is `openxlsx`. Check it out if you are interested.

Chapter 4

Tabulation

4.1 Introduction

R and its add-on packages provide several different tabulation functions with different capabilities. The appropriate function to use depends on your goal. There are at least three different uses for tables.

The first use is to create simple summary statistics that will be used for further calculations in R. For example, a two-by-two table created by the `table` function can be passed to `fisher.test`, which will calculate odds ratios and confidence intervals. The appearance of these tables may, however, be quite basic, as their principal goal is to create new objects for future calculations.

A quite different use of tabulation is to make *production quality* tables for publication. You may want to generate reports from for publication in paper form, or on the World Wide Web. The package `xtable` provides this capability, but it is not covered by this course.

An intermediate use of tabulation functions is to create human-readable tables for discussion within your work-group, but not for publication. The `Epi` package provides a function `stat.table` for this purpose, and this practical is designed to introduce this function.

4.2 The births data

We shall use the births data which concern 500 mothers who had singleton births in a large London hospital. The outcome of interest is the birth weight of the baby, also dichotomised as normal or low birth weight. These data are available in the `Epi` package:

```
library(Epi)
data(births)
names(births)
```

```
## [1] "id"      "bweight" "lowbw"   "gestwks" "preterm" "matage"  "hyp"
## [8] "sex"
```

```
head(births)
```

```
##   id bweight lowbw gestwks preterm matage hyp sex
## 1  1   2974     0   38.52      0     34  0  2
## 2  2   3270     0    NA      NA     30  0  1
## 3  3   2620     0   38.15      0     35  0  2
## 4  4   3751     0   39.80      0     31  0  1
## 5  5   3200     0   38.89      0     33  1  1
## 6  6   3673     0   40.97      0     33  0  2
```

In order to work with this data set we need to transform some of the variables into factors. This is done with the following commands:

```

births$hyp <- factor(births$hyp, labels = c("normal", "hyper"))
births$sex <- factor(births$sex, labels = c("M", "F"))
births$agegrp <-
  cut(
    births$matage,
    breaks = c(20, 25, 30, 35, 40, 45),
    right = FALSE
  )
births$gest4 <-
  cut(
    births$gestwks,
    breaks = c(20, 35, 37, 39, 45),
    right = FALSE
  )

```

Now use `str(births)` to examine the modified data frame. We have transformed the binary variables `hyp` and `sex` into factors with informative labels. This will help when displaying the tables. We have also created grouped variables `agegrp` and `gest4` from the continuous variables `matage` and `gestwks` so that they can be tabulated.

4.3 One-way tables

The simplest table one-way table is created by

```
stat.table(index = sex, data = births)
```

```

## -----
## sex    count()
## -----
## M          264
## F          236
## -----

```

This creates a count of individuals, classified by levels of the factor `sex`. Compare this table with the equivalent one produced by the `table` function. Note that `stat.table` has a `data` argument that allows you to use variables in a data frame without specifying the frame.

You can display several summary statistics in the same table by giving a list of expressions to the `contents` argument:

```

stat.table(
  index = sex,
  contents = list(count(), percent(sex)),
  data = births
)

```

```

## -----
## sex    count() percent(sex)
## -----
## M          264          52.8
## F          236          47.2
## -----

```

Only a limited set of expressions are allowed: see the help page for `stat.table` for details.

You can also calculate marginal tables by specifying `margin=TRUE` in your call to `stat.table`. Do this for the above table. Check that the percentages add up to 100 and the total for `count()` is the same as the number of rows of the data frame `births`.

```
## -----
```

```
## sex      count() percent(sex)
## -----
## M          264         52.8
## F          236         47.2
##
## Total       500        100.0
## -----
```

To see how the mean birth weight changes with `sex`, try

```
stat.table(index = sex, contents = mean(bweight), data = births)
```

```
## -----
## sex    mean(bweight)
## -----
## M          3229.90
## F          3032.83
## -----
```

Add the count to this table. Add also the margin with `margin=TRUE`.

```
## -----
## sex      count() mean(bweight)
## -----
## M          264        3229.90
## F          236        3032.83
##
## Total       500        3136.88
## -----
```

As an alternative to `bweight` we can look at `lowbw` with

```
stat.table(index = sex, contents = percent(lowbw), data = births)
```

```
## -----
## sex    percent(lowbw)
## -----
## M          100.0
## F          100.0
## -----
```

All the percentages are 100! To use the `percent` function the variable `lowbw` must also be in the index, as in

```
stat.table(
  index = list(sex, lowbw),
  contents = percent(lowbw),
  data = births
)
```

```
## -----
##      -----lowbw-----
## sex          0          1
## -----
## M          89.8      10.2
## F          86.0      14.0
## -----
```

The final column is the percentage of babies with low birth weight by different categories of gestation.

- Obtain a table showing the frequency distribution of `gest4`.
- Show how the mean birth weight changes with `gest4`.
- Show how the percentage of low birth weight babies changes with `gest4`.

```
## -----
## gest4      count()
## -----
## [20,35)      31
## [35,37)      32
## [37,39)     167
## [39,45)     260
## -----

## -----
## gest4      mean(bweight)
## -----
## [20,35)     1733.74
## [35,37)     2590.31
## [37,39)     3093.77
## [39,45)     3401.26
## -----

## -----
##              -----gest4-----
## lowbw  [20,35) [35,37) [37,39) [39,45)
## -----
## 0          19.4   59.4   89.2   98.8
## 1          80.6   40.6   10.8    1.2
## -----
```

Another way of obtaining the percentage of low birth weight babies by gestation is to use the ratio function:

```
stat.table(gest4, ratio(lowbw, 1, 100), data = births)
```

```
## -----
## gest4      ratio(lowbw,
##              1, 100)
## -----
## [20,35)      80.65
## [35,37)      40.62
## [37,39)      10.78
## [39,45)       1.15
## -----
```

This only works because `lowbw` is coded 0/1, with 1 for low birth weight.

Tables of odds can be produced in the same way by using `ratio(lowbw, 1-lowbw)`. The `ratio` function is also very useful for making tables of rates with (say) `ratio(D,Y,1000)` where `D` is the number of failures, and `Y` is the follow-up time. We shall return to rates in a later practical.

4.4 Improving the Presentation of Tables

The `stat.table` function provides default column headings based on the `contents` argument, but these are not always very informative. Supply your own column headings using *tagged* lists as the value of the `contents` argument, within a `stat.table` call:

```
stat.table(gest4, contents = list(
  N = count(),
  "(%)" = percent(gest4)
), data = births)
```

```
## -----
## gest4      N      (%)
## -----
```

```
## [20,35)      31      6.3
## [35,37)      32      6.5
## [37,39)     167     34.1
## [39,45)     260     53.1
## -----
```

This improves the readability of the table. It remains to give an informative title to the index variable. You can do this in the same way: instead of giving `gest4` as the `index` argument to `stat.table`, use a named list:

```
stat.table(index = list("Gestation time" = gest4), data = births)
```

```
## -----
## Gestation   count()
## time
## -----
## [20,35)      31
## [35,37)      32
## [37,39)     167
## [39,45)     260
## -----
```

4.5 Two-way Tables

The following call gives a 2×2 table showing the mean birth weight cross-classified by `sex` and `hyp`.

```
stat.table(
  list(sex, hyp),
  contents = mean(bweight),
  data = births
)
```

```
## -----
##      -----hyp-----
## sex    normal  hyper
## -----
## M      3310.75 2814.40
## F      3079.50 2699.72
## -----
```

Add the count to this table and repeat the function call using `margin = TRUE` to calculate the marginal tables.

```
## -----
##      -----hyp-----
## sex    normal  hyper  Total
## -----
## M           221     43    264
##           3310.75 2814.40 3229.90
##
## F           207     29    236
##           3079.50 2699.72 3032.83
##
##
## Total        428     72    500
##           3198.90 2768.21 3136.88
## -----
```

Use `stat.table` with the ratio function to obtain a 2×2 table of percent low birth weight by `sex` and `hyp`.

```
## -----
##      -----hyp-----
```

```
## sex      normal  hyper  Total
## -----
## M          221    43    264
##          3310.75 2814.40 3229.90
##
## F          207    29    236
##          3079.50 2699.72 3032.83
##
##
## Total      428    72    500
##          3198.90 2768.21 3136.88
## -----

## -----hyp-----
## sex      normal  hyper  Total
## -----
## M          221    43    264
##          6.79   27.91  10.23
##
## F          207    29    236
##          12.08   27.59  13.98
##
##
## Total      428    72    500
##          9.35   27.78  12.00
## -----
```

You can have fine-grained control over which margins to calculate by giving a logical vector to the `margin` argument. Use `margin=c(FALSE, TRUE)` to calculate margins over `sex` but not `hyp`. This might not be what you expect, but the `margin` argument indicates which of the index variables are to be **marginalized out**, not which index variables are to remain.

4.6 Printing

Just like every other R function, `stat.table` produces an object that can be saved and printed later, or used for further calculation. You can control the appearance of a table with an explicit call to `print()`

There are two arguments to the print method for `stat.table`. The `width` argument which specifies the minimum column width, and the `digits` argument which controls the number of digits printed after the decimal point. This table

```
odds.tab <-
  stat.table(
    gest4,
    list("odds of low bw" = ratio(lowbw, 1 - lowbw)),
    data = births
  )
print(odds.tab)
```

```
## -----
## gest4      odds of
##           low bw
## -----
## [20,35)    4.17
## [35,37)    0.68
## [37,39)    0.12
## [39,45)    0.01
```

```
## -----
```

shows a table of odds that the baby has low birth weight. Use `width=15` and `digits=3` and see the difference.

```
## -----
## gest4      odds of low bw
## -----
## [20,35)      4.167
## [35,37)      0.684
## [37,39)      0.121
## [39,45)      0.012
## -----
```


Chapter 5

Graphics in R

There are three kinds of plotting functions in R:

- Functions that generate a new plot, e.g. `hist()` and `plot()`.
- Functions that add extra things to an existing plot, e.g. `lines()` and `text()`.
- Functions that allow you to interact with the plot, e.g. `locator()` and `identify()`.

The normal procedure for making a graph in R is to make a fairly simple initial plot and then add on points, lines, text etc., preferably in a script.

5.1 Simple plot on the screen


Load the births data and get an overview of the variables:

```
library(Epi)
data(births)
str(births)
```

```
## 'data.frame':    500 obs. of  8 variables:
## $ id      : num  1 2 3 4 5 6 7 8 9 10 ...
## $ bweight: num  2974 3270 2620 3751 3200 ...
## $ lowbw   : num  0 0 0 0 0 0 0 0 0 0 ...
## $ gestwks: num  38.5 NA 38.2 39.8 38.9 ...
## $ preterm: num  0 NA 0 0 0 0 0 0 0 0 ...
## $ matage  : num  34 30 35 31 33 33 29 37 36 39 ...
## $ hyp     : num  0 0 0 0 1 0 0 0 0 0 ...
## $ sex     : num  2 1 2 1 1 2 2 1 2 1 ...
```

Now look at the birth weight distribution with

```
hist(births$bweight)
```



The histogram can be refined – take a look at the possible options with

```
help(hist)
```

and try some of the options, for example:


```
hist(births$bweight, col = "gray", border = "white")
```



To


look at the relationship between birthweight and gestational weeks, try

```
with(births, plot(gestwks, bweight))
```




You can change the plot-symbol by the option `pch=`. If you want to see all the plot symbols try:

```
plot(1:25, pch = 1:25)
```




- Make a plot of the birth weight versus maternal age with

```
with(births, plot(matage, bweight))
```



Label the axes with


```
with(  
  births,  
  plot(  
    matage,  
    bweight,  
    xlab = "Maternal age",  
    ylab = "Birth weight (g)"  
  )  
)
```



5.2 Colours


There are many colours recognized by R. You can list them all by `colours()` or, equivalently, `colors()` (R allows you to use British or American spelling). To colour the points of birthweight versus gestational weeks, try

```
with(births, plot(gestwks, bweight, pch = 16, col = "green"))
```



This creates a solid mass of colour in the centre of the cluster of points and it is no longer possible to see individual points. You can recover this information by overwriting the points with black circles using the `points()` function.

```
with(births, plot(gestwks, bweight, pch = 16, col = "green"))  
with(births, points(gestwks, bweight, pch = 1))
```



Note: when the number of data points on a scatter plot is large, you may also want to decrease the point size: to get points that are 50% of the original size, add the parameter `cex=0.5` (or another number <1 for different sizes).


5.3 Adding to a plot

The `points()` function just used is one of several functions that add elements to an existing plot. By using these functions, you can create quite complex graphs in small steps.

Suppose we wish to recreate the plot of birthweight *vs* gestational weeks using different colours for male and female babies. To start with an empty plot, with `type='n'` argument.


Then add the points with the `points` function.

```
with(births, plot(gestwks, bweight, type = "n"))
with(
  births,
  points(gestwks[sex == 1], bweight[sex == 1], col = "blue")
)
with(
  births,
  points(gestwks[sex == 2], bweight[sex == 2], col = "red")
)
```



To add a legend explaining the colours, try


```
with(births, plot(gestwks, bweight, type = "n"))
with(
  births,
  points(gestwks[sex == 1], bweight[sex == 1], col = "blue")
)
with(
  births,
  points(gestwks[sex == 2], bweight[sex == 2], col = "red")
)
legend(
  "topleft",
  pch = 1,
  legend = c("Boys", "Girls"),
  col = c("blue", "red")
)
```



which puts the legend in the top left hand corner.

Finally we can add a title to the plot with

```
with(births, plot(gestwks, bweight, type = "n"))
with(
  births,
  points(gestwks[sex == 1], bweight[sex == 1], col = "blue")
)
with(
  births,
  points(gestwks[sex == 2], bweight[sex == 2], col = "red")
)
legend(
  "topleft",
  pch = 1,
  legend = c("Boys", "Girls"),
  col = c("blue", "red")
)
title(
  "Birth weight vs gestational weeks in 500 singleton births"
)
```



5.4 Using indexing for plot elements

One of the most powerful features of R is the possibility to index vectors, not only to get subsets of them, but also for repeating their elements in complex sequences.

Putting separate colours on males and female as above would become very clumsy if we had a 5 level factor instead of sex.

Instead of specifying one color for all points, we may specify a vector of colours of the same length as the `gestwks` and `bweight` vectors. This is rather tedious to do directly, but R allows you to specify an expression anywhere, so we can use the fact that `sex` takes the values 1 and 2, as follows:

First create a colour vector with two colours, and take look at `sex`:

```
c("blue", "red")
```

```
## [1] "blue" "red"
```

```
births$sex
```

```
## [1] 2 1 2 1 1 2 2 1 2 1 1 2 2 1 2 1 2 2 2 1 2 2 2 1 2 1 2 1 2 2 1 1 2 2 1
## [38] 1 2 1 2 1 1 1 1 1 1 1 1 1 1 2 1 2 1 1 1 1 1 2 1 1 2 1 1 2 2 1 2 1 1 1 1
## [75] 1 2 1 1 2 1 2 1 1 2 1 1 1 2 1 1 1 2 1 2 1 1 2 2 2 2 2 2 2 1 1 2 2 2 1
## [112] 2 2 1 2 2 2 2 1 1 2 1 1 2 1 1 2 1 2 1 1 1 1 2 1 2 1 1 1 2 1 2 2 1 1 1 2 2
## [149] 2 2 2 2 1 2 1 1 2 1 1 2 2 1 2 1 1 1 1 2 1 1 2 2 1 1 1 2 2 2 2 1 2 1 2 1
## [186] 1 2 2 1 2 2 1 2 1 2 2 2 2 1 1 1 2 1 1 2 1 1 1 2 2 2 2 2 1 1 2 1 2 2 2 1
## [223] 2 1 1 1 1 2 1 2 2 2 1 2 2 2 1 1 2 1 1 2 1 2 2 2 1 1 2 2 2 2 1 1 1 2 2 1 2
## [260] 2 2 2 1 1 1 1 1 1 2 1 2 1 2 2 2 2 2 1 1 1 2 2 2 1 2 1 2 1 2 1 1 1 2 2 2 2
## [297] 1 2 1 1 2 1 2 1 1 2 1 2 1 2 2 1 1 2 1 1 2 1 1 2 2 1 1 1 1 2 2 2 2 1 1 1 1
## [334] 2 1 1 2 2 1 2 2 1 1 2 1 1 2 1 2 1 1 1 1 2 1 1 2 2 2 1 1 1 1 1 1 2 2 2 2 2
## [371] 1 1 2 1 1 2 1 2 1 1 1 2 2 1 1 2 1 1 1 1 1 1 1 1 1 1 2 2 1 2 2 1 2 2 1 1
```

```
## [408] 2 2 1 1 2 1 1 2 1 2 1 2 2 1 1 2 2 1 2 2 2 2 1 1 2 2 2 2 2 1 2 1 1 2 2 1
## [445] 1 1 2 2 2 2 1 2 2 2 2 1 1 2 1 2 1 1 1 1 2 1 1 1 2 1 2 1 1 1 1 2 2 1 2 1 2
## [482] 2 1 1 2 2 1 1 1 1 2 2 2 1 1 2 1 2 2 1
```


Now see what happens if you index the colour vector by sex:

```
c("blue", "red")[births$sex]
```

```
## [1] "red" "blue" "red" "blue" "blue" "red" "red" "blue" "red" "blue"
## [11] "blue" "red" "red" "blue" "red" "blue" "red" "blue" "red" "red" "red"
## [21] "red" "blue" "red" "red" "red" "blue" "red" "blue" "red" "red" "blue"
## [31] "red" "red" "blue" "blue" "red" "red" "blue" "blue" "red" "blue" "blue"
## [41] "red" "blue" "blue" "blue" "blue" "blue" "blue" "blue" "blue" "blue" "blue"
## [51] "blue" "red" "blue" "red" "blue" "blue" "blue" "blue" "blue" "blue" "blue"
## [61] "red" "blue" "blue" "red" "blue" "blue" "red" "red" "blue" "red" "red"
## [71] "blue" "blue" "blue" "blue" "blue" "red" "blue" "blue" "red" "blue" "blue"
## [81] "red" "blue" "blue" "red" "blue" "blue" "blue" "red" "blue" "blue" "blue"
## [91] "blue" "red" "blue" "red" "blue" "blue" "red" "blue" "red" "red" "red"
## [101] "red" "red" "red" "red" "red" "blue" "blue" "red" "red" "red" "red"
## [111] "blue" "red" "red" "blue" "red" "red" "red" "red" "red" "blue" "blue"
## [121] "red" "blue" "blue" "red" "blue" "blue" "red" "blue" "red" "red" "blue"
## [131] "blue" "blue" "blue" "red" "blue" "red" "blue" "blue" "blue" "blue" "red"
## [141] "blue" "red" "red" "blue" "blue" "blue" "red" "red" "red" "red" "red"
## [151] "red" "red" "blue" "red" "blue" "blue" "red" "blue" "blue" "red" "red"
## [161] "red" "blue" "red" "blue" "blue" "blue" "blue" "blue" "blue" "red" "blue"
## [171] "blue" "red" "red" "blue" "blue" "blue" "red" "red" "red" "red" "red"
## [181] "blue" "red" "blue" "red" "blue" "blue" "red" "red" "red" "blue" "red"
## [191] "red" "blue" "red" "blue" "red" "red" "red" "red" "red" "blue" "blue"
## [201] "blue" "red" "blue" "blue" "red" "blue" "blue" "blue" "blue" "blue" "red"
## [211] "red" "red" "red" "red" "blue" "blue" "red" "blue" "red" "red" "red"
## [221] "red" "blue" "red" "blue" "blue" "blue" "blue" "red" "blue" "red" "red"
## [231] "red" "red" "blue" "red" "red" "red" "blue" "blue" "red" "blue" "blue"
## [241] "blue" "red" "blue" "red" "red" "red" "blue" "blue" "red" "red" "red"
## [251] "red" "red" "blue" "blue" "blue" "red" "red" "blue" "red" "red" "red"
## [261] "red" "red" "blue" "blue" "blue" "blue" "blue" "blue" "blue" "red" "blue"
## [271] "red" "blue" "red" "red" "red" "red" "red" "red" "blue" "blue" "blue"
## [281] "red" "red" "red" "blue" "red" "blue" "red" "blue" "red" "blue" "blue"
## [291] "blue" "blue" "red" "red" "red" "red" "blue" "red" "blue" "blue" "blue"
## [301] "red" "blue" "red" "blue" "blue" "red" "blue" "red" "blue" "red" "red"
## [311] "red" "blue" "blue" "red" "blue" "blue" "red" "blue" "blue" "blue" "red"
## [321] "red" "blue" "blue" "blue" "blue" "red" "red" "red" "red" "red" "blue"
## [331] "blue" "blue" "blue" "red" "blue" "blue" "red" "red" "blue" "red" "red"
## [341] "red" "blue" "blue" "red" "blue" "blue" "red" "blue" "red" "blue" "blue"
## [351] "blue" "blue" "blue" "red" "blue" "blue" "red" "red" "red" "red" "blue"
## [361] "blue" "blue" "blue" "blue" "blue" "red" "red" "red" "red" "red" "red"
## [371] "blue" "blue" "red" "blue" "blue" "red" "blue" "red" "blue" "blue" "blue"
## [381] "blue" "red" "red" "blue" "blue" "red" "blue" "blue" "blue" "blue" "red"
## [391] "blue" "blue" "blue" "blue" "blue" "blue" "blue" "red" "red" "red" "blue"
## [401] "red" "red" "blue" "red" "red" "blue" "blue" "red" "red" "red" "blue"
## [411] "blue" "red" "blue" "blue" "red" "blue" "red" "blue" "red" "red" "red"
## [421] "blue" "blue" "red" "red" "blue" "red" "red" "red" "red" "red" "blue"
## [431] "blue" "red" "red" "red" "red" "red" "red" "red" "blue" "red" "blue"
## [441] "blue" "red" "red" "blue" "blue" "blue" "red" "red" "red" "red" "red"
## [451] "blue" "red" "red" "red" "red" "blue" "blue" "red" "blue" "red" "red"
## [461] "blue" "blue" "blue" "blue" "red" "blue" "blue" "blue" "blue" "red" "blue"
## [471] "red" "blue" "blue" "blue" "blue" "red" "red" "blue" "red" "blue" "blue"
## [481] "red" "red" "blue" "blue" "red" "red" "blue" "blue" "blue" "blue" "blue"
## [491] "red" "red" "red" "blue" "blue" "red" "blue" "red" "red" "red" "blue"
```

For every occurrence of a 1 in `sex` you get "blue", and for every occurrence of 2 you get "red", so the result is a long vector of "blue"s and "red"s corresponding to the males and females. This can now be used in the plot:

```
with(
  births,
  plot(gestwks, bweight, pch = 16, col = c("blue", "red")[sex])
)
```




The same trick can be used if we want to have a separate symbol for mothers over 40 say. We first generate the indexing variable:

```
births$oldmum <- (births$matage >= 40) + 1
```

Note we add 1 because (`matage >= 40`) generates a logic variable, so by adding 1 we get a numeric variable with values 1 and 2, suitable for indexing:

```
with(
  births,
  plot(
    gestwks,
    bweight,
    pch = c(16, 3)[oldmum],
    col = c("blue", "red")[sex]
  )
)
```




so

where `oldmum` is 1 we get `pch=16` (a dot) and where `oldmum` is 2 we get `pch=3` (a cross).

R will accept any kind of complexity in the indexing as long as the result is a valid index, so you don't need to create the variable `oldmum`, you can create it on the fly:

```
with(
  births,
  plot(
    gestwks,
    bweight,
    pch = c(16, 3)[(matage >= 40) + 1],
    col = c("blue", "red")[sex]
  )
)
```



5.5 Generating colours

R has functions that generate a vector of colours for you. For example,


```
rainbow(4)
```

```
## [1] "#FF0000" "#80FF00" "#00FFFF" "#8000FF"
```

produces a vector with 4 colours (not immediately human readable, though). There are a few other functions that generates other sequences of colours, type `?rainbow` to see them. The `color` function (or `colour` function if you prefer) returns a vector of the colour names that R knows about. These names can also be used to specify colours.

Gray-tones are produced by the function `gray` (or `grey`), which takes a numerical argument between 0 and 1; `gray(0)` is black and `gray(1)` is white. Try:

```
plot(0:10, pch = 16, cex = 3, col = gray(0:10 / 10))
points(0:10, pch = 1, cex = 3)
```



5.6 Saving your graphs for use in other documents

If you need to use the plot in a report or presentation, you can save it in a graphics file. Once you have generated the script (sequence of R commands) that produce the graph (and it looks ok on screen), you can start a non-interactive graphics device and then re-run the script. Instead of appearing on the screen, the plot will now be written directly to a file. After the plot has been completed you will need to close the device again in order to be able to access the file. Try:

```
pdf(file = "bweight_gwks.pdf", height = 4, width = 4)
with(births, plot(gestwks, bweight, col = c("blue", "red")[sex]))
legend(
  "topleft",
  pch = 1,
  legend = c("Boys", "Girls"),
  col = c("blue", "red")
)
dev.off()
```

This will give you a portable document file `bweight_gwks.pdf` with a graph which is 4 inches tall and 4 inches wide.

Instead of *pdf*, other formats can be used (*jpg*, *png*, *tiff*, ...). See `help(Devices)` for the available options.

In window-based environments (R GUI for Windows, R-Studio) you may also use the menu (**File**→**Save as ...** or **Export**) to save the active graph as a file and even copy-paste may work (from R graphics window to Word, for instance) – however, writing it manually into the file is recommended for reproducibility purposes (in case you need to redraw your graph with some modifications).

5.6.1 The ‘par()’ command

It is possible to manipulate any element in a graph, by using the graphics options. These are collected on the help page of `par()`. For example, if you want axis labels always to be horizontal, use the command `par(las=1)`. This will be in effect until a new graphics device is opened.

Look at the typewriter-version of the help-page with

```
help(par)
```

or better, use the the html-version through `Help → Html help → Packages → graphics → P → par`.

It is a good idea to take a print of this (having set the text size to *smallest* because it is long) and carry it with you at any time to read in buses, cinema queues, during boring lectures etc. Don't despair, few R-users can understand what all the options are for.

`par()` can also be used to ask about the current plot, for example `par("usr")` will give you the exact extent of the axes in the current plot.

If you want more plots on a single page you can use the command

```
par(mfrow = c(2, 3))
```

This will give you a layout of 2 rows by 3 columns for the next 6 graphs you produce. The plots will appear by row, i.e. in the top row first. If you want the plots to appear columnwise, use `par(mfcol=c(2,3))` (you still get 2 rows by 3 columns).

To restore the layout to a single plot per page use

```
par(mfrow = c(1, 1))
```


If you want a more detailed control over the layout of multiple graphs on a single page look at `?layout`.

5.7 Interacting with a plot

The `locator()` function allows you to interact with the plot using the mouse. Typing `locator(1)` shifts you to the graphics window and waits for one click of the left mouse button. When you click, it will return the corresponding coordinates.

You can use `locator()` inside other graphics functions to position graphical elements exactly where you want them. Recreate the birth-weight plot,

```
with(births, plot(gestwks, bweight, col = c("blue", "red")[sex]))
```



and then add the legend where you wish it to appear by typing

```
legend(
  locator(1),
  pch = 1,
  legend = c("Boys", "Girls"),
  col = c("blue", "red")
)
```

The `identify()` function allows you to find out which records in the data correspond to points on the graph. Try

```
with(births, identify(gestwks, bweight))
```

When you click the left mouse button, a label will appear on the graph identifying the row number of the nearest point in the data frame `births`. If there is no point nearby, R will print a warning message on the console instead. To end the interaction with the graphics window, right click the mouse: the `identify` function returns a vector of identified points.

- Use `identify()` to find which records correspond to the smallest and largest number of gestational weeks and view the corresponding records:

```
with(births, births[identify(gestwks, bweight), ])
```


Chapter 6

Analysis of hazard rates, their ratios and differences and binary regression

This exercise is *very* prescriptive, so you should make an effort to really understand everything you type into R. Consult the relevant slides of the lecture on *Poisson and Binary regression ...*

6.1 Hand calculations for a single rate

Let λ be the true **hazard rate** or theoretical incidence rate of a given outcome event. Its estimator is the empirical **incidence rate** $\hat{\lambda} = D/Y = \text{no. cases/person-years}$. Recall that the standard error of the empirical rate is $SE(\hat{\lambda}) = \hat{\lambda}/\sqrt{D}$.

The simplest approximate 95% confidence interval (CI) for λ is given by

$$\hat{\lambda} \pm 1.96 \times SE(\hat{\lambda})$$

- Suppose 15 outcome events are observed during 5532 person-years in a given study cohort. Let's use R as a simple desk calculator to estimate the underlying hazard rate λ (in 1000 person-years; therefore 5.532) and to get the first version of an approximate confidence interval:

```
library(Epi)
options(digits = 4) # to cut down decimal points in the output

D <- 15
Y <- 5.532 # thousands of years!
rate <- D / Y
SE.rate <- rate / sqrt(D)
c(rate, SE.rate, rate + c(-1.96, 1.96) * SE.rate)
```

6.2 Poisson model for a single rate with logarithmic link

You are able to estimate the hazard rate λ and compute its CI with a **Poisson regression model**, as described in the relevant slides in the lecture handout.

Poisson regression is a **generalized linear model** in which the **family**, *i.e.* the distribution of the response variable, is assumed to be the Poisson distribution. The most commonly applied **link function** in Poisson regression is the natural logarithm; log for short.

- A family object `poisreg`, a modified version of the original `poisson` family object, is available in package `Epi`. When using this, the response is defined as a *matrix* of two columns: numbers of cases D and person-years Y , these being combined into a matrix by `cbind(D,Y)`. No specification of `offset` is needed.

```
mreg <- glm(cbind(D, Y) ~ 1, family = poisreg(link = log))
ci.exp(mreg)
```

- If you want confidence interval for log rate

```
mreg <- glm(cbind(D, Y) ~ 1, family = poisreg(link = log))
ci.lin(mreg)[, c(1, 5, 6)]
```

In this course we endorse the use of family `poisreg` because of its advantages in more general settings.

6.3 Poisson model for a single rate with identity link

The approach leaning on having the number of cases D as the response and $\log(Y)$ as an offset, is limited only to models with log link. A major advantage of the `poisreg` family is that it allows a straightforward use of the *identity* link, too. With this link the response variable is the same, but the parameters to be directly estimated are now the rates itself and their differences, not the log-rates and their differences as with the log link.

- Fit a Poisson model with identity link to our simple data, and use `ci.lin()` to produce the estimate and the confidence interval for the hazard rate from this model:

```
mid <- glm(cbind(D, Y) ~ 1, family = poisreg(link = "identity"))
ci.lin(mid)
ci.lin(mid)[, c(1, 5, 6)]
```

How is the coefficient of this model interpreted? Verify that you actually get the same rate estimate and CI as in section 1.6.1, item 1.

6.4 Poisson model assuming the same rate for several periods

Now, suppose the events and person years are collected over three distinct periods.

- Read in the data and compute period-specific rates

```
Dx <- c(3, 7, 5)
Yx <- c(1.412, 2.783, 1.337)
Px <- 1:3
rates <- Dx / Yx
rates
```

- Using these data, fit the same model with log link as in section 1.6.2, assuming a common single hazard λ for the separate periods. Compare the result from the previous ones

```
m3 <- glm(cbind(Dx, Yx) ~ 1, family = poisreg(link = log))
ci.exp(m3)
```

- Now test whether the rates are the same in the three periods: Try to fit a model with the period as a factor in the model:

```
mp <- glm(cbind(Dx, Yx) ~ factor(Px), family = poisreg(link = log))
ci.exp(mp)
```

Compare the goodness-of-fit of the two models using `anova()` with the argument `test="Chisq"`:

```
anova(m3, mp, test = "Chisq")
```

Compare the test statistic to the deviance of the model `mp`. – What is the deviance indicating?

6.5 Analysis of rate ratio

We now switch to comparison of two rates λ_1 and λ_0 , i.e. the hazard in an exposed group vs. that in an unexposed one.

Consider first estimation of the **hazard ratio** or the underlying *true* rate ratio $\rho = \lambda_1/\lambda_0$ between the groups. Suppose we have pertinent empirical data (cases and person-times) from both groups, (D_1, Y_1) and (D_0, Y_0) . The point estimate of ρ is the empirical **incidence rate ratio**

$$\hat{\rho} = RR = \frac{\hat{\lambda}_1}{\hat{\lambda}_0} = \frac{D_1/Y_1}{D_0/Y_0}$$

Suppose you have 15 events during 5532 person-years in an unexposed group and 28 events during 4783 person-years in an exposed group:

- Calculate the incidence rates in the two groups, their ratio, and the CI of the true hazard ratio ρ by direct application of the above formulae:

```
D0 <- 15
D1 <- 28
Y0 <- 5.532
Y1 <- 4.783
```

- Now achieve this using a Poisson model. For that we first combine the group-specific numbers into pertinent vectors and specify a factor to represent the contrast between the exposed and the unexposed group

```
D <- c(D0, D1)
Y <- c(Y0, Y1)
expos <- 0:1
mm <- glm(cbind(D, Y) ~ factor(expos), family = poisreg(link = log))
```

What do the parameters mean in this model?

- You can extract the estimation results for exponentiated parameters in two ways, as before:

```
ci.exp(mm)
ci.lin(mm, Exp = TRUE)[, 5:7]
```

6.6 Analysis of rate difference

For the **hazard difference** $\delta = \lambda_1 - \lambda_0$, the natural estimator is the **incidence rate difference**

$$\hat{\delta} = \hat{\lambda}_1 - \hat{\lambda}_0 = D_1/Y_1 - D_0/Y_0 = RD.$$

Its variance is just the sum of the variances of the two rates

$$var(RD) = var(\hat{\lambda}_1) + var(\hat{\lambda}_0) = D_1/Y_1^2 + D_0/Y_0^2$$

- Use this formula to compute the point estimate of the rate difference λ and a 95% confidence interval for it:

```
R0 <- D0 / Y0
R1 <- D1 / Y1
RD <- diff(D / Y)
SED <- sqrt(sum(D / Y^2))
c(R1, R0, RD, SED, RD + c(-1, 1) * 1.96 * SED)
```

- Verify that this is the confidence interval you get when you fit an additive model (obtained by identity link) with exposure as a factor:

```
ma <- glm(cbind(D, Y) ~ factor(expos),
  family = poisreg(link = identity)
)
ci.lin(ma)[, c(1, 5, 6)]
```

6.7 Binary regression

Explore the factors associated with risk of low birth weight in 500 singleton births in a London Hospital. Indicator (lowbw) for birth weight less than 2500 g. Data available from the Epi package. Factors of interest are maternal hypertension (hyp), mother's age at birth over 35 years and sex of the baby.

Load the Epi package and the data set and look at its content

```
library(dplyr)
library(Epi)
data(births)
str(births)
```

- Because all variables are numeric we need first to do a little housekeeping. Two of them are directly converted into factors, and categorical versions are created of two continuous variables by function cut().

```
births$hyp <- factor(births$hyp, labels = c("normal", "hyper"))
births$sex <- factor(births$sex, labels = c("M", "F"))
births$gest4 <- cut(births$gestwks,
  breaks = c(20, 35, 37, 39, 45), right = FALSE
)
births$maged <- ifelse(births$matage < 35, 0, 1)
```

- Cross tabulate (dplyr) counts of children by hypertension and low birth weight. calculate (mutate) proportions of low birth weight children by hypertension.

```
births %>%
  count(hyp, lowbw) %>%
  group_by(hyp) %>% # now required with changes to dplyr::count()
  mutate(prop = prop.table(n))
```

- Estimate relative risk of low birth weight for mothers with hypertension compared to those without using binary regression.

```
m <- glm(lowbw ~ hyp, family = binomial(link = log), data = births)
ci.exp(m)
```

- Adjust relative risk of low birth and hypertension with the sex of children

```
m <- glm(lowbw ~ sex + hyp, family = binomial(link = log), data = births)
ci.exp(m)
```

- Adjust relative risk of low birth and hypertension with the sex of children and mother being over 35 years.

```
m <- glm(lowbw ~ maged + sex + hyp, family = binomial(link = log), data = births)
ci.exp(m)
```

6.8 Optional/Homework: Hand calculations and calculations using matrix tools

NB. This subsection requires some familiarity with matrix algebra. Do this only after you have done the other exercises of this session.

First some basic hand calculations.

- Suppose 15 outcome events are observed during 5532 person-years in a given study cohort. Let's use R as a simple desk calculator to estimate the underlying hazard rate λ (in 1000 person-years; therefore 5.532) and to get the first version of an approximate confidence interval:

```
library(Epi)
options(digits = 4) # to cut down decimal points in the output
```

```
D <- 15
Y <- 5.532 # thousands of years!
rate <- D / Y
SE.rate <- rate / sqrt(D)
c(rate, SE.rate, rate + c(-1.96, 1.96) * SE.rate)
```

- Compute now the approximate confidence interval using the method based on log-transformation and compare the result with that in the previous item.

```
SE.logr <- 1 / sqrt(D)
EF <- exp(1.96 * SE.logr)
c(log(rate), SE.logr)
c(rate, EF, rate / EF, rate * EF)
```

- Calculate the incidence rates in the two groups, their ratio, and the CI of the true hazard ratio ρ by direct application of the above formulae:

```
D0 <- 15
D1 <- 28
Y0 <- 5.532
Y1 <- 4.783
R1 <- D1 / Y1
R0 <- D0 / Y0
RR <- R1 / R0
SE.lrr <- sqrt(1 / D0 + 1 / D1)
EF <- exp(1.96 * SE.lrr)
c(R1, R0, RR, RR / EF, RR * EF)
```

- Explore the function `ci.mat()`, which lets you use matrix multiplication (operator `'%*%'` in R) to produce a confidence interval from an estimate and its standard error (or CIs from whole columns of estimates and SEs):

```
ci.mat
ci.mat()
```

As you see, this function returns a 2×3 matrix (2 rows, 3 columns) containing familiar numbers.

- When you combine the single rate and its standard error into a row vector of length 2, i.e. a 1×2 matrix, and multiply this by the 2×3 matrix above, the computation returns a 1×3 matrix containing the point estimate and the confidence limit.

Apply this method to the single rate calculations in 1.6.1, first creating the 1×2 matrix and then performing the matrix multiplication.

```
rateandSE <- c(rate, SE.rate)
rateandSE
rateandSE %*% ci.mat()
```

- When the confidence interval is based on the log-rate and its standard error, the result is obtained by appropriate application of the `exp`-function on the pertinent matrix product

```
lograndSE <- c(log(rate), SE.logr)
lograndSE
exp(lograndSE %*% ci.mat())
```

- For computing the rate ratio and its CI as in 1.6.5, matrix multiplication with `ci.mat()` should give the same result as there:

```
exp(c(log(RR), SE.lrr) %*% ci.mat())
```

- The main argument in function `ci.mat()` is `alpha`, which sets the confidence level: $1 - \alpha$. The default value is `alpha = 0.05`, corresponding to the level $1 - 0.05 = 95\%$. If you wish to get the confidence interval for the rate ratio at the 90% level ($= 1 - 0.1$), for instance, you may proceed as follows:

```
ci.mat(alpha = 0.1)
exp(c(log(RR), SE.lrr) %*% ci.mat(alpha = 0.1))
```

- Now achieve this using a Poisson model. For that we first combine the group-specific numbers into pertinent vectors and specify a factor to represent the contrast between the exposed and the unexposed group

```
D <- c(D0, D1)
Y <- c(Y0, Y1)
expos <- 0:1
```

- Look again to the model used to analyse the rate ratio in. Often one would like to get simultaneously both the rates and the ratio between them. This can be achieved in one go using the *contrast matrix* argument `ctr.mat` to `ci.lin()` or `ci.exp()`. Try:

```
CM <- rbind(c(1, 0), c(1, 1), c(0, 1))
rownames(CM) <- c("rate 0", "rate 1", "RR 1 vs. 0")
CM
mm <- glm(D ~ factor(expos),
  family = poisson(link = log), offset = log(Y)
)
ci.exp(mm, ctr.mat = CM)
```

- Use the same machinery to the additive model to get the rates and the rate-difference in one go. Note that the annotation of the resulting estimates are via the column-names of the contrast matrix.

```
rownames(CM) <- c("rate 0", "rate 1", "RD 1 vs. 0")
ma <- glm(cbind(D, Y) ~ factor(expos),
  family = poisreg(link = identity)
)
ci.lin(ma, ctr.mat = CM)[, c(1, 5, 6)]
```

Chapter 7

Estimation of effects: simple and more complex

This exercise deals with analysis of metric and binary response variables. We start with simple estimation of effects of a binary, categorical or a numeric explanatory variable, the explanatory or exposure variable of interest. Then evaluation of potential modification and/or confounding by other variables is considered by stratification by and adjustment/control for these variables. For such tasks we utilize functions `lm()` and `glm()` which can be used for more general linear and generalized linear models. Finally, more complex spline modelling for the effect of a numeric exposure variable is illustrated.

7.1 Response and explanatory variables

Identifying the *response* or *outcome variable* correctly is the key to analysis. The main types are:

- Metric or continuous (a measurement with units).
- Binary (“yes” vs. “no”, coded 1/0), or proportion.
- Failure in person-time, or incidence rate.

All these response variable are numeric.

Variables on which the response may depend are called *explanatory variables* or *regressors*. They can be categorical factors or numeric variables. A further important aspect of explanatory variables is the role they will play in the analysis.

- Primary role: exposure.
- Secondary role: confounder and/or effect-measure modifier.

The word **effect** is used here as a general term referring to ways of contrasting or comparing the expected values of the response variable at different levels of an explanatory variable. The main comparative measures or effect measures are:

- Differences in means for a metric response.
- Ratios of odds for a binary response.
- Ratios of rates for a failure or count response.

Other kinds of *contrasts* between exposure groups include (a) ratios of geometric means for positive-valued metric outcomes, (b) differences and ratios between proportions (risk difference and risk ratio), and (c) differences between incidence or mortality rates.

Note that in spite of using the causally loaded word *effect*, we treat *outcome regression* modelling here primarily with descriptive or predictive aims in mind. Traditionally, these types of models have also been used to estimate *causal effects* of exposure variables from the pertinent regression coefficients. More serious causal analysis is introduced in the lecture and practical on Tuesday morning, and modern approaches to estimate causal effects will be considered on Thursday afternoon.

7.2 Data set *births*

We shall use the `births` data to illustrate different aspects in estimating effects of various exposures on a metric response variable `bweight` = birth weight, recorded in grams.

1. Load the packages needed in this exercise and the data set, and look at its content

```
library(Epi)
library(mgcv)
data(births)
str(births)

## 'data.frame':    500 obs. of  8 variables:
## $ id      : num  1 2 3 4 5 6 7 8 9 10 ...
## $ bweight: num  2974 3270 2620 3751 3200 ...
## $ lowbw   : num  0 0 0 0 0 0 0 0 0 0 ...
## $ gestwks: num  38.5 NA 38.2 39.8 38.9 ...
## $ preterm: num  0 NA 0 0 0 0 0 0 0 0 ...
## $ matage  : num  34 30 35 31 33 33 29 37 36 39 ...
## $ hyp     : num  0 0 0 0 1 0 0 0 0 0 ...
## $ sex     : num  2 1 2 1 1 2 2 1 2 1 ...
```

2. We perform similar housekeeping tasks as in a previous exercise.

```
births$hyp <- factor(births$hyp, labels = c("normal", "hyper"))
births$sex <- factor(births$sex, labels = c("M", "F"))
births$maged <- cut(births$matage, breaks = c(22, 35, 44), right = FALSE)
births$gest4 <- cut(births$gestwks,
  breaks = c(20, 35, 37, 39, 45), right = FALSE)
```

3. Have a look at univariate summaries of the different variables in the data; especially the location and dispersion of the distribution of `bweight`.

```
summary(births)

##           id           bweight           lowbw           gestwks           preterm
## Min.      : 1.0      Min.      : 628      Min.      :0.00      Min.      :24.69      Min.      :0.0000
## 1st Qu.:125.8      1st Qu.:2862      1st Qu.:0.00      1st Qu.:37.94      1st Qu.:0.0000
## Median :250.5      Median :3188      Median :0.00      Median :39.12      Median :0.0000
## Mean    :250.5      Mean    :3137      Mean    :0.12      Mean    :38.72      Mean    :0.1286
## 3rd Qu.:375.2      3rd Qu.:3551      3rd Qu.:0.00      3rd Qu.:40.09      3rd Qu.:0.0000
## Max.    :500.0      Max.    :4553      Max.    :1.00      Max.    :43.16      Max.    :1.0000
##                                     NAs      :10      NAs      :10
##           matage           hyp           sex           maged           gest4
## Min.      :23.00      normal:428      M:264      [22,35):270      [20,35): 31
## 1st Qu.:31.00      hyper : 72      F:236      [35,44):230      [35,37): 32
## Median :34.00                                     [37,39):167
## Mean    :34.03                                     [39,45):260
## 3rd Qu.:37.00                                     NAs      : 10
## Max.    :43.00
##

with(births, sd(bweight))

## [1] 637.4515
```

7.3 Simple estimation with `lm()` and `glm()`

We are ready to analyze the effect of maternal hypertension `hyp` on `bweight`. A binary explanatory variable, like `hyp`, leads to an elementary two-group comparison of group means for a metric response.

1. Comparison of two groups is commonly done by the conventional t -test and the associated confidence interval.

```
with(births, t.test(bweight ~ hyp, var.equal = TRUE))
```

```
##
## Two Sample t-test
##
## data: bweight by hyp
## t = 5.455, df = 498, p-value = 7.729e-08
## alternative hypothesis: true difference in means between group normal and group hyper is not equal to 0
## 95 percent confidence interval:
## 275.5707 585.8210
## sample estimates:
## mean in group normal mean in group hyper
## 3198.904 2768.208
```

The P -value refers to the test of the null hypothesis that there is no effect of `hyp` on birth weight (somewhat implausible null hypothesis in itself!). However, `t.test()` does not provide the point estimate for the effect of `hyp`; only the test result and a confidence interval. – The estimated effect of `hyp` on birth weight, measured as a difference in means between hypertensive and normotensive mothers, is $3199 - 2768 = 431$ g.

2. The same task can easily be performed by `lm()` or by `glm()`. The main argument in both is the *model formula*, the left hand side being the response variable and the right hand side after `~` defines the explanatory variables and their joint effects on the response. Here the only explanatory variable is the binary factor `hyp`. With `glm()` one specifies the `family`, i.e. the assumed distribution of the response variable. However, in case you use `lm()`, this argument is not needed, because `lm()` fits only models for metric responses assuming Gaussian distribution.

```
m1 <- glm(bweight ~ hyp, family = gaussian, data = births)
summary(m1)
```

```
##
## Call:
## glm(formula = bweight ~ hyp, family = gaussian, data = births)
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept)  3198.90      29.96 106.768 < 2e-16 ***
## hyphyper    -430.70      78.95  -5.455 7.73e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 384203.2)
##
## Null deviance: 202765853 on 499 degrees of freedom
## Residual deviance: 191333183 on 498 degrees of freedom
## AIC: 7852.4
##
## Number of Fisher Scoring iterations: 2
```

3. Note the amount of output that `summary()` method produces. The point estimate plus confidence limits can, though, be concisely obtained by function `ci.lin()` found in `Epi` package.

```
round(ci.lin(m1)[, c(1, 5, 6)], 1)
```

```
##             Estimate  2.5% 97.5%
## (Intercept)  3198.9 3140.2 3257.6
## hyphyper    -430.7 -585.4 -275.9
```

7.4 Stratified effects, and interaction or effect-measure modification

We shall now examine whether and to what extent the *effect* of **hyp** on **bweight**, i.e. the mean difference between hypertensive and normotensive mothers, varies by **sex** without assigning causal interpretation to the estimated contrasts.

1. The following *interaction plot* shows how the mean **bweight** depends jointly on **hyp** and **gest4**

```
par(mfrow = c(1, 1))
with(births, interaction.plot(sex, hyp, bweight))
```

At face value it appears that the mean difference in **bweight** between hypertensive and normotensive mothers is somewhat bigger in boys than in girls.

2. Let us get numerical values for the mean differences in the two levels of **sex**. Stratified estimation of effects can be done by `lm()` as follows:

```
m3 <- lm(bweight ~ sex / hyp, data = births)
round(ci.lin(m3)[, c(1, 5, 6)], 1)
```

```
##              Estimate    2.5%   97.5%
## (Intercept)    3310.7  3230.1  3391.4
## sexF          -231.2  -347.2  -115.3
## sexM:hyphyper -496.4  -696.1  -296.6
## sexF:hyphyper -379.8  -617.4  -142.2
```

The estimated effects of **hyp** in the two strata defined by **sex** thus are -496 g in boys and -380 g among girls. The error margins of the two estimates are quite wide, though.

3. An equivalent model with an explicit *product term* or *interaction term* between **sex** and **hyp** is fitted as follows:

```
m3I <- lm(bweight ~ sex + hyp + sex:hyp, data = births)
round(ci.lin(m3I)[, c(1, 4, 5, 6)], 2)
```

```
##              Estimate      P    2.5%   97.5%
## (Intercept)    3310.75  0.00  3230.14  3391.35
## sexF          -231.25  0.00  -347.15  -115.35
## hyphyper       -496.35  0.00  -696.07  -296.63
## sexF:hyphyper   116.58  0.46  -193.80   426.96
```

From this output you would find a familiar estimate -231 g for girls vs. boys among normotensive mothers and the estimate -496 g contrasting hypertensive and normotensive mothers in the reference class of **sex**, i.e. among boys. The remaining coefficient is the estimate of the interaction effect such that $116.6 = -379.8 - (-496.4)$ g describes the contrast in the effect of **hyp** on **bweight** between girls and boys.

The *P*-value 0.46 as well as the wide confidence interval about zero of this interaction parameter suggest good compatibility of the data with the null hypothesis of no interaction between **hyp** and **sex**. Thus, there is insufficient evidence against the possibility of *effect(-measure) modification* by **sex** on the effect of **hyp**. On the other hand, this test is not very sensitive given the small sample size. Thus, in spite of obtaining a “non-significant” result, the possibility of a real effect-measure modification cannot be ignored based on these data only.

7.5 Controlling or adjusting for the effect of hyp for sex

The estimated effects of **hyp**: -496 in boys and -380 in girls, look quite similar (and the *P*-value against no interaction was quite large, too). Therefore, we may now proceed to estimate the overall effect of **hyp** *controlling for* – or *adjusting for* – **sex**.

1. Adjustment is done by adding **sex** to the model formula:

```
m4 <- lm(bweight ~ sex + hyp, data = births)
ci.lin(m4)[, c(1, 5, 6)]
```

```
##           Estimate      2.5%      97.5%
## (Intercept) 3302.8845 3225.0823 3380.6867
## sexF        -214.9931 -322.4614 -107.5249
## hyphyper    -448.0817 -600.8911 -295.2723
```

The estimated effect of `hyp` on `bweight` adjusted for `sex` is thus -448 g, which is a weighted average of the sex-specific estimates. It is slightly different from the unadjusted estimate -431 g, indicating that there was no essential confounding by `sex` in the simple comparison of means. Note also, that the model being fitted makes the assumption that the estimated effect is the same for boys and girls.

Many people go straight ahead and control for variables which are likely to confound the effect of exposure without bothering to stratify first, but often it is useful to examine the possibility of effect-measure modification before that.

7.6 Numeric exposure, simple linear regression and checking assumptions

If we wished to study the effect of gestation time on the baby's birth weight then `gestwks` is a numeric exposure variable.

1. Assuming that the relationship of the response with `gestwks` is roughly linear (for a continuous response), % or log-linear (for a binary or failure rate response) we can estimate the linear effect of `gestwks` with `lm()` as follows:

```
m5 <- lm(bweight ~ gestwks, data = births)
ci.lin(m5)[, c(1, 5, 6)]
```

```
##           Estimate      2.5%      97.5%
## (Intercept) -4489.1398 -5157.2891 -3820.9905
## gestwks      196.9726  179.7482   214.1971
```

We have fitted a simple linear regression model and obtained estimates of the two regression coefficient: `intercept` and `slope`. The linear effect of `gestwks` is thus estimated by the slope coefficient, which is 197 g per each additional week of gestation.

At this stage it will be best to make some visual check concerning our model assumptions using `plot()`. In particular, when the main argument for the *generic function* `plot()` is a fitted `lm` object, it will provide you some common diagnostic graphs.

2. To check whether `bweight` goes up linearly with `gestwks` try

```
with(births, plot(gestwks, bweight))
abline(m5)
```

3. Moreover, take a look at the basic diagnostic plots for the fitted model.

```
par(mfrow = c(2, 2))
plot(m5)
```

What can you say about the agreement with data of the assumptions of the simple linear regression model, like linearity of the systematic dependence, homoskedasticity and normality of the error terms?

7.7 Penalized spline model

We shall now continue the analysis such that the apparently curved effect of `gestwks` is modelled by a *penalized spline*, based on the recommendations of Martyn in his lecture today.

You cannot fit a penalized spline model with `lm()` or `glm()`. Instead, function `gam()` in package `mgcv` can be used for this purpose. Make sure that you have loaded this package.

1. When calling `gam()`, the model formula contains expression '`s(X)`' for any explanatory variable `X`, for which you wish to fit a smooth function

```
mPs <- mgcv::gam(bweight ~ s(gestwks), data = births)
summary(mPs)

##
## Family: gaussian
## Link function: identity
##
## Formula:
## bweight ~ s(gestwks)
##
## Parametric coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  3138.01      20.11     156   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##              edf Ref.df      F p-value
## s(gestwks)  3.321  4.189 124.7   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.516   Deviance explained = 51.9%
## GCV = 1.9995e+05   Scale est. = 1.9819e+05   n = 490
```

From the output given by `summary()` you find that the estimated intercept is equal to the overall mean birth weight in the data. The estimated residual variance is given by `Scale est.` or from subobject `sig2` of the fitted `gam` object. Taking square root you will obtain the estimated residual standard deviation: 445.2 g.

```
mPs$sig2

## [1] 198186
sqrt(mPs$sig2)

## [1] 445.1808
```

The degrees of freedom in this model are not computed as simply as in previous models, and they typically are not integer-valued. However, the fitted spline seems to consume only a little more degrees of freedom as an 3rd degree polynomial model would take.

2. A graphical presentation of the fitted curve together with the confidence and prediction intervals is more informative. Let us first write a short function script to facilitate the task. We utilize function `matshade()` in `Epi`, which creates shaded areas, and function `matlines()` which draws lines joining the pertinent end points over the x -values for which the predictions are computed.

```
plotFitPredInt <- function(xval, fit, pred, ...) {
  matshade(xval, fit, lwd = 2, alpha = 0.2)
  matshade(xval, pred, lwd = 2, alpha = 0.2)
  matlines(xval, fit, lty = 1, lwd = c(3, 2, 2), col = c("black", "blue", "blue"))
  matlines(xval, pred, lty = 1, lwd = c(3, 2, 2), col = c("black", "brown", "brown"))
}
```

3. Finally, create a vector of x -values and compute the fitted/predicted values as well as the interval limits at these points from the fitted model object utilizing function `predict()`. This function creates a matrix of three columns: (1) fitted/predicted values, (2) lower limits, (3) upper limits and make the graph:

```

nd <- data.frame(gestwks = seq(24, 45, by = 0.25))
pr.Ps <- predict(mPs, newdata = nd, se.fit = TRUE)
str(pr.Ps) # with se.fit=TRUE, only two columns: fitted value and its SE

## List of 2
## $ fit : num [1:85(1d)] 350 385 421 456 491 ...
## ..- attr(*, "dimnames")=List of 1
## .. ..$ : chr [1:85] "1" "2" "3" "4" ...
## $ se.fit: num [1:85(1d)] 324 309 293 278 264 ...
## ..- attr(*, "dimnames")=List of 1
## .. ..$ : chr [1:85] "1" "2" "3" "4" ...

fit.Ps <- cbind(
  pr.Ps$fit,
  pr.Ps$fit - 2 * pr.Ps$se.fit,
  pr.Ps$fit + 2 * pr.Ps$se.fit
)
pred.Ps <- cbind(
  pr.Ps$fit, # must add residual variance to se.fit^2
  pr.Ps$fit - 2 * sqrt(pr.Ps$se.fit^2 + mPs$sig2),
  pr.Ps$fit + 2 * sqrt(pr.Ps$se.fit^2 + mPs$sig2)
)
par(mfrow = c(1, 1))
with(births, plot(bweight ~ gestwks,
  xlim = c(24, 45),
  cex.axis = 1.5, cex.lab = 1.5
))
plotFitPredInt(nd$gestwks, fit.Ps, pred.Ps)

```

Compare this with the graph on slide 20 of the lecture we had. Are you happy with the end result?

7.8 Analysis of binary outcomes

Instead of investigating the distribution and determinants of birth weight as such, it is common in perinatal epidemiology to consider occurrence of low birth weight; whether birth weight is < 2.5 kg or not. Variable `lowbw` with values 1 and 0 in the `births` data represents that dichotomy. Some analyses on `lowbw` were already conducted in a previous practical. Here we illustrate further aspects of effect estimation and modelling binary outcome.

1. We start with simple tabulation of the prevalence of `lowbw` by maternal hypertension

```

stat.table(
  index = list(hyp, lowbw),
  contents = list(count(), percent(lowbw)),
  margins = TRUE, data = births
)

## -----
##          -----lowbw-----
## hyp           0           1     Total
## -----
## normal        388         40      428
##               90.7        9.3    100.0
##
## hyper         52          20       72
##               72.2       27.8    100.0
##
##
##

```

```
## Total      440      60      500
##           88.0     12.0    100.0
## -----
```

It seems that the prevalence for hypertensive mothers is about 18 percent points higher, or about three times as high as that for normotensive mothers

2. The three comparative measures of prevalences can be estimated by `glm()` with different link functions:

```
binRD <- glm(lowbw ~ hyp, family = binomial(link = "identity"), data = births)
round(ci.lin(binRD)[, c(1, 2, 5:6)], 3)
```

```
##           Estimate StdErr  2.5% 97.5%
## (Intercept)   0.093  0.014 0.066 0.121
## hyphyper      0.184  0.055 0.077 0.291
```

```
binRR <- glm(lowbw ~ hyp, family = binomial(link = "log"), data = births)
round(ci.lin(binRR, Exp = TRUE)[, c(1, 2, 5:7)], 3)
```

```
##           Estimate StdErr exp(Est.)  2.5% 97.5%
## (Intercept)  -2.370  0.151    0.093 0.070 0.126
## hyphyper      1.089  0.242    2.972 1.848 4.780
```

```
binOR <- glm(lowbw ~ hyp, family = binomial(link = "logit"), data = births)
round(ci.lin(binOR, Exp = TRUE)[, c(1, 2, 5:7)], 3)
```

```
##           Estimate StdErr exp(Est.)  2.5% 97.5%
## (Intercept)  -2.272  0.166    0.103 0.074 0.143
## hyphyper      1.317  0.311    3.731 2.027 6.865
```

Check that these results were quite compatible with the “about” estimates given in the previous item. How well is the odds ratio approximating the risk ratio here?

3. The prevalence of low birth weight is expected to be inversely related to gestational age (weeks), as is evident from simple tabulation

```
stat.table(
  index = list(gest4, lowbw),
  contents = list(count(), percent(lowbw)),
  margins = TRUE, data = births
)
```

```
## -----
##           -----lowbw-----
## gest4      0      1      Total
## -----
## [20,35)      6      25      31
##           19.4    80.6   100.0
##
## [35,37)     19      13      32
##           59.4    40.6   100.0
##
## [37,39)    149      18     167
##           89.2    10.8   100.0
##
## [39,45)    257       3     260
##           98.8     1.2   100.0
##
## Total      440      60      500
##           88.0     12.0   100.0
## -----
```

4. Let's jump right away to spline modelling of this relationship

```
binm1 <- mgcv::gam(lowbw ~ s(gestwks), family = binomial(link = "logit"), data = births)
summary(binm1)
```

```
##
## Family: binomial
## Link function: logit
##
## Formula:
## lowbw ~ s(gestwks)
##
## Parametric coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -2.8665      0.2364  -12.12   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##              edf Ref.df Chi.sq p-value
## s(gestwks)  1.01  1.021  68.86  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.425   Deviance explained = 42.9%
## UBRE = -0.57194   Scale est. = 1          n = 490
```

```
plot(binm1)
```

Inspect the figure. Would you agree, that the logit of the prevalence of outcome is almost linearly dependent on `gestwks`?

5. Encouraged by the result of the previous item, we continue the analysis with `glm()` and assuming logit-linearity

```
binm2 <- glm(lowbw ~ I(gestwks - 40), family = binomial(link = "logit"), data = births)
round(ci.lin(binm2, Exp = TRUE)[, c(1, 2, 5:7)], 3)
```

```
##              Estimate StdErr exp(Est.)  2.5% 97.5%
## (Intercept)    -4.011  0.338      0.018 0.009 0.035
## I(gestwks - 40)  -0.896  0.108      0.408 0.330 0.505
```

Inspect the results. How do you interpret the estimated coefficients and their exponentiated values?

6. Instead of fitted logits, it can be more informative to plot the fitted prevalences against `gestwks`, in which we utilize the previously created data frame `nd`

```
predm2 <- predict(binm2, newdata = nd, type = "response")
plot(nd$gestwks, predm2, type = "l")
```

The curve seems to cover practically the whole range of the outcome probability scale with a relatively steep slope between 33 to 37 weeks.

Chapter 8

Poisson regression & analysis of curved effects

This exercise deals with modelling incidence rates using Poisson regression. Our special interest is in estimating and reporting curved effects of continuous explanatory variables on the hazard rate, i.e. the theoretical incidence rate

We analyse the `testisDK` data found in the `Epi` package. It contains the numbers of cases of testis cancer and mid-year populations (person-years) in 1-year age groups in Denmark during 1943-96. In this analysis age and calendar time are first treated as categorical but finally, a penalized spline model is fitted.

8.1 Testis cancer: Data input and housekeeping

1. Load the packages and the data set, and inspect its structure:

```
library(Epi)
library(mgcv)

## Loading required package: nlme
## This is mgcv 1.9-4. For overview type '?mgcv'.
data(testisDK)
str(testisDK)

## 'data.frame':    4860 obs. of  4 variables:
## $ A: num  0 1 2 3 4 5 6 7 8 9 ...
## $ P: num  1943 1943 1943 1943 1943 ...
## $ D: num  1 1 0 1 0 0 0 0 0 0 ...
## $ Y: num  39650 36943 34588 33267 32614 ...

summary(testisDK)

##           A           P           D           Y
## Min.      : 0.0      Min.    :1943      Min.    : 0.000      Min.    : 471.7
## 1st Qu.:22.0      1st Qu.:1956      1st Qu.: 0.000      1st Qu.:18482.2
## Median :44.5      Median :1970      Median : 1.000      Median :28636.0
## Mean     :44.5      Mean     :1970      Mean     : 1.812      Mean     :26239.8
## 3rd Qu.:67.0      3rd Qu.:1983      3rd Qu.: 2.000      3rd Qu.:36785.5
## Max.     :89.0      Max.     :1996      Max.     :17.000      Max.     :47226.8

head(testisDK)

##      A      P D      Y
```

```
## 1 0 1943 1 39649.50
## 2 1 1943 1 36942.83
## 3 2 1943 0 34588.33
## 4 3 1943 1 33267.00
## 5 4 1943 0 32614.00
## 6 5 1943 0 32020.33
```

2. There are nearly 5000 observations from 90 one-year age groups and 54 calendar years. To get a clearer picture of what's going on, we do some housekeeping. The age range will be limited to 15-79 years, and age and period are both categorized into 5-year intervals – following to the time-honoured practice in epidemiology.

```
tdk <- subset(testisDK, A > 14 & A < 80)
tdk$Age <- cut(tdk$A, br = 5 * (3:16), include.lowest = TRUE, right = FALSE)
nAge <- length(levels(tdk$Age))
tdk$Per <- cut(tdk$P,
  br = seq(1943, 1998, by = 5),
  include.lowest = TRUE, right = FALSE)
nPer <- length(levels(tdk$Per))
```

8.2 Some descriptive analysis

Computation and tabulation of incidence rates

1. Tabulate numbers of cases and person-years, and compute the incidence rates (per 100,000 y) in each 5 y × 5 y cell using `stat.table()`. Take a look at the structure of the thus created object

```
tab <- stat.table(
  index = list(Age, Per),
  contents = list(
    D = sum(D),
    Y = sum(Y / 1000),
    rate = ratio(D, Y, 10^5)
  ),
  margins = TRUE,
  data = tdk
)
str(tab)
```

```
## 'stat.table' num [1:3, 1:14, 1:12] 10 773.81 1.29 30 813.02 ...
## - attr(*, "dimnames")=List of 3
## ..$ contents: Named chr [1:3] "D" "Y" "rate"
## ..$- attr(*, "names")= chr [1:3] "D" "Y" "rate"
## ..$ Age : chr [1:14] "[15,20)" "[20,25)" "[25,30)" "[30,35)" ...
## ..$ Per : chr [1:12] "[1943,1948)" "[1948,1953)" "[1953,1958)" "[1958,1963)" ...
## - attr(*, "table.fun")= chr [1:3] "sum" "sum" "ratio"
```

The table is too wide to be readable as such. A graphical presentation is more informative.

2. From the saved table object `tab` you can plot an age-incidence curve for each period separately, after you have checked the structure of the table, so that you know the relevant dimensions in it. There is a function `rateplot()` in `Epi` that does default plotting of tables of rates (see the help page of `rateplot`)

```
str(tab)

## 'stat.table' num [1:3, 1:14, 1:12] 10 773.81 1.29 30 813.02 ...
## - attr(*, "dimnames")=List of 3
## ..$ contents: Named chr [1:3] "D" "Y" "rate"
## ..$- attr(*, "names")= chr [1:3] "D" "Y" "rate"
## ..$ Age : chr [1:14] "[15,20)" "[20,25)" "[25,30)" "[30,35)" ...
```

```
## ..$ Per      : chr [1:12] "[1943,1948)" "[1948,1953)" "[1953,1958)" "[1958,1963)" ...
## - attr(*, "table.fun")= chr [1:3] "sum" "sum" "ratio"

par(mfrow = c(1, 1))
rateplot(
  rates = tab[3, 1:nAge, 1:nPer], which = "ap", ylim = c(1, 30),
  age = seq(15, 75, 5), per = seq(1943, 1993, 5),
  col = heat.colors(16), ann = TRUE
)
```

What can you conclude about the trend in age-specific incidence rates over calendar time? What about the effect of age; is there any common pattern in the age-incidence curves across the periods?

8.3 Age and period as categorical factors

We shall first fit a Poisson regression model with log link on age and period model in the traditional way, in which both factors are treated as categorical. The model is additive on the log-rate scale. It is useful to scale the person-years to be expressed in 10^5 y.

1. In fitting the model we utilize the `poisreg` family object found in package `Epi`.

```
tdk$Y <- tdk$Y / 100000
mCat <- glm(cbind(D, Y) ~ Age + Per,
  family = poisreg(link = log), data = tdk )
round(ci.exp(mCat), 2)
```

```
##               exp(Est.) 2.5% 97.5%
## (Intercept)      1.47 1.26  1.72
## Age[20,25)       3.13 2.75  3.56
## Age[25,30)       4.90 4.33  5.54
## Age[30,35)       5.50 4.87  6.22
## Age[35,40)       4.78 4.22  5.42
## Age[40,45)       3.66 3.22  4.16
## Age[45,50)       2.60 2.27  2.97
## Age[50,55)       1.94 1.68  2.25
## Age[55,60)       1.47 1.25  1.72
## Age[60,65)       0.98 0.82  1.18
## Age[65,70)       0.92 0.76  1.12
## Age[70,75)       0.90 0.73  1.12
## Age[75,80]       0.86 0.67  1.11
## Per[1948,1953)   1.12 0.96  1.30
## Per[1953,1958)   1.30 1.13  1.50
## Per[1958,1963)   1.53 1.33  1.76
## Per[1963,1968)   1.68 1.47  1.92
## Per[1968,1973)   1.98 1.74  2.25
## Per[1973,1978)   2.33 2.05  2.64
## Per[1978,1983)   2.66 2.35  3.01
## Per[1983,1988)   2.83 2.50  3.20
## Per[1988,1993)   3.08 2.73  3.47
## Per[1993,1998]   3.31 2.93  3.74
```

What do the estimated rate ratios tell about the age and period effects?

2. A graphical inspection of point estimates and confidence intervals can be obtained as follows. In the beginning it is useful to define shorthands for the pertinent mid-age and mid-period values of the different intervals

```
aMid <- seq(17.5, 77.5, by = 5)
pMid <- seq(1945, 1995, by = 5)
par(mfrow = c(1, 2))
```

```
matplot(aMid, rbind(c(1,1,1), ci.exp(mCat)[2:13, ]), type = "o", pch = 16,
  log = "y", cex.lab = 1.5, cex.axis = 1.5, col= c("black", "blue", "blue"),
  xlab = "Age (years)", ylab = "Rate ratio" )
matplot(pMid, rbind(c(1,1,1), ci.exp(mCat)[14:23, ]), type = "o", pch = 16,
  log = "y", cex.lab = 1.5, cex.axis = 1.5, col=c("black", "blue", "blue"),
  xlab = "Calendar year - 1900", ylab = "Rate ratio" )
```

3. In the fitted model the reference category for each factor was the first one. As age is the dominating factor, it may be more informative to remove the intercept from the model. As a consequence the age effects describe fitted rates at the reference level of the period factor. For the latter one could choose the middle period 1968-72 using `Relevel()`.

```
tdk$Per70 <- Relevel(tdk$Per, ref = 6)
mCat2 <- glm(cbind(D, Y) ~ -1 + Age + Per70,
  family = poisreg(link = log), data = tdk )
round(ci.exp(mCat2), 2)
```

##	exp(Est.)	2.5%	97.5%
## Age[15,20)	2.91	2.55	3.33
## Age[20,25)	9.12	8.31	10.01
## Age[25,30)	14.28	13.11	15.55
## Age[30,35)	16.03	14.72	17.46
## Age[35,40)	13.94	12.76	15.23
## Age[40,45)	10.66	9.71	11.71
## Age[45,50)	7.57	6.83	8.39
## Age[50,55)	5.67	5.05	6.36
## Age[55,60)	4.28	3.75	4.88
## Age[60,65)	2.85	2.43	3.35
## Age[65,70)	2.68	2.25	3.19
## Age[70,75)	2.63	2.16	3.20
## Age[75,80]	2.51	1.98	3.18
## Per70[1943,1948)	0.51	0.44	0.58
## Per70[1948,1953)	0.57	0.50	0.64
## Per70[1953,1958)	0.66	0.58	0.74
## Per70[1958,1963)	0.77	0.69	0.87
## Per70[1963,1968)	0.85	0.76	0.95
## Per70[1973,1978)	1.18	1.07	1.30
## Per70[1978,1983)	1.35	1.22	1.48
## Per70[1983,1988)	1.43	1.30	1.57
## Per70[1988,1993)	1.56	1.42	1.70
## Per70[1993,1998]	1.67	1.53	1.84

4. Let us also plot estimates from the latter model, too.

```
par(mfrow = c(1, 2))
matplot(aMid, rbind(c(1,1,1), ci.exp(mCat2)[2:13, ]), type = "o", pch = 16,
  log = "y", cex.lab = 1.5, cex.axis = 1.5, col=c("black", "blue", "blue"),
  xlab = "Age (years)", ylab = "Rate" )
matplot(pMid, rbind(ci.exp(mCat2)[14:18, ], c(1,1,1), ci.exp(mCat2)[19:23, ]),
  type = "o", pch = 16, log = "y", cex.lab = 1.5, cex.axis = 1.5,
  col=c("black", "blue", "blue"),
  xlab = "Calendar year - 1900", ylab = "Rate ratio" )
abline(h = 1, col = "gray")
```

8.4 Generalized additive model with penalized splines

It is obvious that the age effect on the log-rate scale is highly non-linear. Yet, it is less clear whether the true period effect deviates from linearity. Nevertheless, there are good reasons to try fitting smooth continuous functions for both time scales.

1. As the next task we fit a generalized additive model for the log-rate on continuous age and period applying penalized splines with default settings of function `gam()` in package `mgcv`. In this fitting an *optimal* value for the penalty parameter is chosen based on an AIC-like criterion known as UBRE ('Un-Biased Risk Estimator')

```
library(mgcv)
mPen <- mgcv::gam(cbind(D, Y) ~ s(A) + s(P),
  family = poisreg(link = log), data = tdk
)
summary(mPen)
```

```
##
## Family: poisson
## Link function: log
##
## Formula:
## cbind(D, Y) ~ s(A) + s(P)
##
## Parametric coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  1.70960    0.01793   95.33  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##              edf Ref.df Chi.sq p-value
## s(A)  8.143    8.765   2560  <2e-16 ***
## s(P)  3.046    3.790   1054  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.598   Deviance explained = 53.6%
## UBRE = 0.082051   Scale est. = 1           n = 3510
```

The summary is quite brief, and the only estimated coefficient is the intercept, which sets the baseline level for the log-rates, against which the relative age effects and period effects will be contrasted. On the rate scale the baseline level 5.53 per 100000 y is obtained by `exp(1.7096)`.

2. See also the default plot for the fitted curves (solid lines) describing the age and the period effects which are interpreted as contrasts to the baseline level on the log-rate scale.

```
par(mfrow = c(1, 2))
plot(mPen, se=2, seWithMean = TRUE)
```

The dashed lines describe the approximate 95% confidence band for the pertinent curve. One could get the impression that year 1968 would be some kind of reference value for the period effect, like period 1968-72 chosen as the reference in the categorical model previously fitted. This is not the case, however, because `gam()` by default parametrizes the spline effects such that the reference level, at which the spline effect is nominally zero, is the overall *grand mean* value of the log-rate in the data. This corresponds to the principle of *sum contrasts* (`contr.sum`) for categorical explanatory factors.

From the summary you will also find that the degrees of freedom value required for the age effect is nearly the same as the default dimension $k - 1 = 9$ of the part of the model matrix (or basis) initially allocated for each smooth function. (Here k refers to the relevant argument that determines the basis dimension when specifying

a smooth term by `s()` in the model formula). On the other hand the period effect takes just about 3 df.

3. It is a good idea to do some diagnostic checking of the fitted model

```
par(mfrow = c(2, 2))
gam.check(mPen)

##
## Method: UBRE   Optimizer: outer newton
## full convergence after 7 iterations.
## Gradient range [-9.387549e-10,1.362e-06]
## (score 0.0820511 & scale 1).
## Hessian positive definite, eigenvalue range [0.0002209238,0.0003823997].
## Model rank = 19 / 19
##
## Basis dimension (k) checking results. Low p-value (k-index<1) may
## indicate that k is too low, especially if edf is close to k'.
##
##      k'   edf k-index p-value
## s(A) 9.00 8.14   0.93  0.005 **
## s(P) 9.00 3.05   0.95  0.095 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

The four diagnostic plots are analogous to some of those used in the context of linear models for Gaussian responses, but not all of them may be as easy to interpret. – Pay attention to the note given in the printed output about the value of `k`.

4. Let us refit the model but now with an increased `k` for age:

```
mPen2 <- mgcv::gam(cbind(D, Y) ~ s(A, k = 20) + s(P),
  family = poisreg(link = log), data = tdk
)
summary(mPen2)

##
## Family: poisson
## Link function: log
##
## Formula:
## cbind(D, Y) ~ s(A, k = 20) + s(P)
##
## Parametric coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  1.70863    0.01795   95.17  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##      edf Ref.df Chi.sq p-value
## s(A) 11.132 13.406  2553  <2e-16 ***
## s(P)  3.045  3.788  1054  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.599   Deviance explained = 53.7%
## UBRE = 0.081809   Scale est. = 1         n = 3510
par(mfrow = c(2, 2))
gam.check(mPen2)
```

```
##
## Method: UBRE   Optimizer: outer newton
## full convergence after 6 iterations.
## Gradient range [-2.396066e-12,2.901858e-09]
## (score 0.08180917 & scale 1).
## Hessian positive definite, eigenvalue range [0.00022158,0.0009322215].
## Model rank = 29 / 29
##
## Basis dimension (k) checking results. Low p-value (k-index<1) may
## indicate that k is too low, especially if edf is close to k'.
##
##      k'   edf k-index p-value
## s(A) 19.00 11.13   0.93   0.01 **
## s(P)  9.00  3.05   0.95   0.11
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

With this choice of k the df value for age became about 11, which is well below $k - 1 = 19$. Let us plot the fitted curves from this fitting, too

```
par(mfrow = c(1, 2))
plot(mPen2, seWithMean = TRUE)
abline(v = 1968, h = 0, lty = 3)
```

There does not seem to have happened any essential changes from the previously fitted curves, so maybe 8 df could, after all, be quite enough for the age effect.

- Graphical presentation of the effects using `plot.gam()` can be improved. For instance, we may present the age effect to describe the *mean* incidence rates by age, averaged over the whole time span of 54 years. This is obtained by adding the estimated intercept to the estimated smooth curve for the age effect and showing the antilogarithms of the ordinates of the curve. For that purpose we need to extract the intercept and modify the labels of the y -axis accordingly. The estimated period curve can also be expressed in terms of relative incidence rates in relation to the fitted baseline rate, as determined by the model intercept.

```
par(mfrow = c(1, 2))
icpt <- coef(mPen2)[1] # estimated intercept
plot(mPen2,
     seWithMean = TRUE, select = 1, rug = FALSE,
     yaxt = "n", ylim = c(log(1), log(20)) - icpt,
     xlab = "Age (y)", ylab = "Mean rate (/100000 y)"
)
axis(2, at = log(c(1, 2, 5, 10, 20)) - icpt, labels = c(1, 2, 5, 10, 20))
plot(mPen2,
     seWithMean = TRUE, select = 2, rug = FALSE,
     yaxt = "n", ylim = c(log(0.4), log(2)),
     xlab = "Calendar year", ylab = "Relative rate"
)
axis(2, at = log(c(0.5, 0.75, 1, 1.5, 2)), labels = c(0.5, 0.75, 1, 1.5, 2))
abline(v = 1968, h = 0, lty = 3)
```

Homework You could continue the analysis of these data by fitting an age-cohort model as an alternative to the age-period model, as well as an age-cohort-period model utilizing function `apc.fit()` in `Epi`. See (<http://bendixcarstensen.com/APC/>) for details.

Chapter 9

Causal inference

9.1 Proper adjustment for confounding in regression models

The first exercise of this session will ask you to simulate some data according to pre-specified causal structure (don't take the particular example too seriously) and see how you should adjust the analysis to obtain correct estimates of the causal effects.

Suppose one is interested in the effect of beer-drinking on body weight. Let's *assume* that in addition to the potential effect of beer on weight, the following is true in reality:

- Beer-drinking has an effect on the body weight.
- Men drink more beer than women
- Men have higher body weight than women
- People with higher body weight tend to have higher blood pressure
- Beer-drinking increases blood pressure

The task is to simulate a dataset in accordance with this model, and subsequently analyse it to see, whether the results would allow us to conclude the true association structure.

- Sketch a DAG (not necessarily with R) to see, how should one generate the data
- Suppose the actual effect sizes are following:
- The probability of beer-drinking is 0.2 for females and 0.7 for males
- Men weigh on average 10kg more than women.
- One kg difference in body weight corresponds in average to 0.5mmHg difference in (systolic) blood pressures.
- Beer-drinking increases blood pressure by 10mmHg in average.
- Beer-drinking has **no** effect on body weight.

The R commands to generate the data are:

```
bdat <- data.frame(sex = c(rep(0, 500), rep(1, 500)))  
# a data frame with 500 females, 500 males  
bdat$beer <- rbinom(1000, 1, 0.2 + 0.5 * bdat$sex)  
bdat$weight <- 60 + 10 * bdat$sex + rnorm(1000, 0, 7)  
bdat$bp <-  
  110 + 0.5 * bdat$weight + 10 * bdat$beer + rnorm(1000, 0, 10)
```

- Now fit the following models for body weight as dependent variable and beer-drinking as independent variable. Look, what is the estimated effect size:
- Unadjusted (just simple linear regression)

- Adjusted for sex
- Adjusted for sex and blood pressure
- What would be the conclusions on the effect of beer on weight, based on the three models? Do they agree? Which (if any) of the models gives an unbiased estimate of the actual causal effect of interest?
- How can the answer be seen from the graph?
- Now change the data-generation algorithm so, that in fact beer-drinking does increase the body weight by 2kg. Look, what are the conclusions in the above models now. Thus the data is generated as before, but the weight variable is computed as:

```
bdat$weight <-
60 + 10 * bdat$sex + 2 * bdat$beer + rnorm(1000, 0, 7)
```

- Suppose one is interested in the effect of beer-drinking on blood pressure instead, and is fitting a) an unadjusted model for blood pressure, with beer as an only covariate; b) a model with beer and sex as covariates. Would either a) or b) give an unbiased estimate for the effect? (You may double-check whether the simulated data is consistent with your answer).

9.2 DAG tools in the package *dagitty*

There is a software *DAGitty* (<http://www.dagitty.net/>) and also an R package *dagitty* that can be helpful in dealing with DAGs. Let's try to get the answer to the previous exercise using this package.

```
if (!("dagitty" %in% installed.packages())){
  install.packages("dagitty")
}
library(dagitty)
```

```
##
## Attaching package: 'dagitty'
## The following object is masked from 'package:Epi':
##
## paths
```

Let's recreate the graph on the lecture slide 28 (but omitting the direct causal effect of interest, $C \rightarrow D$):

```
g <- dagitty("dag {
  C <- S -> Y -> U -> D
  C -> Z <- Y
  Z -> D
  C <- X -> D
  C -> Q
  W -> D
}")
plot(g)
```


To get a more similar look as on the slide, we must supply the coordinates (x increases from left to right, y from top to bottom):

```
coordinates(g) <-
list(
  x =
    c(
      S = 1, C = 1, Q = 1, Y = 2, Z = 2,
      X = 2, U = 3, D = 3, W = 3
    ),
  y =
    c(
```

```

    U = 1, Y = 1, S = 1, Z = 2, C = 3,
    D = 3, X = 4, W = 4, Q = 4
  )
)
plot(g)

```



Let's look at all possible paths from C to D :

```
paths(g, "C", "D")
```

As you see, one path contains a collider and is therefore a *closed* path and the others are *open*.

Let's identify the minimal sets of variables needed to adjust the model for D for, to obtain an unbiased estimate of the effect of C . You can specify, whether you want to estimate direct or total effect of C :

```

adjustmentSets(
  g, exposure = "C", outcome = "D", effect = "direct"
)
adjustmentSets(
  g, exposure = "C", outcome = "D", effect = "total"
)

```

Thus, for total effect estimation one should adjust for X and either Y or S , whereas for direct effect estimation, one would also need to adjust for Z .

You can verify that, these are the variables that will block all open paths from C to D .

Now try to do the *beer-weight* exercise using *dagitty*:

- Create the DAG and plot it
- What are the paths from WEIGHT to BEER?
- Will you get the same recommendation for the adjustment variable selection as you found before?

9.3 Identifying the true DAG for the data

The following code creates three DAGs

```


par(mfrow=c(1,3))
g1 <- dagitty("dag {
  U -> Z -> Y

```

```

    U -> X
    W -> Y
    Q -> W -> X
    Q -> Y
  })
g2 <- dagitty("dag {
  U -> Z -> Y -> W
  U -> X -> W
  Q -> W
  Q -> Y
}")
g3 <- dagitty("dag {
  U -> Z -> Y
  U -> X -> W -> Y
  Q -> W
  Q -> Y
}")
coord <-
  list( x = c(X=1, U = 1.3, W = 2, Z = 2.3, Q = 2.7, Y=3),
        y = c(U=1, Z=1.3, X=2, Y=2, W=2.7, Q=3)
  )
coordinates(g1)<-coordinates(g2)<-coordinates(g3)<-coord
plot(g1)
title("(a)")
plot(g2)
title("(b)")
plot(g3)
title("(c)")

```



Now, the following script generates a dataset:

```

source("data/gendata.r")
head(dat)

```

```

##      Q      W      X      Z      Y
## 1 -0.397  1.056  0.199 -0.372 -4.198
## 2  0.460  2.393 -1.647 -3.004 -10.667
## 3  0.164 -0.364  0.205 -0.852  -1.830
## 4  1.584  3.285 -9.941  2.752  -4.121
## 5  1.048  1.637  0.783 -2.047  -4.378
## 6  1.895  3.272 -3.900 -2.313  -9.871

```

Exercise: Using linear models, try to identify, which of the three DAGs has been used to generate the data.

```
adjustmentSets(g1,"X","Y", effect="direct")
```

```
## { W, Z }
```

```
## { U, W }
```

```
summary(lm(Y~X+W+Z,data=dat))
```

```
##
```

```
## Call:
```

```
## lm(formula = Y ~ X + W + Z, data = dat)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
## -3.9888 -0.7441  0.0096  0.7657  3.5928
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.033168   0.024551   1.351   0.177
## X            -0.001691   0.014877  -0.114   0.909
## W            -2.599607   0.031987 -81.270 <2e-16 ***
## Z             1.014166   0.021197  47.846 <2e-16 ***
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
```

```
## Residual standard error: 1.097 on 1996 degrees of freedom
```

```
## Multiple R-squared:  0.9686, Adjusted R-squared:  0.9686
```

```
## F-statistic: 2.054e+04 on 3 and 1996 DF,  p-value: < 2.2e-16
```

```
# no X effect: this DAG is possible
```

```
adjustmentSets(g2,"X","Y", effect="direct")
```

```
## { Z }
```

```
## { U }
```

```
summary(lm(Y~X+Z,data=dat))
```

```
##
```

```
## Call:
```

```
## lm(formula = Y ~ X + Z, data = dat)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
## -7.480 -1.568 -0.107  1.580  8.730
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.10067    0.05092   1.977  0.0482 *
## X            1.13137    0.01077 105.040 <2e-16 ***
## Z            2.38032    0.02680  88.830 <2e-16 ***
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
```

```
## Residual standard error: 2.277 on 1997 degrees of freedom
```

```
## Multiple R-squared:  0.8648, Adjusted R-squared:  0.8647
```

```
## F-statistic: 6387 on 2 and 1997 DF,  p-value: < 2.2e-16
```

X effect is significant: this DAG is not likely

```
adjustmentSets(g3,"X","Y", effect="direct")
```

```
## { Q, W, Z }
```

```
## { Q, U, W }
```

```
summary(lm(Y~X+Q+W+Z,data=dat))
```

```
##
```

```
## Call:
```

```
## lm(formula = Y ~ X + Q + W + Z, data = dat)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
## -3.7210 -0.6425  0.0031  0.6596  3.1985
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.019576   0.022168   0.883   0.377
## X            -0.006038   0.013429  -0.450   0.653
## Q             1.071881   0.050242  21.335 <2e-16 ***
## W            -3.033392   0.035312 -85.903 <2e-16 ***
## Z             1.003544   0.019138  52.437 <2e-16 ***
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
```

```
## Residual standard error: 0.9903 on 1995 degrees of freedom
```

```
## Multiple R-squared:  0.9745, Adjusted R-squared:  0.9744
```

```
## F-statistic: 1.902e+04 on 4 and 1995 DF, p-value: < 2.2e-16
```

no X effect: this DAG is possible

```
adjustmentSets(g1,"Z","W")
```

```
## {}
```

```
summary(lm(W~Z,data=dat))
```

```
##
```

```
## Call:
```

```
## lm(formula = W ~ Z, data = dat)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
## -9.5057 -1.5539  0.0044  1.4534  8.0072
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.02154    0.04919  -0.438   0.661
## Z            0.01908    0.02238   0.852   0.394
```

```
##
```

```
## Residual standard error: 2.2 on 1998 degrees of freedom
```

```
## Multiple R-squared:  0.0003635, Adjusted R-squared: -0.0001368
```

```
## F-statistic: 0.7266 on 1 and 1998 DF, p-value: 0.3941
```

no Z effect: this DAG is possible

```
adjustmentSets(g3,"Z","W")
```

```
## { X }
## { U }
summary(lm(W~Z+X,data=dat))

##
## Call:
## lm(formula = W ~ Z + X, data = dat)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.5471 -0.4981  0.0016  0.5336  2.3622
##
## Coefficients:
##              Estimate Std. Error  t value Pr(>|t|)
## (Intercept) -0.025965   0.017165   -1.513   0.131
## Z            -0.525522   0.009033  -58.178 <2e-16 ***
## X            -0.435859   0.003631 -120.044 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.7676 on 1997 degrees of freedom
## Multiple R-squared:  0.8783, Adjusted R-squared:  0.8782
## F-statistic: 7208 on 2 and 1997 DF,  p-value: < 2.2e-16
# significant Z effect: this DAG is possible
```

If you have made your decision, you may check the script ‘gendata.r’ to see, whether your guess was right.

9.4 Instrumental variables estimation: Mendelian randomization

Suppose you want to estimate the effect of Body Mass Index (BMI) on blood glucose level (associated with the risk of diabetes). Let’s conduct a simulation study to verify that when the exposure-outcome association is confounded, but there is a valid instrument (genotype), one obtains an unbiased estimate of the causal effect.

- Start by generating the genotype variable as $\text{Binomial}(2,p)$, with $p = 0.2$ (and look at the resulting genotype frequencies):

```
n <- 10000
mrdat <- data.frame(G = rbinom(n, 2, 0.2))
table(mrdat$G)
```

- Also generate the confounder variable U

```
mrdat$U <- rnorm(n)
```

- Generate a continuous (normally distributed) exposure variable BMI so that it depends on G and U . Check with linear regression, whether there is enough power to get significant parameter estimates. For instance:

```
mrdat$BMI <- with(mrdat, 25 + 0.7 * G + 2 * U + rnorm(n))
```

- Finally generate Y (“Blood glucose level”) so that it depends on BMI and U (but not on G).

```
mrdat$Y <-
  with(mrdat, 3 + 0.1 * BMI - 1.5 * U + rnorm(n, 0, 0.5))
```

- Verify, that simple regression model for Y , with BMI as a covariate, results in a biased estimate of the causal effect (parameter estimate is different from what was generated)

How different is the estimate from 0.1?

- Estimate a regression model for Y with two covariates, G and BMI . Do you see a significant effect of G ? Could you explain analytically, why one may see a significant parameter estimate for G there?
- Find an IV (instrumental variables) estimate, using G as an instrument, by following the algorithm in the lecture notes (use two linear models and find a ratio of the parameter estimates). Does the estimate get closer to the generated effect size?

```
mgx <- lm(BMI ~ G, data = mrdat)
ci.lin(mgx) # check the instrument effect
bgx <- mgx$coef[2] # save the 2nd coefficient (coef of G)
mgy <- lm(Y ~ G, data = mrdat)
ci.lin(mgy)
bgy <- mgy$coef[2]
causeff <- bgy / bgx
causeff # closer to 0.1?
```

- A proper simulation study would require the analysis to be run several times, to see the extent of variability in the parameter estimates. A simple way to do it here would be using a `for`-loop. Modify the code as follows (exactly the same commands as executed so far, adding a few lines of code to the beginning and to the end):

```
n <- 10000
# initializing simulations:
# 30 simulations (change it, if you want more):
nsim <- 30
mr <- rep(NA, nsim) # empty vector for the outcome parameters
for (i in 1:nsim) { # start the loop
  ## Exactly the same commands as before:
  mrdat <- data.frame(G = rbinom(n, 2, 0.2))
  mrdat$U <- rnorm(n)
  mrdat$BMI <-
    with(mrdat, 25 + 0.7 * G + 2 * U + rnorm(n))
  mrdat$Y <-
    with(mrdat, 3 + 0.1 * BMI - 1.5 * U + rnorm(n, 0, 0.5))
  mgx <- lm(BMI ~ G, data = mrdat)
  bgx <- mgx$coef[2]
  mgy <- lm(Y ~ G, data = mrdat)
  bgy <- mgy$coef[2]
  # Save the i'th parameter estimate:
  mr[i] <- bgy / bgx
} # end the loop
```

Now look at the distribution of the parameter estimate:

```
summary(mr)
```

- (optional) Change the code of simulations so that the assumptions are violated: add a weak direct effect of the genotype G to the equation that generates Y :

```
mrdat$Y <-
  with(
    mrdat,
    3 + 0.1 * BMI - 1.5 * U + 0.05 * G + rnorm(n, 0, 0.5)
  )
```

Repeat the simulation study to see, what is the bias in the average estimated causal effect of BMI on Y .

- (optional) Using library `sem` and function `tsls`, one can obtain a two-stage least squares estimate for the causal effect and also the proper standard error. Do you get the same estimate as before?

```
if (!("sem" %in% installed.packages())) install.packages("sem")
library(sem)
summary(tsls(Y ~ BMI, ~G, data = mrdat))
```

(There are also several other R packages for IV estimation and Mendelian Randomization (*MendelianRandomization* for instance))


9.5 Why are simulation exercises useful for causal inference?

If we simulate the data, we know the data-generating mechanism and the *true* causal effects. So this is a way to check, whether an analysis approach will lead to estimates that correspond to what is generated. One could expect to see similar phenomena in real data analysis, if the data-generation mechanism is similar to what was used in simulations.

Chapter 10

Graphics with ggplot2

The plot below is from randomized study of the effect of Tamoxifen treatment on bone mineral metabolism in a group of patients who were treated for breast cancer (Kristensen et al, 1994). It is reproduced with permission from the author Peter Dalgaard who has also shared the original data.



The purpose of this exercise is to show you how to build a similar plot in R using the `ggplot2` package.

The outcome of the study is the percentage change in serum alkaline phosphate compared with the baseline when the participants first entered the study. The first measurement, on the left of the plot, is zero by definition. Then as we move from left to right, the plot shows the relative changes for the two groups of women: one receiving the anti-cancer drug tamoxifen; the other receiving standard treatment which is referred to as the *control* group. Each woman was randomly allocated to one of these groups on entry to the study. The women returned for follow-up visits until the end of the study at 24 months.

Take a moment to study the graph and try to identify the different graphical elements that we will need to reproduce. Write down a description of these elements and compare it to my summary below.

- The trajectory of the outcome variable is represented by points connected by straight lines. One point represents the mean value for one follow-up visit in one of the groups.
- Separate trajectories are plotted for the treatment and control groups.
- Vertical error bars give a visual impression of the uncertainty associated with each point.
- A horizontal line at zero is drawn as a reference so that we can see the size of the change relative to baseline.
- The x-axis has tick marks at follow-up visits, which are irregularly spaced. Visits occurred every 3 months for the first year, then every 6 months for the second year.
- The y-axis extends beyond the range of the data from -40% to +30%.
- A table underneath the x-axis shows the number of women in the treatment and control groups at each visit. These numbers go down as the study progresses. This phenomenon is extremely common in medical studies and is called “loss to follow-up”.

These are the important graphical elements. You may have spotted some other features of the graph, but we will focus on trying to incorporate the features listed above.

10.1 Data

The data are available in the file `alkfos.csv`. This is a text file in comma-separated variable format. It can be read into R using the `read.csv()` function.

```
alkfos <- read.csv("./data/alkfos.csv")
```

There are 14 rows in the data frame `alkfos`: one for each combination of the 7 visits and the 2 groups. The variables in the data frame are:

- **time** The time since randomization for each follow-up visit.
- **mean** The outcome variable representing the mean percentage change in serum alkaline phosphate.
- **available** The number of women still participating in the study.
- **treat** The treatment group: a character vector taking values “tamoxifen” for the treatment group and “control” for the control group.
- **sem** The standard error of the mean used to calculate the length of the error bars.

10.2 Building the plot

10.2.1 The ggplot2 package

In this exercise, we will be using the R package `ggplot2`. The “gg” stands for “Grammar of Graphics”, which is a conceptual framework for understanding statistical graphics created by Leland Wilkinson.

The core function in the `ggplot2` package is `ggplot()`. This function requires two arguments:

- **data** is the name of the data frame that we will take the variables from.
- **mapping** is the *aesthetic mapping* between variables in the data frame and graphical elements that we want to plot. In this example:
 - y-axis is the mean of the percent change in alkaline phosphate represented by the variable `mean`;
 - x-axis is the visit time represented by the variable `time`.

We make the mapping in a call to the function `aes()`.

The code below creates a basic plot and assigns it to the object `p0`.

```
library(ggplot2, quietly=TRUE)
p0 <- ggplot(data=alkfos, mapping=aes(x=time, y=mean))
```

Running the code creates the plot but does not display it. This is an important difference between base graphics and `ggplot` graphics:


- Base graphics functions do not return an R object but will always display a plot on the current graphics device.
- GGplot graphics functions return an R graphical object, which will only be displayed if we ask for it.

There are two ways to display a graphical object in R:

1. If we have saved the result to an object, call the `print()` or `show()` functions.
2. If we have not saved the result, the graphical object will be displayed in the same way that a regular R object is printed to the screen.

The code chunk below displays the object `p0` created in the chunk above.

```
show(p0)
```



You may find the results a little disappointing. The `ggplot()` function has drawn x- and y-axes that span the range of the data, but otherwise has drawn no points or lines. This is another important difference between `ggplot2` and base graphics. Whereas base plotting functions will usually create a sensible default plot, `ggplot()` makes no assumptions about what you want.

We will recreate the plot in *layers*. Layers are added to a `ggplot` object by calling additional functions and adding their return value to the basic plot using the `+` symbol.

10.2.2 Geometries


The most important layers are *geometries*, which describe how to display the data in the plot. Geometry functions in `ggplot2` all start with the prefix `geom_`.

10.2.2.1 Points and lines

We want to plot points representing the mean of the outcome at each visit. This is done with the function `geom_point()`. We also want to plot lines between successive points in time. This is done with the function

`geom_line()`. The code chunk below adds point and line geometries to our basic plot `p0` and saves the result to the new object `p1`.


```
p1 <- p0 + geom_point() + geom_line()
show(p1)
```



Again, this is not quite what we want. The line geometry is mixing up observations from the treatment and control groups, but we want separate lines for each group. To obtain this we need to add an additional aesthetic mapping inside the call to `ggplot()`. The aesthetic `group` gives the name of the variable in the data frame that distinguishes observations from different groups.

Modify the code chunk below so that it includes the `group` aesthetic to recreate plot `p1` correctly.

```
p1 <- ggplot(data=alkfos, mapping=aes(x=time, y=mean, group=treat)) +
  geom_point() + geom_line()
show(p1)
```



10.2.2.2 Error bars


Now we can add the error bars. These can be added using the `linerrange` geometry. This geometry needs some new aesthetic mappings. It will inherit the `x` mapping from the call to `ggplot()`. However, it needs new mappings `ymin` and `ymax` giving the lower and upper points of the error bar on the y-axis. We could go back and modify the aesthetic mappings in the call to `ggplot()`. However, we can also supply aesthetic mappings inside the call to a geometry function. It can be useful to do this when an aesthetic is only used by one geometry (which is the case here).

Modify the code chunk below so that the `linerrange` geometry has its own aesthetic mapping for `ymin` and `ymax`.

- `ymin` is the mean minus the standard error
- `ymax` is the mean plus the standard error

Note that the error bars are *not confidence intervals*. In statistical graphics it is very common to use plus/minus the standard error to represent the uncertainty in the data without making any statistical claims about coverage.

```
p2 <- p1 + geom_linerrange(mapping=aes(ymin=mean-sem, ymax=mean+sem))
show(p2)
```




Compare the y-axes for plots `p1` and `p2`. Note that when we add the error bars, the y-axis automatically expanded the limits of the y-axis to include the highest and lowest values on the error bars.

10.2.2.3 Horizontal reference line

The last geometry we want to add is the horizontal reference line at zero. This is added using the `hline` geometry. We must supply the argument `yintercept` which determines where the horizontal line intersects the y-axis.

We should distinguish the reference line from the lines used to plot the data, otherwise it may distract the viewer. We could do this by changing the colour of the reference line. But here we will use a thinner line using the optional argument `linewidth` to the `hline` geometry. The default value of `linewidth` is 1. A value less than 1 draws a thinner line; a value greater than 1 draws a thicker line. Choose a suitable value for your plot.

```
p3 <- p2 + geom_hline(yintercept=0, linewidth=0.1)
show(p3)
```



10.2.3 Scales


Next we want to fix the x- and y-axes. The `ggplot()` function will provide default axes for you, which are often good enough. But in this case we want custom axes. These are provided by scale functions. In fact axes are just one kind of scale. In the grammar of graphics framework, a scale is a way to annotate any quantity that represents a dimension of the data. This also includes shapes, colours, and sizes.

Both axes are continuous. Therefore we use the `scale_x_continuous()` and `scale_y_continuous()` functions to create custom axes. These scale functions have the same arguments, but they affect the x- and y-axes, respectively. The arguments you want to supply are

- **breaks** A numeric vector of positions on the axis where you want tick marks.
- **limits** A numeric vector of length 2 giving the minimum and maximum values to be displayed on the axis.


Modify the code chunk below to create the custom axes:

```
p4 <- p3 + scale_x_continuous(breaks=c(0,3,6,9,12,18,24)) +
  scale_y_continuous(breaks=seq(from=-40, to=30, by=5),
                    limits = c(-40, 30))
show(p4)
```



We should also supply informative axis labels. The default axis labels are taken from the names of the variables that we used in the aesthetic mappings for `x` and `y`. Use the `xlab()` and `ylab()` functions to provide informative axis labels.

```
p5 <- p4 + xlab("Months after randomization") +  
  ylab("Percent change in serum alkaline phosphate")  
show(p5)
```




10.2.4 Themes

Our plot contains most of the same graphical elements as the reference plot. But it does not look the same. The original plot was drawn on a white background. But our plot has a grey background with white major and minor grid lines.

Themes control many aspects of the appearance of the plot that do not depend on the data. The default theme, which we have been using so far, is called `theme_gray()`. Use the `help()` function to look up `theme_gray()` and you will see the help page documents the other complete themes provided by the ggplot2 package.

- Modify the chunk below to use different themes until you find one that most closely matches the original plot in the file bentK.pdf.
- Find the parameter that changes to the font size and use it to reduce the font until the y-axis label fits.


```
p6 <- p5 + theme_classic(base_size=9)
show(p6)
```



10.2.5 Putting everything together

In this exercise, you have constructed the plot in steps, saving and displaying the plot in separate chunks to allow you to see the changes as we add each layer of the plot. In practice, we do not work like this. All the instructions for creating the plot should be together. Fill in the chunk below so that the plot is recreated in a single chunk.

```
p <- ggplot(alkfos, mapping=aes(x=time, y=mean, groups=treat)) +
  geom_point(mapping=aes(shape=treat)) +
  geom_line() +
  geom_linerange(mapping=aes(ymin=mean-sem, ymax=mean+sem)) +
  scale_x_continuous(breaks=c(0,3,6,9,12,18,24)) +
  scale_y_continuous(breaks=seq(from=-40, to=30, by=5),
                     limits = c(-40, 30)) +
  xlab("Months after randomization") +
  ylab("Percent change in serum alkaline phosphate") +
  geom_hline(yintercept=0, linewidth=0.1) +
  theme_classic(base_size=9)
show(p)
```



There is still one thing missing from our reproduction. In the original plot, the treatment and control groups were distinguished by different plotting symbols, whereas in our plot they both have circular marks. We can change this by adding the aesthetic mapping **shape**.

Modify the chunk above so that the aesthetic **shape** is mapped onto the variable **treat**. You can do this inside the call to `ggplot()` or in the call to `geom_point()`.

Note that adding the shape aesthetic automatically adds a legend to the plot, telling us which shape corresponds to which group.

10.3 Self study

Congratulations, you have mostly recreated the reference figure from Kristensen et al (1994). However, have you really learned anything? The best way to find out is to come back in a week's time and try to recreate the reference plot without looking at this exercise sheet.

You may find it useful to download the ggplot cheatsheet. You can also use R help to find out more about the functions in the ggplot2 package. Finally, you can refer to the book *Ggplot2: elegant graphics for data analysis* by Hadley Wickham, creator of the ggplot2 package.

10.4 Adding the table (advanced topic)

There is one important element missing from the graph. This is the table of the number of participants in each group. This is an advanced topic so instead of explaining each step I am just going to give you the solution. Do not be concerned if you do not understand this part.

Firstly we create a new graphical object to represent the table. This object has its own aesthetic mappings, geometries and scales. I will skip a detailed description of this step except to note that some of the arguments below will suppress graphical features (e.g. `NULL` and `element_blank()`)

```

tab <- ggplot(data=alkfos,
              mapping=aes(x=time, y=treat, label=available)) +
  geom_text(size=2) + xlab(NULL) + ylab(NULL) +
  scale_x_continuous(breaks=NULL) +
  theme_bw(base_size=9) +
  theme(panel.grid=element_blank())
tab

```




Next use the `plot_grid()` function from the `cowplot` package to stack the displayed objects vertically (`align="v"`), in a single column (`ncol=1`, `nrow=2`), with the left and right ends of the x-axes aligned (`axis="lr"`) and most of the space taken up by the plot (`rel_heights=c(5,1)`).

```

library(cowplot)
plot_grid(plotlist=list(p, tab), align="v", axis="lr",
          ncol=1, nrow=2, rel_heights=c(5,1))

```



10.5 References

Kristensen B, Ejlersen B, Dalgaard P, Larsen L, Holmegaard SN, Transbøl I, Mouridsen HT. Tamoxifen and bone metabolism in postmenopausal low-risk breast cancer patients: a randomized study. *J Clin Oncol*. 1994 May;**12**(5):992-7. doi: 10.1200/JCO.1994.12.5.992. PMID: 8164053.

Chapter 11

Survival analysis with competing risks: Oral cancer patients

11.1 Description of the data

File `oralca2.txt`, that you may access from a url address to be given in the practical, contains data from 338 patients having an oral squamous cell carcinoma diagnosed and treated in one tertiary level oncological clinic in Finland since 1985, followed-up for mortality until 31 December 2008.

The dataset contains the following variables:

Variable	Description
<code>sex</code>	sex, a factor with categories; 1 = "Female", 2 = "Male"
<code>age</code>	age (years) at the date of diagnosing the cancer
<code>stage</code>	TNM stage of the tumour (factor): 1 = "I", ..., 4 = "IV", 5 = "unkn"
<code>time</code>	follow-up time (in years) since diagnosis until death or censoring
<code>event</code>	event ending the follow-up (numeric): 0 = censoring alive, 1 = death from oral cancer, 2 = death from other causes.

11.2 Loading the packages and the data

- Load the R packages `Epi`, and `survival` needed in this exercise.

```
library(Epi)
library(survival)
cB8 <- c("#000000", "#E69F00", "#56B4E9", "#009E73",
        "#F0E442", "#0072B2", "#D55E00", "#CC79A7") #colors chosen
options(digits=3)
```

- Read the datafile `oralca2.txt` from a website, whose precise address will be given in the practical, into an R data frame named `orca`. Look at the head, structure and the summary of the data frame. Using function `table()` count the numbers of censorings as well as deaths from oral cancer and other causes, respectively, from the `event` variable.

```
orca <- read.table(file = 'https://raw.githubusercontent.com/SPE-R/SPE/master/pracs-book/data/oralca2.txt')
head(orca)
```

```
##      sex age stage time event
## 1  Male 65.4  unkn 5.081     0
## 2 Female 83.1   III 0.419     1
## 3  Male 52.6   II  7.915     2
```

```
## 4   Male 77.1    I 2.480    2
## 5   Male 80.3   IV 2.500    1
## 6 Female 82.6   IV 0.167    2

orca$stage<-as.factor(orca$stage)
orca$sex<-factor(orca$sex,levels=c("Male","Female"))
str(orca)

## 'data.frame':   338 obs. of  5 variables:
## $ sex : Factor w/ 2 levels "Male","Female": 1 2 1 1 1 2 1 1 2 1 ...
## $ age : num  65.4 83.1 52.6 77.1 80.3 ...
## $ stage: Factor w/ 5 levels "I","II","III",...: 5 3 2 1 4 4 2 5 4 2 ...
## $ time : num  5.081 0.419 7.915 2.48 2.5 ...
## $ event: int   0 1 2 2 1 2 0 1 1 0 ...

summary(orca)
```

```
##      sex      age      stage      time      event
## Male :186   Min. :15.2   I :50   Min. : 0.085   Min. :0.000
## Female:152 1st Qu.:53.2   II:77  1st Qu.: 1.333   1st Qu.:0.000
##           Median :64.9   III:72 Median : 3.869   Median :1.000
##           Mean   :63.5   IV :68  Mean   : 5.662   Mean   :0.994
##           3rd Qu.:74.3  unkn:71 3rd Qu.: 8.417   3rd Qu.:2.000
##           Max.   :92.2           Max.   :23.258   Max.   :2.000
```

11.3 Total mortality: Kaplan–Meier analyses

- We start our analysis of total mortality pooling the two causes of death into a single outcome. First, construct a *survival object* `orca$suob` from the event variable and the follow-up time using function `Surv()`. Look at the structure and summary of `suob`.

```
suob <- Surv(orca$time, 1 * (orca$event > 0))
str(suob)

## 'Surv' num [1:338, 1:2] 5.081+ 0.419 7.915 2.480 2.500 0.167 5.925+ 1.503 13.333 7.666
## - attr(*, "dimnames")=List of 2
## ..$ : NULL
## ..$ : chr [1:2] "time" "status"
## - attr(*, "type")= chr "right"

summary(suob)
```

```
##      time      status
## Min. : 0.085   Min. :0.000
## 1st Qu.: 1.333   1st Qu.:0.000
## Median : 3.869   Median :1.000
## Mean   : 5.662   Mean   :0.678
## 3rd Qu.: 8.417   3rd Qu.:1.000
## Max.   :23.258   Max.   :1.000
```

- Create a `survfit` object `s.all`, which does the default calculations for a Kaplan–Meier analysis of the overall (marginal) survival curve.

```
s.all <- survfit(suob ~ 1, data = orca)
```

See the structure of this object and apply `print()` method on it, too. Look at the results; what do you find?

```
s.all

## Call: survfit(formula = suob ~ 1, data = orca)
##
##      n events median 0.95LCL 0.95UCL
```

```
## [1,] 338    229    5.42    4.33    6.92
```


```
str(s.all)
```

```
## List of 17
## $ n      : int 338
## $ time   : num [1:251] 0.085 0.162 0.167 0.17 0.246 0.249 0.252 0.329 0.334 0.413 ...
## $ n.risk  : num [1:251] 338 336 334 330 328 327 326 323 322 321 ...
## $ n.event : num [1:251] 2 2 4 2 1 1 3 1 1 1 ...
## $ n.censor : num [1:251] 0 0 0 0 0 0 0 0 0 0 ...
## $ surv     : num [1:251] 0.994 0.988 0.976 0.97 0.967 ...
## $ std.err  : num [1:251] 0.0042 0.00595 0.00847 0.0095 0.00998 ...
## $ cumhaz   : num [1:251] 0.00592 0.01187 0.02385 0.02991 0.03295 ...
## $ std.chaz : num [1:251] 0.00418 0.00593 0.00843 0.00946 0.00994 ...
## $ type     : chr "right"
## $ logse    : logi TRUE
## $ conf.int : num 0.95
## $ conf.type: chr "log"
## $ lower    : num [1:251] 0.986 0.977 0.96 0.953 0.949 ...
## $ upper    : num [1:251] 1 1 0.993 0.989 0.987 ...
## $ t0       : num 0
## $ call     : language survfit(formula = suob ~ 1, data = orca)
## - attr(*, "class")= chr "survfit"
```


- The `summary` method for a `survfit` object would return a lengthy life table. However, the `plot` method with default arguments offers the Kaplan–Meier curve for a conventional illustration of the survival experience in the whole patient group. Alternatively, instead of graphing survival proportions, one can draw a curve describing their complements: the cumulative mortality proportions. This curve is drawn together with the survival curve as the result of the second command line below.

```
plot(s.all,main="KM estimate of the survival
and cum. mortality proportions",
xlab="years", ylab="Survival")
```

KM estimate of the survival and cum. mortality proportions



```
plot(s.all)
lines(s.all, fun = "event", mark.time = F, conf.int = FALSE)
```




The effect of option `mark.time=F` is to omit marking the times when censorings occurred.

11.4 Total mortality by stage

Tumour stage is an important prognostic factor in cancer survival.


- stage I: cancer hasn't spread to lymph nodes or other tissue (local),
stage II: cancer has grown but not spread
stage III: cancer has grown larger but not spread
stage IV: cancer has spread to other organs or areas
- Plot separate cumulative mortality curves for the different stage groups marking them with different colors, the order which you may define yourself. Also find the median survival time for each stage.

```
s.stg <- survfit(suob ~ stage, data = orca)
col5 <- cB8[1:5]
plot(s.stg, col = col5, fun = "event", main="Cum. mortality by stage", mark.time = FALSE)
legend(15, 0.6, title="stage", legend=levels(orca$stage), col = col5, lty=rep(1,5))
```



- Create now two parallel plots of which the first one describes the cumulative hazards and the second one graphs the log-cumulative hazards against log-time for the different stages. Compare the two presentations with each other and with the one in the previous item.

```
par(mfrow = c(1, 2))
plot(s.stg, col = col5, fun = "cumhaz", main = "Cumulative hazards")
legend(1, 3.5, title="stage", legend=levels(orca$stage), col = col5, lty=rep(1,5), cex=0.8)
plot(s.stg, col = col5, fun = "cloglog", main = "cloglog: log cum.haz")
legend(3, -2, title="stage", legend=levels(orca$stage), col = col5, lty=rep(1,5), cex=0.7)
```



- If the survival times were *exponentially* distributed in a given (sub)population the corresponding cloglog-curve should follow an approximately linear pattern. Could this be the case here in the different stages?
- Also, if the survival distributions of the different subpopulations would obey the *proportional hazards* model, the vertical distance between the cloglog-curves should be approximately constant over the time axis. Do these curves indicate serious deviation from the proportional hazards assumption?
- In the lecture handouts it was observed that the crude contrast between males and females in total mortality appears unclear, but the age-adjustment in the Cox model provided a more expected hazard ratio estimate. We shall examine the confounding by age somewhat closer. First categorize the continuous age variable into, say, three categories by function `cut()` using suitable breakpoints, like 55 and 75 years, and cross-tabulate sex and age group:

```
orca$agegr <- cut(orca$age, br = c(0, 55, 75, 95))
stat.table(list(sex, agegr), list(count(), percent(agegr)),
  margins = TRUE,
  data = orca
)
```


```
## -----
##          -----agegr-----
## sex      (0,55] (55,75] (75,95]  Total
## -----
## Male           71      86      29    186
##                38.2    46.2    15.6  100.0
##
## Female          29      74      49    152
##                19.1    48.7    32.2  100.0
##
##
## Total          100     160     78    338
```

```
##          29.6    47.3    23.1    100.0
## -----
```

Male patients are clearly younger than females in these data.

Now, plot Kaplan–Meier curves jointly classified by sex and age.

```
s.agrx <- survfit(suob ~ agegr + sex, data=orca)
par(mfrow=c(1,1))
plot(s.agrx, fun="event", main="Cumulative mortality (KM) by age and sex", xlab="Time since oral cancer dia",
      col=rep(c(cB8[7], cB8[6]),3), lty=c(2,2, 1,1, 5,5),
      xaxs="i", yaxs="i")
legend(12,0.35, legend=c("(0,55] Female ", "(0,55] Male",
                        "(55,75] Female ", "(55,75] Male",
                        "(75,95] Female ", "(75,95] Male" ),
      col=rep(c(cB8[7], cB8[6]),3), lty=c(2,2, 1,1, 5,5), cex=0.65)
```



In each age band the mortality curve for males is on a higher level than that for females.

11.5 Event-specific cumulative mortality curves

We move on to analysing cumulative mortalities for the two causes of death separately, first overall and then by prognostic factors.

- Use the `survfit`-function in `survival` package with option `type="mstate"`.

```
library(survival)
cif1 <- survfit(Surv(time, event, type = "mstate") ~ 1,
  data = orca
)
```

```
## Warning in Surv(time, event, type = "mstate"): type= 'mstate' is deprecated,
## use a factor variable as status
```

```
str(cif1)
```

```
## List of 23
## $ n          : int 338
## $ time       : num [1:251] 0.085 0.162 0.167 0.17 0.246 0.249 0.252 0.329 0.334 0.413 ...
## $ n.risk     : num [1:251, 1:3] 338 336 334 330 328 327 326 323 322 321 ...
## ..- attr(*, "dimnames")=List of 2
## .. ..$ : NULL
## .. ..$ : chr [1:3] "(s0)" "1" "2"
## $ n.event    : num [1:251, 1:3] 0 0 0 0 0 0 0 0 0 0 ...
## ..- attr(*, "dimnames")=List of 2
## .. ..$ : NULL
## .. ..$ : chr [1:3] "(s0)" "1" "2"
## $ n.censor   : num [1:251, 1:3] 0 0 0 0 0 0 0 0 0 0 ...
## ..- attr(*, "dimnames")=List of 2
## .. ..$ : NULL
## .. ..$ : chr [1:3] "(s0)" "1" "2"
## $ pstate     : num [1:251, 1:3] 0.994 0.988 0.976 0.97 0.967 ...
## ..- attr(*, "dimnames")=List of 2
## .. ..$ : NULL
## .. ..$ : chr [1:3] "(s0)" "1" "2"
## $ n.transition: num [1:251, 1:2] 2 2 2 1 1 0 2 1 1 1 ...
## ..- attr(*, "dimnames")=List of 2
## .. ..$ : NULL
## .. ..$ : chr [1:2] "1:2" "1:3"
## $ n.id       : int 338
## $ cumhaz     : num [1:251, 1:2] 0.00592 0.01187 0.01786 0.02089 0.02394 ...
## ..- attr(*, "dimnames")=List of 2
## .. ..$ : NULL
## .. ..$ : chr [1:2] "1:2" "1:3"
## $ p0         : Named num [1:3] 1 0 0
## ..- attr(*, "names")= chr [1:3] "(s0)" "1" "2"
## $ std.err    : num [1:251, 1:3] 0.00417 0.00588 0.00827 0.00922 0.00965 ...
## $ std.chaz   : num [1:251, 1:2] 0.00417 0.00592 0.00727 0.00787 0.00844 ...
## $ std.auc    : num [1:251, 1:3] 0 0.000321 0.000343 0.000356 0.000934 ...
## $ logse     : logi FALSE
## $ transitions : 'table' int [1:3, 1:3] 122 0 0 107 0 0 109 0 0
## ..- attr(*, "dimnames")=List of 2
## .. ..$ from: chr [1:3] "(s0)" "1" "2"
## .. ..$ to  : chr [1:3] "1" "2" "(censored)"
## $ lower      : num [1:251, 1:3] 0.986 0.977 0.96 0.953 0.949 ...
## ..- attr(*, "dimnames")=List of 2
## .. ..$ : NULL
## .. ..$ : chr [1:3] "(s0)" "1" "2"
## $ upper      : num [1:251, 1:3] 1 1 0.993 0.989 0.987 ...
## ..- attr(*, "dimnames")=List of 2
## .. ..$ : NULL
## .. ..$ : chr [1:3] "(s0)" "1" "2"
## $ conf.type  : chr "log"
## $ conf.int   : num 0.95
## $ states    : chr [1:3] "(s0)" "1" "2"
## $ type      : chr "mright"
## $ t0        : num 0
## $ call      : language survfit(formula = Surv(time, event, type = "mstate") ~ 1, data = orca)
```

```
## - attr(*, "class")= chr [1:2] "survfitms" "survfit"
```


- One could apply here the plot method of the survfit object to plot the cumulative incidences for each cause. However, we suggest that you use instead a simple function `plotCIF()` found in the `Epi` package. The main arguments are

data	data frame created by function <code>survfit()</code>
event	indicator for the event: values 1 or 2.

Other arguments are like in the ordinary `plot()` function.

- Draw two parallel plots describing the overall cumulative incidence curves for both causes of death

```
par(mfrow = c(1, 2))
plotCIF(cif1, 1, main = "Cancer death", xlab = "Time since oral cancer diagnosis (years)")
plotCIF(cif1, 2, main = "Other deaths", xlab = "Time since oral cancer diagnosis (years)")
```



- Compute the estimated cumulative incidences by stage for both causes of death. Now you have to add variable `stage` to `survfit`-function.

See the structure of the resulting object, in which you should observe `strata` variable containing the stage grouping variable. Plot the pertinent curves in two parallel graphs. Cut the *y*-axis for a more efficient graphical presentation

```
col5 <- col5
cif2 <- survfit(Surv(time, event, type = "mstate") ~ stage,
  data = orca
)
```

```
## Warning in Surv(time, event, type = "mstate"): type= 'mstate' is deprecated,
## use a factor variable as status
```


```
str(cif2)
```

```
## List of 24
## $ n          : int [1:5] 50 77 72 68 71
## $ time       : num [1:307] 0.17 0.498 0.665 0.832 1.166 ...
## $ n.risk     : num [1:307, 1:3] 50 49 48 47 46 45 44 43 42 41 ...
## ..- attr(*, "dimnames")=List of 2
## .. ..$ : NULL
## .. ..$ : chr [1:3] "(s0)" "1" "2"
## $ n.event    : num [1:307, 1:3] 0 0 0 0 0 0 0 0 0 0 ...
## ..- attr(*, "dimnames")=List of 2
## .. ..$ : NULL
## .. ..$ : chr [1:3] "(s0)" "1" "2"
## $ n.censor   : num [1:307, 1:3] 0 0 0 0 0 0 1 0 0 0 ...
## ..- attr(*, "dimnames")=List of 2
## .. ..$ : NULL
## .. ..$ : chr [1:3] "(s0)" "1" "2"
## $ pstate     : num [1:307, 1:3] 0.98 0.96 0.94 0.92 0.9 ...
## ..- attr(*, "dimnames")=List of 2
## .. ..$ : NULL
## .. ..$ : chr [1:3] "(s0)" "1" "2"
## $ n.transition: num [1:307, 1:2] 0 1 1 0 1 1 0 1 1 0 ...
## ..- attr(*, "dimnames")=List of 2
## .. ..$ : NULL
## .. ..$ : chr [1:2] "1:2" "1:3"
## $ n.id       : int [1:5] 50 77 72 68 71
## $ cumhaz     : num [1:307, 1:2] 0 0.0204 0.0412 0.0412 0.063 ...
## ..- attr(*, "dimnames")=List of 2
## .. ..$ : NULL
## .. ..$ : chr [1:2] "1:2" "1:3"
## $ p0         : num [1:5, 1:3] 1 1 1 1 0 0 0 0 0 ...
## ..- attr(*, "dimnames")=List of 2
## .. ..$ : chr [1:5] "stage=I" "stage=II" "stage=III" "stage=IV" ...
## .. ..$ : chr [1:3] "(s0)" "1" "2"
## $ strata     : Named int [1:5] 49 75 62 58 63
## ..- attr(*, "names")= chr [1:5] "stage=I" "stage=II" "stage=III" "stage=IV" ...
## $ std.err    : num [1:307, 1:3] 0.0198 0.0277 0.0336 0.0384 0.0424 ...
## $ std.chaz   : num [1:307, 1:2] 0 0.0202 0.0289 0.0289 0.036 ...
## $ std.auc    : num [1:307, 1:3] 0 0.00649 0.01028 0.01484 0.02589 ...
## $ logse     : logi FALSE
## $ transitions : 'table' int [1:3, 1:3] 122 0 0 107 0 0 109 0 0
## ..- attr(*, "dimnames")=List of 2
## .. ..$ from: chr [1:3] "(s0)" "1" "2"
## .. ..$ to  : chr [1:3] "1" "2" "(censored)"
## $ lower      : num [1:307, 1:3] 0.942 0.907 0.876 0.848 0.821 ...
## ..- attr(*, "dimnames")=List of 2
## .. ..$ : NULL
## .. ..$ : chr [1:3] "(s0)" "1" "2"
## $ upper      : num [1:307, 1:3] 1 1 1 0.998 0.987 ...
## ..- attr(*, "dimnames")=List of 2
## .. ..$ : NULL
## .. ..$ : chr [1:3] "(s0)" "1" "2"
## $ conf.type  : chr "log"
## $ conf.int   : num 0.95
## $ states     : chr [1:3] "(s0)" "1" "2"
## $ type       : chr "mright"
```

```
## $ t0          : num 0
## $ call        : language survfit(formula = Surv(time, event, type = "mstate") ~ stage, data = orca)
## - attr(*, "class")= chr [1:2] "survfitms" "survfit"

par(mfrow = c(1, 2))
plotCIF(cif2, 1,
  main = "Cancer death by stage",
  col = col5, ylim = c(0, 0.7)
)
plotCIF(cif2, 2,
  main = "Other deaths by stage",
  col = col5, ylim = c(0, 0.7)
)


legend(3, 0.7, title="stage", legend=levels(orca$stage), col = col5, lty=rep(1,5), cex=0.7)
```



Compare the two plots. What would you conclude about the effect of stage on the two causes of death?

- Using another function `stackedCIF()` in `Epi` you can put the two cumulative incidence curves in one graph but stacked upon one another such that the lower curve is for the cancer deaths and the upper curve is for total mortality, and the vertical difference between the two curves describes the cumulative mortality from other causes. You can also add some colours for the different zones:

```
par(mfrow=c(1,1),xaxs="i", yaxs="i") # make plot start 0,0
stackedCIF(cif1,xlim=c(0,20),
  col = c("black"),
  fill=c(cB8[6],cB8[8],cB8[2])) #choosing some colors
text( 10, 0.10, "Oral ca death ", pos = 4)
text( 10, 0.5, " Other death ", pos = 4)
text( 10, 0.80, " Alive ", pos = 4)
```



11.6 Regression modelling of overall mortality.

- Fit the semiparametric proportional hazards regression model, a.k.a. the Cox model, on all deaths including sex, age and stage as covariates. Use function `coxph()` in package `survival`. It is often useful to center and scale continuous covariates like `age` here. The estimated rate ratios and their confidence intervals can also here be displayed by applying `ci.lin()` on the fitted model object.

```
options(show.signif.stars = FALSE)
m1 <- coxph(suob ~ sex + I((age - 65) / 10) + stage, data = orca)
summary(m1)
```

Call:
coxph(formula = suob ~ sex + I((age - 65)/10) + stage, data = orca)

n= 338, number of events= 229
##

	coef	exp(coef)	se(coef)	z	Pr(> z)
## sexFemale	-0.3514	0.7037	0.1414	-2.49	0.01295
## I((age - 65)/10)	0.4160	1.5159	0.0564	7.37	1.6e-13
## stageII	0.0349	1.0355	0.2467	0.14	0.88742
## stageIII	0.3454	1.4126	0.2457	1.41	0.15971
## stageIV	0.8854	2.4240	0.2427	3.65	0.00026
## stageunkn	0.5844	1.7939	0.2512	2.33	0.02002

##

	exp(coef)	exp(-coef)	lower .95	upper .95
## sexFemale	0.704	1.421	0.533	0.928
## I((age - 65)/10)	1.516	0.660	1.357	1.693
## stageII	1.036	0.966	0.639	1.679
## stageIII	1.413	0.708	0.873	2.286

```
## stageIV          2.424      0.413      1.506      3.901
## stageunkn        1.794      0.557      1.096      2.935
##
## Concordance= 0.674 (se = 0.02 )
## Likelihood ratio test= 86.8 on 6 df, p=<2e-16
## Wald test          = 80.5 on 6 df, p=3e-15
## Score (logrank) test = 82.9 on 6 df, p=9e-16
round(ci.exp(m1),3)
```

```
##              exp(Est.)  2.5% 97.5%
## sexFemale      0.704 0.533 0.928
## I((age - 65)/10) 1.516 1.357 1.693
## stageII        1.036 0.639 1.679
## stageIII       1.413 0.873 2.286
## stageIV        2.424 1.506 3.901
## stageunkn      1.794 1.096 2.935
```

Look at the results. What are the main findings?

- Check whether the data are sufficiently consistent with the assumption of proportional hazards with respect to each of the variables separately as well as globally, using the `cox.zph()` function.

```
cox.zph(m1)
```

```
##              chisq df    p
## sex          0.275  1 0.60
## I((age - 65)/10) 1.706  1 0.19
## stage        3.648  4 0.46
## GLOBAL       5.144  6 0.53
```

- No evidence against proportionality assumption could apparently be found. Moreover, no difference can be observed between stages I and II in the estimates. On the other hand, the group with stage unknown is a complex mixture of patients from various true stages. Therefore, it may be prudent to exclude these subjects from the data and to pool the first two stage groups into one. After that fit a model in the reduced data with the new stage variable.

```
orca2 <- subset(orca, stage != "unkn")
orca2$st3 <- Relevel(orca2$stage, list(1:2, 3, 4:5))
levels(orca2$st3) <- c("I-II", "III", "IV")
m2 <- coxph(Surv(orca2$time, 1 * (orca2$event > 0)) ~ sex + I((age - 65) / 10) + st3, data = orca2)
summary(m2)
```

```
## Call:
## coxph(formula = Surv(orca2$time, 1 * (orca2$event > 0)) ~ sex +
##       I((age - 65)/10) + st3, data = orca2)
##
##      n= 267, number of events= 184
##
##              coef exp(coef) se(coef)      z Pr(>|z|)
## sexFemale      -0.2839    0.7528  0.1571 -1.81   0.071
## I((age - 65)/10)  0.3810    1.4637  0.0622  6.13 8.8e-10
## st3III          0.3117    1.3657  0.1827  1.71   0.088
## st3IV           0.8713    2.3900  0.1786  4.88 1.1e-06
##
##              exp(coef) exp(-coef) lower .95 upper .95
## sexFemale          0.753    1.328    0.553    1.02
## I((age - 65)/10)    1.464    0.683    1.296    1.65
## st3III              1.366    0.732    0.955    1.95
## st3IV               2.390    0.418    1.684    3.39
```

```
##
## Concordance= 0.673 (se = 0.022 )
## Likelihood ratio test= 64.8 on 4 df, p=3e-13
## Wald test = 62.3 on 4 df, p=1e-12
## Score (logrank) test = 64.3 on 4 df, p=4e-13
#m2 <- update(m1, . ~ . - stage + st3, data = orca2) #do not work
round(ci.exp(m2), 3)
```

```
##                exp(Est.)  2.5% 97.5%
## sexFemale         0.753 0.553  1.02
## I((age - 65)/10)   1.464 1.296  1.65
## st3III            1.366 0.955  1.95
## st3IV             2.390 1.684  3.39
```

- Plot the predicted cumulative mortality curves by stage, jointly stratified by sex and age, focusing only on 40 and 80 year old patients, respectively, based on the fitted model m2. You need to create a new artificial data frame containing the desired values for the covariates.

```
newd <- data.frame(
  sex = c(rep("Male", 6), rep("Female", 6)),
  age = rep(c(rep(40, 3), rep(80, 3)), 2),
  st3 = rep(levels(orca2$st3), 4)
)
newd
```

```
##      sex age st3
## 1   Male  40 I-II
## 2   Male  40 III
## 3   Male  40 IV
## 4   Male  80 I-II
## 5   Male  80 III
## 6   Male  80 IV
## 7 Female  40 I-II
## 8 Female  40 III
## 9 Female  40 IV
## 10 Female 80 I-II
## 11 Female 80 III
## 12 Female 80 IV
```

```
col3 <- cB8[1:3] #pre-setting color palette
leg<-levels(interaction(levels(factor(orca2$sex)),levels(orca2$st3))) #legend labels by sex and stage
```

```
## Warning in ay * as.numeric(nx) + ax: longer object length is not a multiple of
## shorter object length
```


```
par(mfrow = c(1, 2))
plot(
  survfit(m2, newdata = subset(newd, sex == "Male" & age == 40)),
  col = col3,
  lty= 1,
  fun = "event", mark.time = FALSE,
  main="Cum. mortality for M and F \n age 40"
)
lines(
  survfit(m2, newdata = subset(newd, sex == "Female" & age == 40)),
  col = col3, lty=c(2,2,2),fun = "event", mark.time = FALSE
)
legend(0, 0.95, title="stage",legend=leg,
  col = c(col3[1],col3[1],col3[2],col3[2],col3[3],col3[3]),
```

```


, lty=c(2,1,2,1,2,1), cex=0.6)
plot(
  survfit(m2, newdata = subset(newd, sex == "Male" & age == 80)),
  ylim = c(0, 1), col = col3, fun = "event", mark.time = FALSE,
  main="Cum. mortality for M and F \n age 80"
)
lines(
  survfit(m2, newdata = subset(newd, sex == "Female" & age == 80) ),
  col = col3, fun = "event", lty = 2, mark.time = FALSE
)
legend(10, 0.5, title="stage", legend=leg,
  col = c(col3[1], col3[1], col3[2], col3[2], col3[3], col3[3]),
  , lty=c(2,1,2,1,2,1), cex=0.7)

```

**Cum. mortality for M and F
age 40**



**Cum. mortality for M and F
age 80**



11.7 Modelling event-specific hazards

- Fit the Cox model for the cause-specific hazard of cancer deaths with the same covariates as above. In this case only cancer deaths are counted as events and deaths from other causes are included into censorings.

```

m2haz1 <-
  coxph(
    Surv(time, event == 1) ~ sex + I((age - 65) / 10) + st3,
    data = orca2
  )
round(ci.exp(m2haz1), 4)

```

```

##                exp(Est.)  2.5% 97.5%
## sexFemale        0.983 0.642  1.51

```

```
## I((age - 65)/10)      1.426 1.204  1.69
## st3III               1.514 0.901  2.54
## st3IV                3.181 1.985  5.10
```

```
cox.zph(m2haz1)
```

```
##                chisq df      p
## sex            0.0724  1 0.788
## I((age - 65)/10) 5.7620  1 0.016
## st3            3.4195  2 0.181
## GLOBAL         9.1216  4 0.058
```

Compare the results with those of model `m2`. What are the major differences?

- Fit a similar model for deaths from other causes and compare the results.

```
m2haz2 <-
  coxph(
    Surv(time, event == 2) ~ sex + I((age - 65) / 10) + st3,
    data = orca2
  )
round(ci.exp(m2haz2), 4)
```

```
##                exp(Est.)  2.5% 97.5%
## sexFemale              0.552 0.352 0.868
## I((age - 65)/10)       1.488 1.249 1.772
## st3III                 1.230 0.749 2.021
## st3IV                  1.641 0.952 2.827
```

```
cox.zph(m2haz2)
```

```
##                chisq df      p
## sex            0.4071  1 0.52
## I((age - 65)/10) 0.2435  1 0.62
## st3            0.0443  2 0.98
## GLOBAL         0.6444  4 0.96
```

11.8 Lexis object with multi-state set-up

Before entering to multi-state analyses, it might be instructive to apply some Lexis tools to illustrate the competing-risks set-up. More detailed explanation of these tools will be given by Bendix later.

- Form a Lexis object from the data frame and print a summary of it. We shall name the main (and only) time axis in this object as `stime`.

```
orca.lex <- Lexis(
  exit = list(stime = time),
  exit.status = factor(event,
    labels = c("Alive", "Oral ca. death", "Other death")
  ),
  data = orca
)
```

```
## NOTE: entry.status has been set to "Alive" for all.
## NOTE: entry is assumed to be 0 on the stime timescale.
```


```
summary(orca.lex)
```

```
##
## Transitions:
##      To
```

```
## From      Alive Oral ca. death Other death Records: Events: Risk time: Persons:
##   Alive    109              122         107      338      229      1914      338
```

- Draw a box diagram of the two-state set-up of competing transitions. Run first the following command line

```
boxes(orca.lex, boxpos = TRUE)
```



Now, move the cursor to the point in the graphics window, at which you wish to put the box for *Alive*, and click. Next, move the cursor to the point at which you wish to have the box for *Oral ca. death*, and click. Finally, do the same with the box for *Other death*. If you are not happy with the outcome, run the command line again and repeat the necessary mouse moves and clicks.

Chapter 12

Time-splitting, time-scales and SMR

This exercise is about mortality among Danish Diabetes patients. It is based on the dataset `DMLate`, a random sample of 10,000 patients from the Danish Diabetes Register (scrambled dates), all with date of diagnosis after 1994.

Start by loading the relevant packages:

```
library(Epi)
library(popEpi)
library(mgcv, quietly = TRUE)
```

This is `mgcv` 1.9-4. For overview type `'?mgcv'`.

```
library(tidyverse, quietly = TRUE)
```

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr      1.2.1      v readr      2.2.0
v forcats    1.0.1      v stringr    1.6.0
v ggplot2    4.0.3      v tibble     3.3.1
v lubridate  1.9.5      v tidyr      1.3.2
v purrr      1.2.2

-- Conflicts ----- tidyverse_conflicts() --
x dplyr::collapse()   masks nlme::collapse()
x dplyr::filter()     masks stats::filter()
x lubridate::is.Date() masks popEpi::is.Date()
x dplyr::lag()         masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

Then load the data and take a look at the data:

```
data(DMLate)
str(DMLate)
```

```
'data.frame':  10000 obs. of  7 variables:
 $ sex   : Factor w/ 2 levels "M","F": 2 1 2 2 1 2 1 1 2 1 ...
 $ dobth: num  1940 1939 1918 1965 1933 ...
 $ dodm  : num  1999 2003 2005 2009 2009 ...
 $ dodth: num  NA NA NA NA NA ...
 $ dooad: num  NA 2007 NA NA NA ...
 $ doins: num  NA NA NA NA NA NA NA NA NA NA ...
 $ dox   : num  2010 2010 2010 2010 2010 ...
```

You can get a more detailed explanation of the data by referring to the help page:

?DMLate

1. Set up the dataset as a `Lexis` object with age, calendar time and duration of diabetes as timescales, and date of death as event. Make sure that you know what each of the arguments to `Lexis` mean:

```
LL <- Lexis(entry = list(A = dodm - dobth,
                        P = dodm,
                        dur = 0),
            exit = list(P = dox),
            exit.status = factor(!is.na(dodth),
                                labels = c("Alive", "Dead")),
            data = DMLate)
```

NOTE: `entry.status` has been set to "Alive" for all.

NOTE: Dropping 4 rows with duration of follow up < tol

```
summary(LL)
```

Transitions:

```

      To
From   Alive Dead Records: Events: Risk time: Persons:
  Alive 7497 2499      9996      2499  54273.27      9996
```

Take a look at the first few lines of the resulting dataset, for example using `head()`.

2. Get an overview of the mortality by using `stat.table` to tabulate no. deaths, person-years (`lex.dur`) and the crude mortality rate by sex. Try:

```
stat.table(sex,
            list(D = sum(lex.Xst == "Dead"),
                 Y = sum(lex.dur),
                 rate = ratio(lex.Xst == "Dead",
                              lex.dur,
                              1000)),
            margins = TRUE,
            data = LL)
```

```
-----
sex          D          Y    rate
-----
M          1343.00 27614.21   48.63
F          1156.00 26659.05   43.36

Total      2499.00 54273.27   46.04
-----
```

a less versatile alternative is xtabs:

```
xtabs(cbind(D = lex.Xst == "Dead",
            Y = lex.dur)
      ~ sex,
      data = LL)
```

```
sex      D      Y
M 1343.00 27614.21
F 1156.00 26659.05
```

3. If we want to assess how mortality depends on age, calendar time and duration or how it relates to population mortality, we should in principle split the follow-up along all three time scales. In practice it is sufficient to split it along one of the time-scales and then use the value of each of the time-scales at the

left endpoint of the intervals.

Use `splitLexis` (or `splitMulti` from the `popEpi` package) to split the follow-up along the age-axis in suitable intervals (here set to 1/2 year, but really immaterial as long as it is small):

```
SL <- splitLexis(LL,
  breaks = seq(0, 125, 1 / 2),
  time.scale = "A")
summary(LL)
```

Transitions:

```
      To
From   Alive Dead Records: Events: Risk time: Persons:
  Alive 7497 2499      9996      2499  54273.27      9996
```

```
summary(SL)
```

Transitions:

```
      To
From   Alive Dead Records: Events: Risk time: Persons:
  Alive 115974 2499  118473      2499  54273.27      9996
```

How many records are now in the dataset? How many person-years? Compare to the original `Lexis`-dataset.

12.1 Age-specific mortality

1. Now estimate age-specific mortality curves for men and women separately, using splines as implemented in `gam`. We use `k = 20` to be sure to catch any irregularities by age.

```
r.m <- mgcv::gam(cbind(lex.Xst == "Dead", lex.dur) ~ s(A, k = 10),
  family = poisreg,
  data = subset(SL, sex == "M"))
summary(r.m)
```

Family: poisson

Link function: log

Formula:

```
cbind(lex.Xst == "Dead", lex.dur) ~ s(A, k = 10)
```

Parametric coefficients:

```
      Estimate Std. Error z value Pr(>|z|)
(Intercept) -3.53023      0.04998  -70.64  <2e-16
```

Approximate significance of smooth terms:

```
      edf Ref.df Chi.sq p-value
s(A)  4.341  5.326  1017  <2e-16
```

```
R-sq.(adj) = 0.00313 Deviance explained = 9.13%
```

```
UBRE = -0.80407 Scale est. = 1 n = 60347
```

Make sure you understand all the components on this modeling statement. Fit the same model for women.

There is a convenient wrapper for this, exploiting the `Lexis` structure of data, but which does not have an update

```
r.m <- gamLexis(subset(SL, sex == "M"), ~ s(A, k = 10))
```

```
mgcv::gam Poisson analysis of Lexis object subset(SL, sex == "M") with log link:
Rates for the transition:
Alive->Dead
```

```
r.f <- gamLexis(subset(SL, sex == "F"), ~ s(A, k = 10))
```

```
mgcv::gam Poisson analysis of Lexis object subset(SL, sex == "F") with log link:
Rates for the transition:
Alive->Dead
```

2. Now, extract the estimated rates by using the wrapper function `ci.pred` that computes predicted rates and confidence limits for these.


`glm.Lexis` and `gamLexis` use the `poisreg` family that will return the rates in the (inverse) units in which the person-years were given; that is the units in which `lex.dur` is recorded—in this case 1 year.

```
nd <- data.frame(A = seq(20, 90, 0.5))
p.m <- ci.pred(r.m, newdata = nd)
p.f <- ci.pred(r.f, newdata = nd)
str(p.m)
```

```
num [1:141, 1:3] 0.00132 0.00137 0.00142 0.00146 0.00152 ...
- attr(*, "dimnames")=List of 2
..$ : chr [1:141] "1" "2" "3" "4" ...
..$ : chr [1:3] "Estimate" "2.5%" "97.5%"
```


3. Plot the predicted rates for men and women together - using for example `matplot` or `matshade`.

```
par(mar = c(3.5,3.5,1,1),
    mgp = c(3,1,0) / 1.6,
    las = 1,
    bty = "n",
    lend = "butt")
matplot(nd$A, cbind(p.m, p.f) * 1000,
        type = "l",
        col = rep(c("blue", "red"), each = 3),
        lwd = c(3, 1, 1),
        lty = 1,
        log = "y", yaxt = "n",
        xlab = "Age",
        ylab = "Mortality per 1000 PY")
axis(side = 2,
     at = ll <- outer( c(5, 10, 20), -1:1, function(x,y) x * 10^y),
     labels = ll)
```



The `axis` statement is a trick to get the labels on the y -axis without “.0”s. `matshade` assumes that curves are added to an existing plot, so the `plot=TRUE` must be supplied to create a new plot:

```
par(mar = c(3.5,3.5,1,1),
    mgp = c(3,1,0) / 1.6,
    las = 1,
    bty = "n",
    lend = "butt")
matshade(nd$A, cbind(p.m, p.f) * 1000, plot = TRUE,
        type = "l",
        col = c("blue", "red"),
        lwd = 3,
        lty = 1,
        log = "y", yaxt = "n",
        xlab = "Age",
        ylab = "Mortality per 1000 PY")
axis(side = 2,
     at = ll <- outer( c(5, 10, 20), -1:1, function(x,y) x * 10^y),
     labels = ll)
```



12.2 Further time scales: period and duration

4. We now want to model the mortality rates among diabetes patients also including current date and duration of diabetes, using penalized splines. Use the argument `bs = "cr"` to `s()` to get cubic splines instead of thin plate ("`tp`") splines which is the default.

As before specify the model exploiting the `Lexis` class of the dataset, try:

```
Mcr <- gamLexis(subset(SL, sex == "M"),
  ~ s(A, bs = "cr", k = 10) +
    s(P, bs = "cr", k = 10) +
    s(dur, bs = "cr", k = 10))
```

mgcv::gam Poisson analysis of Lexis object subset(SL, sex == "M") with log link:

Rates for the transition:

Alive->Dead

```
summary(Mcr)
```

Family: poisson

Link function: log

Formula:

```
cbind(trt(Lx$lex.Cst, Lx$lex.Xst) %in% trnam, Lx$lex.dur) ~ s(A,
  bs = "cr", k = 10) + s(P, bs = "cr", k = 10) + s(dur, bs = "cr",
  k = 10)
```

Parametric coefficients:

```
Estimate Std. Error z value Pr(>|z|)
```

```
(Intercept) -3.54074    0.04938   -71.7    <2e-16
```

Approximate significance of smooth terms:

```
      edf Ref.df Chi.sq p-value
s(A)   3.645  4.517 1013.20 < 2e-16
s(P)   1.024  1.048   17.58 3.48e-05
s(dur) 7.586  8.384   74.46 < 2e-16
```

```
R-sq.(adj) = 0.00333    Deviance explained = 9.87%
UBRE = -0.8054    Scale est. = 1          n = 60347
```

Fit the same model for women as well. Are the models reasonably fitting?

```
Fcr <- gamLexis(subset(SL, sex == "F"),
  ~ s(A, bs = "cr", k = 10) +
    s(P, bs = "cr", k = 10) +
    s(dur, bs = "cr", k = 10))
```

mgcv::gam Poisson analysis of Lexis object subset(SL, sex == "F") with log link:

Rates for the transition:

Alive->Dead

```
summary(Fcr)
```

Family: poisson

Link function: log

Formula:

```
cbind(trt(Lx$lex.Cst, Lx$lex.Xst) %in% trnam, Lx$lex.dur) ~ s(A,
  bs = "cr", k = 10) + s(P, bs = "cr", k = 10) + s(dur, bs = "cr",
  k = 10)
```

Parametric coefficients:

```
      Estimate Std. Error z value Pr(>|z|)
(Intercept) -3.78483    0.05808  -65.17    <2e-16
```


Approximate significance of smooth terms:

```
      edf Ref.df Chi.sq p-value
s(A)   2.667  3.366  988.49 < 2e-16
s(P)   1.904  2.391   20.08 0.000136
s(dur) 5.973  6.972   38.98 < 2e-16
```

```
R-sq.(adj) = 0.00417    Deviance explained = 11.1%
UBRE = -0.82405    Scale est. = 1          n = 58126
```

- Plot the estimated effects, using the default plot method for `gam` objects. Remember that there are three effects estimated, so it is useful set up a multi-panel display, and for the sake of comparability to set `ylim` to the same for men and women:

```
par(mfcol = c(3, 2))
plot(Mcr, ylim = c(-3, 3), col = "blue", lwd = 2)
plot(Fcr, ylim = c(-3, 3), col = "red", lwd = 2)
```



What is the absolute scale for these effects?

6. Compare the fit of the naive model with just age and the three-factor models, using `anova`, e.g.:

```
anova(Mcr, r.m, test = "Chisq")
```

Analysis of Deviance Table

```
Model 1: cbind(trt(Lx$lex.Cst, Lx$lex.Xst) %in% trnam, Lx$lex.dur) ~ s(A,
  bs = "cr", k = 10) + s(P, bs = "cr", k = 10) + s(dur, bs = "cr",
  k = 10)
```

```
Model 2: cbind(trt(Lx$lex.Cst, Lx$lex.Xst) %in% trnam, Lx$lex.dur) ~ s(A,
  k = 10)
```

	Resid. Df	Resid. Dev	Df	Deviance	Pr(>Chi)
1	60332	11717			
2	60341	11813	-8.6238	-95.935	< 2.2e-16

```
anova(Fcr, r.f, test = "Chisq")
```

Analysis of Deviance Table

```
Model 1: cbind(trt(Lx$lex.Cst, Lx$lex.Xst) %in% trnam, Lx$lex.dur) ~ s(A,
  bs = "cr", k = 10) + s(P, bs = "cr", k = 10) + s(dur, bs = "cr",
  k = 10)
```

```
Model 2: cbind(trt(Lx$lex.Cst, Lx$lex.Xst) %in% trnam, Lx$lex.dur) ~ s(A,
  k = 10)
```

	Resid. Df	Resid. Dev	Df	Deviance	Pr(>Chi)
1	58112	10204			
2	58122	10268	-9.4572	-63.366	4.812e-10

What do you conclude?

7. The model we fitted has three time-scales: current age, current date and current duration of diabetes,

so the effects that we report are not immediately interpretable, as they are (as in any kind of multiple regressions) to be interpreted as *all else equal* which they are not, as the three time scales advance simultaneously at the same pace.

The reporting should therefore more naturally be showing predicted rates on the mortality scale as a function of age, but showing the mortality for persons diagnosed in different ages, using separate displays for separate years of diagnosis.

This is most easily done using the `ci.pred` function with the `newdata =` argument. So a person diagnosed in age 50 in 1995 will have a mortality measured in cases per 1000 PY as:

```
pts <- seq(0, 12, 1/4) # times since entry (date of diabetes)
nd <- data.frame(A = 50 + pts,
                 P = 1995 + pts,
                 dur = pts)
m.pr <- ci.pred(Mcr, newdata = nd)
```

Note that because we used `gamLexis` which uses the `poisreg` family we need not specify `lex.dur` as a variable in the prediction data frame `nd`. Predictions will be rates in the same units as `lex.dur` (well, the inverse).

Now take a look at the result from the `ci.pred` statement and construct prediction of mortality for men and women diagnosed in a range of ages, say 50, 60, 70, and plot these together in the same graph:

```
cbind(nd, ci.pred(Mcr, newdata = nd))[1:10,]
```

8. From the figure with the three effects it seems that the duration effect is over-modeled, so refit constraining the d.f. to 5:

```
Mcr <- gamLexis(subset(SL, sex == "M"),
               ~ s(A, bs = "cr", k = 10) +
                 s(P, bs = "cr", k = 10) +
                 s(dur, bs = "cr", k = 5))
```


```
mgcv::gam Poisson analysis of Lexis object subset(SL, sex == "M") with log link:
Rates for the transition:
Alive->Dead
```

```
Fcr <- gamLexis(subset(SL, sex == "F"),
               ~ s(A, bs = "cr", k = 10) +
                 s(P, bs = "cr", k = 10) +
                 s(dur, bs = "cr", k = 5))
```

```
mgcv::gam Poisson analysis of Lexis object subset(SL, sex == "F") with log link:
Rates for the transition:
Alive->Dead
```

Plot the estimated rates from the revised models.

```
par(mfcol = c(3, 2))
plot(Mcr, ylim = c(-3, 3), col = "blue", lwd = 2)
plot(Fcr, ylim = c(-3, 3), col = "red", lwd = 2)
```



What do you conclude from the plots?

12.3 SMR

The SMR is the **Standardized Mortality Ratio**, which is the mortality rate-ratio between the diabetes patients and the general population (that is, persons without diabetes). In real studies we would subtract the deaths and the person-years among the diabetes patients from those of the general population, but since we do not have access to these, we make the comparison to the general population at large, i.e. also including the diabetes patients.

So we now want to include the population mortality rates as a fixed variable in the split dataset; for each record in the split dataset we attach the value of the population mortality for the relevant sex, and and calendar time.

This can be achieved in two ways: Either we just use the current split of follow-up time and allocate the population mortality rates for some suitably chosen (mid-)point of the follow-up in each, or we make a second split by date, so that follow-up in the diabetes patients is in the same classification of age and data as the population mortality table.

12. We will use the former approach, using the dataset split in 6 month intervals, and then include as an extra variable the population mortality as available from the data set `M.dk`.

First create the variables in the diabetes dataset that we need for matching with the population mortality data, that is sex and age and date at the midpoint of each of the intervals (or rather at a point 3 months after the left endpoint of the interval; recall we split the follow-up in 6 month intervals).

We need to have variables of the same type when we merge, so we must transform the sex variable in `M.dk` to a factor, and must for each follow-up interval in the SL data have an age and a period variable that can be used in merging with the population data.

```
str(SL)
```

```
Classes 'Lexis' and 'data.frame':  118473 obs. of  14 variables:
 $ lex.id : int  1 1 1 1 1 1 1 1 1 1 ...
```

```

$ A      : num  58.7 59 59.5 60 60.5 ...
$ P      : num  1999 1999 2000 2000 2001 ...
$ dur    : num  0 0.339 0.839 1.339 1.839 ...
$ lex.dur: num  0.339 0.5 0.5 0.5 0.5 ...
$ lex.Cst: Factor w/ 2 levels "Alive","Dead": 1 1 1 1 1 1 1 1 1 ...
$ lex.Xst: Factor w/ 2 levels "Alive","Dead": 1 1 1 1 1 1 1 1 1 ...
$ sex    : Factor w/ 2 levels "M","F": 2 2 2 2 2 2 2 2 2 ...
$ dobth  : num  1940 1940 1940 1940 1940 ...
$ dodm   : num  1999 1999 1999 1999 1999 ...
$ dodth  : num  NA NA NA NA NA NA NA NA NA NA ...
$ dooad  : num  NA NA NA NA NA NA NA NA NA NA ...
$ doins  : num  NA NA NA NA NA NA NA NA NA NA ...
$ dox    : num  2010 2010 2010 2010 2010 ...
- attr(*, "breaks")=List of 3
..$ A   : num [1:251] 0 0.5 1 1.5 2 2.5 3 3.5 4 4.5 ...
..$ P   : NULL
..$ dur : NULL
- attr(*, "time.scales")= chr [1:3] "A" "P" "dur"
- attr(*, "time.since")= chr [1:3] "" "" ""

```

```

SL$Am <- floor(SL$A + 0.25) # data were split in .5 year intervals
SL$Pm <- floor(SL$P + 0.25)
data(M.dk)
str(M.dk)

```

```

'data.frame': 7800 obs. of 6 variables:
 $ A      : num  0 0 0 0 0 0 0 0 0 0 ...
 $ sex    : num  1 2 1 2 1 2 1 2 1 2 ...
 $ P      : num  1974 1974 1975 1975 1976 ...
 $ D      : num  459 303 435 311 405 258 332 205 312 233 ...
 $ Y      : num  35963 34383 36099 34652 34965 ...
 $ rate   : num  12.76 8.81 12.05 8.97 11.58 ...
- attr(*, "Contents")= chr "Number of deaths and risk time in Denmark"

```

```

M.dk <- transform(M.dk,
                  Am = A,
                  Pm = P,
                  sex = factor(sex, labels = c("M", "F")))
str(M.dk)

```

```

'data.frame': 7800 obs. of 8 variables:
 $ A      : num  0 0 0 0 0 0 0 0 0 0 ...
 $ sex    : Factor w/ 2 levels "M","F": 1 2 1 2 1 2 1 2 1 2 ...
 $ P      : num  1974 1974 1975 1975 1976 ...
 $ D      : num  459 303 435 311 405 258 332 205 312 233 ...
 $ Y      : num  35963 34383 36099 34652 34965 ...
 $ rate   : num  12.76 8.81 12.05 8.97 11.58 ...
 $ Am     : num  0 0 0 0 0 0 0 0 0 0 ...
 $ Pm     : num  1974 1974 1975 1975 1976 ...

```

Then match the rates from M.dk into SL; sex, Am and Pm are the common variables, and therefore the match is on these variables:

```

SLr <- merge(SL,
             M.dk[, c("sex", "Am", "Pm", "rate")])
dim(SLr)

```

```
dim(SLr)
```

```
[1] 118454      17
```

This merge (remember to `?merge!`) only takes rows that have information from both datasets, hence the slightly fewer rows in `SLr` than in `SL`.

- Compute the expected number of deaths as the person-time multiplied (`lex.dur`) by the corresponding population rate, and put it in a new variable, `E`, say (`Expected`). Use `stat.table` to make a table of observed, expected and the ratio (SMR) by age (suitably grouped, look for `cut`) and sex.
13. Fit a poisson model with sex as the explanatory variable and `E` as denominator to derive the SMR (and c.i.). Some of the population mortality rates are 0, so you need to exclude those records from the analysis.

```
mmsr <- glm(cbind((lex.Xst == "Dead"), E) ~ sex - 1,
            family = poisreg,
            data = subset(SLr, E > 0))
ci.exp(mmsr)
```

```
      exp(Est.)      2.5%      97.5%
sexM  1.685699  1.597881  1.778344
sexF  1.541922  1.455441  1.633541
```

Do you recognize the numbers?

- We can assess the ratios of SMRs between men and women by using the `ctr.mat` argument which should be a matrix:

```
(CM <- rbind(M = c(1, 0),
             W = c(0, 1),
             "M/F" = c(1, -1)))
```

```
      [,1] [,2]
M         1   0
W         0   1
M/F       1  -1
```

```
round(ci.exp(mmsr, ctr.mat = CM), 2)
```

```
      exp(Est.) 2.5% 97.5%
M         1.69 1.60  1.78
W         1.54 1.46  1.63
M/F       1.09 1.01  1.18
```

What do you conclude about the SMRs among men and women?

What do you conclude about the mortality rates among men and women?

12.4 SMR modeling

14. Now model the SMR using age and date of diagnosis and diabetes duration as explanatory variables, including the expected-number instead of the person-years, using separate models for men and women.

Note that you cannot use `gamLexis` from the code you used for fitting models for the rates, you need to use `gam` with the `poisreg` family. And remember to exclude those units where no deaths in the population occur (that is where the population mortality rate is 0).

```
Mmsr <- gam(cbind(lex.Xst == "Dead", E)
            ~ s( A, bs = "cr", k = 5) +
              s( P, bs = "cr", k = 5) +
              s(dur, bs = "cr", k = 5),
            family = poisreg,
```


```
data = subset(SLr, E > 0 & sex == "M"))
ci.exp(Msmr)
```

	exp(Est.)	2.5%	97.5%
(Intercept)	2.1667053	2.0001046	2.3471833
s(A).1	0.5960527	0.5181962	0.6856068
s(A).2	0.8927331	0.8650901	0.9212594
s(A).3	0.3663506	0.2818154	0.4762435
s(A).4	0.4587803	0.3763049	0.5593320
s(P).1	0.9893048	0.9732835	1.0055898
s(P).2	0.9608720	0.9085181	1.0162428
s(P).3	0.9027294	0.7825209	1.0414040
s(P).4	0.9173557	0.8132938	1.0347323
s(dur).1	0.6581828	0.5769629	0.7508362
s(dur).2	0.8446645	0.7613599	0.9370838
s(dur).3	0.7830215	0.6916313	0.8864877
s(dur).4	1.8176744	1.1041338	2.9923368

```
Fsmr <- update(Msmr, data = subset(SLr, E > 0 & sex == "F"))
```

Plot the estimated smooth effects for both men and women using e.g. `plot.gam`. What do you see?

```
par(mfcol = c(3, 2))
plot(Msmr, ylim = c(-3, 3), col = "blue", lwd = 2)
plot(Fsmr, ylim = c(-3, 3), col = "red", lwd = 2)
```



15. Plot the predicted SMRs from the models for men and women diagnosed in ages 50, 60 and 70 as you did for the rates. What do you see?

```
par(mfrow = c(1,1))
n50 <- nd
```


```

n60 <- mutate(n50, A = A + 10)
n70 <- mutate(n50, A = A + 20)
head(n70)

      A      P dur
1 70.00 1995.00 0.00
2 70.25 1995.25 0.25
3 70.50 1995.50 0.50
4 70.75 1995.75 0.75
5 71.00 1996.00 1.00
6 71.25 1996.25 1.25

matshade(n50$A, cbind(ci.pred(Msmr, n50),
                      ci.pred(Fsmr, n50)), plot = TRUE,
        col = c("blue", "red"), lwd = 3,
        ylim = c(0.5, 5), log = "y", xlim = c(50, 80))
matshade(n60$A, cbind(ci.pred(Msmr, n60),
                      ci.pred(Fsmr, n60)),
        col = c("blue", "red"), lwd = 3)
matshade(n70$A, cbind(ci.pred(Msmr, n70),
                      ci.pred(Fsmr, n70)),
        col = c("blue", "red"), lwd = 3)
abline(h = 1)
abline(v = 50 + 0:5, lty = 3, col = "gray")

```



Describe the shapes of the curves. What aspects of the shapes are induced by the model ?

16. Try to simplify the model to one with a simple sex effect, separate linear effects of age and date of follow-up for each sex, and a smooth effect of duration common for both sexes, giving an estimate of the change in SMR by age and calendar time.

```

Bsmr <- gam(cbind(lex.Xst == "Dead", E)
  ~ sex / A +
    sex / P +
    s(dur, bs = "cr", k = 5),
  family = poisreg,
  data = subset(SLr, E > 0))
round(ci.exp(Bsmr)[-1,], 3)

```

	exp(Est.)	2.5%	97.5%
sexF	52149.525	0.000	2.375788e+23
sexM:A	0.981	0.977	9.860000e-01
sexF:A	0.980	0.975	9.860000e-01
sexM:P	0.987	0.972	1.002000e+00
sexF:P	0.981	0.965	9.980000e-01
s(dur).1	0.663	0.601	7.320000e-01
s(dur).2	0.861	0.795	9.310000e-01
s(dur).3	0.885	0.805	9.740000e-01
s(dur).4	1.412	0.945	2.111000e+00

How much does SMR change by each year of age? And by each calendar year?


What is the meaning of the `sexF` parameter?

17. Use your previous code to plot the predicted mortality from this model too. Are the predicted SMR curves credible?

```

m50 <- mutate(n50, sex = "M")
f50 <- mutate(n50, sex = "F")
m60 <- mutate(m50, A = A + 10)
f60 <- mutate(f50, A = A + 10)
m70 <- mutate(m50, A = A + 20)
f70 <- mutate(f50, A = A + 20)
matshade(n50$A, cbind(ci.pred(Bsmr, m50),
  ci.pred(Bsmr, f50)), plot = TRUE,
  col = c("blue", "red"), lwd = 3,
  ylim = c(0.5, 5), log = "y", xlim = c(50, 80))
matshade(n60$A, cbind(ci.pred(Bsmr, m60),
  ci.pred(Bsmr, f60)),
  col = c("blue", "red"), lwd = 3)
matshade(n70$A, cbind(ci.pred(Bsmr, m70),
  ci.pred(Bsmr, f70)),
  col = c("blue", "red"), lwd = 3)
abline(h = 1)
abline(h = 1:5, lty = 3, col = "gray")

```



What is your conclusion about SMR for diabetes patients relative to the general popuation?

Chapter 13

Nested case-control study and case-cohort study: Risk factors of coronary heart disease

In this exercise we shall apply both the nested case-control (NCC) design and the case-cohort (CC) design in sampling control subjects from a defined cohort or closed study population. The case group comprises those cohort members who die from coronary heart disease (CHD) during a > 20 years follow-up of the cohort. The risk factors of interest are cigarette smoking, systolic blood pressure, and total cholesterol level.

Our study population is an occupational cohort comprising 1501 men working in blue-collar jobs in one Nordic country. Eligible subjects had no history of coronary heart disease when recruited to the study in the early 1990s. Smoking habits and many other items were inquired at baseline by a questionnaire, and blood pressure was measured by a research nurse, the values being written down on the questionnaire. Serum samples were also taken from the cohort members at the same time and were stored in a freezer. For some reason, the data in the questionnaires were not entered to any computer file, but the questionnaires were kept in a safe storehouse for further purposes. Also, no biochemical analyses were initially performed for the sera collected from the participants. However, dates of birth and dates of entry to the study were recorded in an electronic file.

In 2010 the study was suddenly reactivated by those investigators of the original team who were still alive then. As the first step mortality follow-up of the cohort members was executed by record linkage to the national population register, from which the dates of death and emigration were obtained. Another linkage was performed with the national register of causes of death in order to get the deaths from coronary heart disease identified. As a result a data file `occoh.txt` was completed containing the following variables:

Variable	Description
<code>id</code>	identification number
<code>birth</code>	date of birth
<code>entry</code>	date of recruitment and baseline measurements
<code>exit</code>	date of exit from mortality follow-up
<code>death</code>	indicator for vital status at the end of follow-up; 1, if dead from any cause, and = 0, if alive
<code>chdeath</code>	indicator for death from coronary heart disease; 1, if <i>yes</i> , and 0, if <i>no</i> .

This exercise is divided into five main parts:

1. Description of the study base or the follow-up experience of the whole cohort, identification of the cases and illustrating the risk sets.
2. Nested case-control study within the cohort:
 - (i) selection of controls by risk set or time-matched sampling using function `ccwc()` in package `Epi`,

- (ii) collection of exposure data for cases and controls from the pertinent data base of the whole cohort to the case-control data set using function `merge()`, and
 - (iii) analysis of the case-control data set with stratified Cox model using function `clogit()` in package `survival()`,
3. Case-cohort study within the cohort:
 - (i) selection of a subcohort by simple random sampling from the cohort,
 - (ii) collection of exposure data for subcohort members and cases, and
 - (iii) analysis of the case-cohort data set with Cox model by weighted partial likelihood including appropriate weighting and correction of estimated covariance matrix for the model coefficients using function `cch()` in package `survival()`.
 4. Comparison of results from all previous analyses, also with those from a full cohort design.
 5. Further tasks and homework.

13.1 Reading the cohort data, illustrating the study base and risk sets

- Load the packages `Epi` and `survival`. Read in the cohort data file and name the resulting data frame as `oc`. See its structure and print the univariate summaries.

```
library(Epi)
library(survival)
url <- "https://raw.githubusercontent.com/SPE-R/SPE/master/pracs-book/data"
oc <- read.table(paste(url, "occoh.txt", sep = "/"), header = TRUE)
str(oc)
```

```
## 'data.frame': 1501 obs. of 6 variables:
## $ id : int 1 2 3 4 5 6 7 8 9 10 ...
## $ birth : chr "1943-02-19" "1934-07-06" "1939-03-05" "1939-07-03" ...
## $ entry : chr "1990-08-14" "1990-08-14" "1990-08-14" "1990-08-14" ...
## $ exit : chr "2009-12-31" "2009-12-31" "2009-12-31" "2009-12-31" ...
## $ death : int 0 0 0 0 1 1 1 1 0 0 ...
## $ chdeath: int 0 0 0 0 0 0 0 1 0 0 ...
```

```
summary(oc)
```

```
##      id      birth      entry      exit
## Min.   : 1   Length :1501   Length :1501   Length :1501
## 1st Qu.:376   N.unique :1349   N.unique : 302   N.unique : 290
## Median :751   N.blank  : 0   N.blank  : 0   N.blank  : 0
## Mean   :751   Min.nchar: 10   Min.nchar: 10   Min.nchar: 10
## 3rd Qu.:1126   Max.nchar: 10   Max.nchar: 10   Max.nchar: 10
## Max.   :1501
##      death      chdeath
## Min.   :0.0000   Min.   :0.00000
## 1st Qu.:0.0000   1st Qu.:0.00000
## Median :0.0000   Median :0.00000
## Mean   :0.1972   Mean   :0.07995
## 3rd Qu.:0.0000   3rd Qu.:0.00000
## Max.   :1.0000   Max.   :1.00000
```

- It is convenient to change all the dates into fractional calendar years

```
oc$ybirth <- cal.yr(oc$birth)
oc$yentry <- cal.yr(oc$entry)
oc$yexit <- cal.yr(oc$exit)
```

We shall also compute the age at entry and at exit, respectively, as age will be the main time scale in our analyses.

```
oc$agencyentry <- oc$yentry - oc$ybirth
oc$agencyexit <- oc$yexit - oc$ybirth
```

- As the next step we shall create a `lexis` object from the data frame along the calendar period and age axes, and as the outcome event we specify the coronary death.

```
oc.lex <- Lexis(
  entry = list(
    per = yentry,
    age = yentry - ybirth
  ),
  exit = list(per = yexit),
  exit.status = chdeath,
  id = id, data = oc
)
```

NOTE: entry.status has been set to 0 for all.

```
str(oc.lex)
```

```
## Classes 'Lexis' and 'data.frame': 1501 obs. of 17 variables:
## $ per : 'cal.yr' num 1991 1991 1991 1991 1991 ...
## $ age : 'cal.yr' num 47.5 56.1 51.4 51.1 55.5 ...
## $ lex.dur: 'cal.yr' num 19.4 19.4 19.4 19.4 15.6 ...
## $ lex.Cst: num 0 0 0 0 0 0 0 0 0 0 ...
## $ lex.Xst: int 0 0 0 0 0 0 0 1 0 0 ...
## $ lex.id : int 1 2 3 4 5 6 7 8 9 10 ...
## $ id : int 1 2 3 4 5 6 7 8 9 10 ...
## $ birth : chr "1943-02-19" "1934-07-06" "1939-03-05" "1939-07-03" ...
## $ entry : chr "1990-08-14" "1990-08-14" "1990-08-14" "1990-08-14" ...
## $ exit : chr "2009-12-31" "2009-12-31" "2009-12-31" "2009-12-31" ...
## $ death : int 0 0 0 0 1 1 1 1 0 0 ...
## $ chdeath: int 0 0 0 0 0 0 0 1 0 0 ...
## $ ybirth : 'cal.yr' num 1943 1935 1939 1940 1935 ...
## $ yentry : 'cal.yr' num 1991 1991 1991 1991 1991 ...
## $ yexit : 'cal.yr' num 2010 2010 2010 2010 2006 ...
## $ agencyentry: 'cal.yr' num 47.5 56.1 51.4 51.1 55.5 ...
## $ agencyexit : 'cal.yr' num 66.9 75.5 70.8 70.5 71.1 ...
## - attr(*, "time.scales")= chr [1:2] "per" "age"
## - attr(*, "time.since")= chr [1:2] "" ""
## - attr(*, "breaks")=List of 2
## ..$ per: NULL
## ..$ age: NULL
```


```
summary(oc.lex)
```

```
##
## Transitions:
## To
## From 0 1 Records: Events: Risk time: Persons:
## 0 1381 120 1501 120 25280.91 1501
```

- At this stage it is informative to examine a graphical presentation of the follow-up lines and outcome cases in a conventional Lexis diagram. Make use of the `plot` method for `Lexis` objects. Gray lifelines are drawn and a bullet is put at the exit point of those lifelines that end with the outcome event.

```
par(mfrow = c(1, 1))
plot(oc.lex, xlim = c(1990, 2010), grid = TRUE)
```

```
points(oc.lex, pch = c(NA, 16)[oc.lex$lex.Xst + 1])
```




- As age is here the main time axis, we shall graphically illustrate the **study base**, *i.e.* the follow-up lines and outcome events, only along the age scale, being ordered by age at exit. Vertical lines at those ages when new coronary deaths occur are drawn to identify the pertinent **risk sets**. For that purpose it is useful first to sort the data frame and the **Lexis** object jointly by age at exit & age at entry, and to give a new ID number according to that order.

```
oc.ord <- cbind(ID = 1:1501, oc[order(oc$agexit, oc$agentry), ])
oc.lexord <- Lexis(
  entry = list(age = agentry),
  exit = list(age = agexit),
  exit.status = chdeath,
  id = ID, data = oc.ord
)
```


NOTE: entry.status has been set to 0 for all.

```
plot(oc.lexord, "age")
points(oc.lexord, pch = ifelse(oc.lexord$lex.Xst == 1, 16, NA))
with(
  subset(oc.lexord, lex.Xst == 1),
  abline(v = agexit, lty = 3)
)
```



For a closer look, we now zoom the graphical illustration of the risk sets into event times occurring between 50 to 58 years. – Copy the last four lines from the previous item and add arguments `xlim` and `ylim` to the call of `plot()`.

```
plot(oc.lexord, "age", xlim = c(50, 58), ylim = c(5, 65))
points(
  oc.lexord, "age", pch = ifelse(oc.lexord$lex.Xst == 1, 16, NA)
)
with(
  subset(oc.lexord, lex.Xst == 1),
  abline(v = agexit, lty = 3)
)
```



13.2 Nested case-control study

We shall now employ the strategy of **risk-set sampling** or **time-matched** sampling of controls, *i.e.* we are conducting a *nested case-control study* within the cohort.

- The risk sets are defined according to the age at diagnosis of the case. Further matching is applied for age at entry by 1-year agebands. For this purpose we first generate a categorical variable **agen2** for age at entry

```
oc.lex$agen2 <- cut(oc.lex$agentry, br = seq(40, 62, 1))
```

Matched sampling from risk sets may be carried out using function `ccwc()` found in the **Epi** package. Its main arguments are the times of **entry** and **exit** which specify the time at risk along the main time scale (here age), and the outcome variable to be given in the **fail** argument. The number of controls per case is set to be two, and the additional matching factor is given.

- After setting the RNG seed (with your own number), make a call of this function and see the structure of the resulting data frame **cactrl** containing the cases and the chosen individual controls.

```
set.seed(98623)
cactrl <-
  ccwc(
    entry = agentry, exit = agexit, fail = chdeath,
    controls = 2, match = agen2,
    include = list(id, agentry),
    data = oc.lex, silent = FALSE
  )
```

```
##
```

```
## Sampling risk sets: .....
```

```
str(cactrl)
```

```
## 'data.frame':   360 obs. of  7 variables:
## $ Set      : num  1 1 1 2 2 2 3 3 3 4 ...
## $ Map      : num  8 1423 1 95 381 ...
## $ Time     : num  63.9 63.9 63.9 66.7 66.7 ...
## $ Fail     : num  1 0 0 1 0 0 1 0 0 1 ...
## $ agen2    : Factor w/ 22 levels "(40,41]", "(41,42]", ...: 8 8 8 8 8 8 8 8 8 ...
## $ id       : int   8 1423 1 95 381 106 115 44 1343 504 ...
## $ agentry  : num  47.7 47.9 47.5 47.5 48 ...
```

Check the meaning of the four first columns of the case-control data frame from the help page of function `ccwc()`.

- Now we shall start collecting data on the risk factors for the cases and their matched controls, including determination of the total cholesterol levels from the frozen sera! The storehouse of the risk factor measurements for the whole cohort is file `occoh-Xdata.txt`. It contains values of the following variables.

Variable	Description
id	identification number, the same as in <code>occoh.txt</code>
smok	cigarette smoking with categories; 1: never, 2: former, 3: 1-14/d, 4: 15+/d
sbp	systolic blood pressure (mmHg)
tchol	total cholesterol level (mmol/l)

```
ocX <-
  read.table(
    paste(url, "occoh-Xdata.txt", sep = "/"), header = TRUE
  )
str(ocX)
```

```
## 'data.frame':   1501 obs. of  6 variables:
## $ id       : int   1 2 3 4 5 6 7 8 9 10 ...
## $ birth    : chr  "1943-02-19" "1934-07-06" "1939-03-05" "1939-07-03" ...
## $ entry    : chr  "1990-08-14" "1990-08-14" "1990-08-14" "1990-08-14" ...
## $ smok     : int   4 3 3 1 2 2 1 2 1 1 ...
## $ sbp      : int  130 128 157 102 138 119 155 154 164 124 ...
## $ tchol    : num   7.56 6.55 8.13 5.93 7.92 5.9 7.28 7.43 5.34 6.24 ...
```

- In the next step we collect the values of the risk factors for our cases and controls by merging the case-control data frame and the storehouse file. In this operation we utilize function `merge()` to select columns of two data frames: `cactrl` (all columns) and `ocX` (four columns) and to merge these into a single file (see exercise 1.1, subsection 1.1.8, where `merge()` was introduced). The `id` variable in both files is used as the key to link each individual case or control with his own data on risk factors.

```
oc.ncc <- merge(cactrl, ocX[, c("id", "smok", "tchol", "sbp")],
  by = "id"
)
str(oc.ncc)
```

```
## 'data.frame':   360 obs. of  10 variables:
## $ id       : int   1 8 10 12 14 28 37 41 41 42 ...
## $ Set      : num   1 1 77 9 44 50 5 5 7 9 ...
## $ Map      : num   1 8 10 12 14 28 37 41 41 42 ...
## $ Time     : num  63.9 63.9 73 70.3 75.2 ...
## $ Fail     : num   0 1 0 0 0 0 0 1 0 1 ...
## $ agen2    : Factor w/ 22 levels "(40,41]", "(41,42]", ...: 8 8 20 17 19 19 1 1 1 17 ...
## $ agentry  : num  47.5 47.7 59.5 56.1 58.7 ...
## $ smok     : int   4 2 1 2 4 1 3 4 4 2 ...
```

```
## $ tchol : num 7.56 7.43 6.24 5 3.73 4.56 5.15 6.09 6.09 5.41 ...
## $ sbp : int 130 154 124 148 165 230 116 125 125 156 ...
```

- We shall treat smoking as categorical and total cholesterol and systolic blood pressure as quantitative risk factors, but the values of the latter will be divided by 10 to get more interpretable effect estimates.

Variable	Description
cholgrp	cholesterol class; 1: <5, 2: 5-<6.5, 3: >=6.5
sbpgrp	blood pressure class; 1: <130, 2: 130-<150, 3: 150-<170, 4: >=170

Convert the smoking variable into a factor.

```
oc.ncc$smok <- factor(oc.ncc$smok,
  labels = c("never", "ex", "1-14/d", ">14/d")
)
```

- It is useful to start the analysis of case-control data by simple tabulations by the categorized risk factors. Crude estimates of the rate ratios associated with them, in which matching is ignored, can be obtained as follows. We shall focus on smoking

```
stat.table(
  index = list(smok, Fail),
  contents = list(count(), percent(smok)),
  margins = TRUE,
  data = oc.ncc
)
```

```
## -----
##          -----Fail-----
##   smok          0          1    Total
## -----
## never           85          31     116
##                35.4       25.8    32.2
##
## ex              61          19      80
##                25.4       15.8    22.2
##
## 1-14/d          54          42      96
##                22.5       35.0    26.7
##
## >14/d           40          28      68
##                16.7       23.3    18.9
##
##
## Total          240          120     360
##                100.0      100.0   100.0
## -----
```

```
smok.crncc <- glm(Fail ~ smok, family = binomial, data = oc.ncc)
round(ci.exp(smok.crncc), 3)
```

```
##          exp(Est.)  2.5% 97.5%
## (Intercept)    0.365 0.242 0.550
## smokex         0.854 0.442 1.651
## smok1-14/d     2.133 1.199 3.794
## smok>14/d      1.919 1.018 3.619
```

- A proper analysis takes into account matching that was employed in the selection of controls for each case from the pertinent risk set, further restricted to subjects who were about the same age at entry as the

case was. Also, adjustment for the other risk factors is desirable. In this analysis function `clogit()` in `survival` package is utilized. It is in fact a wrapper of function `coxph()`.

```
m.clogit <- clogit(Fail ~ smok + I(sbp / 10) + tchol +
  strata(Set), data = oc.ncc)
summary(m.clogit)

## Call:
## coxph(formula = Surv(rep(1, 360L), Fail) ~ smok + I(sbp/10) +
##       tchol + strata(Set), data = oc.ncc, method = "exact")
##
## n= 360, number of events= 120
##
##              coef exp(coef) se(coef)      z Pr(>|z|)
## smokex      -0.22100   0.80172  0.35085 -0.630  0.52877
## smok1-14/d   0.59186   1.80734  0.29894  1.980  0.04772 *
## smok>14/d    0.56616   1.76149  0.33326  1.699  0.08935 .
## I(sbp/10)    0.11171   1.11818  0.05453  2.049  0.04050 *
## tchol        0.40091   1.49319  0.11142  3.598  0.00032 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##              exp(coef) exp(-coef) lower .95 upper .95
## smokex          0.8017      1.2473    0.4031    1.595
## smok1-14/d      1.8073      0.5533    1.0060    3.247
## smok>14/d       1.7615      0.5677    0.9167    3.385
## I(sbp/10)       1.1182      0.8943    1.0048    1.244
## tchol           1.4932      0.6697    1.2003    1.858
##
## Concordance= 0.679 (se = 0.046 )
## Likelihood ratio test= 28.81 on 5 df,  p=3e-05
## Wald test            = 24.31 on 5 df,  p=2e-04
## Score (logrank) test = 27.31 on 5 df,  p=5e-05

round(ci.exp(m.clogit), 3)

##              exp(Est.)  2.5% 97.5%
## smokex          0.802 0.403 1.595
## smok1-14/d      1.807 1.006 3.247
## smok>14/d       1.761 0.917 3.385
## I(sbp/10)       1.118 1.005 1.244
## tchol           1.493 1.200 1.858
```

Compare these with the crude estimates obtained above.

13.3 Case-cohort study

Now we start applying the second major outcome-selective sampling strategy for collecting exposure data from a big study population

- The subcohort is selected as a simple random sample ($n = 260$) from the whole cohort. The id-numbers of the individuals that are selected will be stored in vector `subcids`, and `subcind` is an indicator for inclusion to the subcohort.

```
N <- 1501
n <- 260
set.seed(15792)
subcids <- sample(N, n)
oc.lexord$subcind <- 1 * (oc.lexord$id %in% subcids)
```


- We form the data frame `oc.cc` to be used in the subsequent analysis selecting the union of the subcohort members and the case group from the data frame of the full cohort. After that we collect the data of the risk factors from the data storehouse for the subjects in the case-cohort data

```
oc.cc <- subset(oc.lexord, subcind == 1 | chdeath == 1)
oc.cc <- merge(oc.cc, ocX[, c("id", "smok", "tchol", "sbp")],
  by = "id"
)
str(oc.cc)
```

```
## Classes 'Lexis' and 'data.frame': 362 obs. of 21 variables:
## $ id : int 8 10 34 40 41 42 43 45 51 62 ...
## $ age : num 47.7 59.5 51.9 45.2 40.2 ...
## $ lex.dur: num 16.21 19.38 19.36 4.59 19.32 ...
## $ lex.Cst: num 0 0 0 0 0 0 0 0 0 0 ...
## $ lex.Xst: int 1 0 0 0 1 1 1 0 1 1 ...
## $ lex.id : int 430 1483 1074 6 121 987 746 486 1056 663 ...
## $ ID : int 430 1483 1074 6 121 987 746 486 1056 663 ...
## $ birth : chr "1942-11-24" "1931-03-01" "1938-09-14" "1945-06-22" ...
## $ entry : chr "1990-08-15" "1990-08-15" "1990-08-20" "1990-08-21" ...
## $ exit : chr "2006-10-30" "2009-12-31" "2009-12-31" "1995-03-26" ...
## $ death : int 1 0 0 1 1 1 1 0 1 1 ...
## $ chdeath: int 1 0 0 0 1 1 1 0 1 1 ...
## $ ybirth : num 1943 1931 1939 1945 1950 ...
## $ yentry : num 1991 1991 1991 1991 1991 ...
## $ yexit : num 2007 2010 2010 1995 2010 ...
## $ agentry: num 47.7 59.5 51.9 45.2 40.2 ...
## $ agexit : num 63.9 78.8 71.3 49.8 59.5 ...
## $ subcind: num 0 1 1 1 0 0 0 1 1 0 ...
## $ smok : int 2 1 1 4 4 2 3 2 1 2 ...
## $ tchol : num 7.43 6.24 5.8 5.88 6.09 5.41 5.72 5.98 7.12 5.57 ...
## $ sbp : int 154 124 113 141 125 156 128 115 137 173 ...
## - attr(*, "breaks")=List of 1
## ..$ age: NULL
## - attr(*, "time.scales")= chr "age"
## - attr(*, "time.since")= chr ""
```

- We shall now create a graphical illustration of the lifelines contained in the case-cohort data. Lines for the subcohort non-cases are grey without bullet at exit, those for subcohort cases are blue with blue bullet at exit, and for cases outside the subcohort the lines are red and dotted with red bullets at exit.

```
plot(subset(oc.cc, chdeath == 0), "age")
lines(subset(oc.cc, chdeath == 1 & subcind == 1), col = "blue")
lines(subset(oc.cc, chdeath == 1 & subcind == 0), col = "red")
points(subset(oc.cc, chdeath == 1),
  pch = 16,
  col = c("blue", "red")[oc.cc$subcind + 1]
)
```



- Define the categorical smoking variable again.

```
oc.cc$smok <- factor(oc.cc$smok,
  labels = c("never", "ex", "1-14/d", ">14/d")
)
```

A crude estimate of the hazard ratio for the various smoking categories k vs. non-smokers ($k = 1$) can be obtained by tabulating cases (D_k) and person-years (y_k) in the subcohort by smoking and then computing the relevant exposure odds ratio for each category:

$$HR_k^{\text{crude}} = \frac{D_k/D_1}{y_k/y_1}$$

```
sm.cc <- stat.table(
  index = smok,
  contents = list(Cases = sum(lex.Xst), Pyrs = sum(lex.dur)),
  margins = TRUE,
  data = oc.cc
)
print(sm.cc, digits = c(sum = 0, ratio = 1))
```

```
## -----
## smok      Cases   Pyrs
## -----
## never      31    2071
## ex         19    1174
## 1-14/d     42    1399
## >14/d      28     846
## -----
## Total      120    5489
## -----
```

```
HRcc <-
  (sm.cc[1, -5] / sm.cc[1, 1]) / (sm.cc[2, -5] / sm.cc[2, 1])
round(HRcc, 3)
```

```
## never      ex 1-14/d  >14/d
## 1.000  1.081  2.006  2.211
```

Do these estimates resemble those obtained from nested case-control data?

- To estimate the rate ratios associated with smoking and adjusted for the other risk factors we now fit the pertinent Cox model applying the method of *weighted partial likelihood* as presented by Ling & Ying (1993) and Barlow (1994). This analysis can be done using function `cch()` in package `survival` with `method = "LinYing"`

```
oc.cc$survobj <- with(oc.cc, Surv(agentry, agexit, chdeath))
cch.LY <- cch(survobj ~ smok + I(sbp / 10) + tchol,
  stratum = NULL,
  subcoh = ~subcind, id = ~id, cohort.size = N, data = oc.cc,
  method = "LinYing"
)
summary(cch.LY)
```

```
## Case-cohort analysis,x$method, LinYing
## with subcohort of 260 from cohort of 1501
##
## Call: cch(formula = survobj ~ smok + I(sbp/10) + tchol, data = oc.cc,
##      subcoh = ~subcind, id = ~id, stratum = NULL, cohort.size = N,
##      method = "LinYing")
##
## Coefficients:
##              Coef      HR (95%  CI)      p
## smokex      -0.096 0.909 0.457 1.806 0.785
## smok1-14/d   0.769 2.157 1.180 3.943 0.012
## smok>14/d    1.085 2.959 1.532 5.718 0.001
## I(sbp/10)    0.218 1.244 1.119 1.383 0.000
## tchol        0.350 1.419 1.152 1.748 0.001
```

13.4 Full cohort analysis and comparisons

Finally, suppose the investigators after all could afford to collect the data on risk factors from the storehouse for the whole cohort.

- Let us form the data frame corresponding to the full cohort design and convert again smoking to be categorical.

```
oc.full <- merge(oc.lex, ocX[, c("id", "smok", "tchol", "sbp")],
  by.x = "id", by.y = "id"
)
oc.full$smok <- factor(oc.full$smok,
  labels = c("never", "ex", "1-14/d", ">14/d")
)
```

Juts for comparison with the corresponding analysis in case-cohort data perform a similar crude estimation of hazard ratios associated with smoking.

```
sm.coh <- stat.table(
  index = smok,
  contents = list(Cases = sum(lex.Xst), Pyrs = sum(lex.dur)),
  margins = TRUE,
```

```

data = oc.full
)
print(sm.coh, digits = c(sum = 0, ratio = 1))

## -----
##      smok      Cases      Pyrs
## -----
## never          31    10363
## ex             19    4879
## 1-14/d         42    6246
## >14/d          28    3793
##
## Total          120    25281
## -----
HRcoh <-
  (sm.coh[1, -5] / sm.coh[1, 1]) / (sm.coh[2, -5] / sm.coh[2, 1])
round(HRcoh, 3)

## never      ex 1-14/d  >14/d
## 1.000  1.302  2.248  2.468

```

- Fit now the ordinary Cox model to the full cohort. There is no need to employ extra tricks upon the ordinary `coxph()` fit.

```

cox.coh <- coxph(Surv(agentry, agexit, chdeath) ~
  smok + I(sbp / 10) + tchol, data = oc.full)
summary(cox.coh)

## Call:
## coxph(formula = Surv(agentry, agexit, chdeath) ~ smok + I(sbp/10) +
##      tchol, data = oc.full)
##
##      n= 1501, number of events= 120
##
##              coef exp(coef) se(coef)      z Pr(>|z|)
## smokex      0.10955   1.11577  0.29240  0.375 0.707922
## smok1-14/d  0.72567   2.06612  0.23704  3.061 0.002203 **
## smok>14/d   0.95054   2.58711  0.26198  3.628 0.000285 ***
## I(sbp/10)   0.14372   1.15456  0.04096  3.509 0.000450 ***
## tchol       0.26517   1.30366  0.07089  3.740 0.000184 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
##              exp(coef) exp(-coef) lower .95 upper .95
## smokex          1.116    0.8962    0.629    1.979
## smok1-14/d      2.066    0.4840    1.298    3.288
## smok>14/d       2.587    0.3865    1.548    4.323
## I(sbp/10)       1.155    0.8661    1.065    1.251
## tchol           1.304    0.7671    1.135    1.498
##
## Concordance= 0.681 (se = 0.026 )
## Likelihood ratio test= 41.16 on 5 df,  p=9e-08
## Wald test              = 42.05 on 5 df,  p=6e-08
## Score (logrank) test = 43.29 on 5 df,  p=3e-08

```

- Lastly, a comparison of the point estimates and standard errors between the different designs, including variants of analysis for the case-cohort design, can be performed.

```
betas <- cbind(coef(cox.coh), coef(m.clogit), coef(cch.LY))
colnames(betas) <- c("coh", "ncc", "cch.LY")
round(betas, 3)
```

```
##           coh      ncc cch.LY
## smokex      0.110 -0.221 -0.096
## smok1-14/d  0.726  0.592  0.769
## smok>14/d   0.951  0.566  1.085
## I(sbp/10)   0.144  0.112  0.218
## tchol       0.265  0.401  0.350
```

```
SEs <- cbind(
  sqrt(diag(cox.coh$var)),
  sqrt(diag(m.clogit$var)),
  sqrt(diag(cch.LY$var))
)
colnames(SEs) <- colnames(betas)
round(SEs, 3)
```

```
##           coh      ncc cch.LY
## smokex      0.292 0.351  0.350
## smok1-14/d  0.237 0.299  0.308
## smok>14/d   0.262 0.333  0.336
## I(sbp/10)   0.041 0.055  0.054
## tchol       0.071 0.111  0.106
```

You will notice that the point estimates of the coefficients obtained from the full cohort, nested case-control, and case-cohort analyses, respectively, are somewhat variable. However, the standard errors from the NCC and CC analyses should be quite similar when the numbers of cases and non-cases are similar.

13.5 Further exercises and homework

- If you have time, you could run both the NCC study and CC study again but now with a larger control group or subcohort; for example 4 controls per case in NCC and $n = 520$ as the subcohort size in CC. Remember resetting the seed first. Pay attention in the results to how much closer will be the point estimates and the proper SEs to those obtained from the full cohort design.
- Instead of simple linear terms for `sbp` and `tchol` you could try to fit spline models to describe their effects.
- A popular alternative to weighted partial likelihood in the analysis of case-cohort data is the *pseudo-likelihood method* (Prentice 1986), which is based on *late entry* to follow-up of the case subjects not belonging to the subcohort. The way to do this is provided by function `cch()` which you can apply directly to the case-cohort data `oc.cc` as before but now with `method = "Prentice"`. – Try this and compare the results with those obtained by weighted partial likelihood in model `cch.LY`.
- Yet another computational solution for maximizing weighted partial likelihood is provided by a combination of functions `twophase()` and `svycoxph()` of the `survey` package. The approach is illustrated with an example in a vignette *Two-phase designs in epidemiology* by Thomas Lumley (see <http://cran.r-project.org/web/packages/survey/vignettes/epi.pdf>). – You can try this at home and check that you would obtain similar results as with model `cch.LY`.

Chapter 14

Causal inference 2: Model-based estimation of causal estimands


Sources of inspiration: Luque Fernandez, M.A. et al. (2018) *Stat Med* 2018;37(16):2530-2546 and Smith et al. (2022) *Stat Med* 2022;41(2):407-432.

14.1 Introduction

We shall illustrate with simulated data the estimation of causal effects of a binary exposure X when the outcome Y is also binary, and there is a set of four covariates $Z = (Z_1, Z_2, Z_3, Z_4)$. As a background story, we imagine a population of cancer patients, in whom the variables and the assumed marginal distributions of the covariates are

Variable	Description
X	treatment; 1: radiotherapy only, 0: radiotherapy + chemotherapy
Y	death during one year after diagnosis of cancer
Z_1	sex; 0: man, 1: woman; $Z_1 \sim \text{Bern}(0.5)$
Z_2	age group 0: young , 1: old ; $Z_2 \sim \text{Bern}(0.65)$
Z_3	stage of cancer; 4 classes; $Z_3 \sim \text{DiscUnif}(1, \dots, 4)$
Z_4	comorbidity score; 5 classes; $Z_3 \sim \text{DiscUnif}(1, \dots, 5)$

For simplicity, covariates Z_3 and Z_4 are treated as continuous variables in the models. The assumed causal diagram is drawn using `dagitty` and is shown below.



For more generic notation, the probabilities of $Y = 1$ will be expressed as expectations, e.g.

$$E(Y^{X=x}) = P(Y^{X=x} = 1) \quad \text{and} \quad E(Y|X = x, Z = z) = P(Y = 1|X = x, Z = z),$$

where Z is the vector of relevant covariates.

The same principle is applied in expressing the conditional probability of $X = 1$, i.e. being exposed, given $Z = z$:

$$E(X|Z = z) = P(X = 1|Z = z).$$

The fitted or predicted probabilities of $Y = 1$ are denoted as fitted $\hat{Y}$ or predicted values $\tilde{Y}$ of Y with pertinent subscripts and/or superscripts. – Both X and Y are modelled by logistic regression. The expit-function or inverse of the logit function is defined: $\text{expit}(u) = 1/(1 + e^{-u})$, $u \in R$. This is equal to the cumulative distribution function of the standard logistic distribution, the values of which are returned in R by `plogis(u)`. The R function that returns values of the logit-function is `qlogis()`.

The true model assumed for the dependence of exposure X on covariates:

$$E(X|Z_1 = z_1, \dots, Z_4 = z_4) = \text{expit}(-5 + 0.05z_2 + 0.25z_3 + 0.5z_4 + 0.4z_2z_4).$$

The assumed true model for the outcome is

$$E(Y|X = x, Z_1 = z_1, \dots, Z_4 = z_4) = \text{expit}(-1 + x - 0.1z_1 + 0.35z_2 + 0.25z_3 + 0.20z_4 + 0.15z_2z_4)$$

Note that X does not depend on Z_1 , and that in both models there is a product term Z_2Z_4 , the effect of which appears weaker for the outcome model.

14.2 Control of confounding

1. Based on inspection of the causal diagram, can you provide justification for the claim that variables Z_2, Z_3 and Z_4 form a proper subset of the four covariates, which is sufficient to block all backdoor paths between X and Y and thus remove confounding?

2. Even though we have such a minimal sufficient set as indicated in item (a), why it still might be useful to include covariate Z_1 , too, when modelling the outcome?

14.3 Generation of target population and true models

1. Load the necessary packages.

```
library(Epi)

##
## Attaching package: 'Epi'

## The following object is masked from 'package:dagitty':
##
##      paths

library(stdReg)
library(PSweight)
```

2. Define two R-functions, which compute the expected values for the exposure and those for the outcome based on the assumed true exposure model and the true outcome model, respectively.

```
EX <- function(z2, z3, z4) {
  plogis(-5 + 0.05 * z2 + 0.25 * z3 + 0.5 * z4 + 0.4 * z2 * z4)
}
EY <- function(x, z1, z2, z3, z4) {
  plogis(-1 + x - 0.1 * z1 + 0.35 * z2 + 0.25 * z3 +
    0.20 * z4 + 0.15 * z2 * z4)
}
```

3. Define the function for the generation of data by simulating random values from pertinent probability distributions based on the given assumptions.

```
genData <- function(N) {
  z1 <- rbinom(N, size = 1, prob = 0.5) # Bern(0.5)
  z2 <- rbinom(N, size = 1, prob = 0.65) # Bern(0.65)
  z3 <- trunc(runif(N, min = 1, max = 5), digits = 0) # DiscUnif(1,4)
  z4 <- trunc(runif(N, min = 1, max = 6), digits = 0) # DiscUnif(1,5)
  x <- rbinom(N, size = 1, prob = EX(z2, z3, z4))
  y <- rbinom(N, size = 1, prob = EY(x, z1, z2, z3, z4))
  data.frame(z1, z2, z3, z4, x, y)
}
```

4. Generate a data frame dd for a big target population of 500000 subjects

```
N <- 500000
set.seed(7777)
dd <- genData(N)
```

14.4 Factual and counterfactual risks - associational and causal contrasts

1. Compute the factual risks of death for the two exposure groups

$$E(Y|X = x) = P(Y = 1|X = x) = \frac{P(Y = 1 \& X = x)}{P(X = x)}, \quad x = 0, 1,$$

in the whole target population, as well as their associational contrasts: risk difference, risk ratio, and odds ratio. Before that define a useful function

```

Contr <- function(mu1, mu0) {
  RD <- mu1 - mu0
  RR <- mu1 / mu0
  OR <- (mu1 / (1 - mu1)) / (mu0 / (1 - mu0))
  return(c(mu1, mu0, RD = RD, RR = RR, OR = OR))
}
Ey1fact <- with(dd, sum(y == 1 & x == 1) / sum(x == 1))
Ey0fact <- with(dd, sum(y == 1 & x == 0) / sum(x == 0))
round(Contr(Ey1fact, Ey0fact), 4)

```

```

##                RD      RR      OR
## 0.8949 0.6288 0.2661 1.4232 5.0286

```

How much bigger is the risk of death of those factually exposed to radiotherapy only as compared with those receiving chemotherapy, too?

2. Compute now the true **counterfactual** or **potential risks** of death

$$E(Y_i^{X_i=x}) = P(Y_i^{X_i=x} = 1) = \pi_i^{X_i=x}$$

for each individual under the alternative treatments or exposure values $x = 0, 1$ with given covariate values, then the average or overall counterfactual risks $E(Y^{X=1}) = \pi^1$ and $E(Y^{X=0}) = \pi^0$ in the population, and finally the true **marginal causal contrasts** for the effect of X :

$$RD = E(Y^{X=1}) - E(Y^{X=0}), \quad RR = E(Y^{X=1})/E(Y^{X=0}),$$

$$OR = \frac{E(Y^{X=1})/[1 - E(Y^{X=1})]}{E(Y^{X=0})/[1 - E(Y^{X=0})]}$$

```

dd <- transform(dd,
  EY1.ind = EY(x = 1, z1, z2, z3, z4),
  EY0.ind = EY(x = 0, z1, z2, z3, z4)
)
EY1pot <- mean(dd$EY1.ind)
EY0pot <- mean(dd$EY0.ind)
round(Contr(EY1pot, EY0pot), 4)

```

```

##                RD      RR      OR
## 0.8273 0.6530 0.1743 1.2670 2.5462

```

3. Compare the associational contrasts computed in item 4.1 with the causal contrasts in item 4.2. What do you conclude about confoundedness of the associational contrasts?

14.5 Outcome modelling and estimation of causal contrasts by g-formula

As the first approach for estimating causal contrasts of interest we apply the method of **standardization** or **g-formula**. It is based on a hopefully realistic enough model for $E(Y|X = x, Z = z)$, i.e. how the risk of outcome is expected to depend on the exposure variable X and on a sufficient set Z of confounders. The potential or counterfactual risks $E(Y^{X=x})$, $x = 0, 1$, are marginal expectations of the above quantities, standardized over the joint distribution of the confounders Z in the target population.

$$E(Y^{X=x}) = E_Z[E(Y|X = x, Z)] = \int E(Y|X = x, Z = z)dF_Z(z), \quad x = 0, 1.$$

1. Assume now a *slightly misspecified* model $\mathbf{mY}$ for the outcome, which contains only main effect terms of the explanatory variables:

$$\pi_i = E(Y_i|X_i = x_i, Z_{i1} = z_{i1}, \dots, Z_{i4} = z_{i4}) = \text{expit} \left(\beta_0 + \delta x_i + \sum_{j=1}^4 \beta_j z_{ij} \right)$$

Fit this model on the target population using function `glm()`

```
mY <- glm(y ~ x + z1 + z2 + z3 + z4, family = binomial, data = dd)
round(ci.lin(mY, Exp = TRUE)[, c(1, 5)], 3)
```

```
##           Estimate exp(Est.)
## (Intercept)  -1.240    0.289
## x             1.052    2.863
## z1           -0.095    0.909
## z2             0.767    2.153
## z3             0.250    1.284
## z4             0.279    1.322
```

There is not much idea in looking at the standard errors or confidence intervals in such a big population.

- For each subject i , compute the fitted individual risk $\widehat{Y}_i$ as well as the predicted potential (counterfactual) risks $\widetilde{Y}_i^{X_i=x}$ for both exposure levels $x = 0, 1$ separately, keeping the individual values of the Z -variables as they are.

```
dd$yh <- predict(mY, type = "response") # fitted values
dd$yp1 <- predict(mY, newdata = data.frame(
  x = rep(1, N), # x=1
  dd[, c("z1", "z2", "z3", "z4")]
), type = "response")
dd$yp0 <- predict(mY, newdata = data.frame(
  x = rep(0, N), # x=0
  dd[, c("z1", "z2", "z3", "z4")]
), type = "response")
```

- Applying the method of **standardization** or **g-formula** compute now the point estimates

$$\widehat{E}_g(Y^{X=x}) = \frac{1}{n} \sum_{i=1}^n \widetilde{Y}_i^{X_i=x}, \quad x = 0, 1.$$

of the two counterfactual risks $E(Y^{X=1}) = \pi^1$ and $E(Y^{X=0}) = \pi^0$ as well as those of the marginal causal contrasts

```
EY1pot.g <- mean(dd$yp1)
EY0pot.g <- mean(dd$yp0)
round(Contr(EY1pot.g, EY0pot.g), 4)
```

```
##           RD      RR      OR
## 0.8330 0.6508 0.1822 1.2800 2.6763
```

The marginal expectations $E_Z[E(X = x, Z)]$ taken over the joint distribution of the confounders Z are empirically estimated from the actual data representing the target population by simply computing the arithmetic means of the individually predicted values $\widetilde{Y}_i^{X_i=x}$ of the outcome for the two exposure levels.

Compare the estimated contrasts with the true ones in item 4.2 above. How big is the bias due to slight misspecification of the outcome model? Compare in particular the estimate of the marginal OR here with the conditional OR obtained in item 5.1 from the pertinent coefficient in the logistic model. Which one is closer to 1?

- Perform the same calculations using the tools in package **stdReg** (see Sjölander 2016)

```
mY.std <- stdGlm(fit = mY, data = dd, X = "x")
summary(mY.std)
```

```
##
## Formula: y ~ x + z1 + z2 + z3 + z4
## Family: binomial
## Link function: logit
## Exposure: x
```

```
##
## Estimate Std. Error lower 0.95 upper 0.95
## 0 0.651 0.000733 0.649 0.652
## 1 0.833 0.001562 0.830 0.836
round(summary(mY.std, contrast = "difference", reference = 0)$est.table, 4)
```

```
## Estimate Std. Error lower 0.95 upper 0.95
## 0 0.0000 0.0000 0.0000 0.0000
## 1 0.1822 0.0017 0.1788 0.1856
round(summary(mY.std, contrast = "ratio", reference = 0)$est.table, 4)
```

```
## Estimate Std. Error lower 0.95 upper 0.95
## 0 1.00 0.0000 1.0000 1.0000
## 1 1.28 0.0028 1.2745 1.2855
round(summary(mY.std, transform = "odds",
               contrast = "ratio", reference = 0)$est.table, 4)
```

```
## Estimate Std. Error lower 0.95 upper 0.95
## 0 1.0000 0.0000 1.0000 1.0000
## 1 2.6763 0.0315 2.6146 2.7379
```

Check that you got the same point estimates as in the previous item. Again, the confidence intervals are not very meaningful when analysing the data covering the whole big target population. Of course, when applied to real sample data they are relevant. In `stdReg` package, the standard errors are obtained by the multivariate delta method built upon M-estimation and robust sandwich estimator of the pertinent covariance matrix, and approximate confidence intervals are derived from these in the usual way.

14.6 Inverse probability weighting (IPW) by propensity scores

The next method is based on weighting each individual observation by the inverse of the probability of belonging to that particular exposure group, which was realized, this probability being predicted by the determinants of exposure.

1. Fit first a somewhat misspecified model for the exposure including the main effects of the Z -variables only.

$$p_i = E(X_i | Z_{1i} = z_{1i}, \dots, Z_{4i} = z_{4i}) = \text{expit}(\gamma_0 + \gamma_1 z_{1i} + \gamma_2 z_{2i} + \gamma_3 z_{3i} + \gamma_4 z_{4i}), \quad i = 1, \dots, N$$

```
mX <- glm(x ~ z1 + z2 + z3 + z4,
          family = binomial(link = logit), data = dd
)
round(ci.lin(mX, Exp = TRUE)[, c(1, 5)], 4)
```

```
## Estimate exp(Est.)
## (Intercept) -6.3031 0.0018
## z1 -0.0161 0.9840
## z2 1.6260 5.0833
## z3 0.2391 1.2702
## z4 0.8369 2.3093
```

2. Extract the **propensity scores**, i.e. fitted probabilities of belonging to exposure group 1: $PS_i = \hat{p}_i$, and compare their distribution between the two groups.

```
dd$PS <- predict(mX, type = "response")
summary(dd$PS)
```

```
## Min. 1st Qu. Median Mean 3rd Qu. Max.
## 0.005256 0.041532 0.112765 0.172440 0.248554 0.613993
```

```
with(subset(dd, x == 0), plot(density(PS), lty = 2))
with(subset(dd, x == 1), lines(density(PS), lty = 1))
```

How different are the distributions? Are they sufficiently overlapping?

3. Compute the weights

$$W_i = \frac{1}{\text{PS}_i}, \quad \text{when } X_i = 1,$$

$$W_i = \frac{1}{1 - \text{PS}_i}, \quad \text{when } X_i = 0.$$

Look at the sum as well as the distribution summary of the weights in the exposure groups. The sum of weights should be close to N in both groups.

```
dd$w <- ifelse(dd$x == 1, 1 / dd$PS, 1 / (1 - dd$PS))
with(dd, tapply(w, x, sum))
```

```
##          0          1
## 498880.5 565237.5
```

4. Compute now the weighted estimates of the potential or counterfactual risks for both exposure categories

$$\widehat{E}_w(Y^{X=x}) = \frac{\sum_{i=1}^n \mathbf{1}_{\{X_i=x\}} W_i Y_i}{\sum_{i=1}^n \mathbf{1}_{\{X_i=x\}} W_i} = \frac{\sum_{X_i=x} W_i Y_i}{\sum_{X_i=x} W_i}, \quad x = 0, 1,$$

and their causal contrasts, for instance

$$\widehat{\text{RD}}_w = \widehat{E}_w(Y^{X=1}) - \widehat{E}_w(Y^{X=0}) = \frac{\sum_{i=1}^n X_i W_i Y_i}{\sum_{i=1}^n X_i W_i} - \frac{\sum_{i=1}^n (1 - X_i) W_i Y_i}{\sum_{i=1}^n (1 - X_i) W_i}$$

```
EY1pot.w <- sum(dd$x * dd$w * dd$y) / sum(dd$x * dd$w)
EY0pot.w <- sum((1 - dd$x) * dd$w * dd$y) / sum((1 - dd$x) * dd$w)
round(Contr(EY1pot.w, EY0pot.w), 4)
```

```
##          RD          RR          OR
## 0.8037 0.6519 0.1518 1.2329 2.1868
```

These estimates seem to be somewhat downward biased when comparing to true values. Could this be because of omitting the relatively strong product term effect of Z_2 and Z_4 ?

14.7 Improving IPW estimation and using R package PSweight

We now try to improve IPW-estimation by a richer exposure model. In computations we shall utilize the R package `PSweight` (see `PSweight` vignette).

1. First, we compute the propensity scores and weights from a more flexible exposure model, which contains all pairwise product terms of the parents of X . According to the causal diagram, Z_1 is not in that subset, so it is left out. The exposure model is specified and the weights are obtained as follows using function `SumStat()` in `PSweight`:

```
mX2 <- glm(x ~ (z2 + z3 + z4)^2, family = binomial, data = dd)
round(ci.lin(mX2, Exp = TRUE)[, c(1, 5)], 3)
```

```
##          Estimate exp(Est.)
## (Intercept)   -5.085      0.006
## z2             0.143      1.154
## z3             0.243      1.275
## z4             0.527      1.693
## z2:z3          -0.003      0.997
## z2:z4          0.376      1.456
## z3:z4          0.001      1.001
```

```
psw2 <- SumStat(
  ps.formula = mX2$formula, data = dd,
  weight = c("IPW", "treated", "overlap")
)
dd$PS2 <- psw2$propensity[, 2]
dd$w2 <- ifelse(dd$x == 1, 1 / dd$PS2, 1 / (1 - dd$PS2))
plot(density(dd$PS2[dd$x == 0]), lty = 2)
lines(density(dd$PS2[dd$x == 1]), lty = 1)
```

Note that apart from ordinary IPW, other types of weights can also be obtained. These are relevant when estimating other kinds of causal contrasts, like “average treatment effect among the treated” (ATT, see below) and “average treatment effect in the overlap (or equipose) population” (ATO).

2. PSweight includes some useful tools to examine the properties of the distribution and to check the balance of the propensity scores, for instance

```
plot(psw2, type = "balance", metric = "PSD")
```

It is desirable that the horizontal values of these measures for given weights are less than 0.1.

3. Estimation and reporting of the causal contrasts. For relative contrasts, the summary method provides the results on the log-scale; therefore exp-transformation

```
ipw2est <- PSweight(ps.formula = mX2, yname = "y", data = dd, weight = "IPW")
ipw2est
```

```
## Original group value: 0, 1
##
## Point estimate:
## 0.6526, 0.8255
```

```
summary(ipw2est)
```

```
##
## Closed-form inference:
##
## Original group value: 0, 1
##
## Contrast:
##           0 1
## Contrast 1 -1 1
##
##           Estimate Std.Error      lwr      upr Pr(>|z|)
## Contrast 1 0.1729179 0.0025636 0.1678934 0.17794 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
(logRR.ipw2 <- summary(ipw2est, type = "RR"))
```

```
##
## Closed-form inference:
##
## Inference in log scale:
## Original group value: 0, 1
##
## Contrast:
##           0 1
## Contrast 1 -1 1
##
##           Estimate Std.Error      lwr      upr Pr(>|z|)
```

```
## Contrast 1 0.2350512 0.0031789 0.2288207 0.24128 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
round(exp(logRR.ipw2$estimates[c(1, 4, 5)]), 3)
```

```
## [1] 1.265 1.257 1.273
round(exp(summary(ipw2est, type = "OR")$estimates[c(1, 4, 5)]), 3)
```

```
## [1] 2.519 2.434 2.606
```

Compare these with the previous IPW estimate as well as the true values. Have we obtained nearly unbiased results?

The standard errors provided by `PSweight` are by default based on the empirical sandwich covariance matrix and application of delta method as appropriate. Bootstrapping is also possible but is computationally very intensive and is recommended to be used only in relatively small samples.

14.8 Effect of exposure among those exposed

If we are interested in the causal contrasts describing the **effect of exposure among those exposed** (like ATE), the relevant factual and counterfactual risks in that subset are

$$\begin{aligned}\pi_1^1 &= E(Y^{X=1}|X=1) = E(Y|X=1) = \pi_1, \\ \pi_1^0 &= E(Y^{X=0}|X=1) = \sum_{X_i=1} E(Y|X=0, Z=z)P(Z=z|X=1)\end{aligned}$$

We are thus making an “observed vs. expected” comparison, in which the z -specific risks in the unexposed are weighted by the distribution of Z in the exposed subset of the target population. The risks and their contrasts are estimated from the fit of the outcome model:

```
EY1att.g <- mean(subset(dd, x == 1)$yp1)
EY0att.g <- mean(subset(dd, x == 1)$yp0)
round(Contr(EY1att.g, EY0att.g), 4)
```

```
##                RD      RR      OR
## 0.8949 0.7560 0.1389 1.1837 2.7488
```

Compare the results here with those for the whole target population. What do you observe?

2. Have you any guess about the causal effect of exposure among the unexposed; is it bigger or smaller than among the exposed or among the whole population?
3. Incidentally, the true causal contrasts among the exposed based on the true model are similarly obtained from the quantities in item 4.2 above:

```
EY1att <- mean(subset(dd, x == 1)$EY1.ind)
EY0att <- mean(subset(dd, x == 1)$EY0.ind)
round(Contr(EY1att, EY0att), 4)
```

```
##                RD      RR      OR
## 0.8955 0.7680 0.1275 1.1661 2.5896
```

Compare the estimates in the previous item with the true values obtained here.

4. When wishing to estimate the effect of exposure among the exposed using the IPW approach, then the weights are $W_i = 1$ for the exposed and $W_i = PS_i/(1 - PS_i)$ for the unexposed. Call again `PSweight` but with another choice of weight:

```
psatt <- PSweight(ps.formula = mX2, yname = "y", data = dd, weight = "treated")
psatt
```

```
## Original group value: 0, 1
## Treatment group value: 1
##
## Point estimate:
## 0.7667, 0.8949
round(summary(psatt)$estimates[1], 4)

## [1] 0.1282
round(exp(summary(psatt, type = "RR")$estimates[1]), 3)

## [1] 1.167
round(exp(summary(psatt, type = "OR")$estimates[1]), 3)

## [1] 2.592
```

Compare the results here with those obtained by g-formula in item 8.1 and with the true contrasts above.

14.9 Double robust estimation by augmented IPW

Let us attempt to correct the estimates by a **double robust** (DR) approach called **augmented IPW estimation** (AIPW), which combines the g-formula and the IPW approach. The classical AIPW-estimator can be expressed in two ways: either an IPW-corrected g-formula estimator, or a g-corrected IPW-estimator.

$$\begin{aligned}\widehat{E}_a(Y^{X=x}) &= \widehat{E}_g(Y^{X=x}) + \frac{1}{n} \sum_{i=1}^n \mathbf{1}_{\{X_i=x\}} W_i (Y_i - \widetilde{Y}_i^{X_i=x}) \\ &= \widehat{E}_w(Y^{X=x}) + \frac{1}{n} \sum_{i=1}^n (1 - \mathbf{1}_{\{X_i=x\}} W_i) \widetilde{Y}_i^{X_i=x}.\end{aligned}$$

1. We shall first combine the results from the slightly misspecified outcome model with those from the more misspecified exposure model.

```
EY1pot.a <- EY1pot.g + mean( 1*(dd$x==1) * dd$w * (dd$y - dd$yp1) )
EY0pot.a <- EY0pot.g + mean( 1*(dd$x==0) * dd$w * (dd$y - dd$yp0) )
round(Contr(EY1pot.a, EY0pot.a), 4)
```

```
##                RD        RR        OR
## 0.8241 0.6523 0.1718 1.2634 2.4972
```

Compare these results with those obtained by g-formula and by non-augmented IPW method. Was augmentation successful?

2. Let us then look, how close we get when combining the results from the slightly misspecified outcome model with the correct exposure model using the alternative AIPW-formula

```
EY1pot.w2 <- ipw2est$muhat[2]
EY0pot.w2 <- ipw2est$muhat[1]
EY1pot.a2 <- EY1pot.w2 + mean( (1 - 1*(dd$x==1) * dd$w2) * dd$yp1 )
EY0pot.a2 <- EY0pot.w2 + mean( (1 - 1*(dd$x==0) * dd$w2) * dd$yp0 )
round(Contr(EY1pot.a2, EY0pot.a2), 4)
```

```
##          1          0    RD.1    RR.1    OR.1
## 0.8261 0.6526 0.1735 1.2658 2.5285
```

Compare the results with previous ones. How successful was augmentation now?

AIPW-estimates and confidence intervals for the causal contrasts of interest can be obtained, for instance, using `PSweight` by adding the model formula of the outcome model as the value for the argument `out.formula`.

14.10 Double robust, targeted maximum likelihood estimation (TMLE)

We now consider now another double robust approach, known as **targeted maximum likelihood estimation** (TMLE). It also corrects the estimator obtained from the outcome model by elements that are derived from the exposure model. The corrections are, though, not as intuitive as those in AIPW. See Schuler and Rose (2017) for more details.

1. The first step is to utilize the propensity scores obtained above for the correct exposure model and define the “clever covariates”

```
dd$H1 <- dd$x / dd$PS2
dd$H0 <- (1 - dd$x) / (1 - dd$PS2)
```

2. Then, a working model is fitted for the outcome, in which the clever covariates are explanatory variables, but the model also includes the previously fitted linear predictor $\hat{\eta}_i = \text{logit}(\hat{Y}_i)$ from the original outcome model `mY` as an offset term; see item 5.2. Moreover, the intercept is removed.

```
epsmod <- glm(y ~ -1 + H0 + H1 + offset(qlogis(yh)),
  family = binomial(link = logit), data = dd
)
eps <- coef(epsmod)
eps
```

```
##           H0           H1
## 0.007197951 -0.002859654
```

3. The logit-transformed predicted values $\hat{Y}_i^{X_i=1}$ and $\hat{Y}_i^{X_i=0}$ of the potential or counterfactual individual risks from the original outcome model are now corrected by the estimated coefficients of the clever covariates, and the corrected predictions are returned to the original scale.

```
ypred0.H <- plogis(qlogis(dd$yp0) + eps[1] / (1 - dd$PS2))
ypred1.H <- plogis(qlogis(dd$yp1) + eps[2] / dd$PS2)
```

4. Estimates of the causal contrasts:

```
EY0pot.t <- mean(ypred0.H)
EY1pot.t <- mean(ypred1.H)
round(Contr(EY1pot.t, EY0pot.t), 4)
```

```
##           RD          RR          OR
## 0.8246 0.6526 0.1720 1.2635 2.5020
```

Compare these with previous results and with the true values.

14.11 Double robust estimation with AIPW and tmle packages

It may be difficult to specify conventional generalized linear models or even generalized additive models for exposure and outcome which are sufficiently realistic, yet which do not suffer from overfitting. Modern approaches of **statistical learning**, aka “machine learning” provide tools for flexible modelling, which may be used to reduce the risk of misspecification, if thoughtfully applied.

There are a few R packages in which some general algorithmic approaches for supervised learning are implemented for estimating causal parameters. For instance, package **AIPW** (see Zhong et al., 2021) utilizes several learning algorithms for exposure and outcome modelling and then performs AIPW estimation of the parameters of interest coupled with calculation of confidence intervals. Package **tmle** (see Karim and Frank, 2021) performs same tasks but uses the TMLE approach in estimation.

Both AIPW and **tmle** lean on the **SuperLearner** package, which uses multiple learning algorithms (e.g. GLM, GAM, Random Forest, Recursive Partitioning, Gradient Boosting, etc.) for constructing predictions of the counterfactual quantities, and then creates an optimal weighted average of those models, aka an “ensemble”.

These algorithms are computationally highly intensive. Fitting models with only 3 or 4 covariates as in this practical on our target cohort of 500,000 subjects would take hours on an ordinary laptop. With a study population of 5000 it takes several minutes.

Chapter 15

Time-dependent variables and multiple states

The following practical exercise is based on the data from paper:

P Hovind, L Tarnow, P Rossing, B Carstensen, and HH Parving: Improved survival in patients obtaining remission of nephrotic range albuminuria in diabetic nephropathy. *Kidney Int*, **66**(3):1180–1186, Sept 2004.

You can find a .pdf-version of the paper here: <http://BendixCarstensen.com/AdvCoh/papers/Hovind.2004.pdf>

15.1 The renal failure dataset

The dataset `renal.dta` contains data on follow up of 125 patients from Steno Diabetes Center. They enter the study when they are diagnosed with nephrotic range albuminuria (NRA). This is a condition where the levels of albumin in the urine is exceeds a certain level as a sign of kidney disease. The levels may however drop as a consequence of treatment, this is called remission. Patients exit the study at death or kidney failure (dialysis or transplant).

Variable	Description
<code>id</code>	Patient id
<code>sex</code>	1=male, 2=female
<code>dob</code>	Date of birth
<code>doe</code>	Date of entry into the study (2.5 years after NRA)
<code>dor</code>	Date of remission. Missing if no remission has occurred
<code>dox</code>	Date of exit from study
<code>event</code>	Exit status: 1,2,3=event (death, ESRD), 0=censored

1. The dataset is in Stata-format, so you must read the dataset using `read.dta` from the `foreign` package (which is part of the standard R-distribution). At the same time, convert `sex` to a proper factor. Choose where to read the dataset.

```
library(Epi)
library(survival)
library(mgcv)
```

Loading required package: nlme

This is mgcv 1.9-4. For overview type '?mgcv'.

```
library(foreign)
# renal <- read.dta(
```

```
# "https://raw.githubusercontent.com/SPE-R/SPE/master/pracs-book/data/renal.dta")
renal <- read.dta("http://BendixCarstensen.com/SPE/data/renal.dta")
renal$sex <- factor(renal$sex, labels = c("M", "F"))
head(renal)
```

2. Use the `Lexis` function to declare the data as survival data with age, calendar time and time since entry into the study as timescales. Label any event > 0 as *ESRD*, i.e. renal death (death of kidney (transplant or dialysis), or person). Note that you must make sure that the *alive* state (here *NRA*) is the first, as `Lexis` assumes that everyone starts in this state (unless of course `entry.status` is specified):

```
Lr <- Lexis(entry = list(per = doe,
                        age = doe - dob,
                        tfi = 0),
           exit = list(per = dox),
           exit.status = factor(event > 0, labels = c("NRA", "ESRD")),
           data = renal)
```

NOTE: `entry.status` has been set to "NRA" for all.

```
str(Lr)
```

```
Classes 'Lexis' and 'data.frame': 125 obs. of 14 variables:
 $ per : num 1996 1990 1988 1995 1988 ...
 $ age : num 28.1 30.2 25.8 44.5 26.6 ...
 $ tfi : num 0 0 0 0 0 0 0 0 0 0 ...
 $ lex.dur: num 1.08 6.6 5.39 8.75 16.07 ...
 $ lex.Cst: Factor w/ 2 levels "NRA","ESRD": 1 1 1 1 1 1 1 1 1 1 ...
 $ lex.Xst: Factor w/ 2 levels "NRA","ESRD": 2 2 2 1 1 2 2 1 2 1 ...
 $ lex.id : int 1 2 3 4 5 6 7 8 9 10 ...
 $ id : num 17 26 27 33 42 46 47 55 62 64 ...
 $ sex : Factor w/ 2 levels "M","F": 1 2 2 1 2 2 1 1 2 1 ...
 $ dob : num 1968 1959 1962 1951 1961 ...
 $ doe : num 1996 1990 1988 1995 1988 ...
 $ dor : num NA 1990 NA 1996 1997 ...
 $ dox : num 1997 1996 1993 2004 2004 ...
 $ event : num 2 1 3 0 0 2 1 0 2 0 ...
 - attr(*, "time.scales")= chr [1:3] "per" "age" "tfi"
 - attr(*, "time.since")= chr [1:3] "" "" ""
 - attr(*, "breaks")=List of 3
 ..$ per: NULL
 ..$ age: NULL
 ..$ tfi: NULL
```

```
summary(Lr)
```

Transitions:

To


From NRA ESRD Records: Events: Risk time: Persons:

NRA	48	77	125	77	1084.67	125
-----	----	----	-----	----	---------	-----

Make sure you know what the variables in `Lr` stand for.

3. Visualize the follow-up in a Lexis-diagram, by using the `plot` method for `Lexis` objects.

```
plot(Lr, col = "black", lwd = 3)
```



```
subset(Lr, age < 0)
```

```
lex.id    per    age tfi lex.dur lex.Cst lex.Xst id sex    dob    doe dor    dox
88 1989.34 -38.81  0    3.5    NRA    ESRD 586  M 2028.15 1989.34 NA 1992.84
event
1
```


What is wrong here? List the data for the person with negative entry age.

4. Correct the data and make a new plot, for example by:

```
Lr <- transform(Lr, age = ifelse(dob > 2000, age + 100, age),
  dob = ifelse(dob > 2000, dob - 100, dob))
subset(Lr, id == 586)
```

```
lex.id    per    age tfi lex.dur lex.Cst lex.Xst id sex    dob    doe dor    dox
88 1989.34 61.19  0    3.5    NRA    ESRD 586  M 1928.15 1989.34 NA 1992.84
event
1
```

```
plot(Lr, col = "black", lwd = 3)
```



5. Now make a Cox-regression analysis of ESRD occurrence with the variables sex and age at entry into the study, using time since entry to the study as time scale.

```
mc <- coxph(Surv(lex.dur, lex.Xst == "ESRD")
  ~ I(age / 10) + sex, data = Lr)
summary(mc)
```

Call:

```
coxph(formula = Surv(lex.dur, lex.Xst == "ESRD") ~ I(age/10) +
  sex, data = Lr)
```

n= 125, number of events= 77

	coef	exp(coef)	se(coef)	z	Pr(> z)
I(age/10)	0.5514	1.7357	0.1402	3.932	8.43e-05
sexF	-0.1817	0.8338	0.2727	-0.666	0.505

	exp(coef)	exp(-coef)	lower .95	upper .95
I(age/10)	1.7357	0.5761	1.3186	2.285
sexF	0.8338	1.1993	0.4886	1.423

Concordance= 0.612 (se = 0.036)

Likelihood ratio test= 16.07 on 2 df, p=3e-04

Wald test = 16.38 on 2 df, p=3e-04

Score (logrank) test = 16.77 on 2 df, p=2e-04

What is the hazard ratio between males and females? Between two persons who differ 10 years in age at entry?

6. The main focus of the paper was to assess whether the occurrence of remission (return to a lower level of albumin excretion, an indication of kidney recovery) influences mortality. *Remission* is a time-dependent

variable which is initially 0, but takes the value 1 when remission occurs. In order to handle this, each person who sees a remission must have two records:

- One record for the time before remission, where entry is `doe`, exit is `dor`, remission is 0, and event is 0.
- One record for the time after remission, where entry is `dor`, exit is `dox`, remission is 1, and event is 0 or 1 according to whether the person had an event at `dox`.

This is accomplished using the `cutLexis` function on the `Lexis` object, where we introduce a remission state `Rem`. Also use `split.state=TRUE` to have different ESRD states according to whether a person had had remission or not prior to ESRD. The statement to do this is:

```
Lc <- cutLexis(Lr, cut = Lr$dor, # where to cut follow up
               timescale = "per", # what timescale are we referring to
               new.state = "Rem", # name of the new state
               split.state = TRUE) # different states depending on previous
summary(Lc)
```

Transitions:

	To							
From	NRA	Rem	ESRD	ESRD(Rem)	Records	Events	Risk time	Persons
NRA	24	29	69	0	122	98	824.77	122
Rem	0	24	0	8	32	8	259.90	32
Sum	24	53	69	8	154	106	1084.67	125


List the records from a few select persons (choose values for `lex.id`, using for example `subset(Lc, lex.id %in% c(5,7,9))`).

- Now show how the states are connected and the number of transitions between them by using `boxes`. This is an interactive command that requires you to click in the graph window:

```
boxes(Lc)
```

It has a couple of fancy arguments, try:

```
boxes(Lc, boxpos = TRUE, scale.R = 100, show.BE = TRUE, hm = 1.5, wm = 1.5)
```




You may even be tempted to read the help page for `boxes.Lexis` ...

- Plot a Lexis diagram where different coloring is used for different segments of the follow-up. The `plot.Lexis` function draws a line for each record in the dataset, so you can index the coloring by `lex.Cst` and `lex.Xst` as appropriate — indexing by a factor corresponds to indexing by the *index number* of the factor levels, so you must be know which order the factor levels are in:

```
levels(Lc) # names and order of states in lex.Cst and lex.Xst
```

```
[1] "NRA"      "Rem"      "ESRD"     "ESRD(Rem)"
par(mai = c(3, 3, 1, 1) / 4, mgp = c(3, 1, 0) / 1.6)
plot(Lc, col = c("red", "limegreen")[Lc$lex.Cst],
      xlab = "Calendar time", ylab = "Age",
      lwd = 3, grid = 0:20 * 5, las = 1,
      xlim = c(1970, 2010), ylim = c(20, 70),
      xaxs = "i", yaxs = "i")
points(Lc, pch = c(NA, NA, 16, 16)[Lc$lex.Xst],
       col = c("red", "limegreen", "transparent", "transparent")[Lc$lex.Cst])
points(Lc, pch = c(NA, NA, 1, 1)[Lc$lex.Xst],
       col = "black", lwd = 2)
```



9. Make a Cox-regression of mortality rates (i.e. endpoint ESRD or ESRD(Rem)) with sex, age at entry and remission as explanatory variables, using time since entry as timescale, and include `lex.Cst` as time-dependent variable, and indicate that each record represents follow-up from `tfi` to `tfi+lex.dur`. Make sure that you know why what goes where here in the call to `coxph`.

```
(EP <- levels(Lc)[3:4])           # define EndPoint states

[1] "ESRD"      "ESRD(Rem)"

m1 <- coxph(Surv(tfi,               # entry time
               tfi + lex.dur,       # exit time
               lex.Xst %in% EP)     # event
            ~ sex + I((doe - dob - 50) / 10) + # fixed covariates
              (lex.Cst == "Rem"),      # time-dependent variable
            data = Lc)
summary(m1)
```

Call:

```
coxph(formula = Surv(tfi, tfi + lex.dur, lex.Xst %in% EP) ~ sex +
      I((doe - dob - 50)/10) + (lex.Cst == "Rem"), data = Lc)
```

n= 154, number of events= 77

	coef	exp(coef)	se(coef)	z	Pr(> z)
sexF	-0.05534	0.94616	0.27500	-0.201	0.840517
I((doe - dob - 50)/10)	0.52190	1.68522	0.13655	3.822	0.000132
lex.Cst == "Rem"TRUE	-1.26241	0.28297	0.38483	-3.280	0.001036

	exp(coef)	exp(-coef)	lower .95	upper .95
sexF	0.9462	1.0569	0.5519	1.6220
I((doe - dob - 50)/10)	1.6852	0.5934	1.2895	2.2024

```
lex.Cst == "Rem"TRUE      0.2830      3.5339      0.1331      0.6016
```

```
Concordance= 0.664 (se = 0.033 )
Likelihood ratio test= 30.31 on 3 df,  p=1e-06
Wald test              = 27.07 on 3 df,  p=6e-06
Score (logrank) test = 29.41 on 3 df,  p=2e-06
```

What is the effect of of remission on the rate of ESRD?

15.2 Splitting the follow-up time

In order to explore the effect of remission on the rate of ESRD, we split the data further into small pieces of follow-up. To this end we use the function `splitLexis`. The rates can then be modeled using a Poisson-model, and the shape of the effect of the underlying *rates* be explored. Furthermore, we can allow effects of both time since NRA and current age. To this end we will use splines, so we need the `splines` and also the `mgcv` packages.

- Now split the follow-up time every month after entry, and verify that the number of events and risk time is the same as before and after the split:

```
sLc <- splitLexis(Lc, "tfi", breaks = seq(0, 30, 1/12))
summary( Lc)
```

Transitions:

	To							
From	NRA	Rem	ESRD	ESRD(Rem)	Records:	Events:	Risk time:	Persons:
NRA	24	29	69	0	122	98	824.77	122
Rem	0	24	0	8	32	8	259.90	32
Sum	24	53	69	8	154	106	1084.67	125

```
summary(sLc)
```

Transitions:

	To							
From	NRA	Rem	ESRD	ESRD(Rem)	Records:	Events:	Risk time:	Persons:
NRA	9854	29	69	0	9952	98	824.77	122
Rem	0	3139	0	8	3147	8	259.90	32
Sum	9854	3168	69	8	13099	106	1084.67	125

- Now fit the Poisson-model corresponding to the Cox-model we fitted previously. The function `Ns()` produces a model matrix corresponding to a piece-wise cubic function, modeling the baseline hazard explicitly (think of the `Ns` terms as the baseline hazard that is not visible in the Cox-model). You can use the wrapper function `glm.Lexis`

```
mp <- glm.Lexis(sLc,
  ~ Ns(tfi, knots = c(0, 2, 5, 10)) +
    sex + I((doe - dob - 40) / 10) +
    I(lex.Cst == "Rem"))
```

stats::glm Poisson analysis of Lexis object sLc with log link:

Rates for transitions:

NRA->ESRD

Rem->ESRD(Rem)

```
ci.exp(mp)
```

	exp(Est.)	2.5%	97.5%
(Intercept)	0.01664432	0.003956666	0.07001685
Ns(tfi, knots = c(0, 2, 5, 10))1	5.18917655	1.949197027	13.81469029

```

Ns(tfi, knots = c(0, 2, 5, 10))2 34.20004199 1.764818735 662.75524463
Ns(tfi, knots = c(0, 2, 5, 10))3 4.43318269 2.179977108 9.01528219
sexF 0.91751162 0.536258443 1.56981691
I((doe - dob - 40)/10) 1.70082390 1.300813859 2.22384004
I(lex.Cst == "Rem")TRUE 0.27927558 0.131396852 0.59358233

```

How does the effects of sex change from the Cox-model?

12. Try instead using the `gam` function from the `mgcv` package. There is convenience wrapper for this for Lexis objects as well:

```

mx <- gam.Lexis(sLc,
  ~ s(tfi, k = 10) +
    sex + I((doe - dob - 40) / 10) +
    I(lex.Cst == "Rem"))

```

`mgcv::gam` Poisson analysis of Lexis object `sLc` with log link:

Rates for transitions:

NRA->ESRD

Rem->ESRD(Rem)

```
ci.exp(mp, subset = c("Cst", "doe", "sex"))
```

```

exp(Est.) 2.5% 97.5%
I(lex.Cst == "Rem")TRUE 0.2792756 0.1313969 0.5935823
I((doe - dob - 40)/10) 1.7008239 1.3008139 2.2238400
sexF 0.9175116 0.5362584 1.5698169

```

```
ci.exp(mx, subset = c("Cst", "doe", "sex"))
```

```

exp(Est.) 2.5% 97.5%
I(lex.Cst == "Rem")TRUE 0.2784659 0.1309446 0.5921838
I((doe - dob - 40)/10) 1.6992069 1.2995225 2.2218192
sexF 0.9309945 0.5435486 1.5946150

```

We see that there is virtually no difference between the two approaches in terms of the regression parameters.

13. Extract the regression parameters from the models using `ci.exp` and compare with the estimates from the Cox-model:

```
ci.exp(mx, subset = c("sex", "dob", "Cst"), pval = TRUE)
```

```

exp(Est.) 2.5% 97.5% P
sexF 0.9309945 0.5435486 1.5946150 0.7945394004
I((doe - dob - 40)/10) 1.6992069 1.2995225 2.2218192 0.0001066910
I(lex.Cst == "Rem")TRUE 0.2784659 0.1309446 0.5921838 0.0008970863

```

```
ci.exp(m1)
```

```

exp(Est.) 2.5% 97.5%
sexF 0.9461646 0.5519334 1.621985
I((doe - dob - 50)/10) 1.6852196 1.2895097 2.202360
lex.Cst == "Rem"TRUE 0.2829710 0.1330996 0.601599

```

```
round(ci.exp(mp, subset = c("sex", "dob", "Cst")) / ci.exp(m1), 2)
```

```

exp(Est.) 2.5% 97.5%
sexF 0.97 0.97 0.97
I((doe - dob - 40)/10) 1.01 1.01 1.01
I(lex.Cst == "Rem")TRUE 0.99 0.99 0.99


```

How large are the differences in estimated regression parameters?

- ```
plot(mx)
```

- ```
nd <- data.frame(tfi = seq(0, 20, 0.1),
```

```
plot = TRUE,
type = "l", lwd = 3:4, col = c("black", "forestgreen"),
log = "y", xlab = "Time since entry (years)",
ylab = "ESRD rate (per 100 PY) for 50 year old men")
```



Try to overlay with the corresponding prediction from the `glm` model using `Ns`.

Prediction from the multistate model

If we want to make proper statements about the survival and disease probabilities we must know not only how the occurrence of remission influences the rate of death/ESRD, but we must also model the occurrence rate of remission itself.

17. The rates of ESRD were modelled by a Poisson model with effects of age and time since NRA — in the models `mp` and `mx`. But if we want to model whole process we must also model the remission rates transition from NRA to Rem, but the number of events is rather small so we restrict covariates in this model to only time since NRA and sex. Note that only the records that represent follow-up in the NRA state should be used; this is most easily done using the `gam.Lexis` function

```
mr <- gam.Lexis(sLc, ~ s(tfi, k = 10) + sex,
               from = "NRA",
               to = "Rem")
```

`mgcv::gam` Poisson analysis of Lexis object `sLc` with log link:

Rates for the transition:

NRA->Rem

```
summary(mr)
```

Family: poisson

Link function: log

Formula:

```
cbind(trt(Lx$lex.Cst, Lx$lex.Xst) %in% trnam, Lx$lex.dur) ~ s(tfi,
  k = 10) + sex
```

Parametric coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-3.7025	0.2582	-14.339	<2e-16
sexF	0.9579	0.3728	2.569	0.0102

Approximate significance of smooth terms:

	edf	Ref.df	Chi.sq	p-value
s(tfi)	1.013	1.025	0.066	0.813

R-sq.(adj) = -5.65e-06 Deviance explained = 1.65%

UBRE = -0.96024 Scale est. = 1 n = 9952

```
ci.exp(mr, pval = TRUE)
```

	exp(Est.)	2.5%	97.5%	P
(Intercept)	0.02466174	0.01486718	0.04090901	1.254019e-46
sexF	2.60620470	1.25503844	5.41202779	1.019130e-02
s(tfi).1	1.00499489	0.89131271	1.13317662	9.351638e-01
s(tfi).2	0.99623769	0.80778743	1.22865188	9.718940e-01
s(tfi).3	0.99822247	0.91911268	1.08414140	9.663137e-01
s(tfi).4	1.00188999	0.89006998	1.12775801	9.750528e-01
s(tfi).5	0.99842904	0.92280785	1.08024715	9.687920e-01
s(tfi).6	0.99817367	0.90142303	1.10530865	9.719666e-01
s(tfi).7	1.00168704	0.92615949	1.08337380	9.663850e-01
s(tfi).8	0.99448155	0.68445450	1.44493689	9.768400e-01
s(tfi).9	0.94790052	0.63349476	1.41834700	7.946918e-01

What is the remission rate-ratio between men and women?

18. If we want to predict the probability of being in each of the three states using these estimated rates, we may resort to analytical calculations of the probabilities from the estimated rates, which is actually doable in this case, but which will be largely intractable for more complicated models. Alternatively we can *simulate* the life course for a large group of (identical) individuals through a model using the estimated rates. That will give a simulated cohort (in the form of a **Lexis** object), and we can then just count the number of persons in each state at each of a set of time points. This is accomplished using the function **simLexis**. The input to this is the initial status of the persons whose life-course we shall simulate, and the transition rates in suitable form:

- Suppose we want predictions for men aged 50 at NRA. The input is in the form of a **Lexis** object (where **lex.dur** and **lex.Xst** will be ignored). Note that in order to carry over the **time.scales** and the **time.since** attributes, we construct the input object using **subset** to select columns, and **NULL** to select rows (see the example in the help file for **simLexis**):

```
inL <- subset(sLc, select = 1:11)[NULL, ]
str(inL)
```

Classes 'Lexis' and 'data.frame': 0 obs. of 11 variables:

```
$ lex.id : int
$ per    : num
$ age    : num
$ tfi    : num
$ lex.dur: num
$ lex.Cst: Factor w/ 4 levels "NRA","Rem","ESRD",...
$ lex.Xst: Factor w/ 4 levels "NRA","Rem","ESRD",...
$ id     : num
```

```

$ sex      : Factor w/ 2 levels "M","F":
$ dob      : num
$ doe      : num
- attr(*, "time.scales")= chr [1:3] "per" "age" "tfi"
- attr(*, "time.since")= chr [1:3] "" "" ""
- attr(*, "breaks")=List of 3
..$ per: NULL
..$ age: NULL
..$ tfi: num [1:361] 0 0.0833 0.1667 0.25 0.3333 ...
timeScales(inL)

[1] "per" "age" "tfi"
inL[1, "lex.id"] <- 1
inL[1, "per"] <- 2000
inL[1, "age"] <- 50
inL[1, "tfi"] <- 0
inL[1, "lex.Cst"] <- "NRA"
inL[1, "lex.Xst"] <- NA
inL[1, "lex.dur"] <- NA
inL[1, "sex"] <- "M"
inL[1, "dob"] <- 1950
inL[1, "doe"] <- 2000
inL <- rbind(inL, inL)
inL[2, "sex"] <- "F"
inL

lex.id per age tfi lex.dur lex.Cst lex.Xst id sex dob doe
1 2000 50 0 NA NRA <NA> NA M 1950 2000
1 2000 50 0 NA NRA <NA> NA F 1950 2000
str(inL)

```

```

Classes 'Lexis' and 'data.frame': 2 obs. of 11 variables:
 $ lex.id : num 1 1
 $ per : num 2000 2000
 $ age : num 50 50
 $ tfi : num 0 0
 $ lex.dur: num NA NA
 $ lex.Cst: Factor w/ 4 levels "NRA","Rem","ESRD",...: 1 1
 $ lex.Xst: Factor w/ 4 levels "NRA","Rem","ESRD",...: NA NA
 $ id : num NA NA
 $ sex : Factor w/ 2 levels "M","F": 1 2
 $ dob : num 1950 1950
 $ doe : num 2000 2000
- attr(*, "breaks")=List of 3
..$ per: NULL
..$ age: NULL
..$ tfi: num [1:361] 0 0.0833 0.1667 0.25 0.3333 ...
- attr(*, "time.scales")= chr [1:3] "per" "age" "tfi"
- attr(*, "time.since")= chr [1:3] "" "" ""

```

The other input for the simulation is the models for the transitions. This is given as a list with an element for each transient state (that is `NRA` and `Rem`), each of which is again a list with names equal to the states that can be reached from the transient state. The content of the list will be `glm` objects, in this case the models we just fitted, describing the transition rates:

```

Tr <- list("NRA" = list("Rem" = mr,
                        "ESRD" = mx),

```

```
"Rem" = list("ESRD(Rem)" = mx))
```

With this as input we can now generate a cohort, using N=5 to simulate life course of 10 persons (5 for each set of starting values in inL):

```
(iL <- simLexis(Tr, inL, N = 10))
```

lex.id	per	age	tfi	lex.dur	lex.Cst	lex.Xst	id	sex	dob	doe	cens
1	2000.00	50.00	0.00	4.44	NRA	ESRD	NA	M	1950	2000	2020
2	2000.00	50.00	0.00	7.36	NRA	ESRD	NA	M	1950	2000	2020
3	2000.00	50.00	0.00	4.93	NRA	ESRD	NA	M	1950	2000	2020
4	2000.00	50.00	0.00	5.96	NRA	ESRD	NA	M	1950	2000	2020
5	2000.00	50.00	0.00	2.48	NRA	ESRD	NA	M	1950	2000	2020
6	2000.00	50.00	0.00	2.64	NRA	ESRD	NA	M	1950	2000	2020
7	2000.00	50.00	0.00	10.92	NRA	ESRD	NA	M	1950	2000	2020
8	2000.00	50.00	0.00	4.23	NRA	ESRD	NA	M	1950	2000	2020
9	2000.00	50.00	0.00	2.38	NRA	Rem	NA	M	1950	2000	2020
9	2002.38	52.38	2.38	6.71	Rem	ESRD(Rem)	NA	M	1950	2000	2020
10	2000.00	50.00	0.00	2.59	NRA	ESRD	NA	M	1950	2000	2020
11	2000.00	50.00	0.00	6.41	NRA	Rem	NA	F	1950	2000	2020
11	2006.41	56.41	6.41	5.00	Rem	ESRD(Rem)	NA	F	1950	2000	2020
12	2000.00	50.00	0.00	1.76	NRA	Rem	NA	F	1950	2000	2020
12	2001.76	51.76	1.76	18.24	Rem	Rem	NA	F	1950	2000	2020
13	2000.00	50.00	0.00	3.28	NRA	Rem	NA	F	1950	2000	2020
13	2003.28	53.28	3.28	15.06	Rem	ESRD(Rem)	NA	F	1950	2000	2020
14	2000.00	50.00	0.00	0.84	NRA	ESRD	NA	F	1950	2000	2020
15	2000.00	50.00	0.00	12.17	NRA	ESRD	NA	F	1950	2000	2020
16	2000.00	50.00	0.00	7.91	NRA	Rem	NA	F	1950	2000	2020
16	2007.91	57.91	7.91	12.09	Rem	Rem	NA	F	1950	2000	2020
17	2000.00	50.00	0.00	5.00	NRA	Rem	NA	F	1950	2000	2020
17	2005.00	55.00	5.00	5.94	Rem	ESRD(Rem)	NA	F	1950	2000	2020
18	2000.00	50.00	0.00	14.14	NRA	ESRD	NA	F	1950	2000	2020
19	2000.00	50.00	0.00	9.39	NRA	ESRD	NA	F	1950	2000	2020
20	2000.00	50.00	0.00	15.33	NRA	Rem	NA	F	1950	2000	2020
20	2015.33	65.33	15.33	4.67	Rem	Rem	NA	F	1950	2000	2020

```
summary(iL, by = "sex")
```

\$M

Transitions:

	To								
From	NRA	Rem	ESRD	ESRD(Rem)	Records:	Events:	Risk time:	Persons:	
NRA	0	1	9	0	10	10	47.93	10	
Rem	0	0	0	1	1	1	6.71	1	
Sum	0	1	9	1	11	11	54.64	10	

\$F

Transitions:

	To								
From	NRA	Rem	ESRD	ESRD(Rem)	Records:	Events:	Risk time:	Persons:	
NRA	0	6	4	0	10	10	76.24	10	
Rem	0	3	0	3	6	3	61.01	6	
Sum	0	9	4	3	16	13	137.25	10	

What type of object have you got as iL?

19. Now generate the life course of, say, 5,000 persons, and look at the summary. The `system.time` command

is just to tell you how long it took, you may want to start with 500 just to see how long that takes.

```
system.time(sM <- simLexis(Tr, inL, N = 500, t.range = 12))
```

```
   user  system elapsed
2.737   3.457   2.170
```

```
summary(sM, by = "sex")
```

\$M

Transitions:

	To							
From	NRA	Rem	ESRD	ESRD(Rem)	Records	Events	Risk time	Persons
NRA	42	90	368	0	500	458	2666.1	500
Rem	0	48	0	42	90	42	585.0	90
Sum	42	138	368	42	590	500	3251.1	500

\$F

Transitions:

	To							
From	NRA	Rem	ESRD	ESRD(Rem)	Records	Events	Risk time	Persons
NRA	35	171	294	0	500	465	2579.15	500
Rem	0	102	0	69	171	69	1140.51	171
Sum	35	273	294	69	671	534	3719.66	500

Why are there so many ESRD-events in the resulting data set?

20. Now count how many persons are present in each state at each time for the first 10 years after entry (which is at age 50). This can be done by using `nState`. Try:

```
nStm <- nState(subset(sM, sex == "M"), time.scale = "age",
               at = seq(0, 10, 0.1),
               from = 50)
nStf <- nState(subset(sM, sex == "F"), time.scale = "age",
               at = seq(0, 10, 0.1),
               from = 50)
head(nStf, 15)
```

	State			
when	NRA	Rem	ESRD	ESRD(Rem)
50	500	0	0	0
50.1	494	5	1	0
50.2	485	10	5	0
50.3	481	11	8	0
50.4	479	12	9	0
50.5	472	16	12	0
50.6	467	19	14	0
50.7	463	21	16	0
50.8	457	25	18	0
50.9	451	29	20	0
51	445	33	22	0
51.1	441	35	24	0
51.2	436	38	26	0
51.3	430	40	30	0
51.4	426	43	30	1

What is in the object `nStf`?

21. With the counts of persons in each state at the designated time points (in `nStm`), compute the cumulative

fraction over the states, arranged in order given by `perm`:

```
ppm <- pState(nStm, perm = c(2, 1, 3, 4))
ppf <- pState(nStf, perm = c(2, 1, 3, 4))
head(ppf)
```

	State			
when	Rem	NRA	ESRD	ESRD(Rem)
50	0.000	1.000	1	1
50.1	0.010	0.998	1	1
50.2	0.020	0.990	1	1
50.3	0.022	0.984	1	1
50.4	0.024	0.982	1	1
50.5	0.032	0.976	1	1


```
tail(ppf)
```

	State			
when	Rem	NRA	ESRD	ESRD(Rem)
59.5	0.234	0.382	0.904	1
59.6	0.234	0.380	0.902	1
59.7	0.236	0.378	0.902	1
59.8	0.236	0.376	0.902	1
59.9	0.232	0.372	0.898	1
60	0.232	0.370	0.898	1

What do the entries in `ppf` represent?

22. Try to plot the cumulative probabilities using the `plot` method for `pState` objects:


```
plot(ppf)
```



Is this useful?

23. Now try to improve the plot so that it is easier to read, and easier to compare between men and women, for example:

```
par(mfrow = c(1, 2))
# Men
plot(ppm, col = c("limegreen", "red", "#991111", "forestgreen"))
lines(as.numeric(rownames(ppm)), ppm[, "NRA"], lwd = 2)
text(59.5, 0.95, "Men", adj = 1, col = "white", font = 2, cex = 1.2)
axis(side = 4, at = 0:10 / 10)
axis(side = 4, at = 1:99 / 100, labels = NA, tck = -0.01)
# Women
plot(ppf, col = c("limegreen", "red", "#991111", "forestgreen"),
      xlim = c(60, 50)) # inverted x-axis
lines(as.numeric(rownames(ppf)), ppf[, "NRA"], lwd = 2)
text(59.5, 0.95, "Women", adj = 0, col = "white", font = 2, cex = 1.2)
axis(side = 2, at = 0:10 / 10)
axis(side = 2, at = 1:99 / 100, labels = NA, tck = -0.01)
```



What is the 10-year risk of remission for men and women respectively?